

MANCHESTER
1824

COMP27112

Computer
Graphics
and
Image Processing

5: Viewing 1: The camera

Toby.Howard@manchester.ac.uk

The diagram illustrates the camera viewing frustum and coordinate systems. A black digital SLR camera is positioned at the origin of a 3D coordinate system. The camera's optical axis is defined by the vertical red dashed line labeled \hat{U}' . The horizontal axis is labeled \hat{S}' , and the depth axis is labeled \hat{F} . A black cow is located in the scene. A red dashed rectangular box, representing the camera's field of view, is centered on the cow. The vertices of this box are labeled c (top-left), \hat{R} (top-right), \hat{L} (bottom-left), and \hat{S} (bottom-right). The camera is mounted on a black tripod, which is supported by a grey checkered plane representing the ground.

MANCHESTER
1824

Introduction

- In part 1 of our study of Viewing, we'll look at
 - Viewing in 2D
 - Clipping
 - Viewing in 3D
 - Viewing in OpenGL

The diagram illustrates a 3D viewing pipeline. At the top right, a wireframe teapot is labeled "Teapot at (0,0,0), but then rotated by 20° about Y, and then translated by (0,0,0.2);". Above it, a box contains the text "'Camera' at (10,10,10), looking at (0,0,0)". A green vertical line extends from the camera position through the teapot to the eye point. A red line represents the optical axis. A blue line shows the projection of the eye point onto the view plane. A dotted arrow points from the camera box to the teapot. A solid arrow points from the camera box to the eye point. A dashed arrow points from the eye point to the teapot.

2

MANCHESTER
1824

Viewing in 2D

- We'll start with looking at viewing in 2D

The diagram shows a house in world coordinates on the left, represented by a red X-Y coordinate system. A question mark points from this to a monitor on the right. The monitor displays the same house, labeled 'Screen coordinates', with a blue U-V coordinate system. The monitor is a Dell model.

- We need to specify **what** we want to see and **where** we want to see it

3

MANCHESTER
1824

Viewing in 2D

- How do we specify the mapping from our scene to the display screen?

The diagram shows a house in world coordinates on the left, represented by a red X-Y coordinate system. A question mark points from this to a monitor on the right. The monitor displays the same house, labeled 'Screen coordinates', with a blue U-V coordinate system. The monitor is a Dell model.

- We use the **analogy** of photographing the scene with a camera
- In this case the camera is confined to the 2D plane

4

MANCHESTER
1824

Viewing in 2D

- We specify a “window” in world coordinates, and a “viewport” in screen coordinates

World coordinates

Window

X

Y

Viewport

U

V

Screen coordinates

M_{view}

5

MANCHESTER
1824

Window to viewport mapping

- We can find M_{view} easily, in 3 steps
 - M1** : Translate by $(-x_0, -y_0)$ to place the window at the origin
 - M2** : Scale the window to be the same shape as the viewport
 - M3** : Shift to the viewport position

World coordinates

Window

X

Y

Viewport

U

V

Screen coordinates

M_{view}

(x_0, y_0)

(x_1, y_1)

(u_0, v_0)

(u_1, v_1)

6

MANCHESTER
1824

Window to viewport mapping

- **M1** : Translate $(-x_0, -y_0)$
- **M2** : Scale $(u_1-u_0/x_1-x_0, v_1-v_0/y_1-y_0)$
- **M3** : Translate (u_0, v_0)
- $\mathbf{M}_{\text{view}} = \mathbf{M}_3 * \mathbf{M}_2 * \mathbf{M}_1$
- $\mathbf{P}_{\text{screen}} = \mathbf{M}_{\text{view}} * \mathbf{P}_{\text{world}}$

World coordinates

Viewport

Screen coordinates

7

MANCHESTER
1824

Window to viewport mapping

- $\mathbf{M}_{\text{view}} = \mathbf{M}_3 * \mathbf{M}_2 * \mathbf{M}_1$

$$\begin{bmatrix} 1 & 0 & u_0 \\ 0 & 1 & v_0 \\ 0 & 0 & 1 \end{bmatrix} * \begin{bmatrix} \frac{u_1 - u_0}{x_1 - x_0} & 0 & 0 \\ 0 & \frac{v_1 - v_0}{y_1 - y_0} & 0 \\ 0 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -x_0 \\ 0 & 1 & -y_0 \\ 0 & 0 & 1 \end{bmatrix}$$

- So we can now map a world space point $\mathbf{P}_{\text{world}}$ into screen coordinates $\mathbf{P}_{\text{screen}}$ as follows. This is our "view".

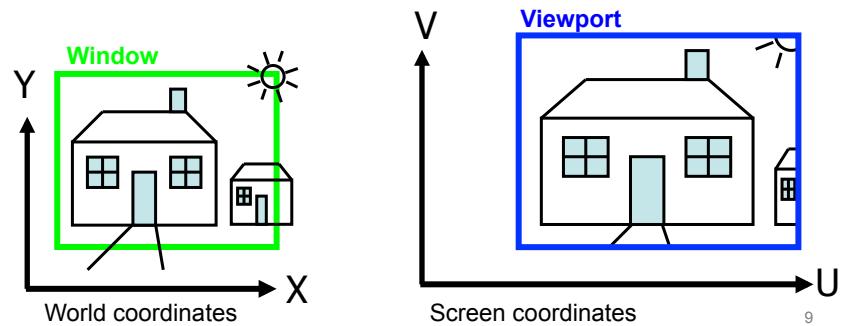
$$\mathbf{P}_{\text{screen}} = \begin{bmatrix} \frac{u_1 - u_0}{x_1 - x_0} & 0 & -x_0 * \frac{u_1 - u_0}{x_1 - x_0} + u_0 \\ 0 & \frac{v_1 - v_0}{y_1 - y_0} & -y_0 * \frac{v_1 - v_0}{y_1 - y_0} + v_0 \\ 0 & 0 & 1 \end{bmatrix} * \mathbf{P}_{\text{world}}$$

Note: We rarely multiply matrices by hand like this. The graphics system will do it for us. So no need for you to remember this complicated matrix.

8

Clipping

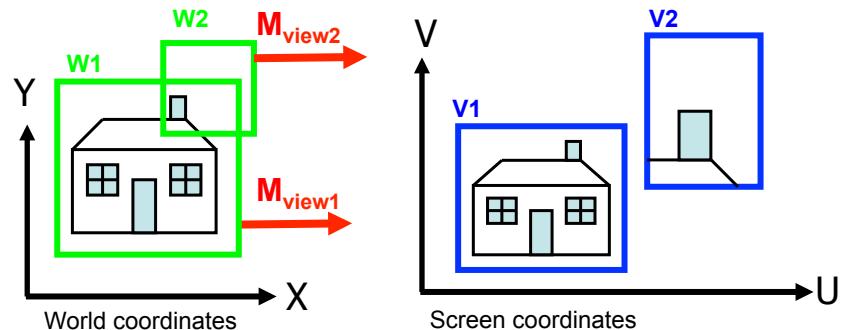
- Normally we will want to **CLIP** against the viewport
 - ...to remove those parts of primitives whose coordinates are outside the window
 - There are standard algorithms for clipping lines and polygons



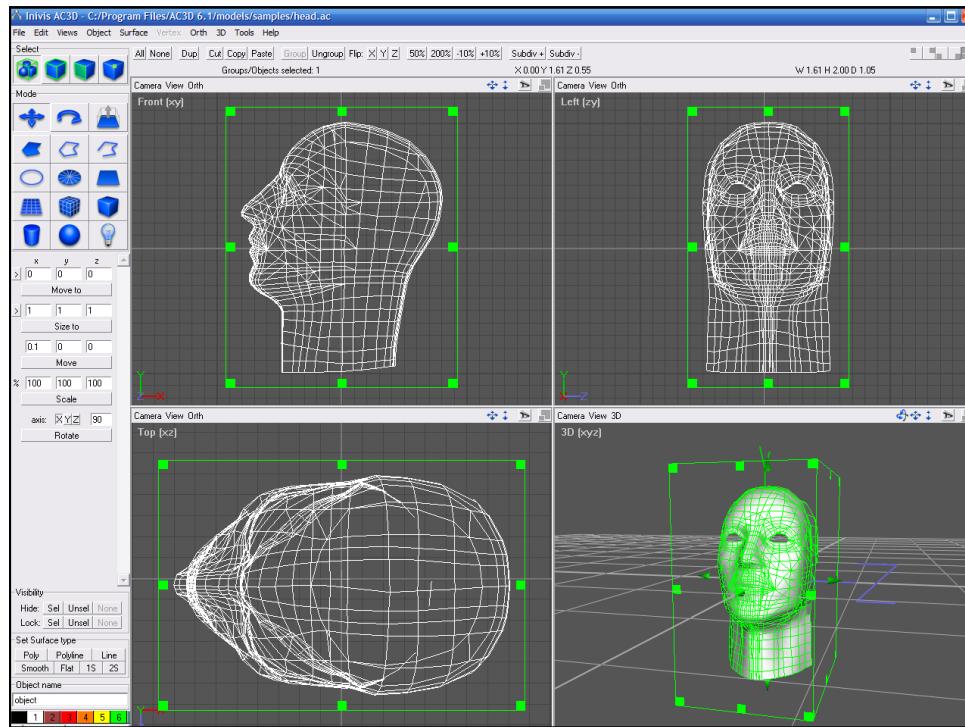
9

Multiple windows and viewports

- Sometimes it's useful to use multiple windows and viewports, to help arrange items on the screen



10



MANCHESTER
1824

Viewing in 2D: summary

- We use the analogy of photographing our scene with a 2D “camera” which can slide in the XY plane
- We compute a **viewing transformation M_{view}**
- $P_{screen} = M_{view} * P_{world}$

MANCHESTER
1824

There is no camera

- Notice what we have done
- We have imagined we have a camera...
- ...but we don't really have one
- So we have simulated a camera by instead **transforming the model**

The diagram shows a house in world coordinates on the left, defined by a green rectangle labeled 'Window'. A red arrow labeled M_{view} points from the window to a blue rectangle on the right labeled 'Screen coordinates'. The screen coordinates also have a vertical axis labeled 'Y' and a horizontal axis labeled 'U'. The house's position is transformed from its original location in world coordinates to a new position on the screen.

13

MANCHESTER
1824

Viewing in 3D

- In 2D graphics, we view our world by mapping from 2D world coordinates to 2D screen coordinates: easy and obvious
- In 3D graphics, in order to view our 3D world, we have to somehow **reduce** our 3D information to 2D information, so that it can be displayed on the 2D display: **not** so easy and obvious
- Here we see a 2D view of an object defined in 3D. It's been projected from 3D to 2D.
- To specify how this view is created, we again use the analogy of "taking a picture using a camera", but this time our "camera" is like a real-world camera: It has a position and orientation in 3D space, and a particular type of lens.

A wireframe rendering of a teapot, showing its internal structure and shape. This represents a 3D object that has been projected onto a 2D plane, similar to how a camera captures a 3D scene.

14

The camera analogy

- The process of transforming a synthetic 3D model into a 2D view is analogous to using a camera in the real world to take 2D pictures of a 3D scene



15

The camera analogy



Step 1: Arrange the scene into the desired composition

16

MANCHESTER
1824

The camera analogy



Step 2: Position and point the camera at the scene

17

MANCHESTER
1824

The camera analogy



Step 3: Choose a camera lens (wide-angle? zoom?)

18

MANCHESTER
1824

The camera analogy

Step 4: Decide the size of the final photograph

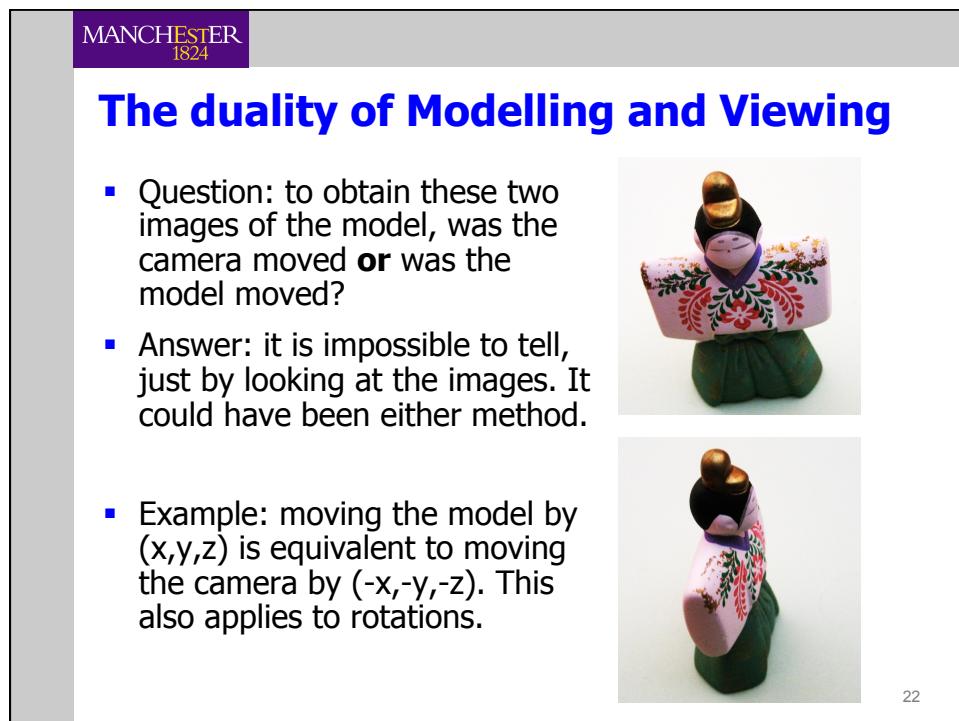
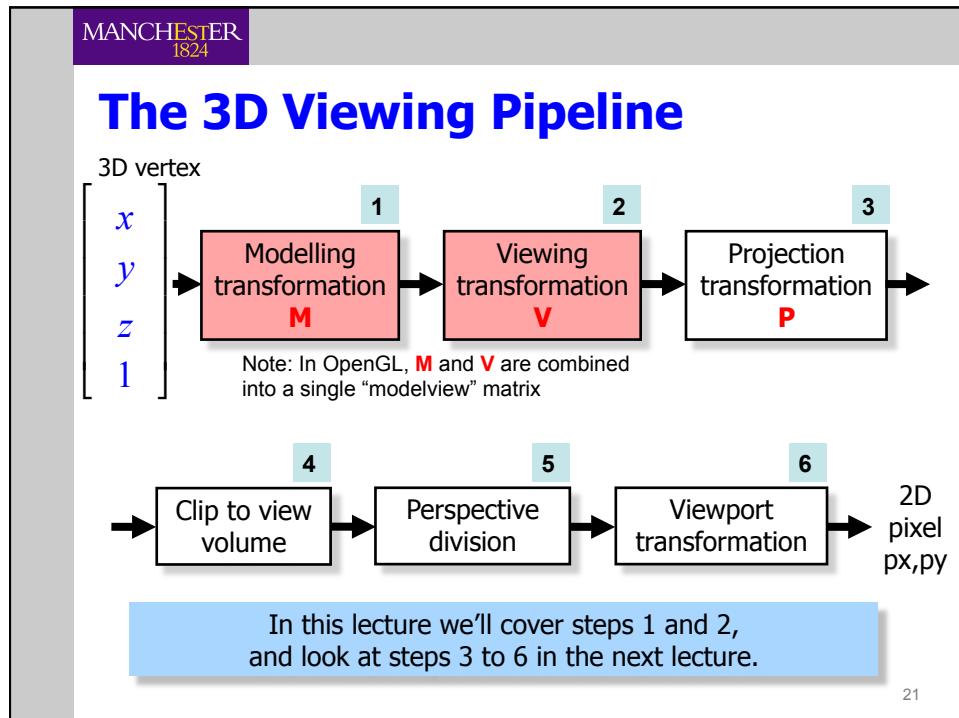
19

MANCHESTER
1824

3D Viewing: the camera analogy

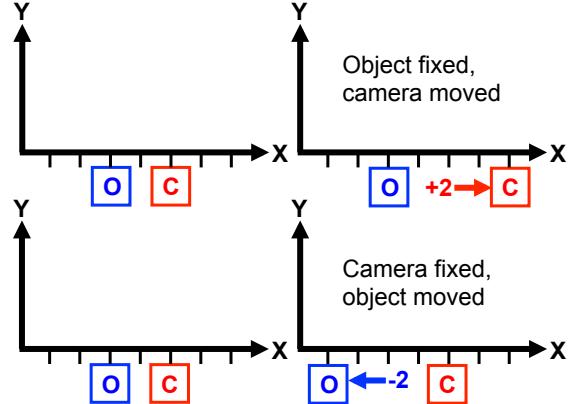
	Real World	Computer Graphics
1	Arrange the scene into the desired composition	Set Modelling Transformation
2	Point the camera at the scene	Set Viewing Transformation
3	Choose the camera lens, or adjust the zoom	Set Projection Transformation
4	Determine the size and shape of the final photograph	Set Viewport Transformation

20



The duality of Modelling and Viewing

- Here, in 2D, we have an object **O** and a camera **C**, both sitting on the X axis



- If we keep the object fixed and move the camera by **+2**, **O** and **C** are now **4** units apart.

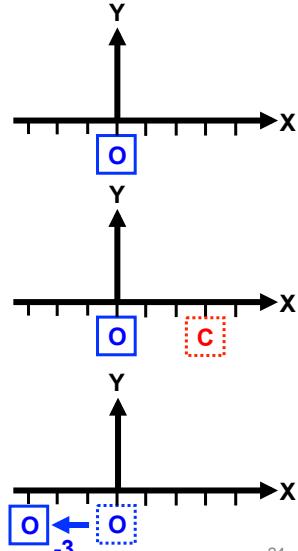
- If we keep the camera fixed and move the object by **-2**, **O** and **C** are now **4** units apart.

- Whether we move **O** or **C**, their **relative positions** will be the same, so the view from the camera will be the same

23

The duality of Modelling and Viewing

- Now for the **KEY IDEA**, which we'll present in 2D
- Imagine we have an object **O** sitting at the origin



- We want to view **O** with a camera **C** located at position **X=3**
- But we don't actually have a camera!

- But we can **SIMULATE** the effect, by instead moving the object by **-3 in X**

24

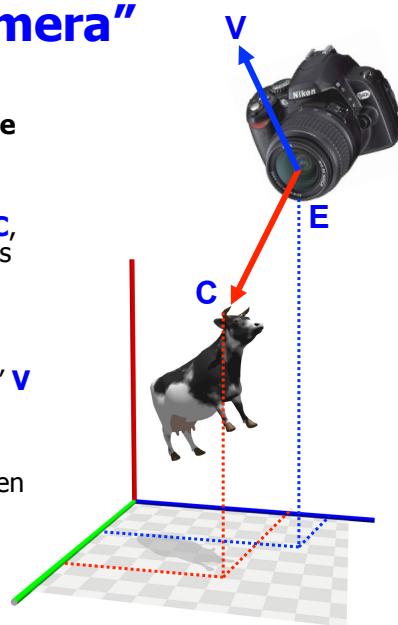
Achieving viewing by modelling

- We've just seen that we can create the same view from a camera at a certain location and orientation, by **instead** transforming the object
- This is exactly what we do in computer graphics
- However, the idea of having a camera is very natural to us, so we **pretend** we really have a camera...
- ...and we express the view we want in terms of "camera location and orientation", but to implement this we actually compute a suitable viewing transformation which we **apply to the object**

25

Specifying the “camera”

- We specify where the “camera” is located in 3D space. This is the **eye point, E**
- We specify a **centre of interest C**, a 3D point at which the “camera” is looking
- We specify the **up** direction of the “camera”, using a “view up vector” **V**
- We can then use **E, C**, and **V** to derive a transformation which, when applied to the **model**, would give the same view as if we really had this “camera”.



MANCHESTER
1824

Defining the camera coordinate system

- We use \mathbf{E} , \mathbf{C} and \mathbf{V} to derive a coordinate system for the "camera"
- We'll call the axes of the camera coordinate system $\hat{\mathbf{s}}$ $\hat{\mathbf{u}}$ $\hat{\mathbf{f}}$

MANCHESTER
1824

Defining the camera coordinate system

- Let's start with $\hat{\mathbf{f}}$. This is straightforward to define, in two steps:
- Step 1. $\mathbf{f} = \mathbf{C} - \mathbf{E}$
- Step 2. The length of \mathbf{f} could be anything, depending on \mathbf{C} and \mathbf{E} , and since we are using \mathbf{f} to specify a direction only, we normalize it to be of length 1, and then call it $\hat{\mathbf{f}}$

Defining the camera coordinate system

- Now we consider the "up" direction of the camera. The user must define this (just like they must define C and E) using a vector \mathbf{V} , called the "view-up" vector.
- Here are some possible values of \mathbf{V} :



$$\mathbf{V} = (0, 1, 0)$$



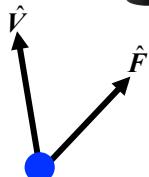
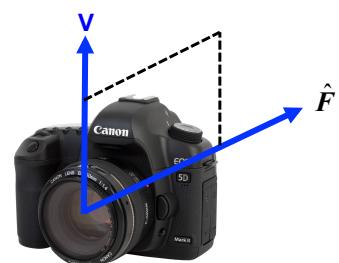
$$\mathbf{V} = (-1, 1, 0)$$



$$\mathbf{V} = (-1, 0, 0)$$

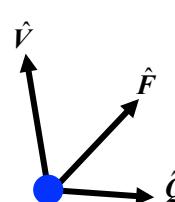
Defining the camera coordinate system

- Let's continue defining the $\hat{\mathbf{S}} \hat{\mathbf{U}} \hat{\mathbf{F}}$ coordinate system of the camera. So far we have one axis, $\hat{\mathbf{F}}$, so let's try and define another axis.
- If \mathbf{V} is the 'up' direction of the camera, you would think that \mathbf{V} would be orthogonal to \mathbf{F}
- OK, let's assume it is, then let's normalise it, to give $\hat{\mathbf{V}}$, and we now have defined 2 of the three axes of the camera.
- We will come back to the red box above, shortly...



30

Defining the camera coordinate system

- Now all that remains is to define the 3rd axis, which needs to be orthogonal to both \hat{V} and \hat{F} .
 - This is straightforward, we just take the cross product of \hat{V} and \hat{F} . Let's call this \hat{Q} .
 - $\hat{Q} = \text{normalise}(\hat{V} \times \hat{F})$
- 
- So now we have established the camera's coordinate system.
 - Yes, but there's a potential problem...

31

Defining the camera coordinate system

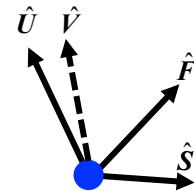
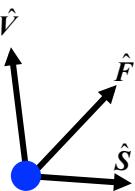
- This system is based on the assumption we stated:

If V is the 'up' direction of the camera, you would think that V would be orthogonal to F
- In other words, we are **assuming** that the user has made sure that V is actually orthogonal to F . To do this, they would first have to work out F , and then define V accordingly.
- Can we rely on all users to always do that? No!
- But if they don't, and instead specify a V that is not orthogonal to F , then the whole derivation of the camera coordinate system will not work, and the axes will point in all kinds of strange directions.
- So what's the solution?

32

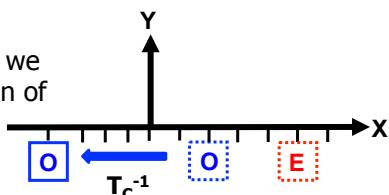
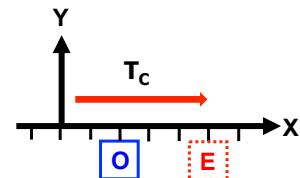
Defining the camera coordinate system

- The solution is to decouple \hat{V} and \hat{F} , by **not making any assumptions** about their relationship
- We first derive \hat{F} as before: $\hat{F} = \text{normalize}(C - E)$
- Then we use \hat{F} and \hat{V} to create \hat{S} orthogonal to them both: $\hat{S} = \text{normalise}(\hat{F} \times \hat{V})$
- Then we use \hat{S} and \hat{F} to create \hat{U} : $\hat{U} = \text{normalise}(\hat{S} \times \hat{F})$ and we **forget all about \hat{V}**
- We've constructed a camera coordinate system $\hat{S} \hat{U} \hat{F}$ that has involved **making no assumptions**, and is guaranteed to be a correct orthogonal coordinate system.



The duality principle applied

- Imagine we have a real camera **C** in the world, at location **E** in the world and rotated by **R**. Call the composite transformation needed to move the camera, from the origin, to where it is, T_c
- Using this camera we photograph an object **O** in the world, and get a view
- We can achieve the **same view** if we compute the inverse transformation of T_c (i.e., T_c^{-1}) ...
- ... and apply this to **O**



34

MANCHESTER
1824

The viewing transformation

- The viewing transformation T_c^{-1} moves & rotates the $\hat{S} \hat{U} \hat{F}$ camera system to align with the origin of the world coordinates system:

The diagram illustrates the viewing transformation between two coordinate systems. On the left, a 'world coordinate system' is shown with three axes: X (blue), Y (red), and Z (green). A point labeled 'E' (the eye) is located at the origin of this system. A camera coordinate system is defined by three unit vectors: \hat{U} (upward, black), \hat{S} (forward, black), and \hat{F}' (right, black). The origin of this system is at point 'E'. A red arrow labeled ' T_c^{-1} translates and rotates' points from the world coordinate system towards the camera coordinate system. The camera coordinate system is also labeled as the "camera" coordinate system.

- KEY IDEA:** If we apply T_c^{-1} to **objects**, we will get the **same view** as if we had a real camera at **E**

35

The viewing transformation in OpenGL

- This OpenGL function takes (**E,C,V**) and computes the viewing transformation we have just seen

```
void gluLookAt (GLdouble eyex,  
                GLdouble eyey,  
                GLdouble eyez,  
                GLdouble centrex,  
                GLdouble centrey,  
                GLdouble centrez,  
                GLdouble upx,  
                GLdouble upy,  
                GLdouble upz);
```

- As we have seen with other OpenGL transformation functions, `gluLookAt()` creates a temporary matrix **T**, and then multiplies the modelview matrix by **T**: $\mathbf{M}_{\text{modelview}} = \mathbf{M}_{\text{modelview}} * \mathbf{T}$

37

The viewing transformation in OpenGL

- The viewing transformation specifies the location and orientation of the “camera” (by in fact transforming the model)
- We incorporate this transformation into the modelview matrix as follows:

```
glMatrixMode(GL_MODELVIEW);  
glLoadIdentity(); // M= identity matrix (I)  
gluLookAt(...stuff...) // M is now I * VIEW
```

- Because we want the viewing transformation to take place **AFTER** any true modelling transformations, we need to “pre-load” the modelview matrix with the viewing transformation...
- ...And then all subsequent modelling transformations will get multiplied into the modelview matrix

38

Modelling and viewing together

```
// First set the viewing transformation
glMatrixMode(GL_MODELVIEW);
glLoadIdentity(); // M= identity matrix (I)
gluLookAt(...stuff...) // M is now I * VIEW

// Now draw a transformed teapot
glTranslatef(tx, ty, tz);
// OpenGL computes temp translation matrix T,
// then sets M= M x T, so now M is (VIEW x T)
glRotatef(theta, 0.0, 1.0, 0.0);
// OpenGL computes temp rotation matrix R,
// then sets M= M x R, so M is now (VIEW x T x R)
glutWireTeapot(1.0);
```

- So all points P will be transformed first by R, then T, then VIEW

39

Example: Setting a view in OpenGL

- Here's a real fragment showing the use of a view transformation and a modelling transformation together
- Note that we also need to set the **projection**, but we'll cover that in the next lecture, so ignore it for now.

```
glMatrixMode(GL_MODELVIEW); // select modelview matrix

glLoadIdentity(); // initialise it

// set the projection (see next lecture)
gluPerspective(...stuff...);

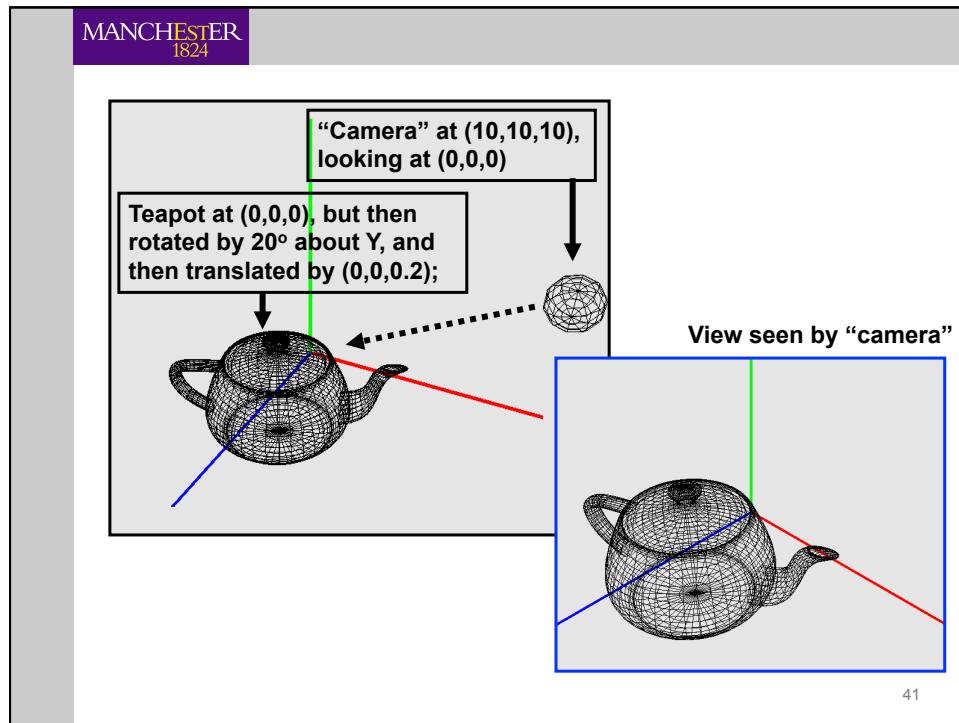
// set the view transformation
gluLookAt(10,10,10, 0,0,0, 0,1,0);

// move/rotate the model however we want
glTranslatef(0.0, 0.0, 0.2);
glRotatef(20.0, 0.0, 1.0, 0.0);

glutWireTeapot(3.0); // draw it
```

See the next slide for a visualisation of this.

40



The screenshot shows a demonstration interface with the "MANCHESTER 1824" logo at the top. The main content area is titled "Demonstrations" in blue.

- World-space view:** Shows a car model with a coordinate system and green lines representing the camera's frustum.
- Screen-space view:** Shows the same car model centered in a square frame.
- World-space view:** Shows a character model with a coordinate system and green lines representing the camera's frustum.
- Screen-space view:** Shows the same character model centered in a square frame.

Below each view are command manipulation windows:

- World-space view:**

```
glTranslatef( 0.00 , 0.00 , 0.00 );
glRotatef( 0.0 , 0.00 , 1.00 , 0.00 );
glScalef( 1.00 , 1.00 , 1.00 );
glBegin( ... );
...
Click on the arguments and move the mouse to modify values.
```
- Screen-space view:**

```
fovy aspect zNear zFar
gluPerspective( 60.0 , 1.00 , 1.0 , 10.0 );
gluLookAt( 0.00 , 0.00 , 2.00 ,
           0.00 , 0.00 , 0.00 ,
           0.00 , 1.00 , 0.00 ); <- up
Click on the arguments and move the mouse to modify values.
```

At the bottom right is the number "42".