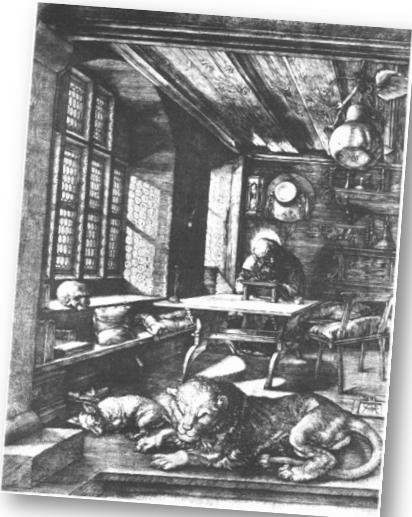


MANCHESTER
1824

COMP27112

Computer Graphics and Image Processing



6: Viewing 2: Projections

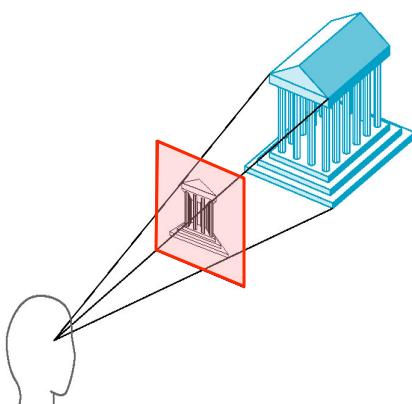
Toby.Howard@manchester.ac.uk

1

MANCHESTER
1824

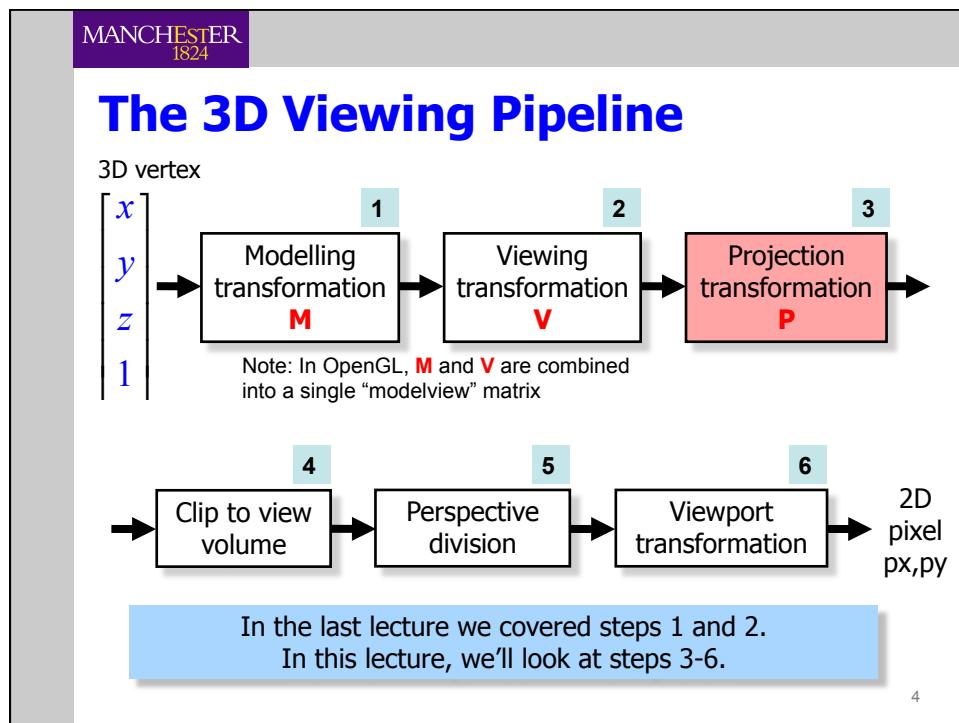
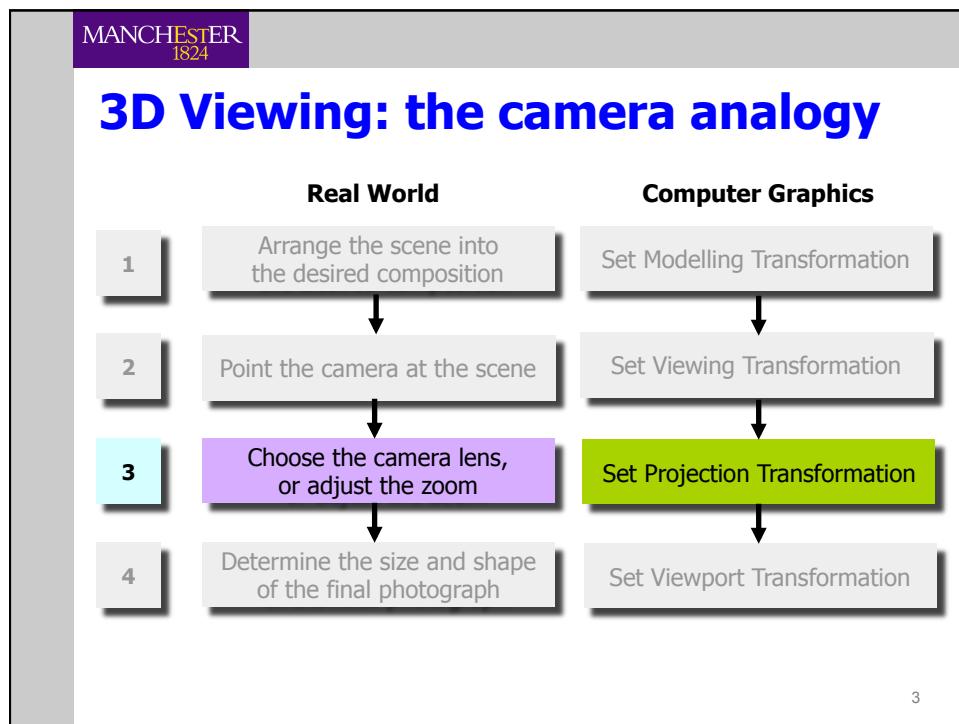
Introduction

- In part 2 of our study of Viewing, we'll look at
 - The theory of geometrical planar projections
 - Classes of projections
 - Perspective
 - Parallel
 - Basic mathematics of projections
 - Matrix representations of projections



Temple figures courtesy: Prof. Ed Angel, University of New Mexico

2

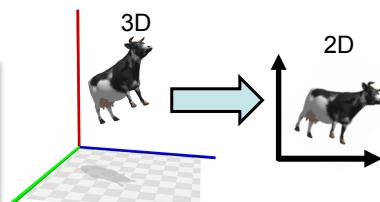


MANCHESTER
1824

Planar geometric projections

- How do we specify the mapping from 3D coordinates to 2D coordinates on the screen?
- We clearly have to somehow convert all the 3D coordinates into suitable 2D coordinates.
- This conversion process is called **projection**.
- There are two classes of planar geometric projection:
 - Parallel projection
 - Perspective projection

Note: We're considering only **planar geometric projections**: the kind that project lines to lines. Other kinds of projections change lines into curves and vice versa (used in cartography, for example), which we will not study.



5

MANCHESTER
1824

Planar Geometric Projections

```
graph TD; Parallel[Parallel] --> Orthographic[Orthographic]; Parallel --> Oblique[Oblique]; Parallel --> Axonometric[Axonometric]; Parallel --> Isometric[Isometric]; Parallel --> Dimetric[Dimetric]; Parallel --> Trimetric[Trimetric]; Parallel --> Cavalier[Cavalier]; Parallel --> Cabinet[Cabinet]; Perspective[Perspective] --> OnePoint[1-point]; Perspective --> TwoPoint[2-point]; Perspective --> ThreePoint[3-point]
```

In Computer Graphics, **perspective** projections are common, for their sense of "realism".

Parallel projections are often used in CAD and Engineering Drawing, because they allow precise measurements to be taken from the image.

MANCHESTER
1824

Parallel projection

The **projection** is the set of points at which the **projectors** **intersect** the projection plane.

Parallel edges remain parallel in the projection. Angles between edges may be **distorted**.

The diagram illustrates parallel projection. A blue 3D rectangular prism represents the '3D object'. It is positioned above a red parallelogram representing the '2D projection plane, oriented in 3D space'. Several dashed lines, labeled 'Projectors', extend from the vertices of the 3D object through the plane to its surface. The intersection points on the plane represent the 'projection' of the object's features. The text notes that parallel edges remain parallel in the projection, while angles between edges may be distorted.

Centre of projection, or “eye point”, is at infinity, so projectors are parallel

7

MANCHESTER
1824

Parallel projection

The diagram illustrates parallel projection using a classical building as the '3D object'. The building is shown in perspective on the left, with its columns and pediment. Projectors, represented by black lines, extend from the building's features through a red parallelogram representing the '2D projection plane, oriented in 3D space' to intersect it on the right. The text notes that the 'Centre of projection, or “eye point”, is at infinity, so projectors are parallel'.

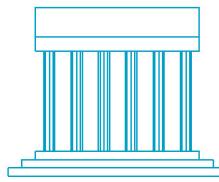
8

Parallel projection: orthographic

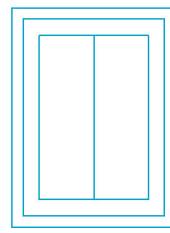
- In orthographic parallel projection, projectors are **perpendicular** to the projection plane
- The projection plane is **parallel** to one of the (X,Y,Z) axes of the 3D object being viewed
- The projected image shows only 2 of the 3 axes
- There is no distortion of lengths or angles, so well-suited for engineering drawing and CAD



Front



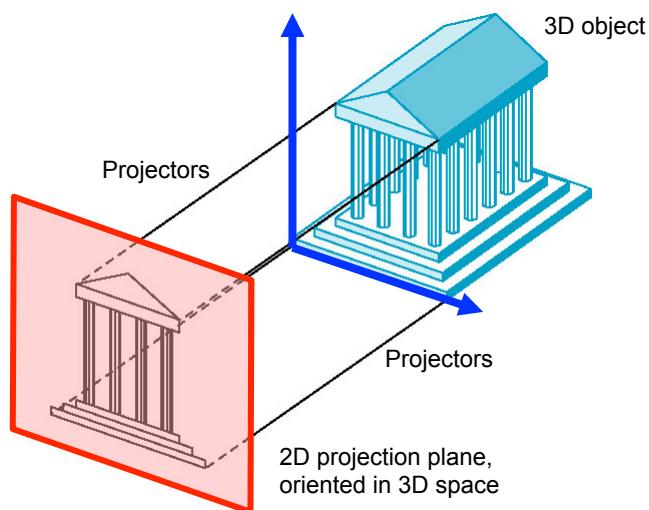
Side



Top

9

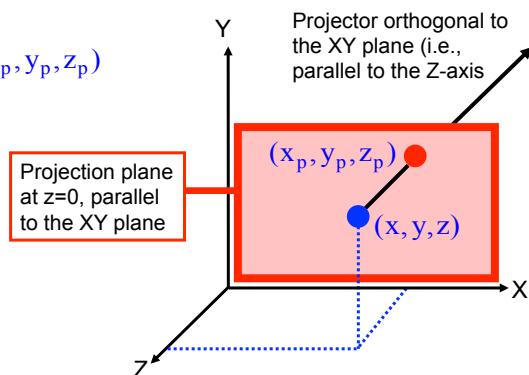
Parallel projection



10

Computing orthographic projection

- Suppose we have a projection plane located at $z=0$, and the plane is parallel to the XY plane
- A 3D point (x, y, z) will project to (x_p, y_p, z_p) on the projection plane
- We want to find (x_p, y_p, z_p)
- By definition:
 - $x_p = x$
 - $y_p = y$
 - $z_p = 0$



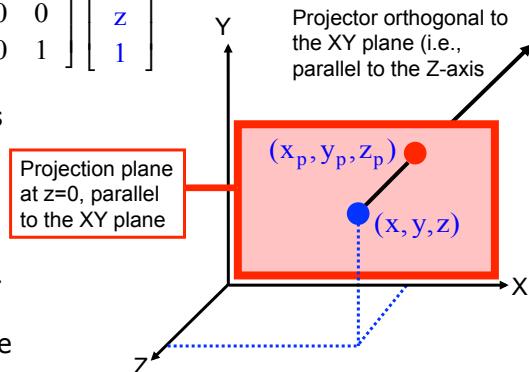
11

Matrix for orthographic projection

- We can express the projection as the matrix

$$\begin{bmatrix} x_p \\ y_p \\ z_p \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

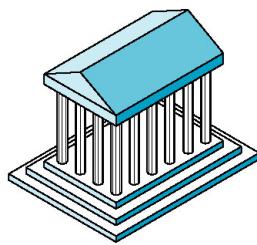
- This is the same as scaling by $(1, 1, 0)$
- Notice this matrix is singular
- We can derive similar matrices for other positions of the projection plane



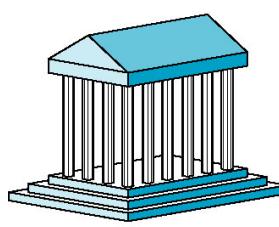
12

Parallel projection: axonometric

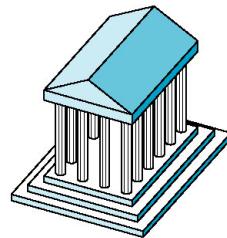
- In axonometric parallel projection, projectors are **perpendicular** to the projection plane
- The projection plane can have **any orientation** relative to the object being viewed
- Now we can see the three axes. There is distortion of lengths or angles, but measurements can still be made



Isometric: projection plane is symmetrical to (X,Y,Z)



Dimetric: projection plane is symmetrical to 2 of (X,Y,Z)

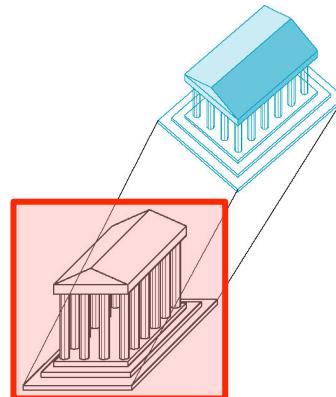


Trimetric: projection plane placed arbitrarily

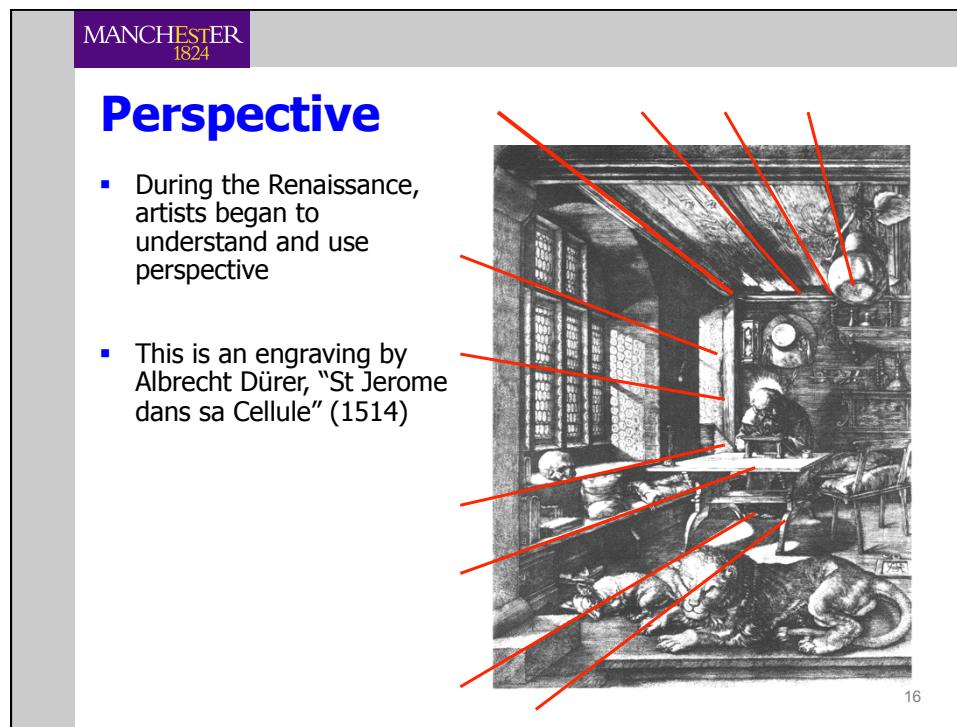
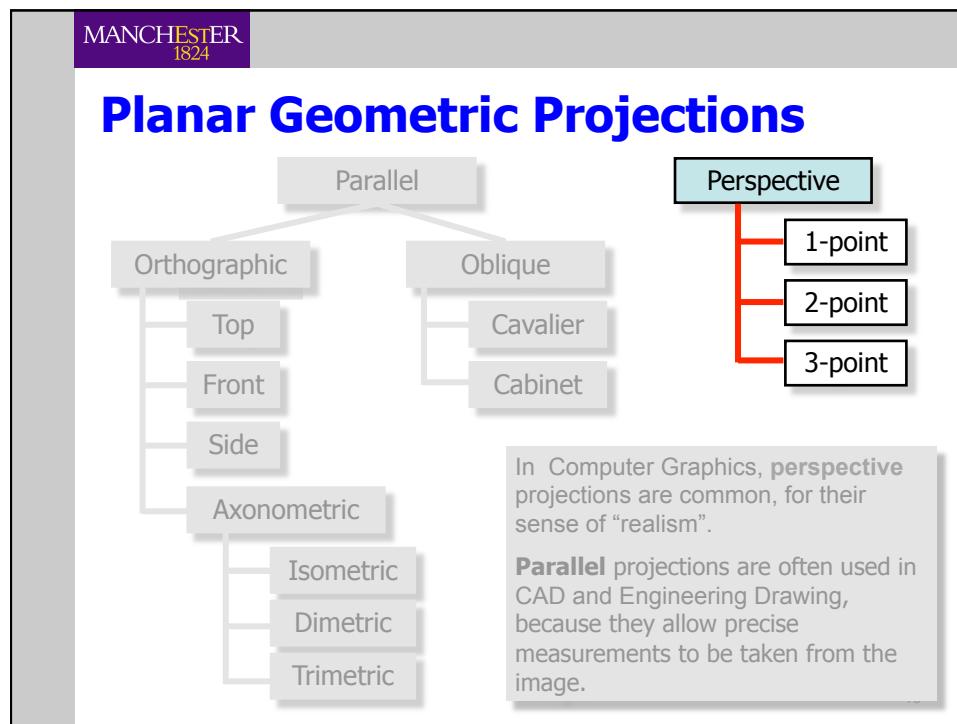
13

Parallel projection: oblique

- This is the most general case of parallel projection
- Projectors can make **any angle** with the projection plane
- The projection plane can have **any orientation** relative to the object being viewed
- We see the three axes. There is distortion of lengths or angles, but measurements can still be made
- We can of course derive matrices to perform **all the parallel projections** we have seen (*but we will not go into the details in this course*)



14



MANCHESTER
1824

Perspective projection

Perspective projection reflects the way we see (eye's lens and retina)

The diagram illustrates perspective projection. A blue 3D object, resembling a classical building, is positioned above a red 2D projection plane. A blue dot on the left, labeled 'Centre of projection, or "eye point"', represents the eye's lens. Lines from this center point converge towards the projection plane, forming a red frame that represents the field of view. The text 'Perspective projection reflects the way we see (eye's lens and retina)' is displayed in a purple box.

3D object

2D projection plane, oriented in 3D space

Centre of projection, or “eye point”

17

MANCHESTER
1824

Perspective projection

The **projection** is the set of points at which the projectors intersect the projection plane.

The diagram illustrates perspective projection. A blue 3D object is positioned above a red 2D projection plane. A blue dot on the left, labeled 'Centre of projection, or "eye point"', represents the eye's lens. Multiple dashed lines, labeled 'Projectors', converge from this center point towards the projection plane, forming a red frame that represents the field of view. The text 'The projection is the set of points at which the projectors intersect the projection plane.' is displayed in a purple box.

3D object

Projectors

2D projection plane, oriented in 3D space

Centre of projection, or “eye point”, is a **point**, so projectors **converge**

18

MANCHESTER
1824

Perspective projection

Objects further away from the COP become smaller. Edges that were parallel may converge. Angles between edges may be **distorted**.

The diagram illustrates perspective projection. A blue 3D cube labeled "3D object" is positioned above a red parallelogram labeled "2D projection plane, oriented in 3D space". Dotted lines represent the projection rays originating from a central point (the Center of Projection, COP) and passing through the vertices of the 3D cube to intersect the 2D plane. The 2D projection of the cube appears smaller and distorted.

19

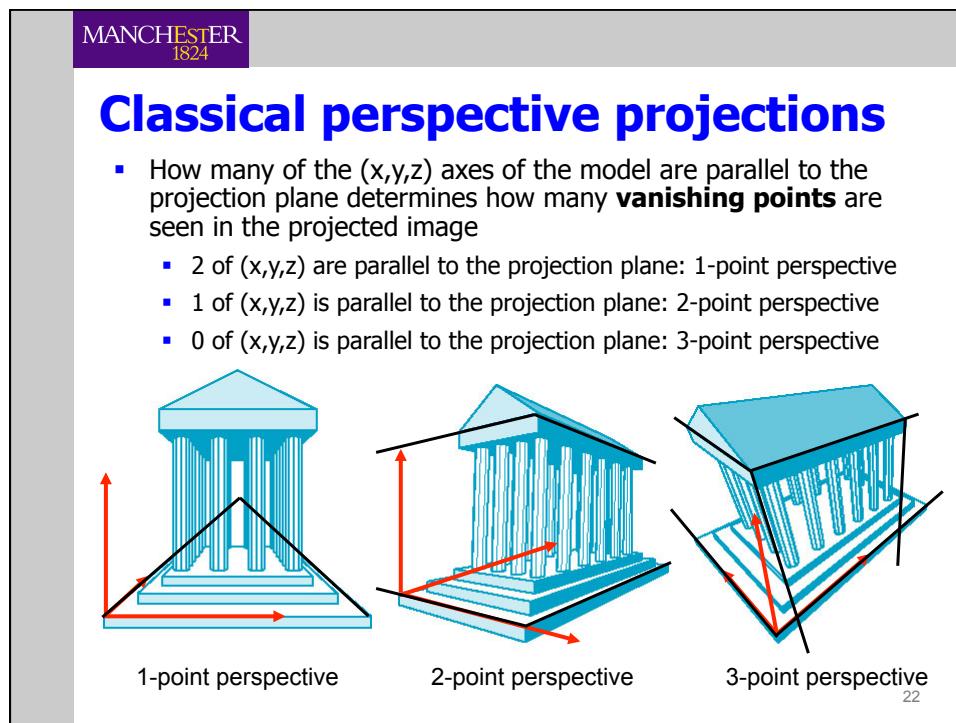
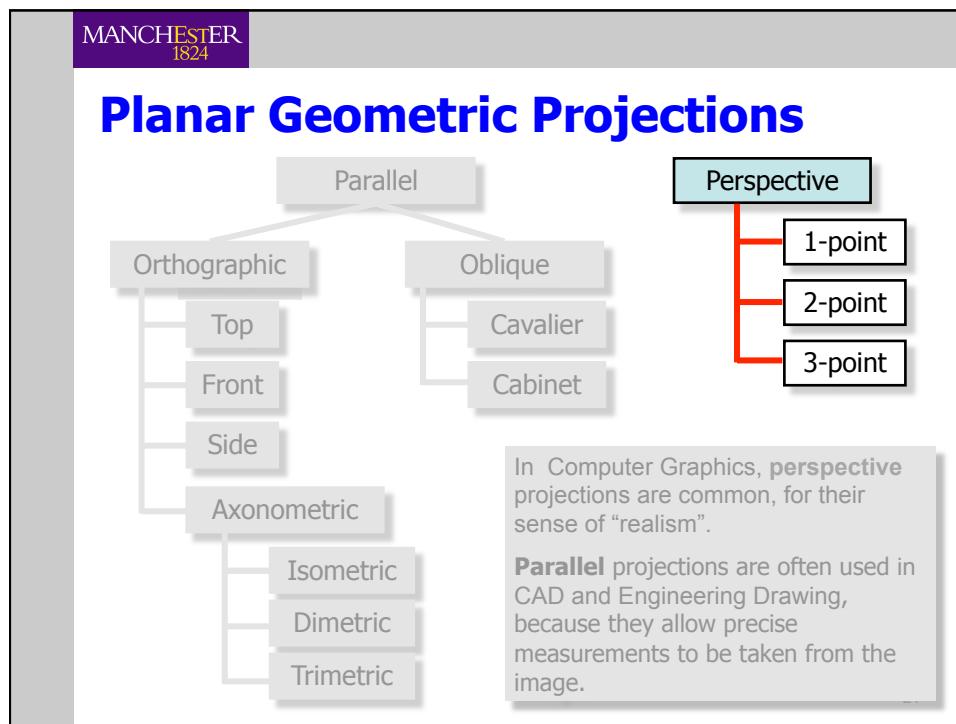
MANCHESTER
1824

Small or far away?

- With perspective projection, distant objects appear smaller in the projected image

The diagram shows a blue dot representing the Center of Projection (COP). From this COP, several red lines (projection rays) extend towards a red parallelogram representing the 2D projection plane. Three black and white cow models are positioned at different distances from the COP: one very close (small), one medium distance (medium size), and one far away (large but appearing small due to perspective projection). The projection rays from the COP pass through the cows and intersect the 2D plane, creating smaller and smaller projections of the cows as they get farther away.

20



22



MANCHESTER
1824

Computing perspective

- Suppose we have a projection plane located at $z=d$, and the plane is parallel to the XY plane
- A point (x, y, z) will project to (x_p, y_p, z_p) on the projection plane
- We want to find (x_p, y_p, z_p)
- By definition, $z_p = d$

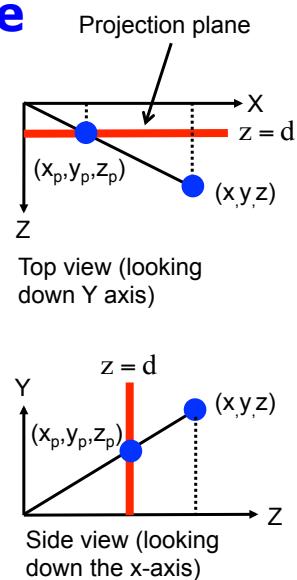
The diagram shows a 3D coordinate system with X, Y, and Z axes. A red rectangle represents the "Projection plane, at $z=d$, parallel to the XY plane". A point (x, y, z) is shown in blue, with its projection (x_p, y_p, z_p) on the red plane also in blue. Dashed lines connect the points, showing the projection process. The Z-axis arrow points towards the viewer, indicating depth.

24

Computing perspective

- We need to find x_p and y_p
- Looking at the top view, looking down the y-axis, by similar triangles, we have:
$$\frac{x_p}{d} = \frac{x}{z} \rightarrow x_p = \frac{dx}{z} \rightarrow x_p = \frac{x}{z/d}$$
- Similarly, in a side view, with our eye at the origin and looking along the x-axis, we have:

$$y_p = \frac{y}{z/d}$$



25

Perspective projection matrix

- We can express this transformation as a 4x4 matrix:
 - If we multiply a point by this matrix, we get the transformed point
- $$\begin{bmatrix} x \\ y \\ z \\ z/d \end{bmatrix} = \begin{bmatrix} x' \\ y' \\ z' \\ w' \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$
- Recall however that we must always end up with a 3D point which has $w = 1$, so we need to divide through all elements by w :
 - This is called **perspective division**

$$\begin{bmatrix} \frac{x}{z/d} \\ \frac{y}{z/d} \\ \frac{z}{z/d} \\ 1 \end{bmatrix} \rightarrow \begin{bmatrix} \frac{x}{z/d} \\ \frac{y}{z/d} \\ d \\ 1 \end{bmatrix}$$

26

Perspective matrices: summary

- We derived the matrix to perform a 1-point projection where the projection plane is parallel to the XY plane.

$$\begin{bmatrix} x' \\ y' \\ z' \\ w' \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d_z & 0 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

We've renamed \mathbf{d} to \mathbf{d}_z just to remind us that it's a z-coordinate

- We can also derive (details omitted) a matrix which expresses the most general case, **3-point projection**, where the projection plane is not parallel to any of the XY, XZ or YZ planes:

$$\begin{bmatrix} x' \\ y' \\ z' \\ w' \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 1/d_x & 1/d_y & 1/d_z & 0 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

The projection plane intersects the X, Y and Z axes at (d_x d_y d_z)

27

Projections: summary

- Parallel orthographic projection onto $z=0$ plane

$$\begin{bmatrix} x_p \\ y_p \\ z_p \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

- 1-point perspective projection onto XY plane

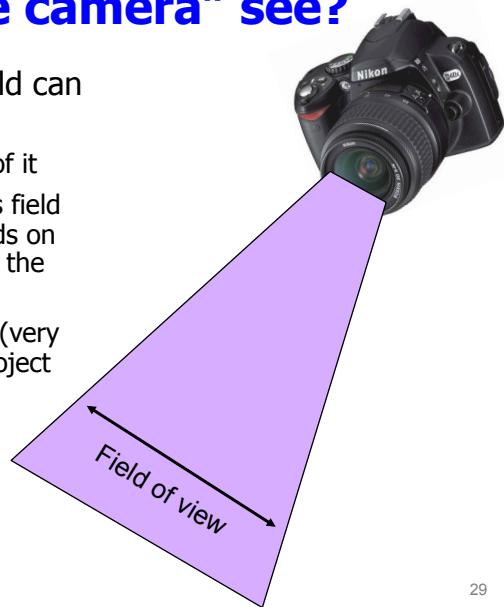
$$\begin{bmatrix} x' \\ y' \\ z' \\ w' \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Then divide by w'
to give 3D point $\begin{bmatrix} x \\ y \\ z \\ d \\ 1 \end{bmatrix}$

28

What can “the camera” see?

- What part of the world can a “camera” see?
 - Only objects in front of it
 - Only objects within its field of view, which depends on its lens (about 90° for the human eye)
 - Only a finite distance (very distant objects will project too small to be seen)



29

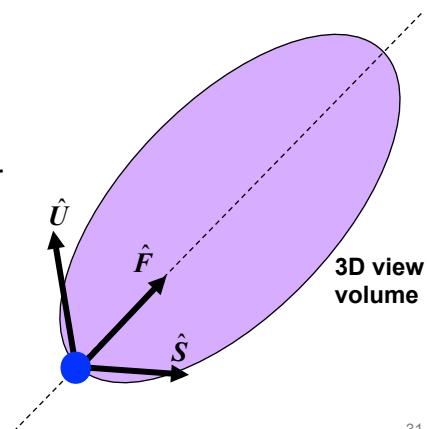
Viewing volumes and clipping

- So, having specified the position and orientation of the camera (see 3D Viewing (1) ...)
- ...we now describe the field of view of the “camera” by defining a 3D “view volume”
- Only those parts of 3D objects which lie inside the view volume are displayed – all objects are **clipped** against the view volume
- We’ll look at clipping shortly, but first we’ll see how the view volumes are defined.

30

Defining the view volume

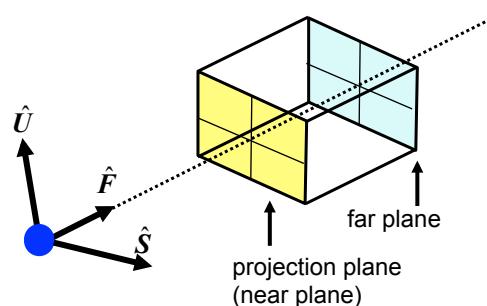
- In **Viewing (1)** we established a coordinate system for the “camera” – with axes $\hat{S} \hat{U} \hat{F}$
- We define a **3D view volume** which is “attached” to the camera
- As we would expect, the shapes of the volumes for parallel and perspective projection are different.



31

View volume for parallel projection

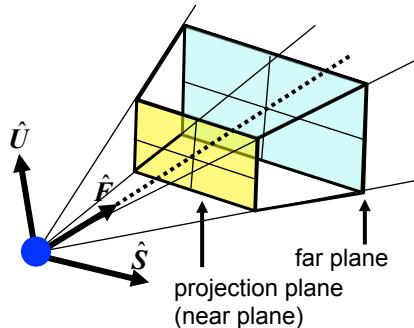
- A view volume is a 3D shape, defined by six planes
- For parallel orthographic projection, it's a **cuboid**
- The cuboid is defined by a near plane (the projection plane) and a far plane, and top/bottom, left/right planes
- The near and far planes are orthogonal to the camera's \hat{F} axis



32

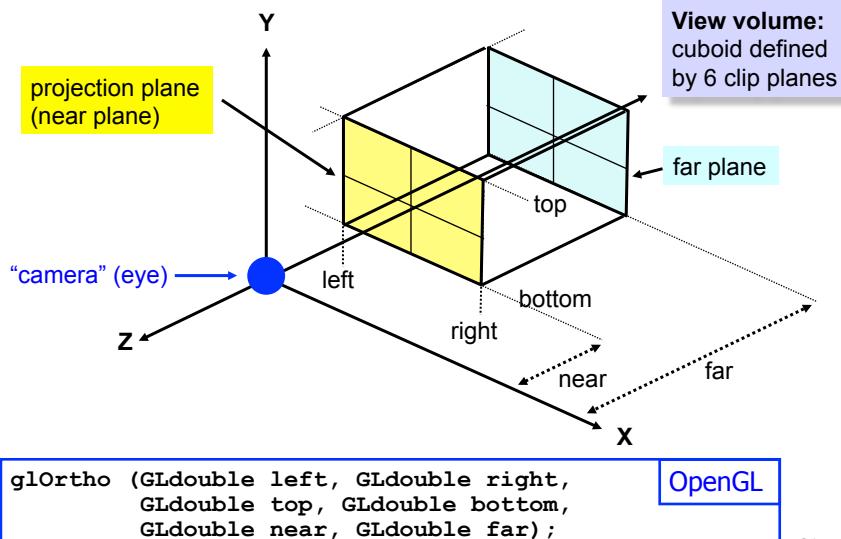
View volume for perspective projection

- For perspective projection, this view volume is a **frustum**, a truncated pyramid
- The frustum is defined by a near plane (the projection plane) and a far plane
- The near and far planes are orthogonal to the camera's \hat{F} axis



33

Orthographic projection in OpenGL



34

MANCHESTER
1824

Perspective projection in OpenGL

projection plane (near plane)

far plane

FOV

"camera" (eye)

near

far

`gluPerspective (GLdouble fovy,
 GLdouble aspect,
 GLdouble near, GLdouble far);`

OpenGL

35

MANCHESTER
1824

Perspective projection in OpenGL

- Suppose the projection plane (near plane) is at $z = -d$, and the eye (or "camera") is at $(0,0,0)$
- Then the (1-point) projection transformation, as we saw in **Viewing (2)** is given by

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix}$$

- Note: this projection transformation is **independent** of the far plane – which is only for clipping

36

Perspective: a problem

- A projection matrix transforms a 3D point to a 3D point with $W \neq 1$
- For example, for a 1-point perspective projection:

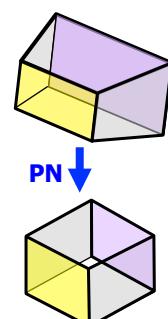
$$\begin{bmatrix} x_p \\ y_p \\ z_p \\ w_p \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} \quad \text{In this case } w_p = z/d$$

- So we need to divide through by w , to get (x_p, y_p, z_p)
- $z_p = d$, which means we have **lost** any idea of the original 3D object's z -coordinates – they've all been set to d .
- This is unhelpful, because we want to **keep** the z depth information, to do hidden-surface removal.

37

Solution: Projection normalization

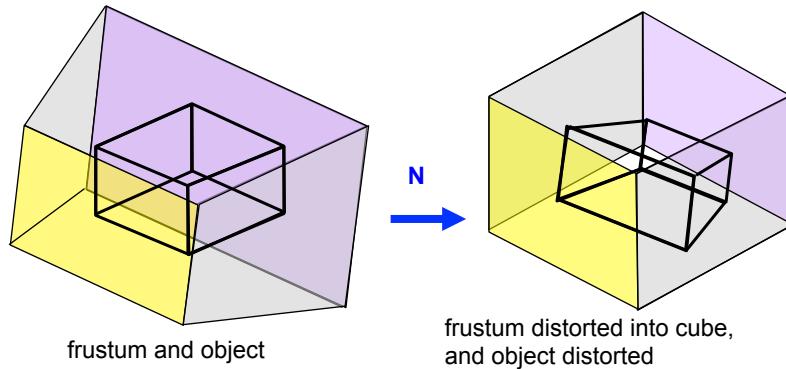
- We would like to have a perspective transformation which preserves depth.
- Suppose we have a particular perspective transformation expressed by frustum F and matrix M
- It can be shown that we can derive a new transformation matrix PN , based on M , that **distorts F into a cube**:
- Transforming our model by PN , and then taking an orthographic projection...
- ...produces **exactly the same result** as performing our original perspective transformation M , with one difference: the z depth values are preserved
- This new technique is called **projection normalization**, and OpenGL creates PN for us automatically.



38

Solution: Projection normalization

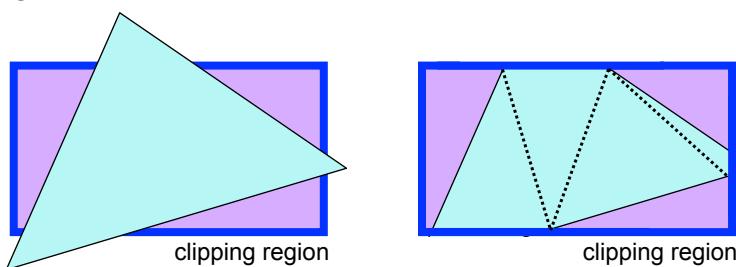
- **PN** distorts the perspective frustum to a unit cube
- An object inside the frustum is also distorted
- The distortion of the object is such that if we then apply an **orthographic projection**, we get the same view as if we had applied a perspective projection to the original undistorted object



39

The clipping operation

- Clipping takes place in the cube produced by projection normalization (easier than clipping against a frustum)
- Clipping is easiest to visualise in 2D
- Here we see polygon being clipped against a rectangular clipping region



- There are established algorithms for clipping in 2D and 3D

40

MANCHESTER
1824

Viewing volume and clipping

```
fovy aspect zNear zFar
gluPerspective( 36.0 , 1.00 , 1.3 , 2.6 );
gluLookAt( 0.00 , 0.00 , 2.00 , <- eye
           0.00 , 0.00 , 0.00 , <- center
           0.00 , 1.00 , 0.00 ); <- up
```

Click on the arguments and move the mouse to modify values.

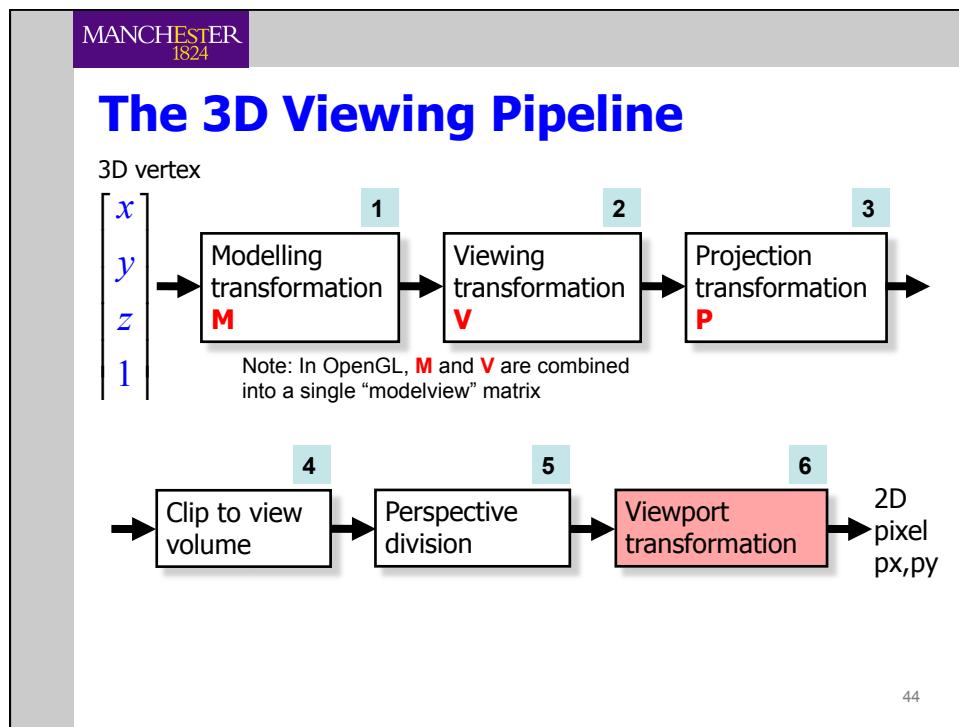
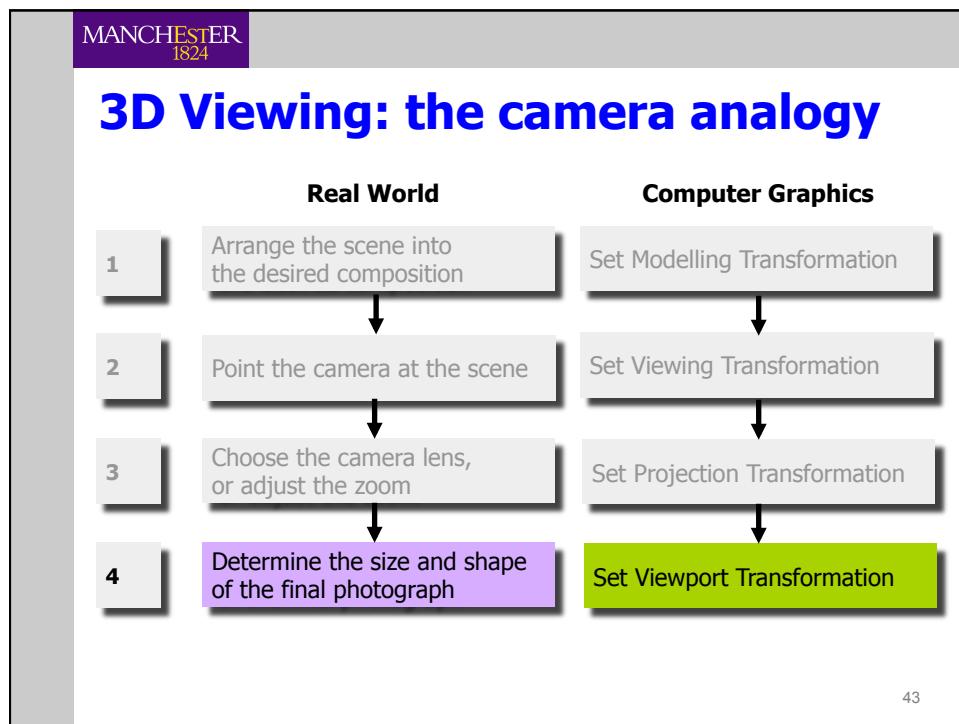
MANCHESTER
1824

Perspective division

- The clipping operation returns a set of (x,y,z,w) vertices defining polygons which are inside the view volume
- OpenGL now performs the **perspective division by w** to convert these (x,y,z,w) values to (x,y,z) 3D points

```
graph LR; A[Projection transformation] --> B[Clip to view volume]; B --> C[Perspective division]
```

42



MANCHESTER
1824

The Viewport Transformation

- After perspective division, we now have 3D coordinates
- OpenGL has scaled these into the range $[-1,+1]$ in X, Y and Z
- In the final step, using only the X and Y values, OpenGL computes a transformation from $[-1,+1]$ to the display screen
 - Either the whole **window**
 - Or a part of it called a **viewport**
- But we retain Z**, to remove hidden surfaces

45

MANCHESTER
1824

Summary of 3D viewing

- We can now summarise the steps of the viewing process
- 1. The modelling transformation arranges objects in our 3D world
- 2. The viewing transformation transforms the world to give the same view as if it were being photographed by a camera
- 3. The projection transformation performs a parallel/perspective projection within limits (the clip planes)
- 4. Those parts of the 3D world outside the clip planes are discarded
- 5. If it's a perspective view, the perspective division "flattens" the image
- 6. The viewport transformation maps the final image to a position in part of the display screen window.

46