

Transport Services and Protocols


Andy Carpenter

(Andy.Carpenter@manchester.ac.uk)

Elements these slides come from Kurose and Ross, authors of "Computer Networking: A Top-down Approach", and are copyright Kurose and Ross


Transport Services and Protocols

- Provides logical communication between applications
 - process-to-process
- Runs in end systems
 - send side: breaks app messages into segments, passes to network layer
 - receiver side: reassembles segments into messages, passes to app layer
- Multiplex/ id end-points
- Protocols available:
 - Internet: TCP, UDP



Transport Service QoS Params

- Provides QoS to applications
- Options
 - Connectionless/connection-orientated
 - reliability
 - flow control
 - congestion control
- Alters underlying network QoS
- Services not available:
 - delay guarantees
 - bandwidth guarantees



End-points (Ports/Sockets)


- Support multiple destinations
- Identify process via 16-bit identifier; port
- Applications use sockets which have associated port
- Well know services allocated ports 0-1024

Port	Mnemonic	Service
7	ECHO	Echo
20	FTP-DATA	FTP (default data)
21	FTP	FTP (control)
53	DOMAIN	Domain name service
80	HTTP	Hypertext Transfer protocol

Centrally
allocated

- Client ports are dynamically allocated by o/s

End-Point: Connecting to a Service



- Packets are sent to ports with attached applications
- Data to other ports will generate an error message

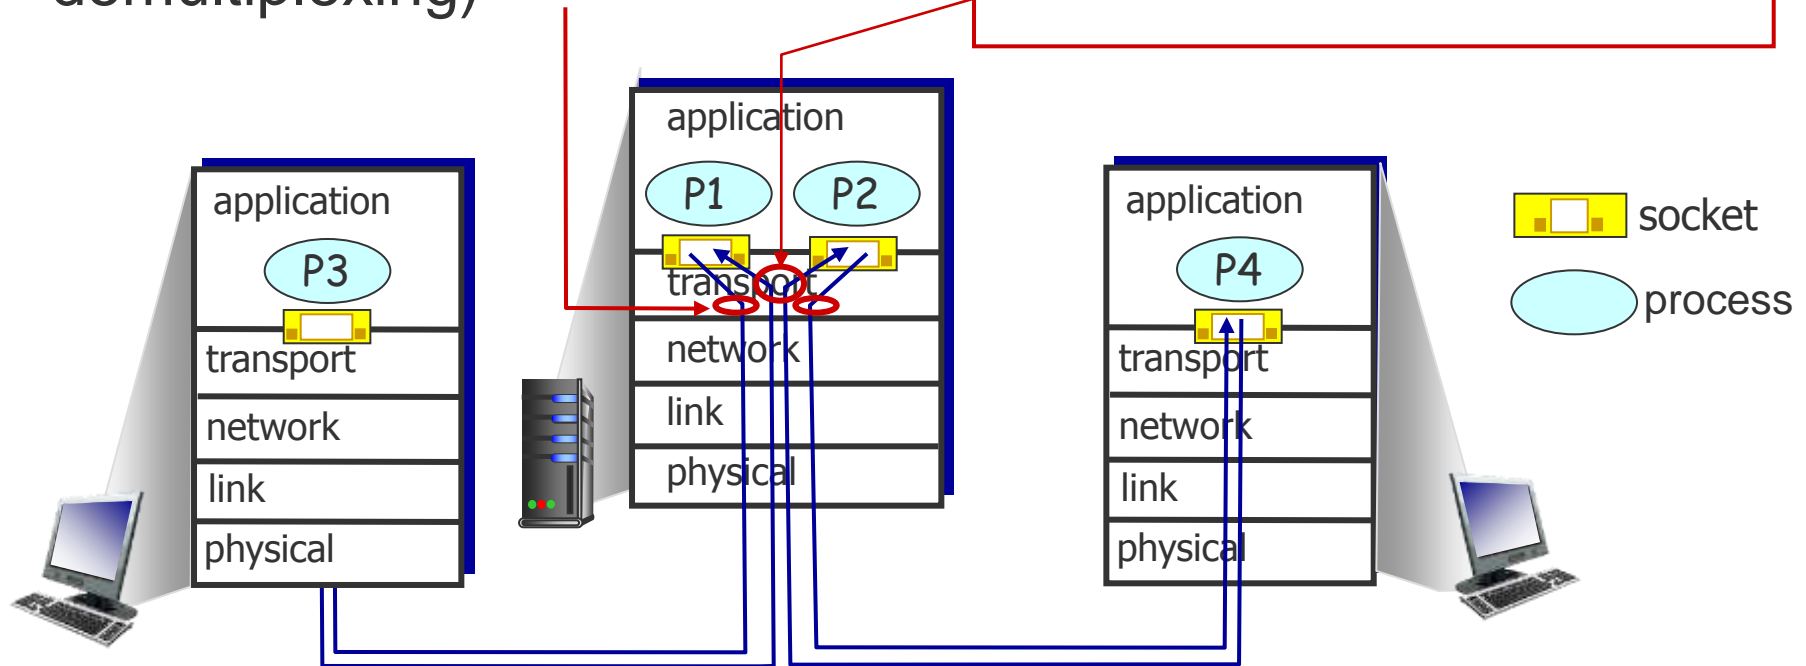
End-Point: Multiplexing

multiplexing at sender:

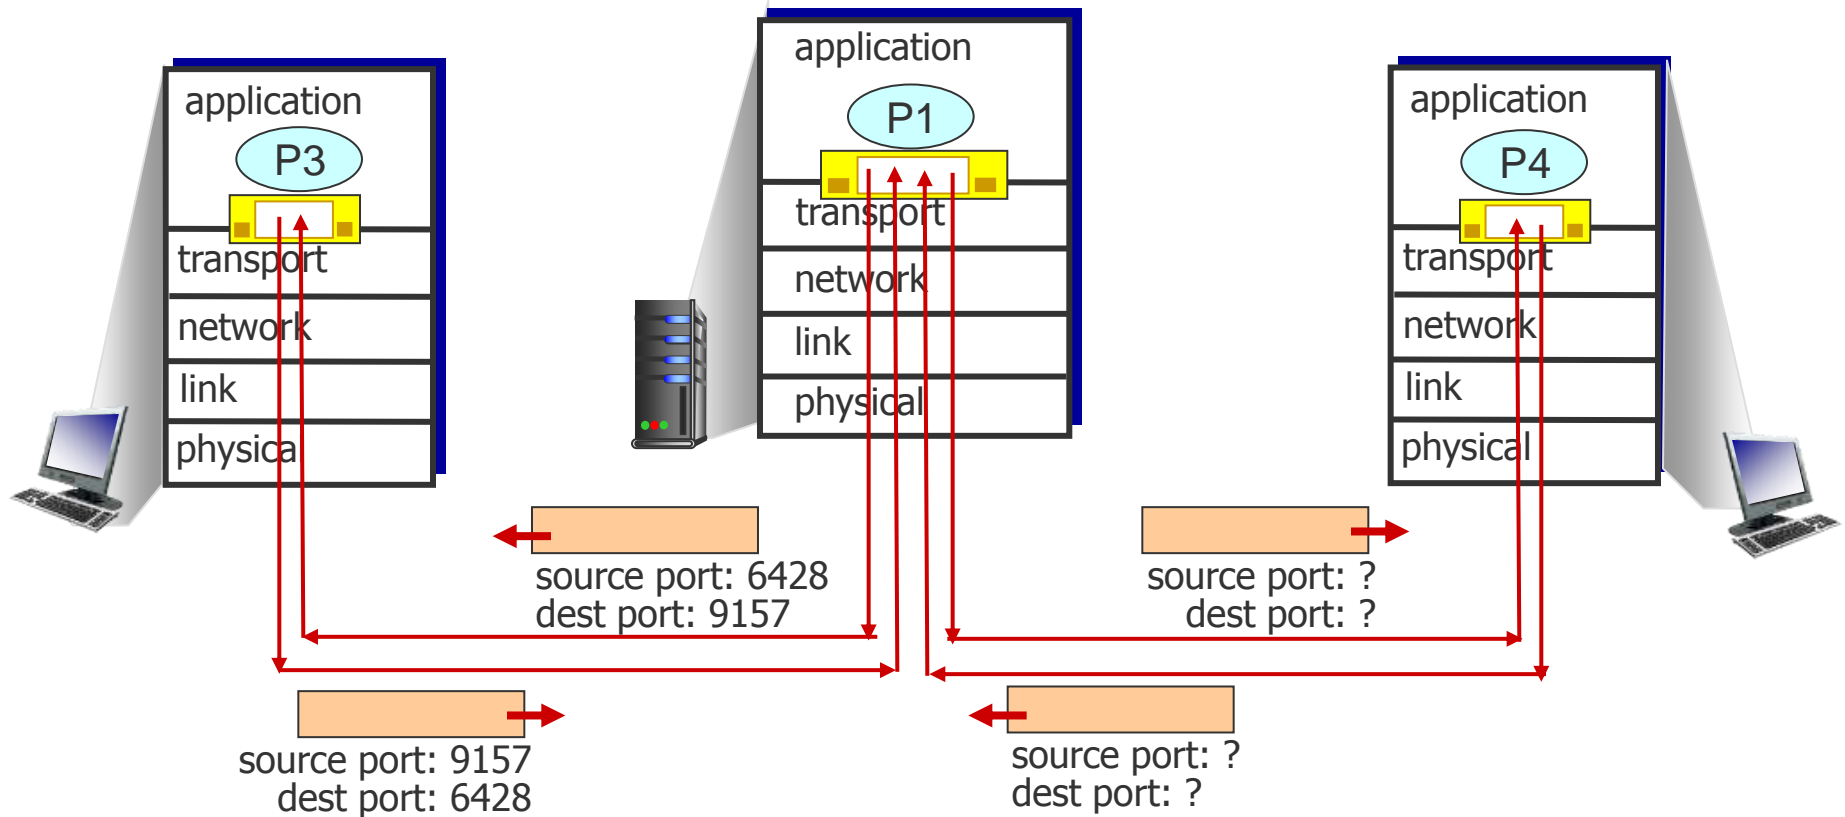
handle data from multiple sockets, add transport header (later used for demultiplexing)

demultiplexing at receiver:

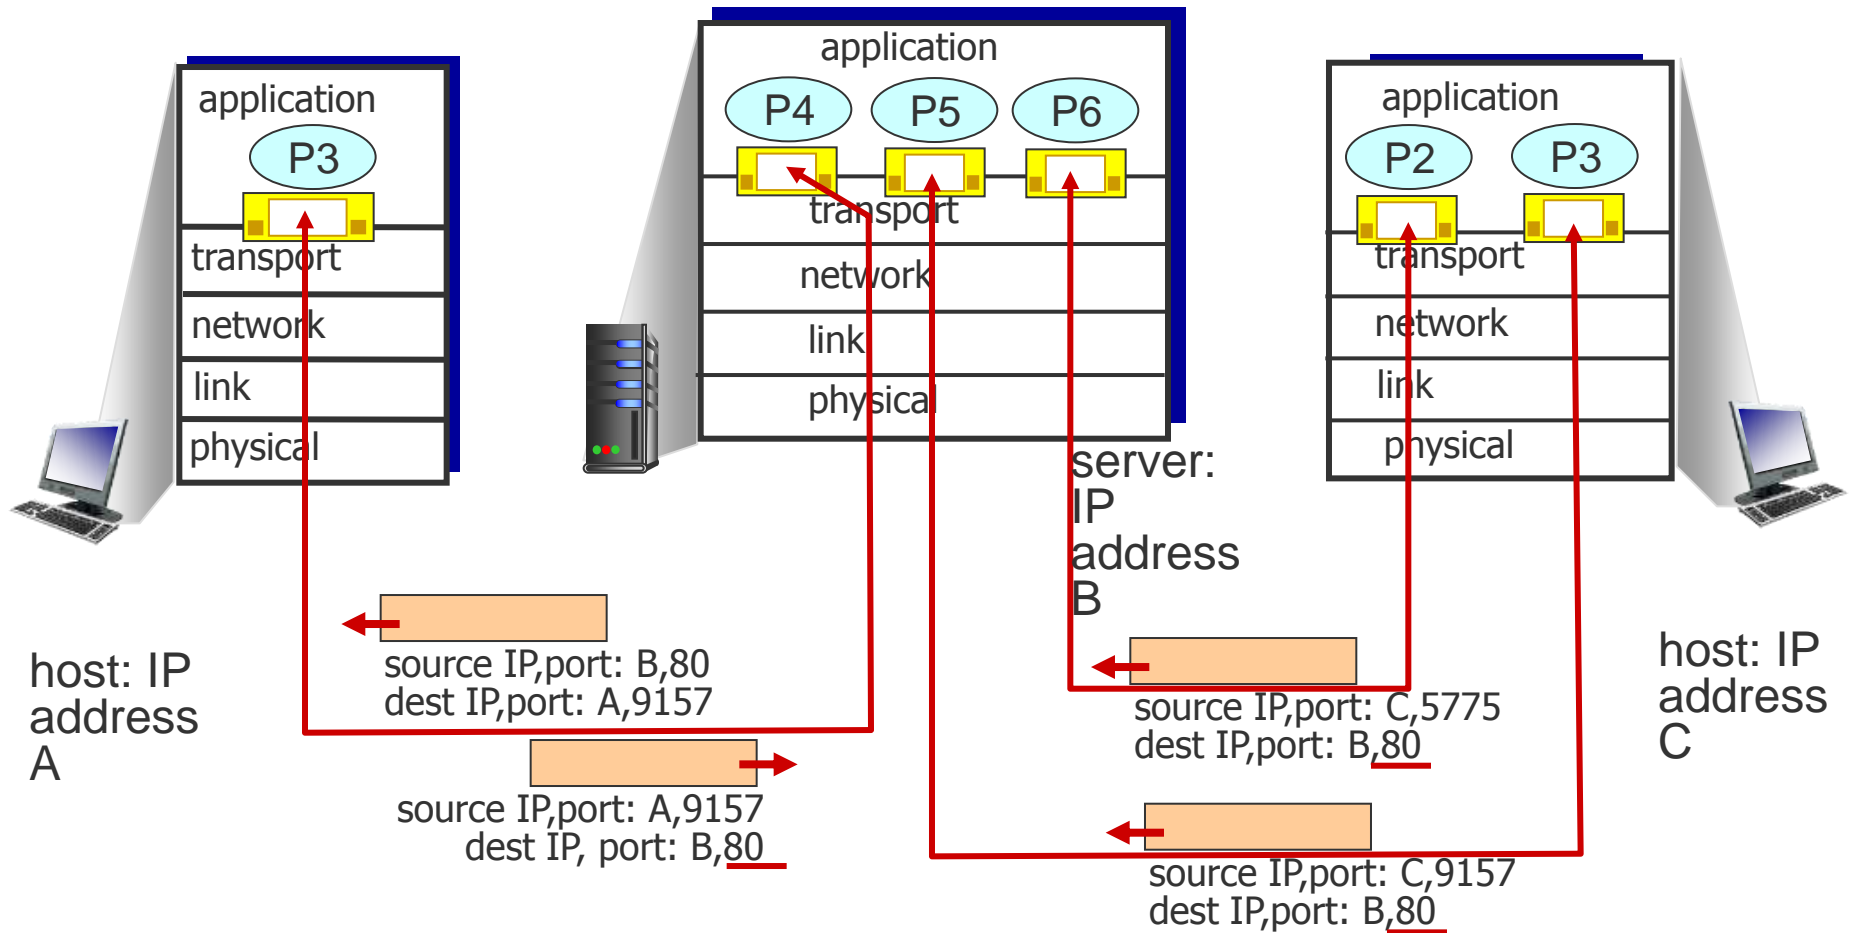
use header info to deliver received segments to correct socket



End-Point: Id (Connectionless)



End-Point: Id (Connection-Oriented)



User Datagram Protocol (UDP)

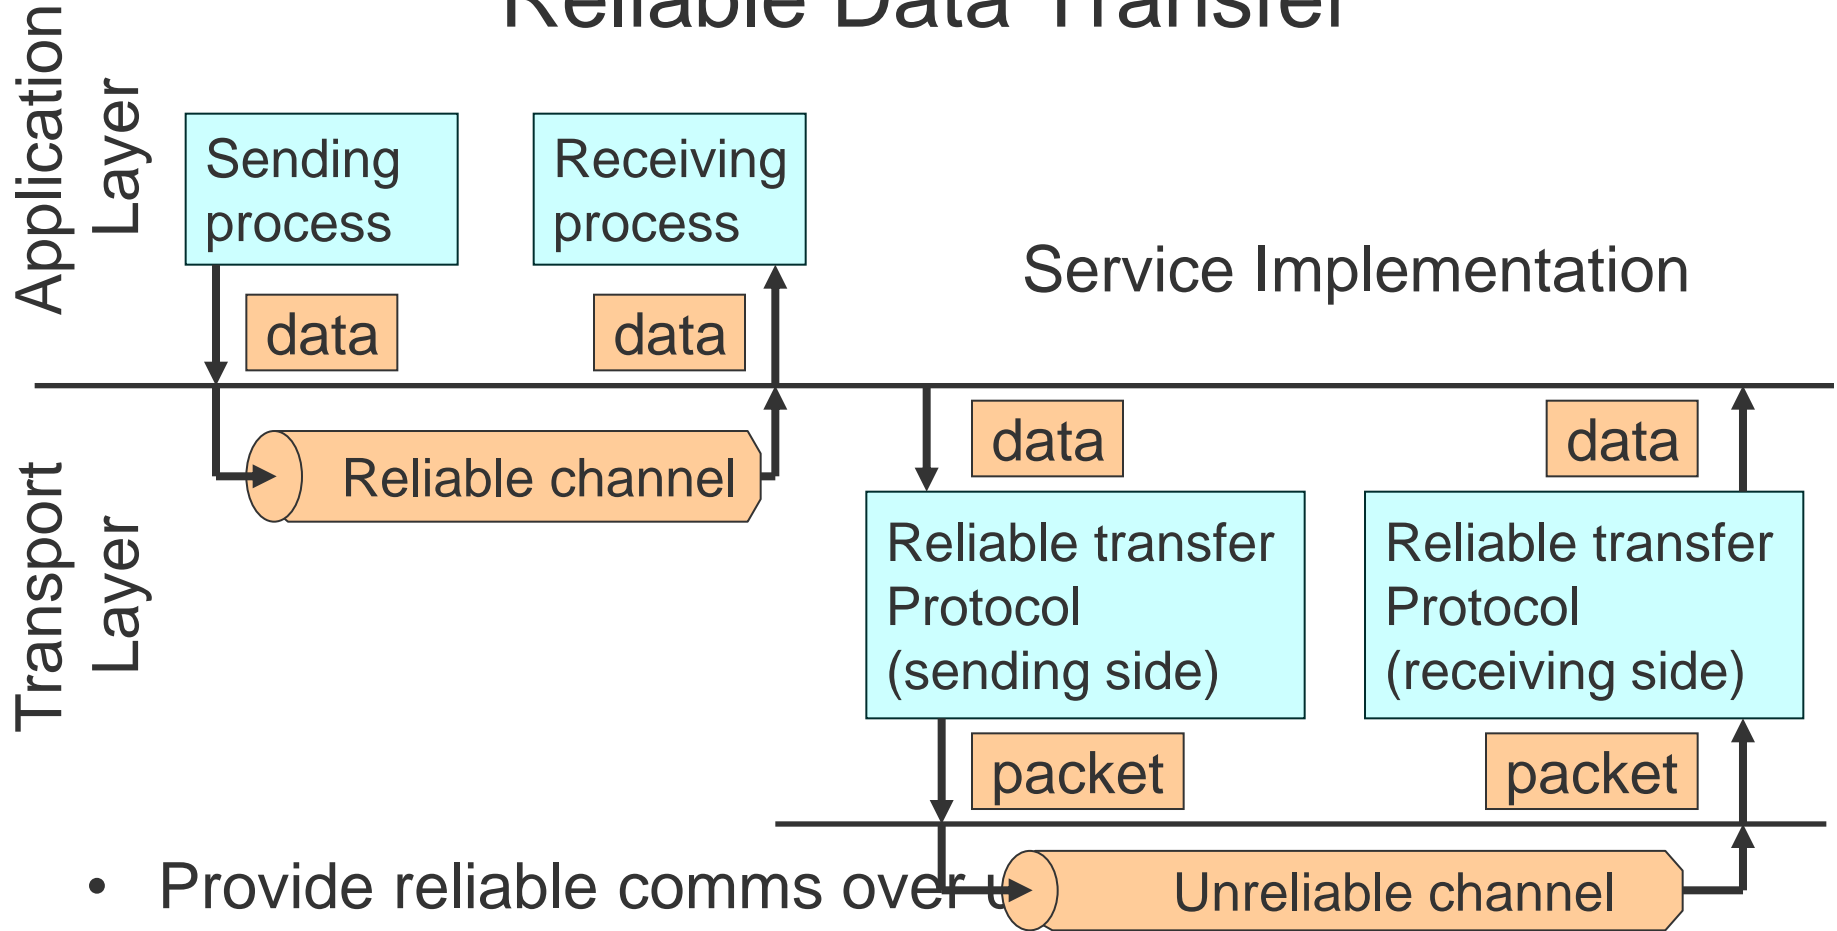
- “no frills”, “best effort” connectionless service, segments may be:
 - lost, reordered
- Segments independent
- Often used for streaming multimedia applications
 - loss tolerant
 - rate sensitive
- Other uses:
 - DNS, SNMP
 - NFS
 - o/s applications

Why is there a UDP?

- no connection setup (adds delay)
- simple: no connection state at sender, receiver
- small segment header
- no congestion control: UDP can blast away as fast as desired

[RFC 768]

Reliable Data Transfer


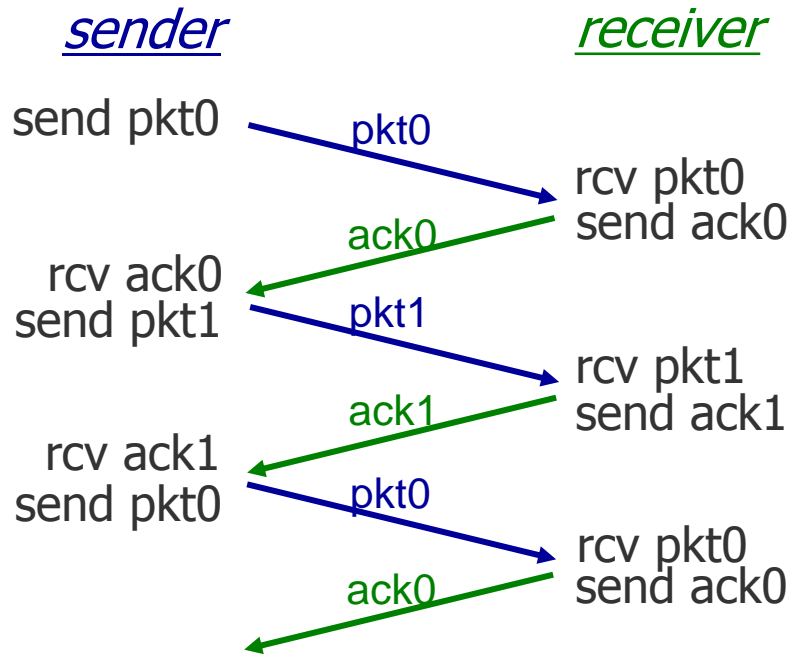


- Provide reliable comms over unreliable channel
- Complexity depends on characteristics of unreliable channel

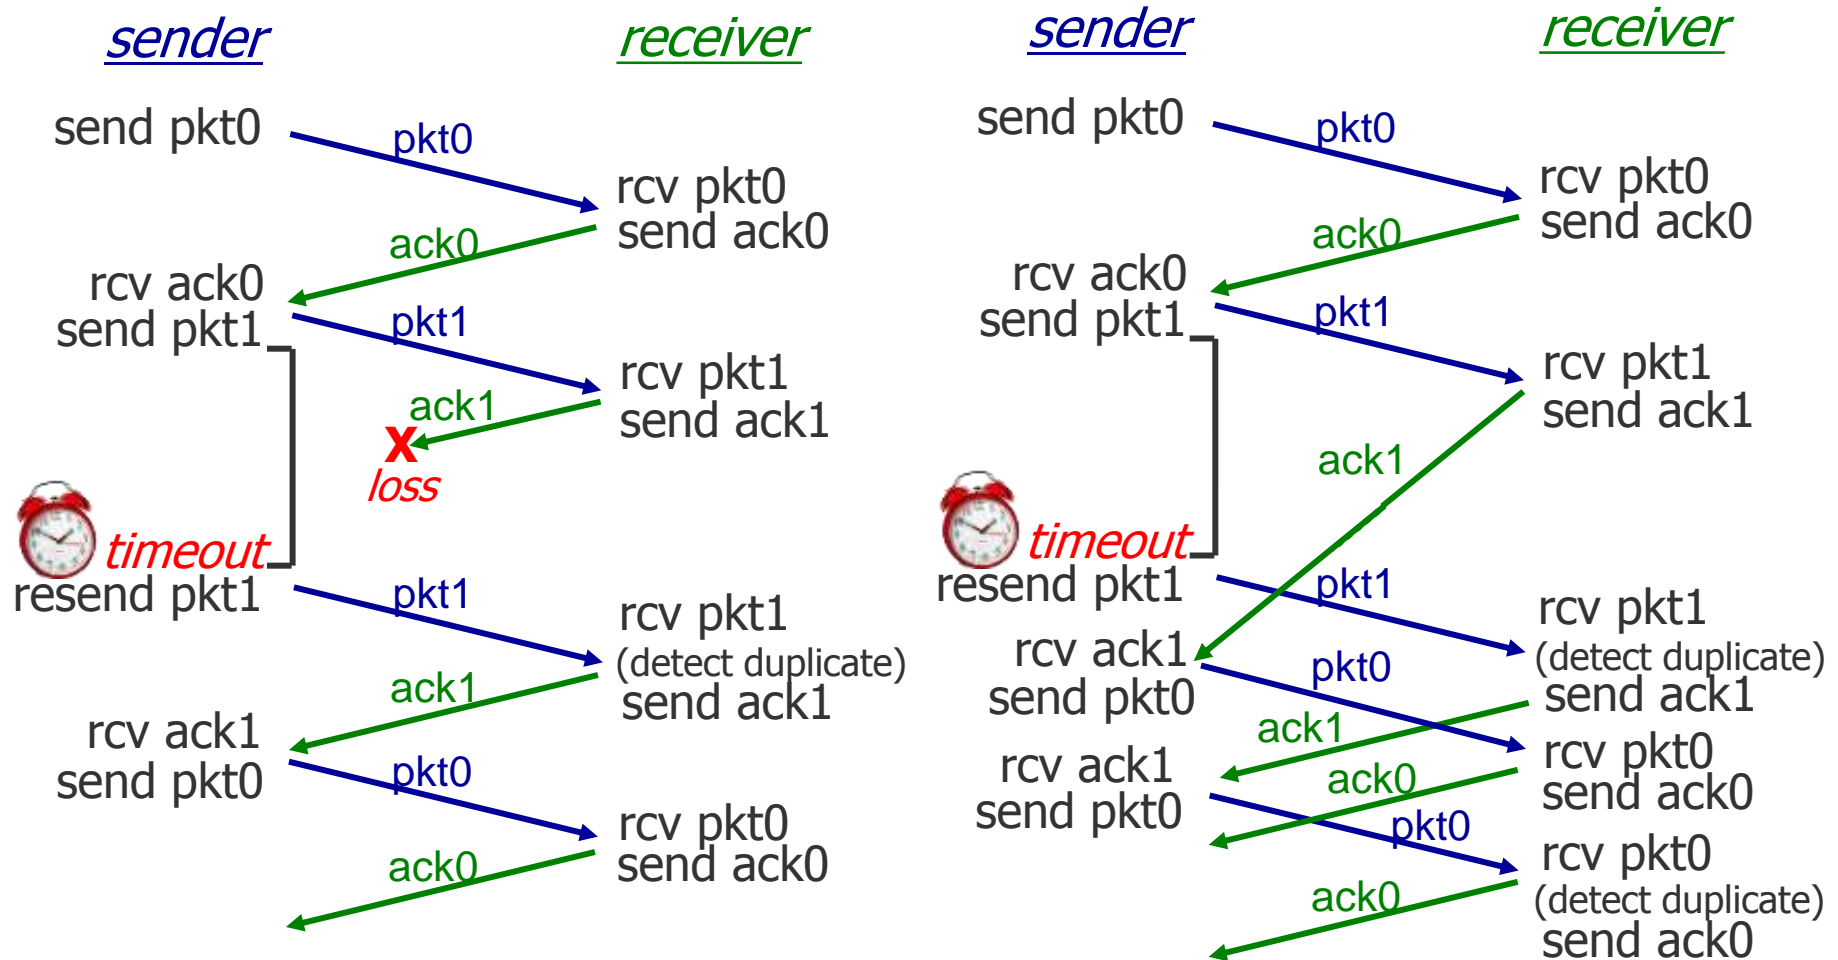
Recovering from Errors

- Recovery uses two mechanisms:
 - acknowledgements and timeouts
- Acknowledgement (**ACK**) is control packet from
 - receiver to transmitter of data packet being ACKed
- Receipt of ACK confirms delivery of data
- If ACK not received within **timeout**:
 - transmitter of data retransmits data; needs copy
- Process called automated repeat request (ARQ)
- ARQ mechanisms: **stop-and-wait, sliding window**


Reliability: Stop-and-Wait



Reliability: Stop-and-Wait




Reliability: S-a-Wait - Utilisation




$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

Network protocols limit use
of physical resources

Reliability: Pipelined Protocols

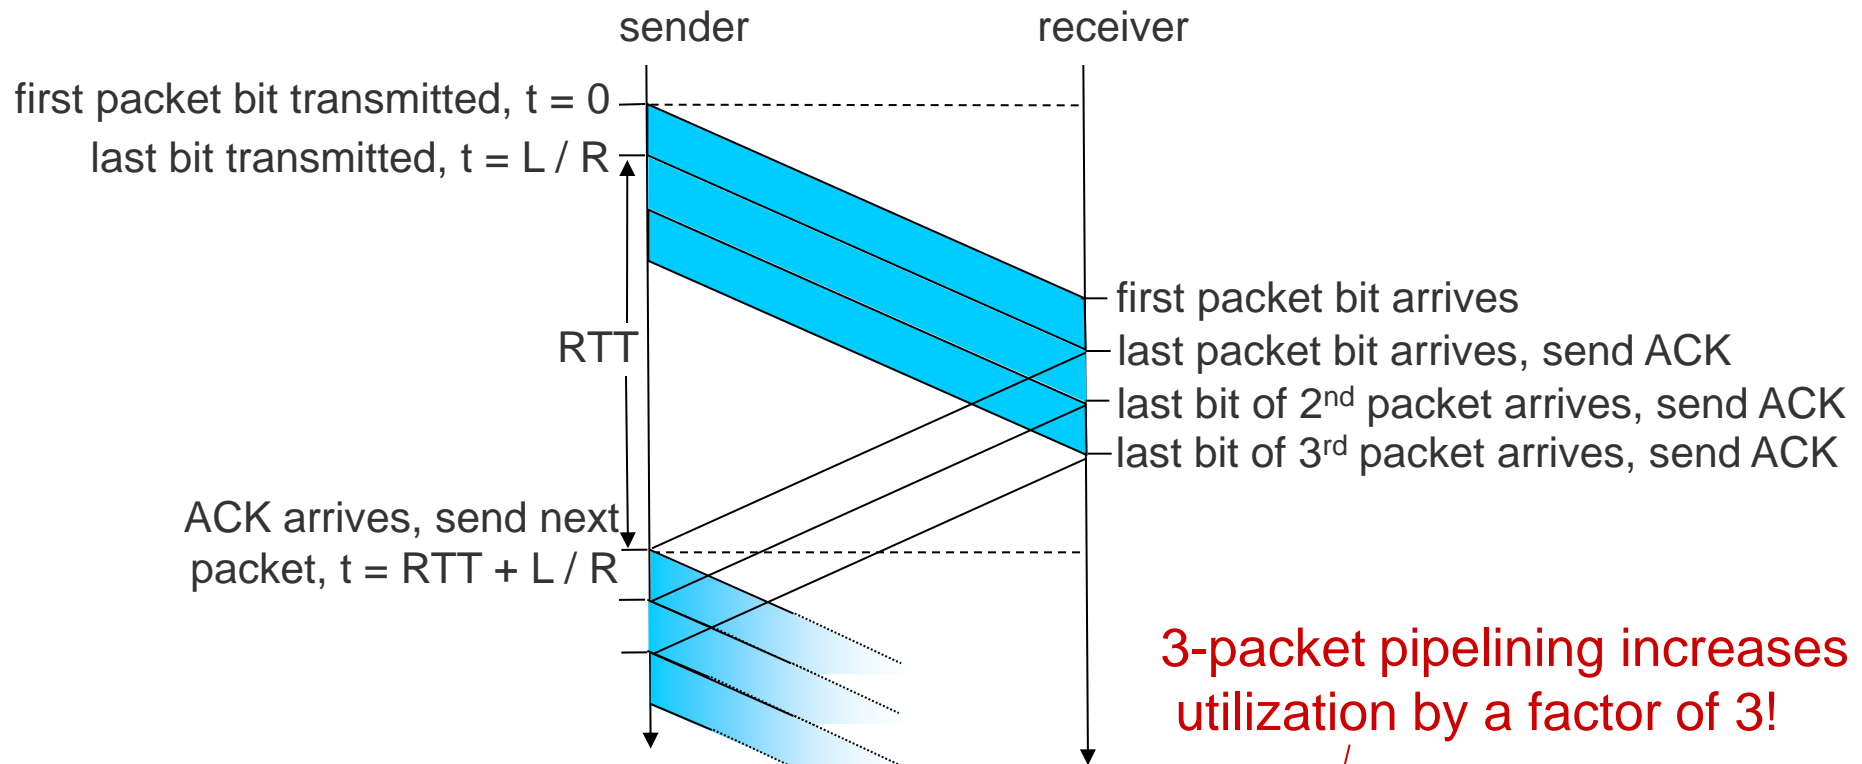


- Allow multiple, “in-flight”, yet-to-be-acknowledged pkts
- Example: received ACK 5, window 6



- Often referred to as sliding window

Reliability: Pipelined Utilisation




$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.00081$$

Reliability: Pipelined Approaches

- What if expecting packet 8 and get packet 9?
- If packet 8 delayed, when arrives can send ACK(9)
- If lost, timeout will expire and will be resent
- **Go-Back-N**, send cumulative ACKs:
 - ACK(n) acknowledges all packets upto n
 - likely that will also get packets 9, 10, ... resent
- **Selective repeat**:
 - explicitly acknowledges all packet
 - only unsuccessfully received packets are resent
- Could **negatively acknowledge** (NACK) packet 8,
 - requests retransmission without timeout expiring

Reliability: Go-Back-N



- k-bit seq # in pkt header
- “window” of up to N , consecutive unack'ed pkts allowed
- ACK(n): ACKs all pkts up to, including seq # n - “cumulative ACK”
- may receive duplicate ACKs (see receiver)
- timer for oldest in-flight pkt
- timeout(n): retransmit packet n and all higher seq # pkts in window

Reliability: Go-Back-N Example (H)

sender window (N=4)

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

sender

send pkt0
send pkt1
send pkt2
send pkt3
(wait)

rcv ack0, send pkt4
rcv ack1, send pkt5

ignore duplicate ACK



pkt 2 timeout

send pkt2
send pkt3
send pkt4
send pkt5

receiver

receive pkt0, send ack0
receive pkt1, send ack1

receive pkt3, discard,
(re)send ack1

receive pkt4, discard,
(re)send ack1

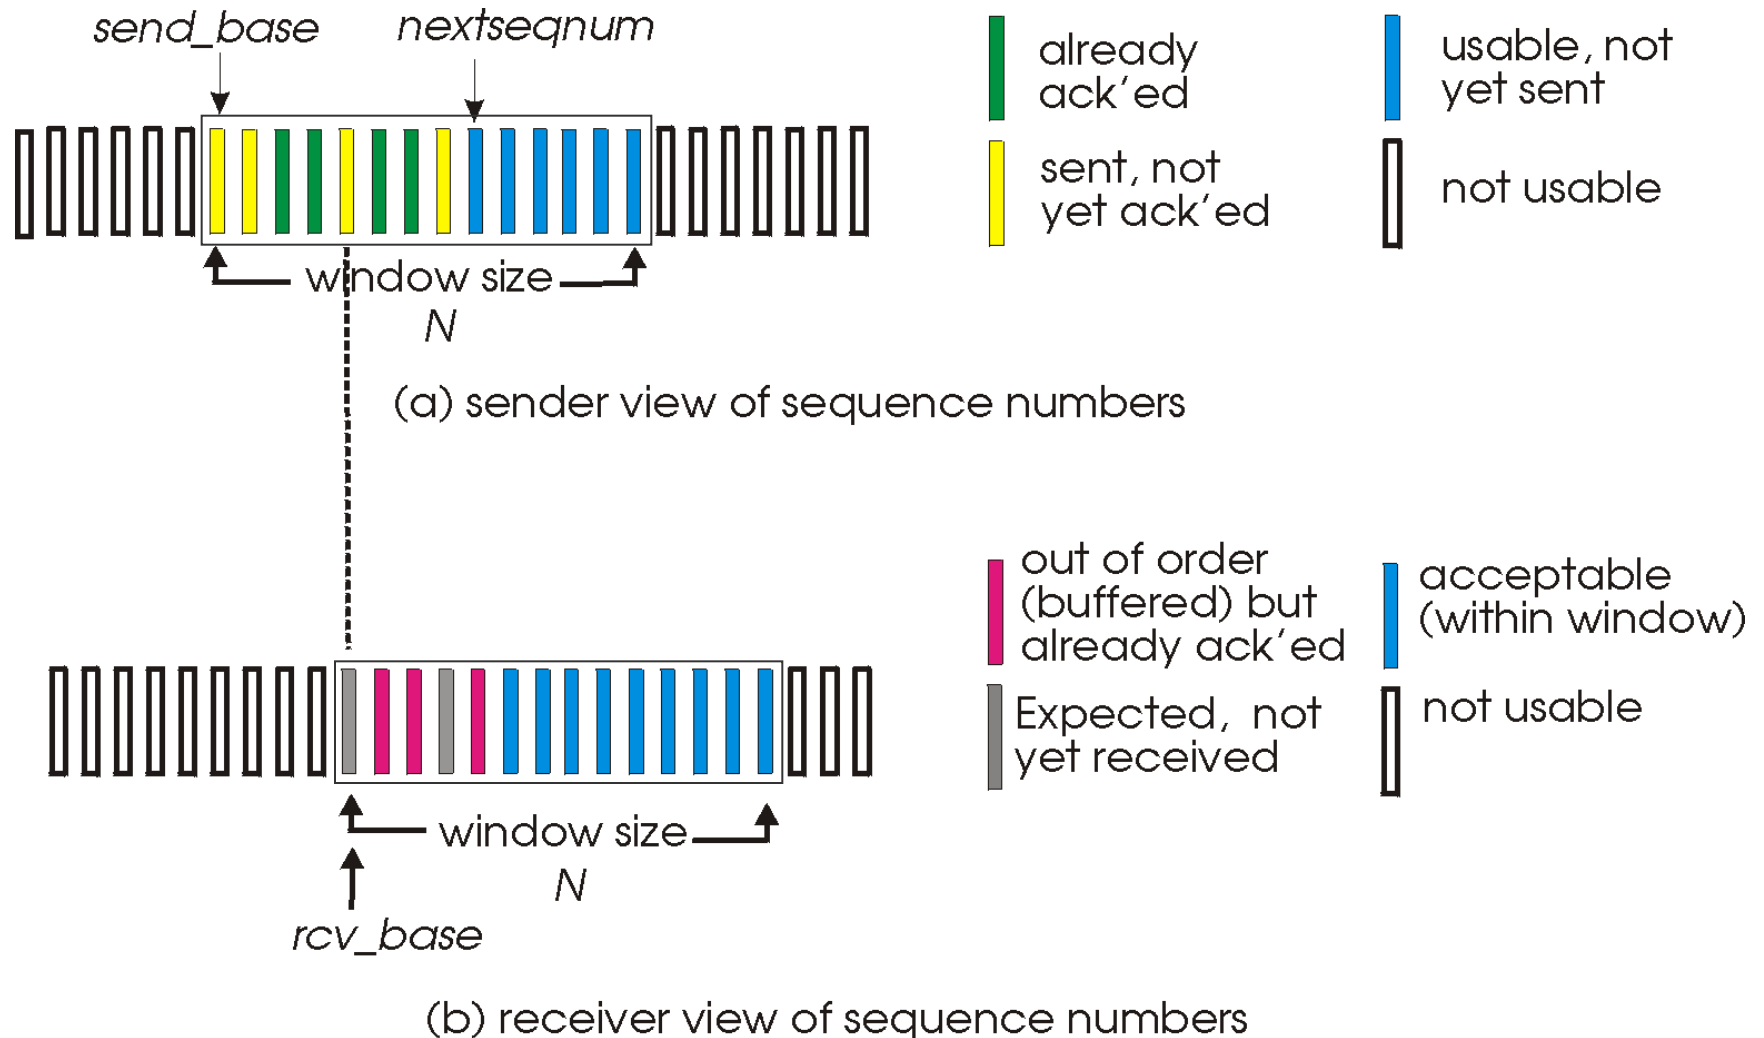
receive pkt5, discard,
(re)send ack1

rcv pkt2, deliver, send ack2
rcv pkt3, deliver, send ack3
rcv pkt4, deliver, send ack4
rcv pkt5, deliver, send ack5

X loss

Transport Layer

Reliability: Selective Repeat



Reliability: Selective Repeat

sender window (N=4)

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
[]

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

sender

send pkt0
send pkt1
send pkt2
send pkt3
(wait)

rcv ack0, send pkt4
rcv ack1, send pkt5

record ack3 arrived



pkt 2 timeout

send pkt2
record ack4 arrived
record ack5 arrived

receiver

receive pkt0, send ack0
receive pkt1, send ack1

receive pkt3, buffer,
send ack3

receive pkt4, buffer,
send ack4

receive pkt5, buffer,
send ack5

rcv pkt2; deliver pkt2,
pkt3, pkt4, pkt5; send ack2

Q: what happens when ack2 arrives?

