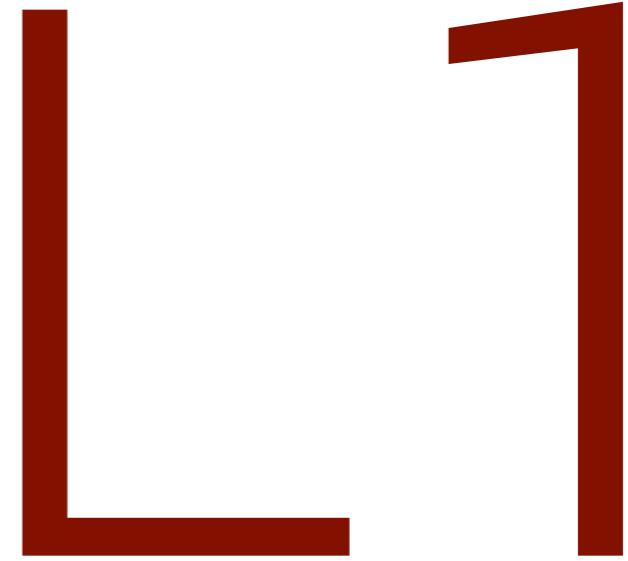


Slides

L1 - L6

Fundamentals of Databases

Alvaro A A Fernandes, SCS, UoM



Introduction to Data Management

Fundamentals of Databases
Alvaro A A Fernandes, SCS, UoM

Readings

X

This lecture relies on your having read the assigned mandatory readings associated with this lecture:

Abraham Silberschatz, Henry F. Korth, and S. Sudarshan. Database System Concepts. 6th Edition. McGraw-Hill, 2006. pp. 1-35.

You can download the material using this link:

<https://contentstore.cla.co.uk/secure/link?id=4151d4df-4e05-e611-80bd-0cc47a6bddeb>

If you fail to study the material, you're likely to find the lecture harder to follow.

More information is available in the Learning Activities Manual.

In This Lecture

x

- We will learn that data is an enterprise asset and that DBMSs are crucial to manage it well.
- We will learn about the importance of adopting different levels of abstraction in designing and implementing databases.
- We will learn that data models lead to a distinction between schemas and instances that enables a logical view of the data.

Database Management Systems

2

- Database management systems (DBMSs) can be seen in the broader context of software engineering strategies.
- Engineering software is hard and costly.
- It is not cost-effective to develop every piece of software from scratch.
- It is not cost-effective to develop every piece of software in a bespoke way.
- We factor out certain functionalities into software products that act as service providers for applications.
 - ▶ We don't build bespoke operating systems (OSs) from scratch.
 - ▶ The services they provide are universally useful and comprise a stable set.

Database Management Systems

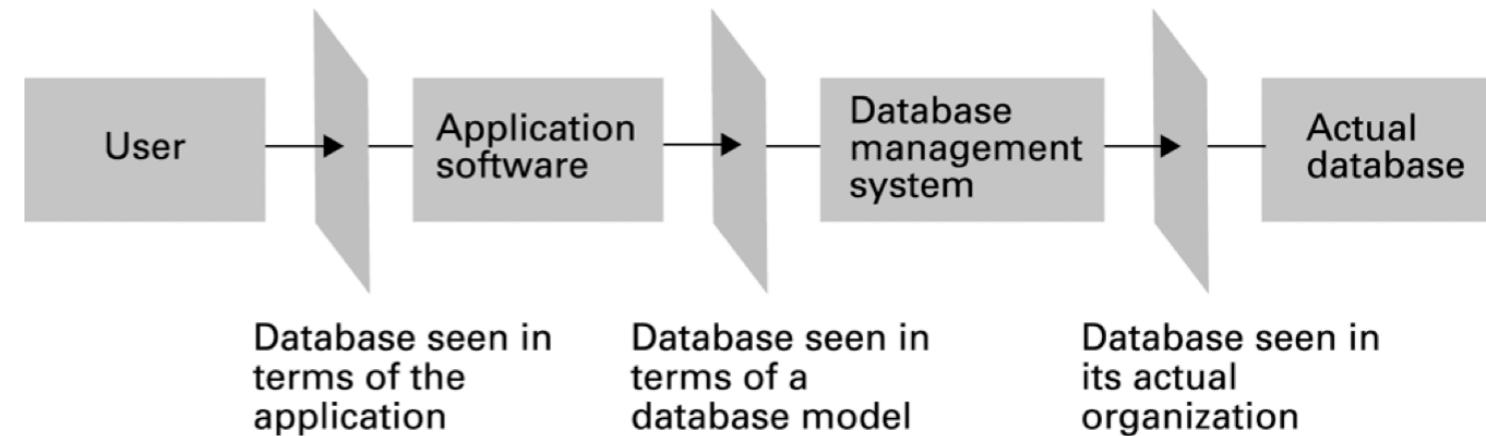
3

- DBMSs, like OSs, factor out functionalities as services that they provide for applications because it has been shown over decades that:
 - ▶ data-related services are universally useful,
 - ▶ there is a stable set of them.
- For that, DBMSs enforce
 - ▶ several principles (such as logical/physical independence)
 - ▶ using advanced techniques (such as transaction and recovery management)
 - ▶ that emerge to software engineers as a design methodology.

How do we view data? How should we manage it?

4

- Data is an asset, as crucial to an organization as capital or human resources.
- A database management systems (DBMS) is a sophisticated tool for managing data intra- and inter-organizationally.
- It exposes services that allow people and applications to have a single, unified view of the organization through its data assets.



Where are DBMSs used?

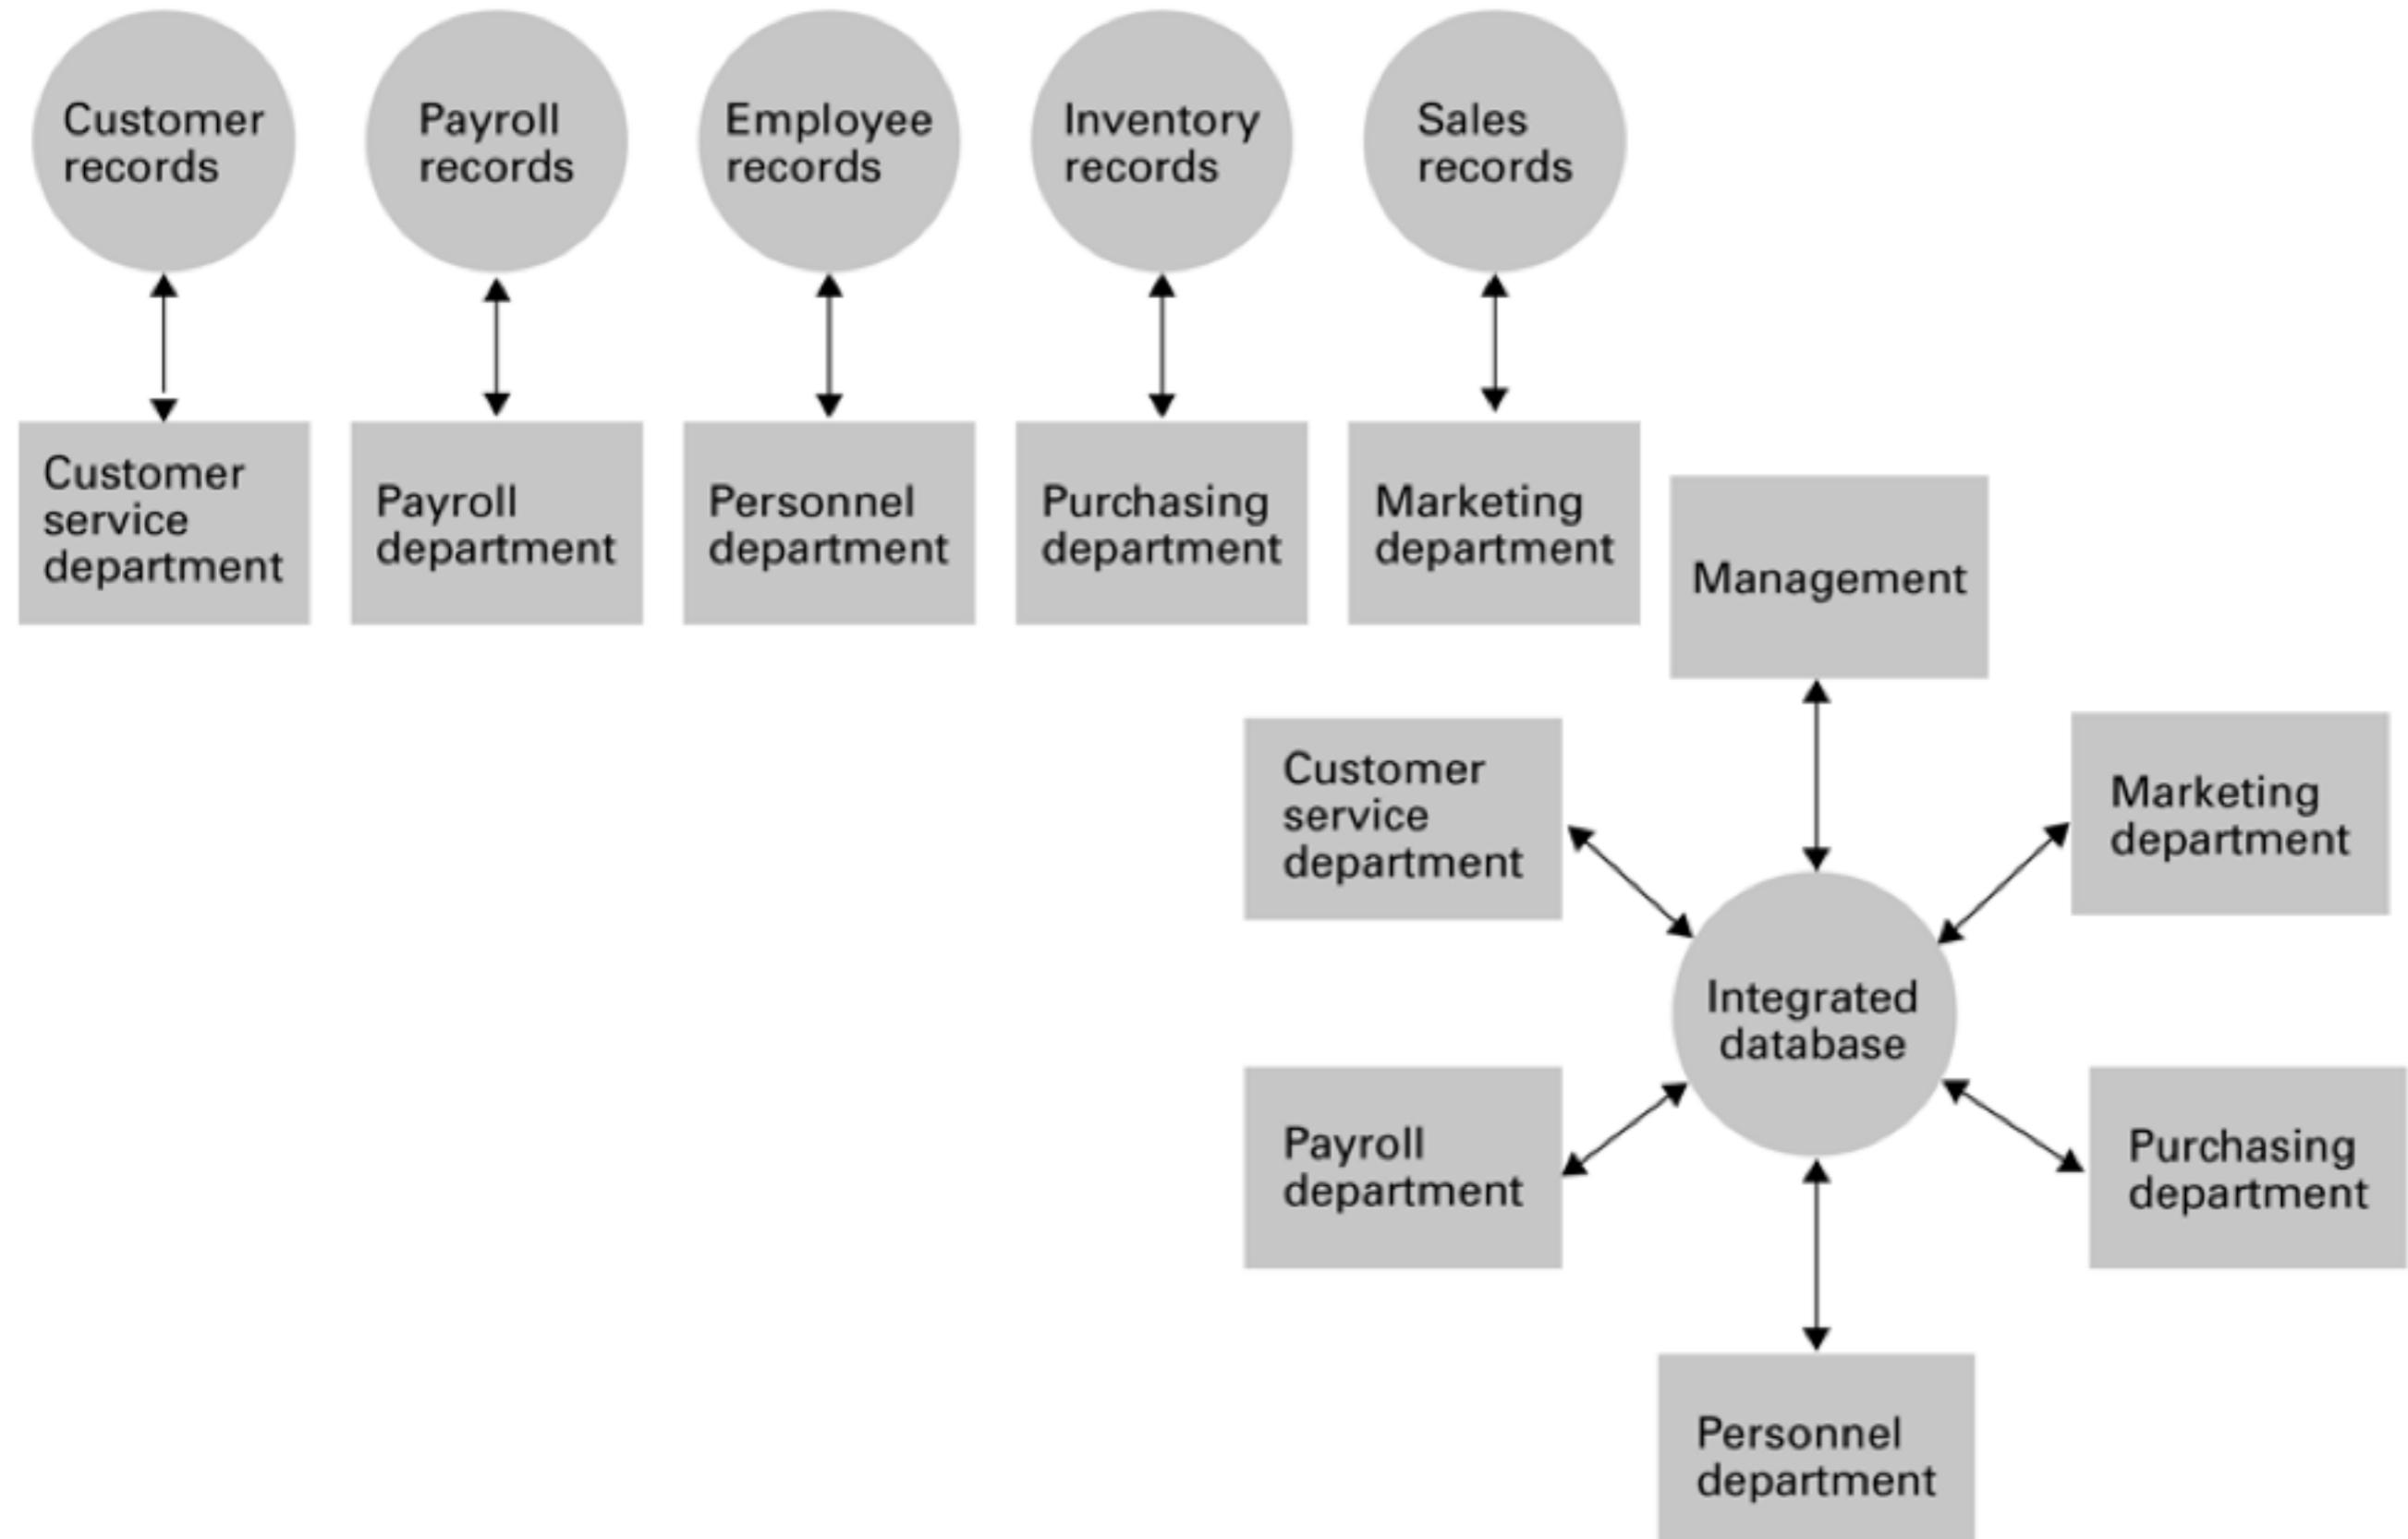
5

- Inside organizations: to support front-end and back-end business processes
- Across organizations: to allow cooperation and coordination



Why DBMSs? Why not use files?

6



Why DBMSs? Why not use files?

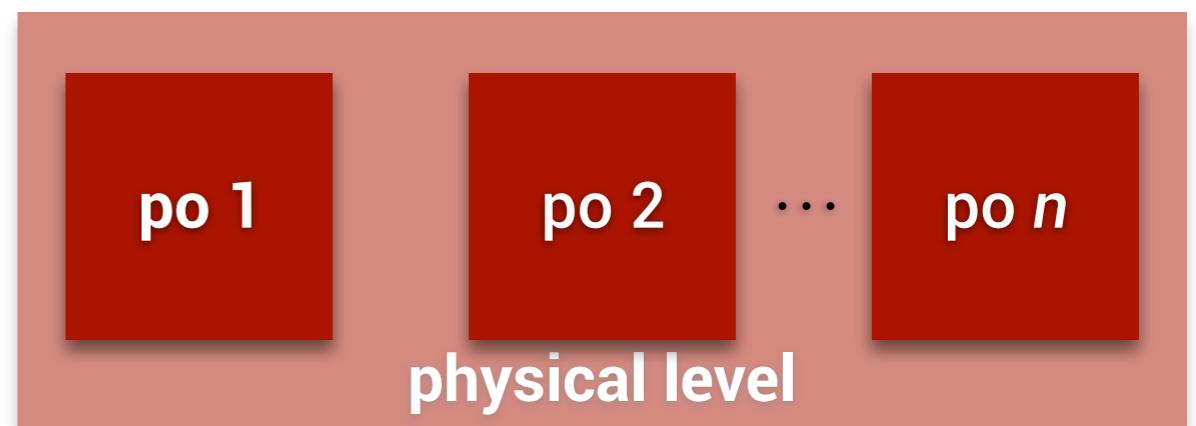
7

- The applications that matter to an organization change faster than the data they use.
- It is better to loosely couple application and data.
- This also allows decoupling a logical view of data from storage technology and formats.
- A DBMS enables such decoupling, so applications can change as needed.
- A DBMS therefore lowers the cost of developing new applications.
- A DBMS offers a set of sophisticated services beyond storage, including:
 - ▶ efficient, scalable querying and updating;
 - ▶ concurrent, fault-tolerant transactions;
 - ▶ privacy, and role-based access and authorization control.

How do different levels of abstraction in DBMSs work?

8

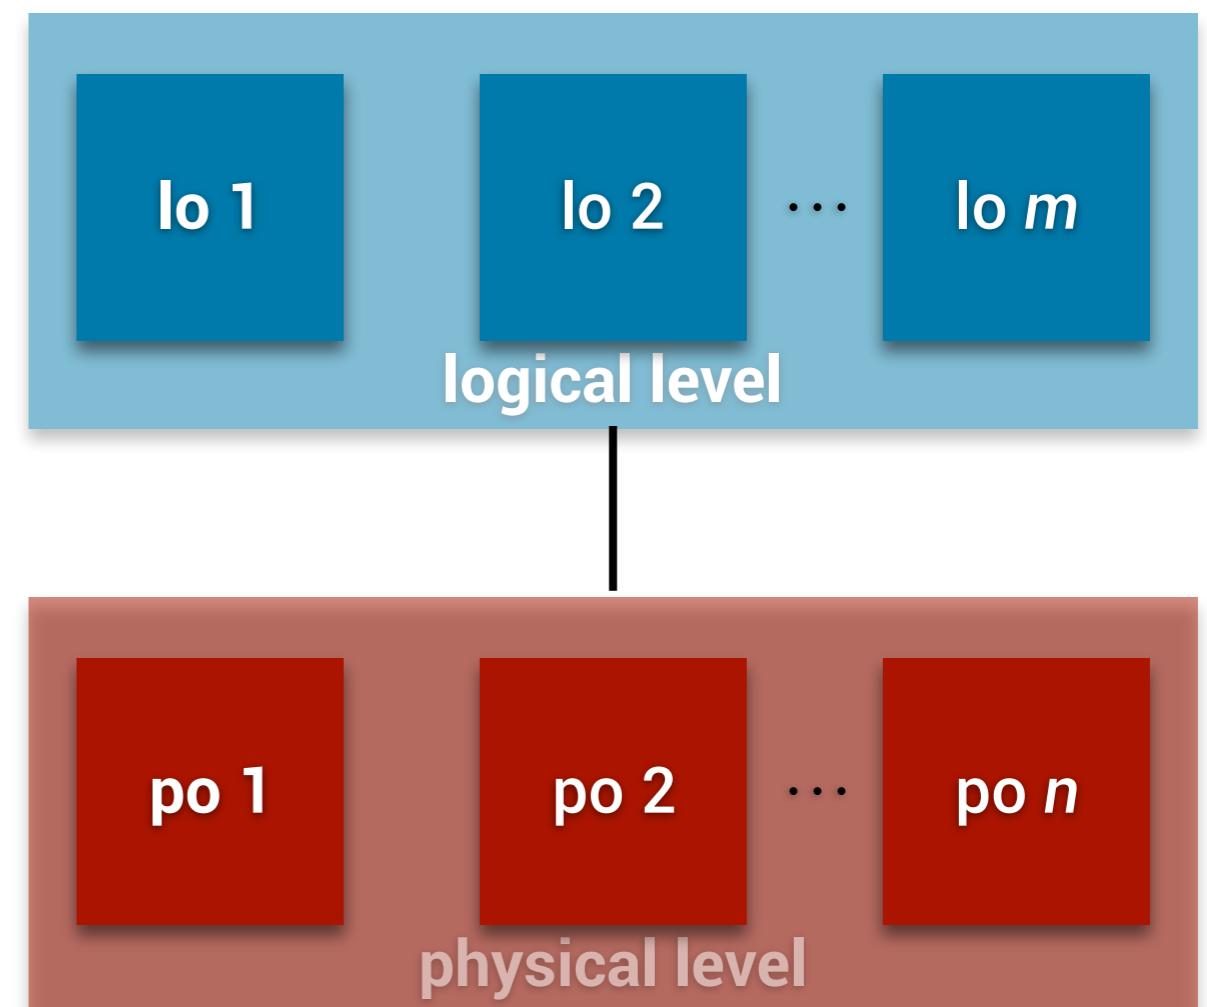
- The physical level
 - ▶ defines formats, records, files, clustering, indexes, compression, replication, redundancy,
 - ▶ is expressed as a storage or **physical model**.



How do different levels of abstraction in DBMSs work?

9

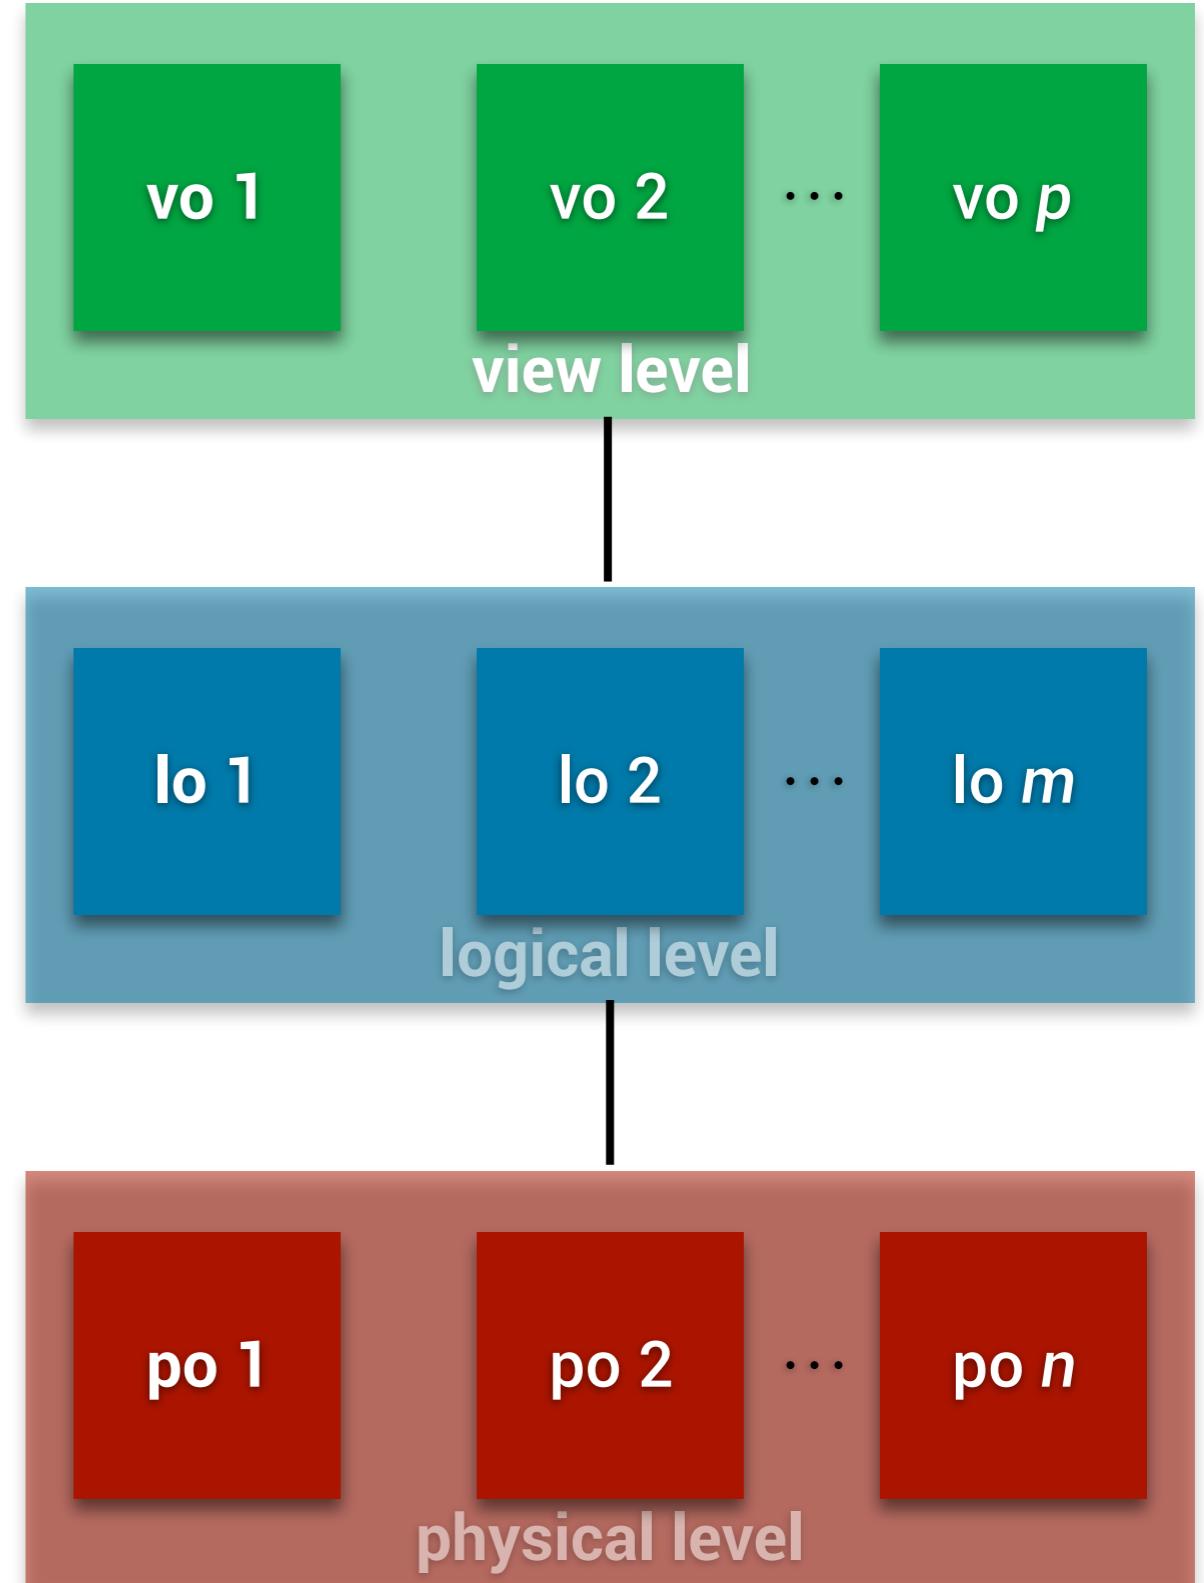
- The logical level
 - ▶ defines the concepts that give the entire enterprise a specific character:
 - supplier, product, stock_unit, customer, order, payment, etc.,
 - employee, stores, banker, accounts_payable, accounts_receivable, etc.;
 - ▶ is expressed by a **logical data model**, possibly derived from a **conceptual model**.



How do different levels of abstraction in DBMSs work?

10

- The view level
 - ▶ defines what software applications work with and partitions functionality:
 - Human resource applications do not see balances in bank account.
 - Point-of-sale applications do not see employee salaries.
 - ▶ is expressed by a **database application model** (comprising queries, views, triggers, etc.).



What is a (logical) data model?

11

- A data model is a formal construct for describing data, their relationship, semantics and constraints.
- The relational model is the most widely-used:
 - ▶ rows of columns are organized in tables that may reference one another, and
 - ▶ if they do, then no dangling allowed.
- It comes with special database languages for interaction, e.g., a relational algebra.
- A DBMS supports one (or more) data models and associated languages.
- The Oracle DBMS, e.g., supports:
 - ▶ relational,
 - ▶ object-relational,
 - ▶ XML-based, and
 - ▶ document-based
 - ▶ using SQL, XQuery, etc.

What are schemas and instances?

12

- Each data model defines:
 - ▶ whether there is a distinct notion of a schema (e.g., recent data models don't),
 - ▶ what the schema is allowed to declare (e.g., which structures and which constraints).
- Then, instances must be valid w.r.t. the schema:
 - ▶ must have compatible structure, and
 - ▶ must satisfy the constraints.
- A schema defines the logical structure of the database, e.g.:
`share(account, customer)`
- An instance is the actual content of the database at a particular point in time, e.g.:
`share(015467, 'Mary Quaint')`
`share(015467, 'Vivienne Eastwood')`
- Schemas are similar to types and variables.
- Instances are similar to variable values (i.e., they change over time).

What different types of database language there are?

13

- There are different kinds of database language:
 - ▶ DDL: Data Definition Language
 - ▶ DML: Data Manipulation Language
 - ▶ (D)QL: Query Language
- Again, the terminology is not crisp:
 - ▶ Sometimes, there is no special DDL.
 - ▶ Often, QL is not distinguished from DML.
- The DDL updates the data dictionary.
- It is a declaration language: it defines the schema (i.e., types and constraints).
- The DML updates the database.
- It is a state-transition language, it changes values (e.g., rows and columns).
- The QL is a pure expression language: evaluation returns values without changing the database instance.

Why are different types of database language needed?

14

- The main advantage of this approach is enforcing **separation of concerns**, e.g.
 - ▶ DDL is enterprise-wide, for all applications.
 - ▶ DML is for reflecting changes caused by events in the life of the enterprise (e.g., a new account).
 - ▶ QL is for retrieving information (e.g., the current average balance).
 - ▶ The lack of assignment makes the QL easier to rewrite for optimization purposes.
- In practice, the distinction is not often enforced.
- SQL \supseteq DDL \cup DML \cup QL
- SQL has become extremely expressive with modern extensions.

Why are different types of database language needed?

15

- The main advantage of this approach is enforcing **separation of concerns**, e.g.
 - ▶ DDL is enterprise-wide, for all applications.
 - ▶ DML is for reflecting changes caused by events in the life of the enterprise (e.g., a new account).
 - ▶ QL is for retrieving information (e.g., the current average balance).
- ▶ The lack of assignment makes the QL easier to rewrite for optimization purposes.
- In practice, the distinction is not often enforced:
 - ▶ SQL \supseteq DDL \cup DML \cup QL
 - ▶ SQL has become extremely expressive with modern extensions.

In the Next Lecture

16

- We'll begin to explore the steps in **database design**, which will take us all of eight lectures to complete.
- The first step is to learn about the approach to conceptual modelling known as **entity-relationship modelling**.

L 2

Designing Databases: The Entity-Relationship Approach

Fundamentals of Databases

Alvaro A A Fernandes, SCS, UoM

Readings

X

This lecture relies on your having read the assigned mandatory readings associated with this lecture:

Ramez Elmasri and Shamkant Navathe. Fundamentals of Database Systems. New International Edition. Pearson, 2013. pp. 201-245.

You can download the material using this link:

<https://contentstore.cla.co.uk/secure/link?id=73ff189a-8463-e611-80c6-005056af4099>

If you fail to study the material, you're likely to find the lecture harder to follow.

More information is available in the Learning Activities Manual.

In the Previous Lecture

x

- We learned that data is an enterprise asset and that DBMSs are crucial to manage it well.
- We learned about the importance of adopting different levels of abstraction in designing and implementing databases.
- We learned that data models lead to a distinction between schemas and instances that enables a logical view of the data.

In This Lecture

X

- What are database applications?
- Why conceptual modelling?
- What are the phases in the design of database applications?
- What is entity-relationship (ER) modelling?

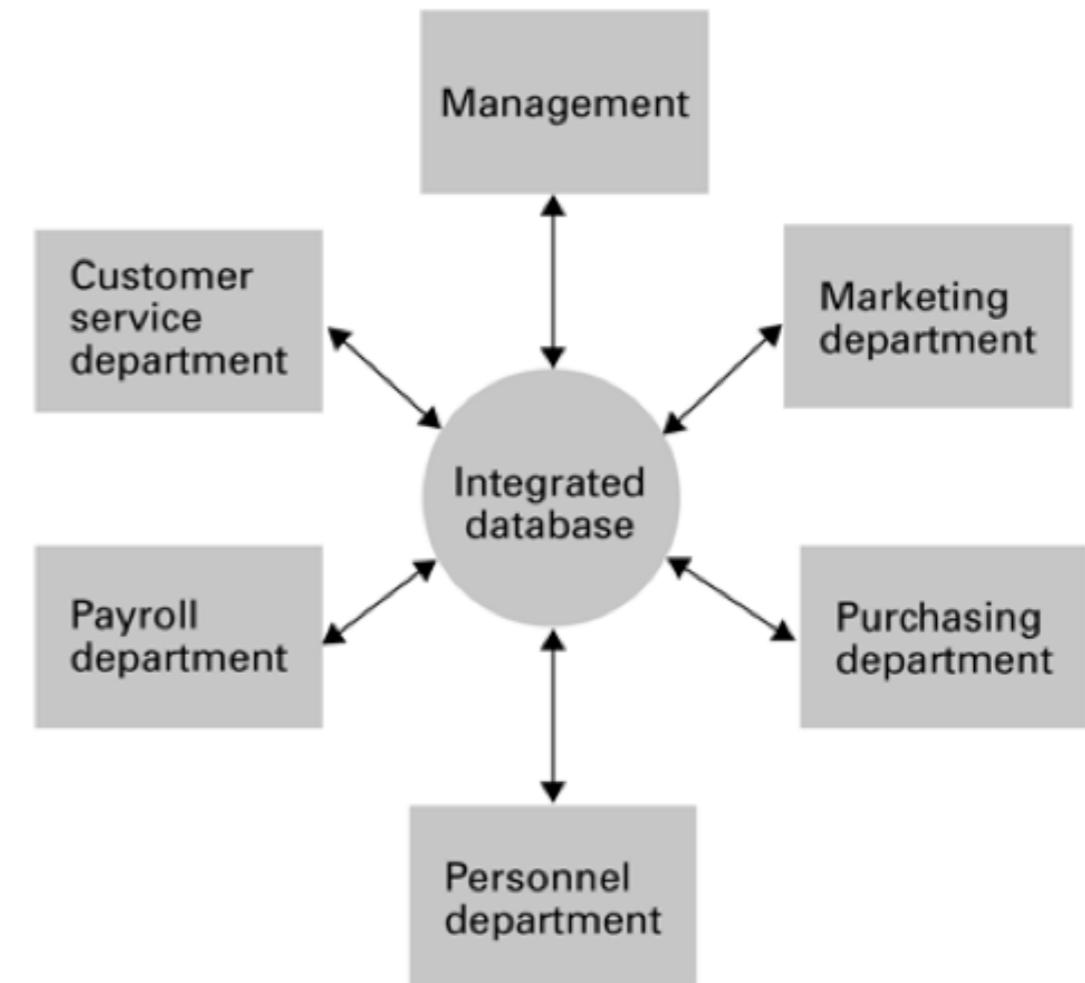
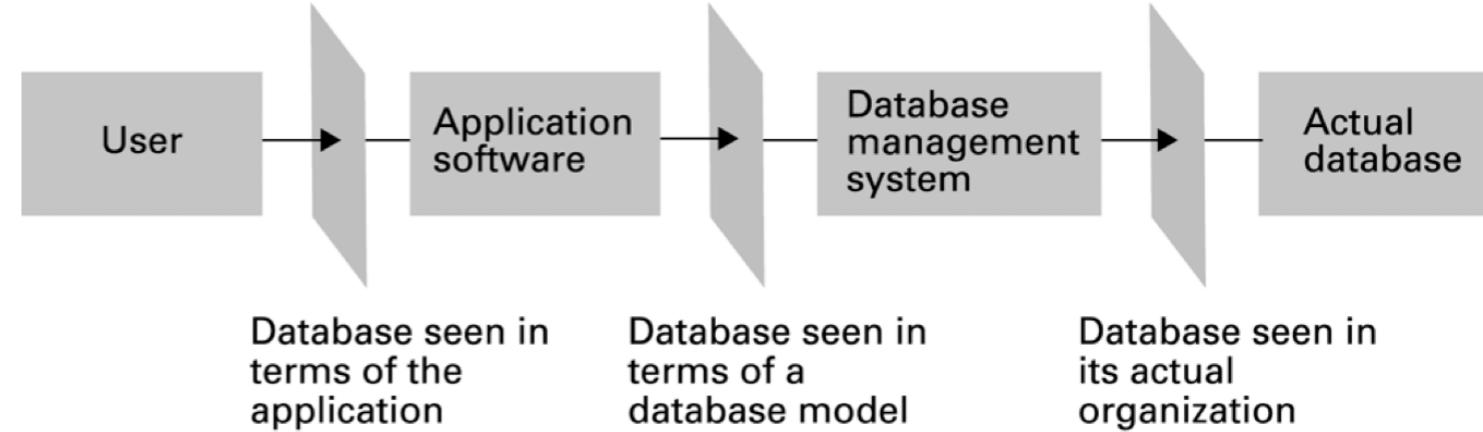
Conceptual Design of Database Applications

18

- **Database applications**

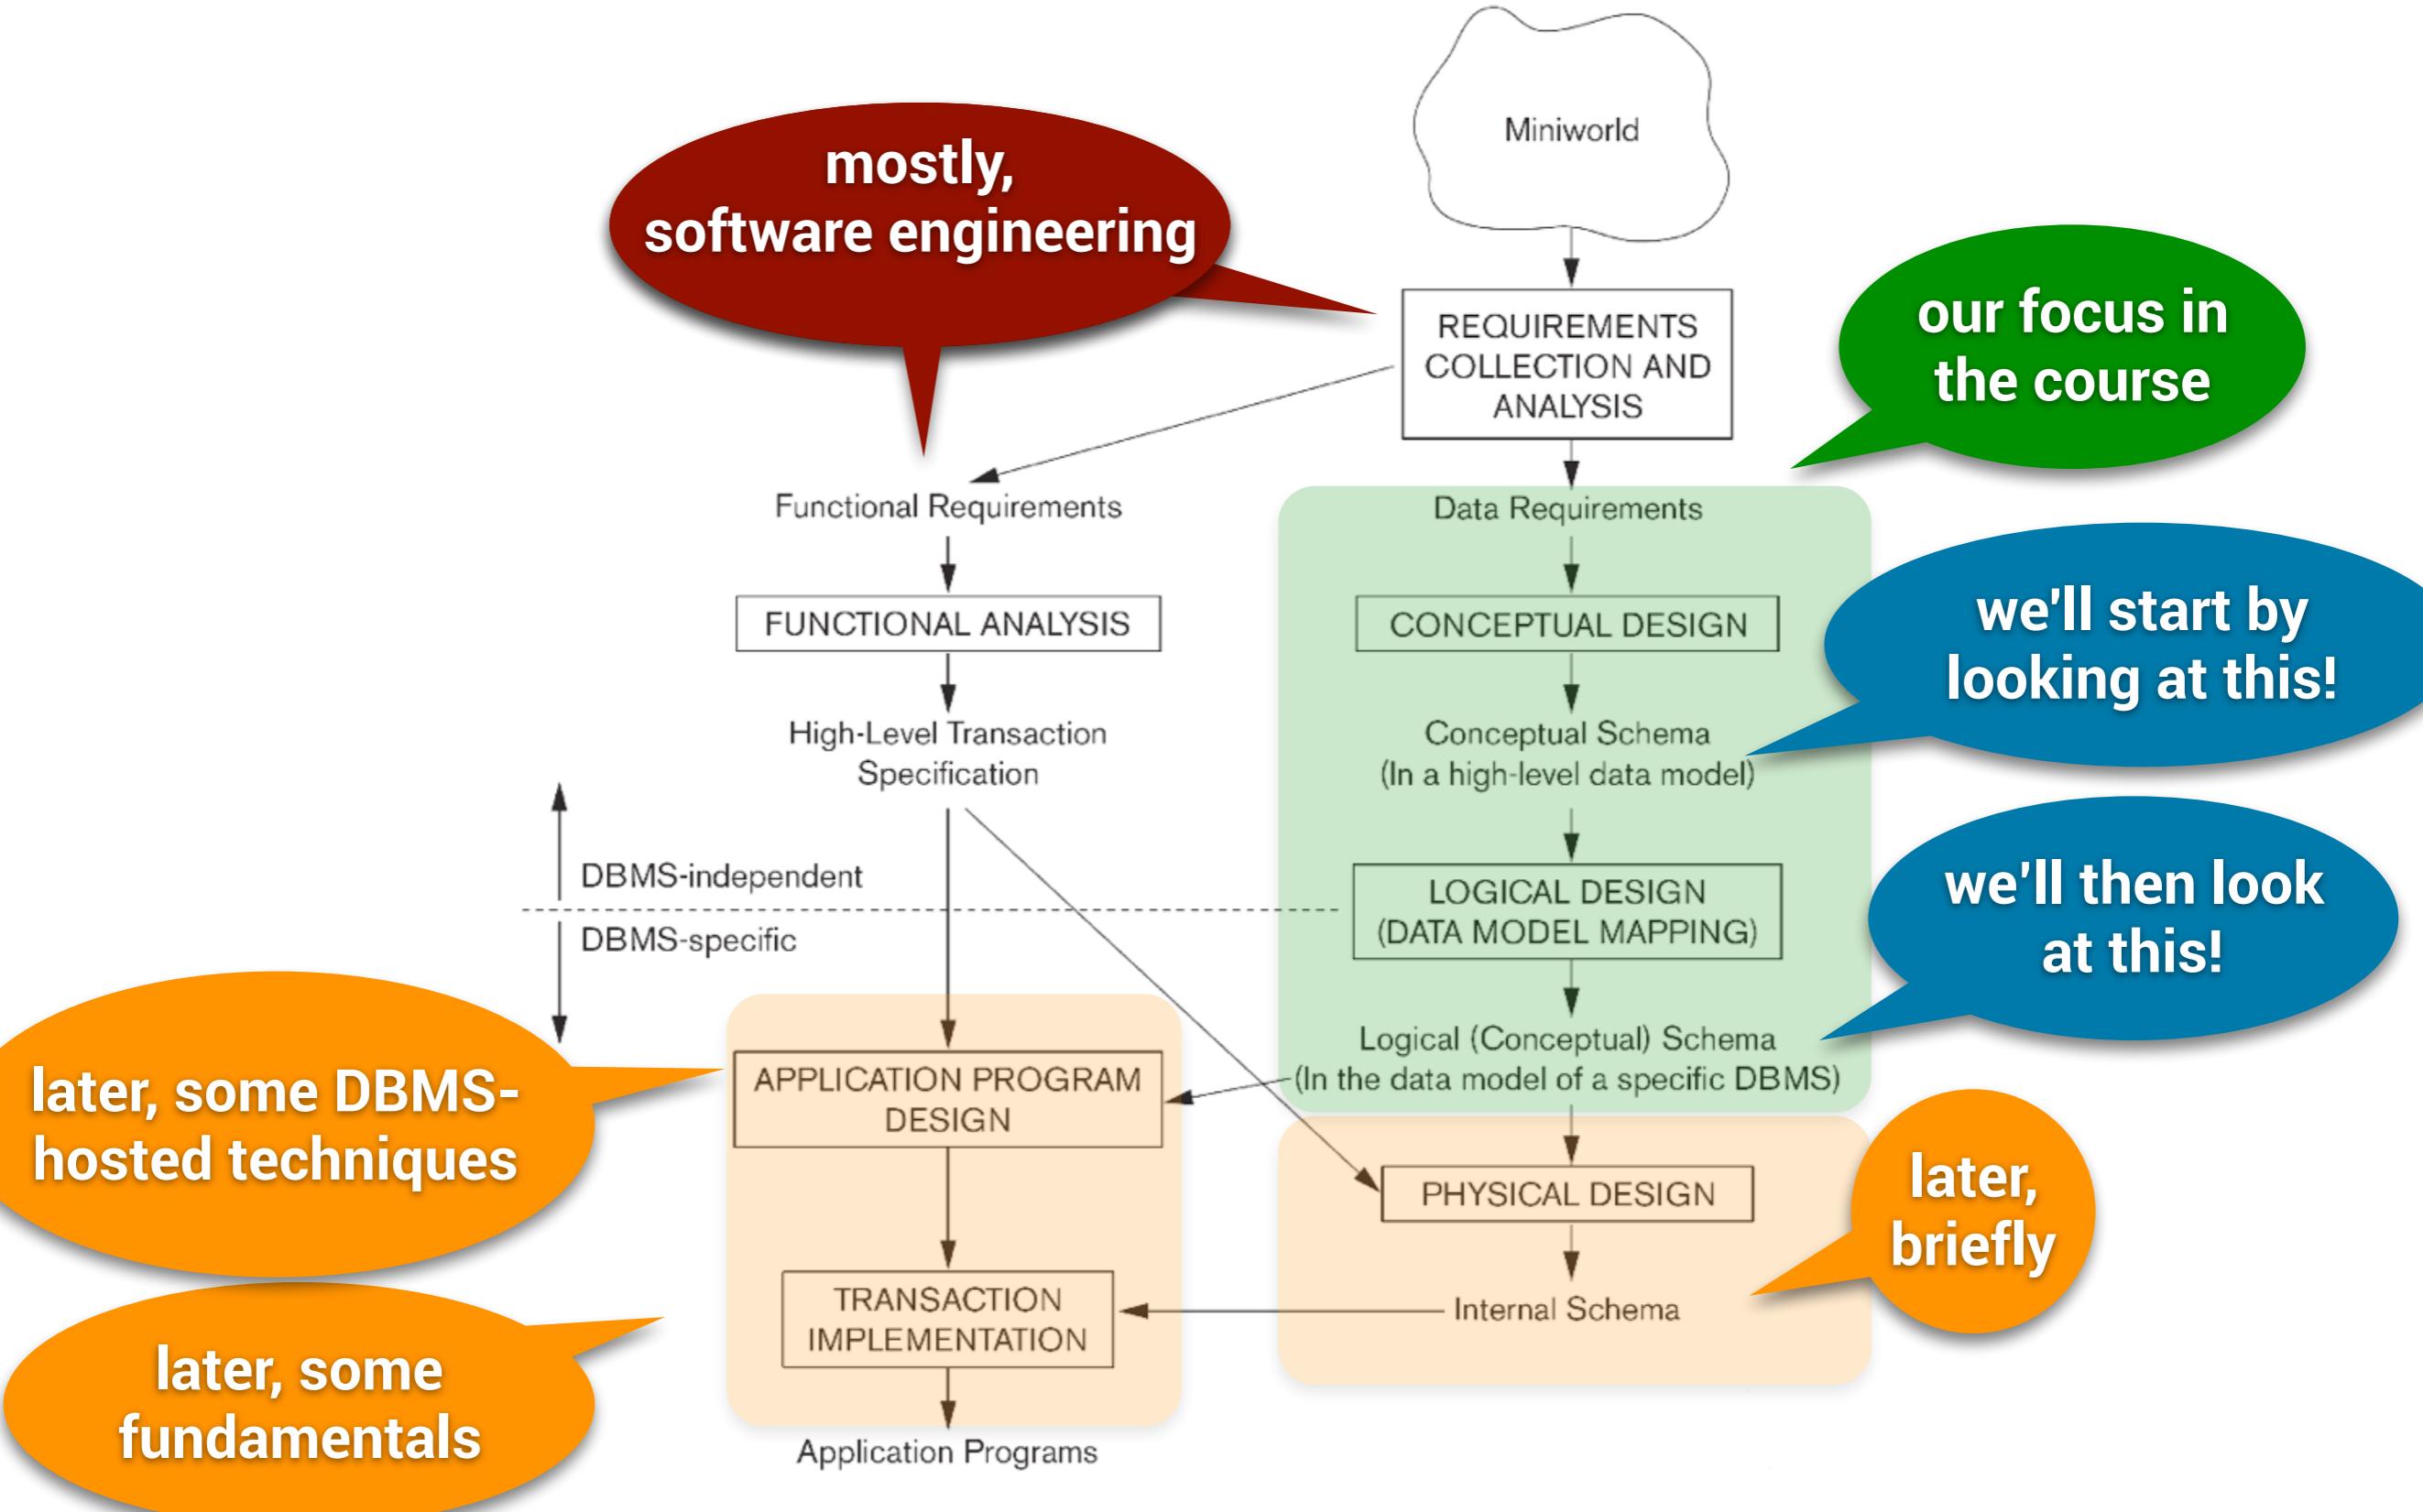
comprise:

- ▶ an integrated database
- ▶ plus application programs that
- ▶ implement queries and updates to that database



Database Application Design Phases

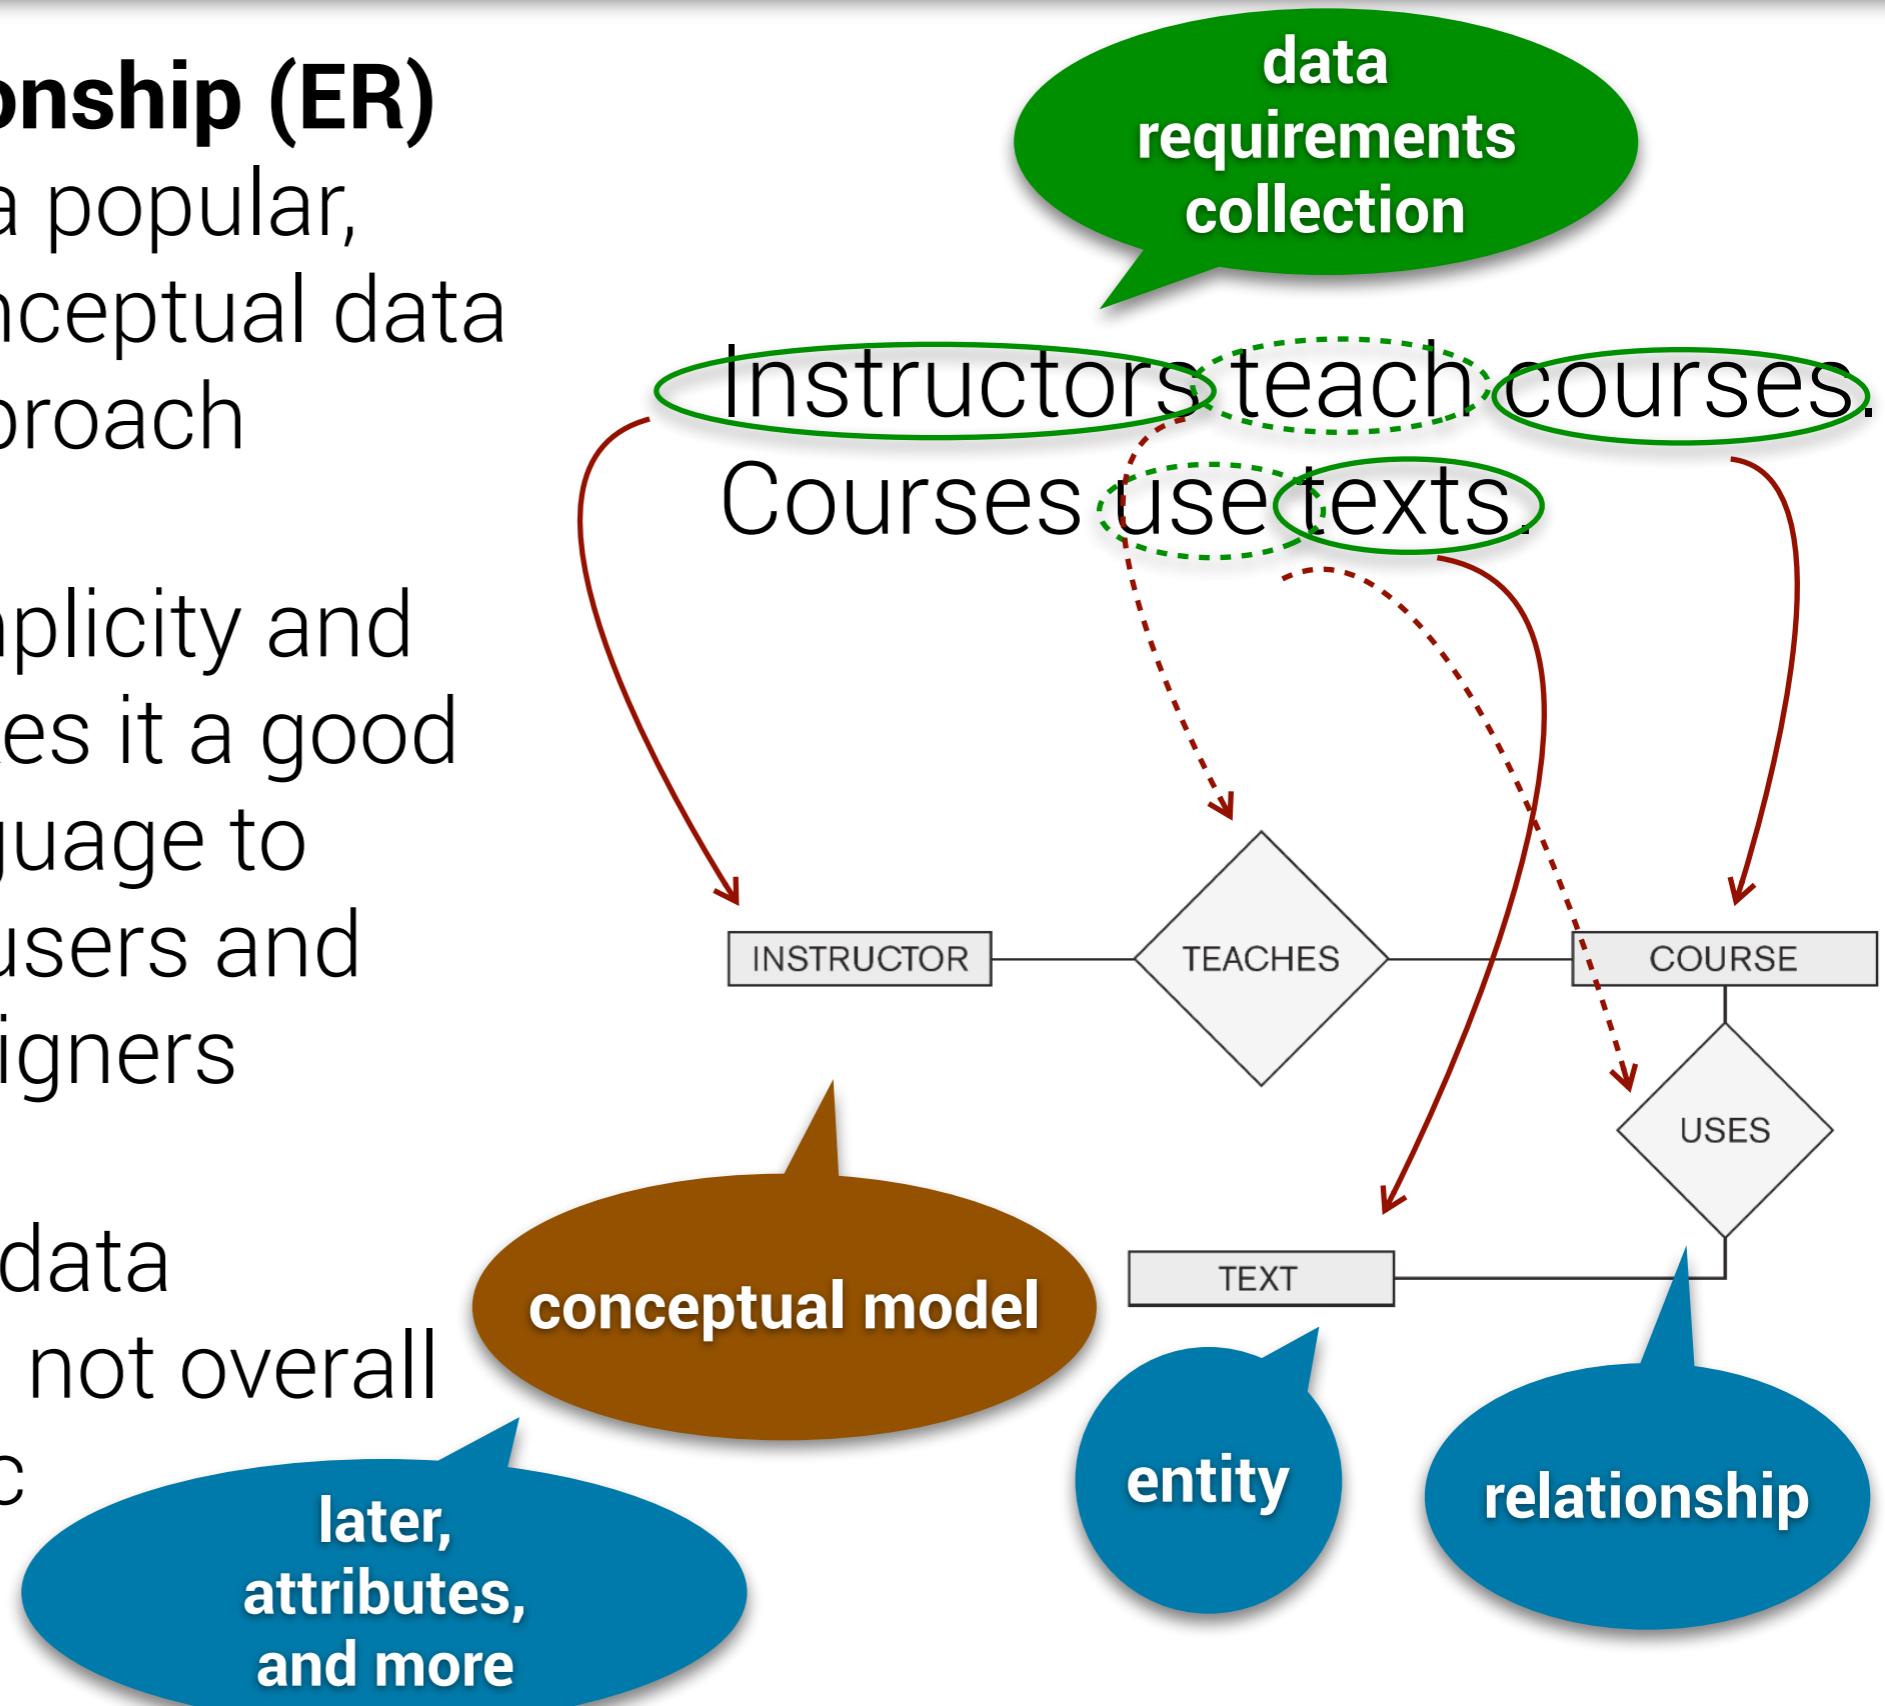
19



Conceptual Design of Database Applications

20

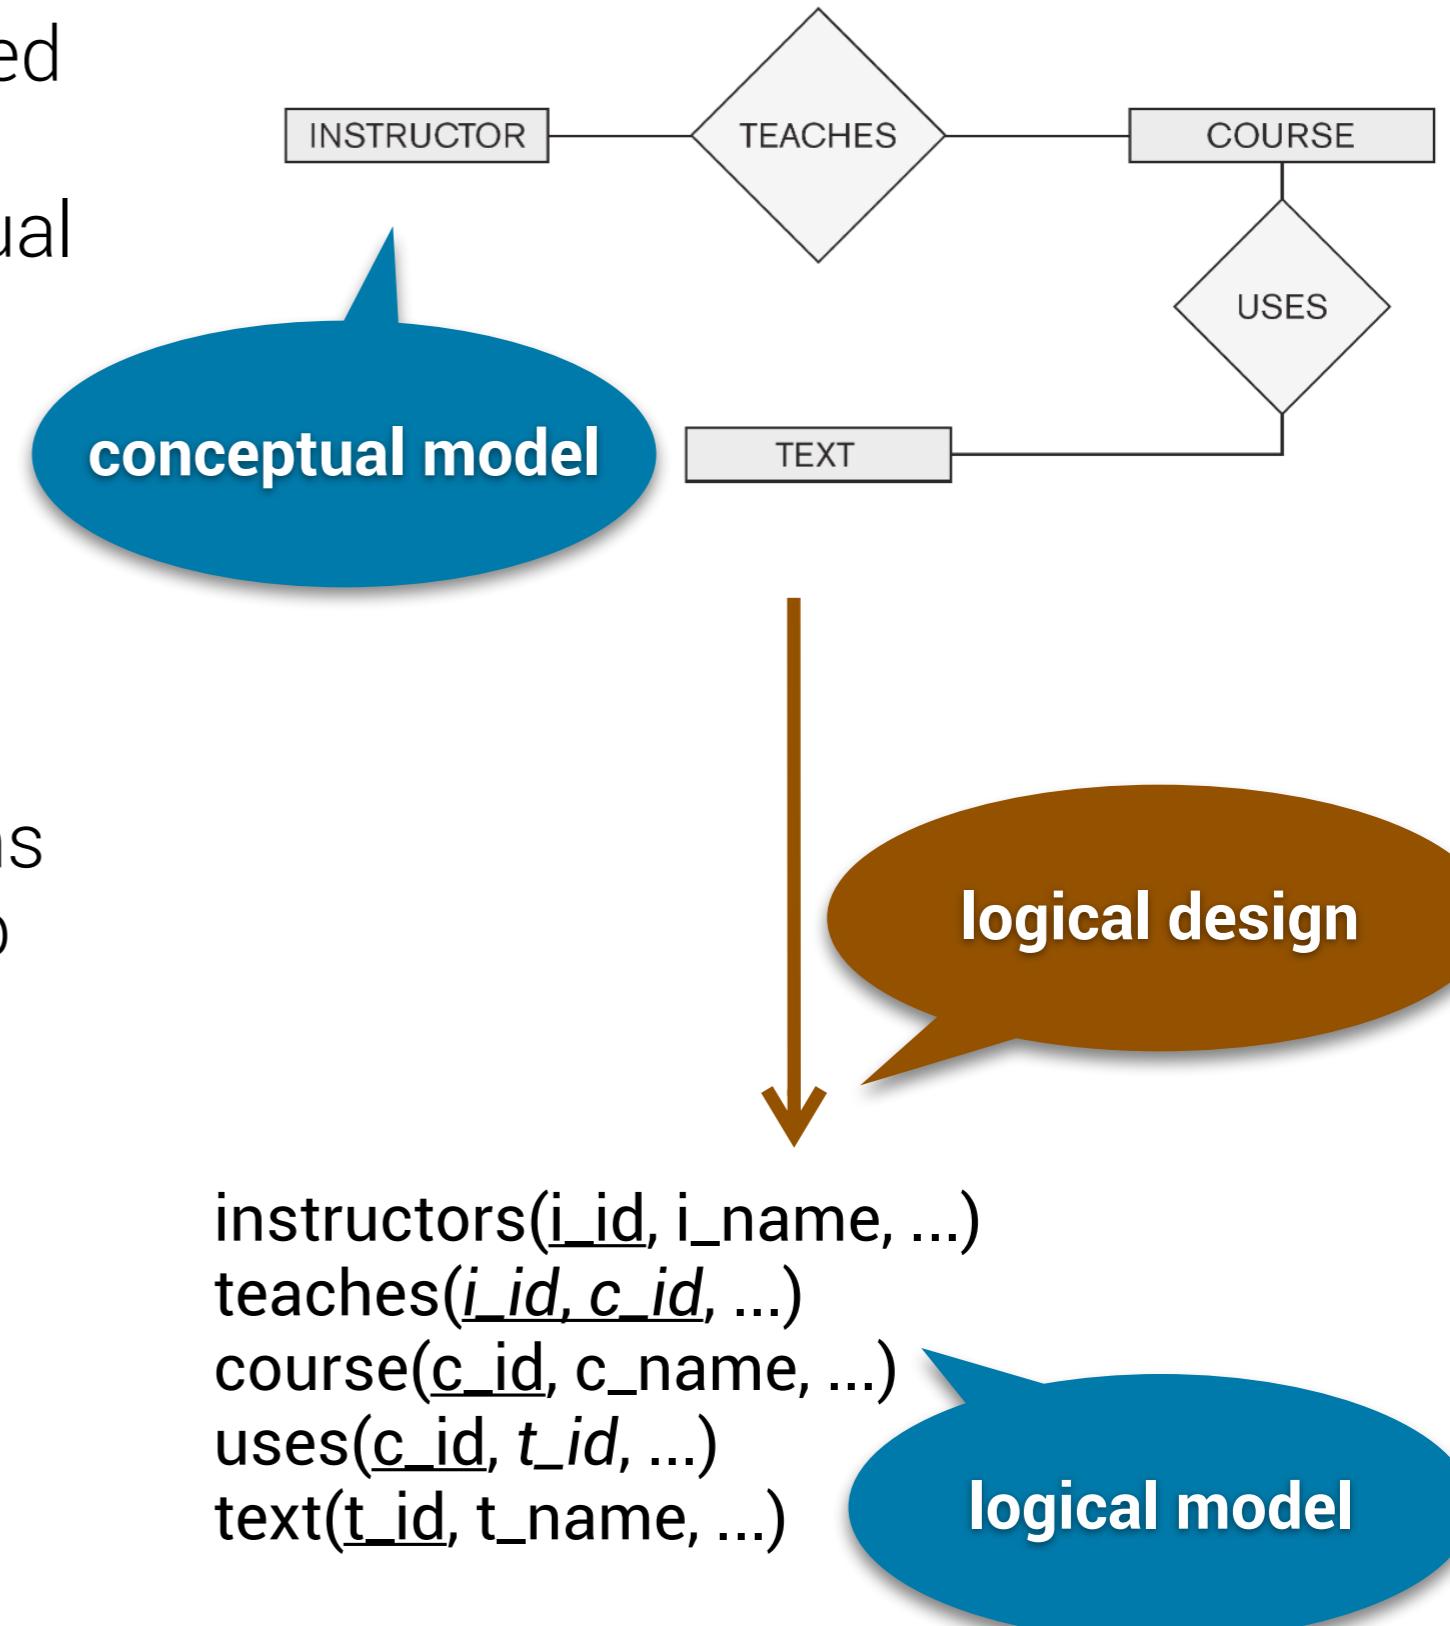
- **Entity-Relationship (ER) Modelling** is a popular, high-level conceptual data modelling approach
- Its formal simplicity and flexibility makes it a good common language to connect end users and database designers
- It focuses on data requirements, not overall business logic



From Conceptual to Logical Models

21

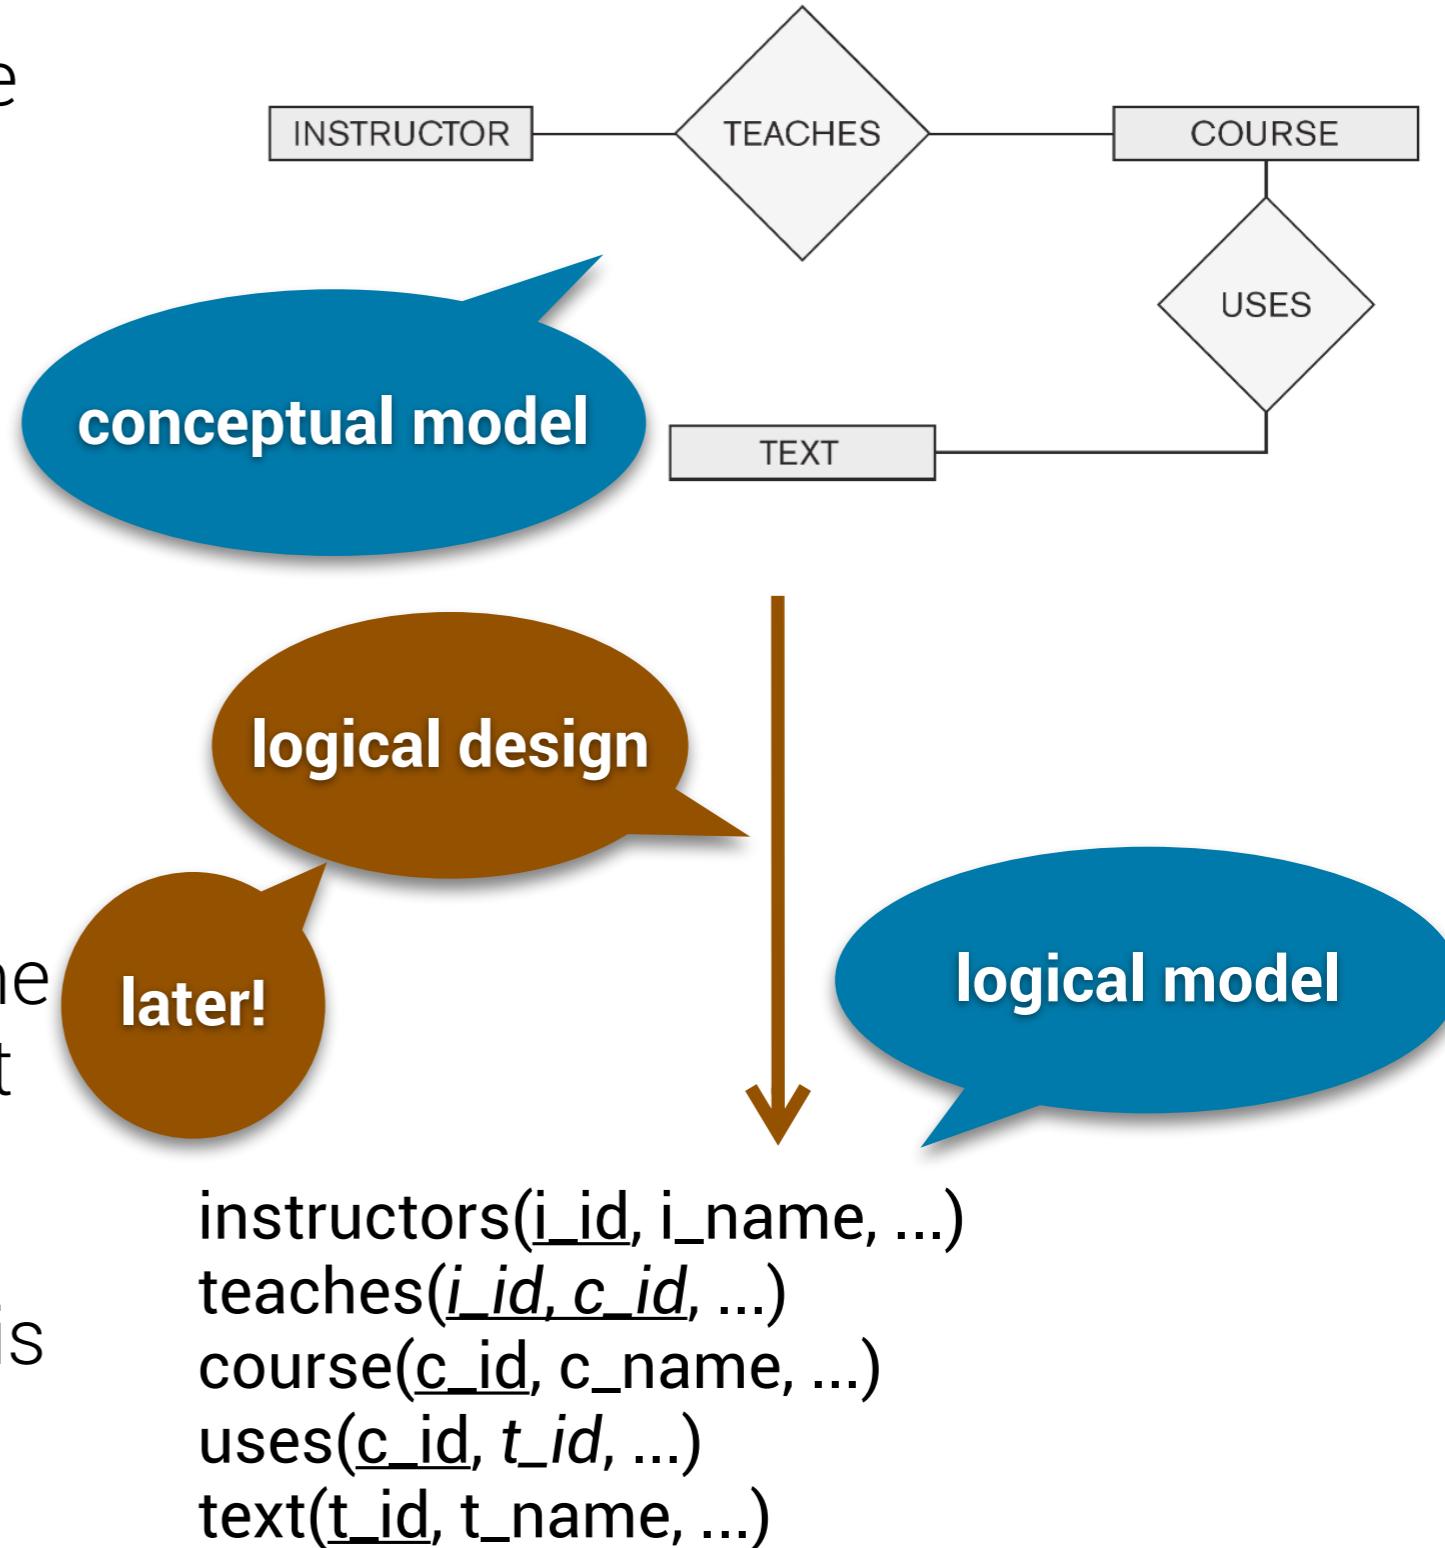
- A **conceptual model** (also called a **conceptual schema**) is the main outcome of the conceptual design phase.
- It is a rigorous, but high-level, description of the data requirements
- It includes detailed descriptions of the entity types, relationship types, and of important constraints on the latter
- It aims to make the mapping into a more rigorous logical model fairly systematic.



From Conceptual to Logical Models

22

- A conceptual model is suitable to being mapped into an **(implementation) logical model or schema**
- In our case, the conceptual model is an **ER model**
- By a **logical model or schema** we mean one (e.g., relational) that is directly supported by the DBMS we intend to implement the database application on
- In our case, the logical model is a **relational model**



ER Modelling: Basic Constructs

23

- The basic constructs in ER modelling capture the data requirements in the form of:
 - ▶ Entity types
 - ▶ Attribute types
 - ▶ Relationship types

similar to
relations

similar to
attributes

but less
constrained

how are they
captured then?

no direct
counterpart in the
relational model

ER Modelling: Attribute Types

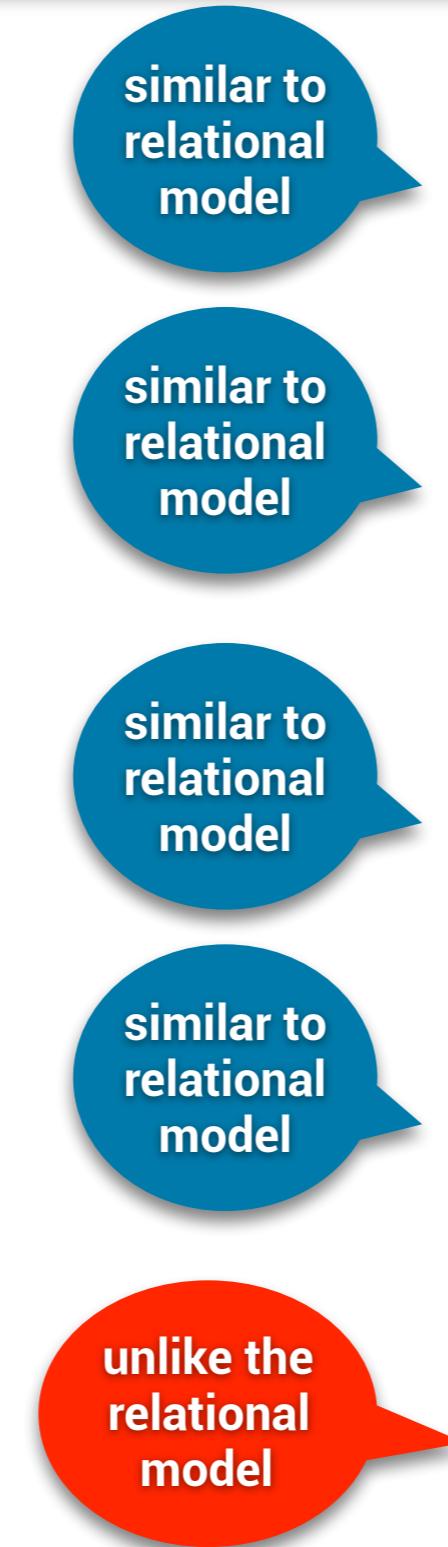
24



- Attribute types can be:
 - ▶ **Composite v. simple (or atomic)**
 - ▶ **Single-valued v. multivalued**
 - ▶ **Stored v. derived**
 - ▶ **NULL-valued**
 - ▶ **Complex-valued** (i.e., arbitrarily-nested composite/multivalued)
- Try not to confuse the notion of a 'composite' attribute type with that of a 'complex' one

ER Modelling: Keys

25



- In ER modelling, a **key** (or **uniqueness**) **constraint** holds.
- Entities must be uniquely identifiable in every possible entity set of an entity type.
- Keys are attributes whose values are distinct for each individual entity in any entity set.
- There can be more than one key in an entity type.
- An entity type that has no natural or assigned key is said to be **weak**.

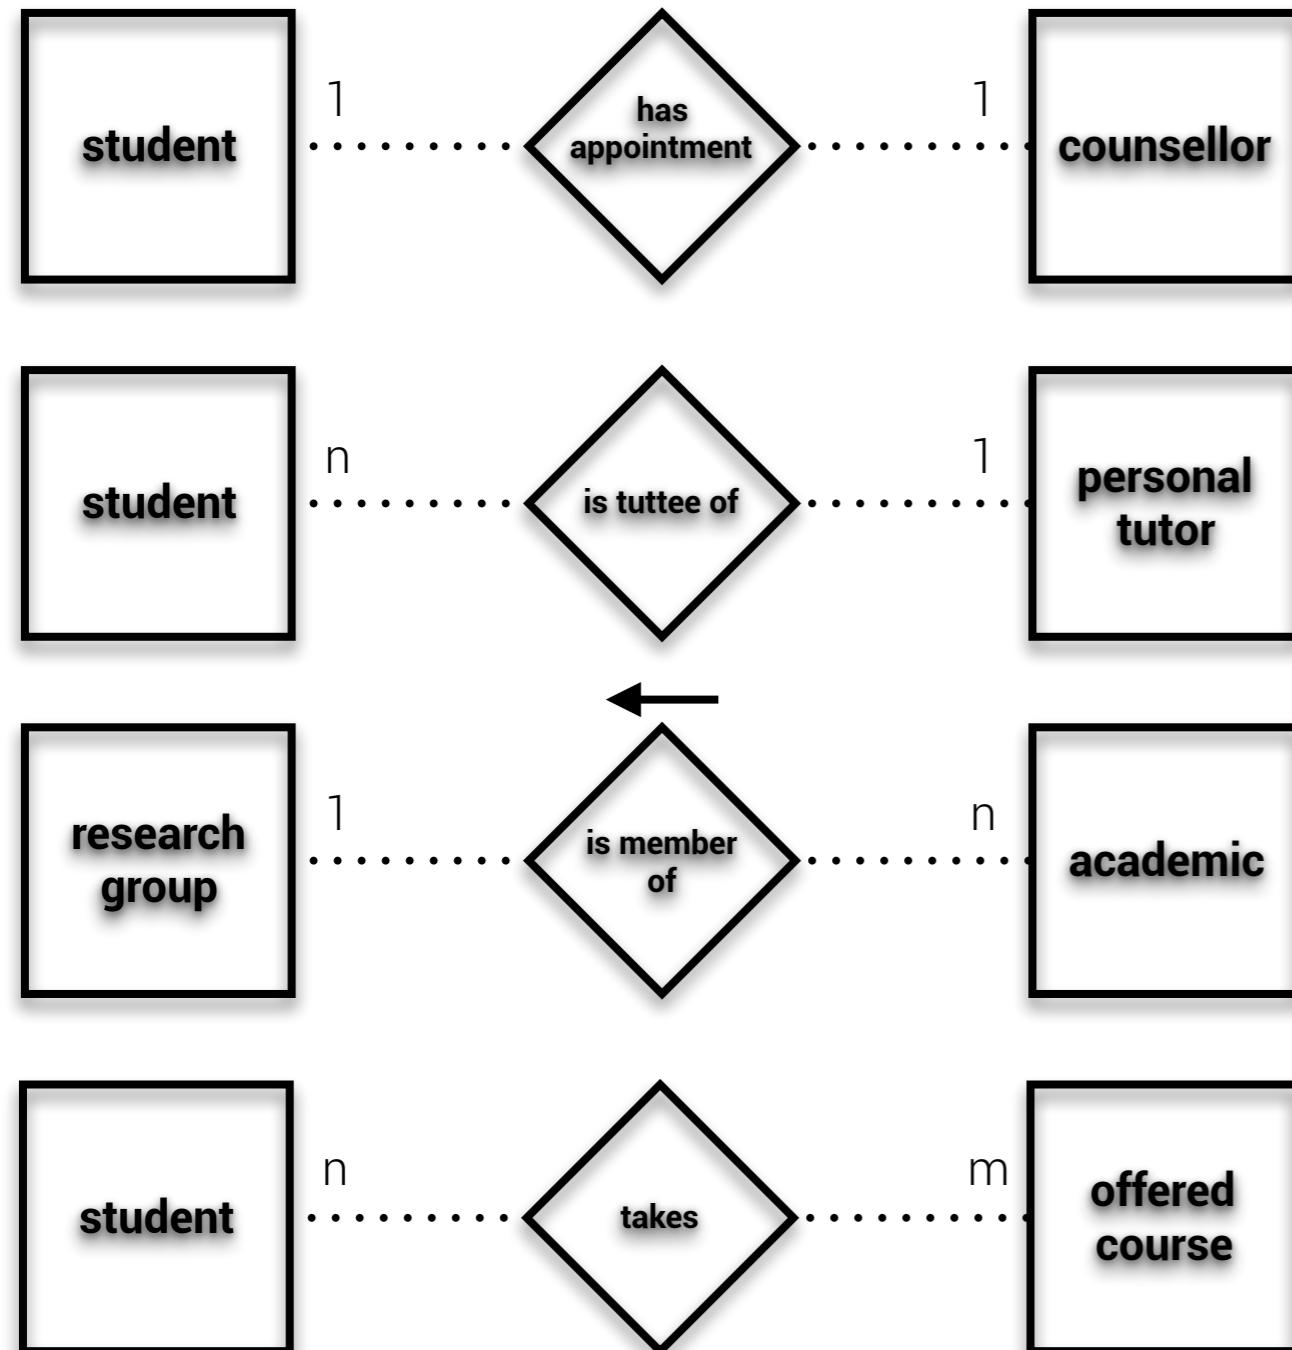
ER Modelling: A Word of Caution on Keys

26

- Try not to confuse the **ER** notion of '**key**' with the **relational** ones of '**primary key**' and '**foreign key**'.
- The notion of key in ER modelling is related (but not identical) to the notion of a primary/foreign key in a relational database.
- It is best compared to the relational notion of a *candidate key*, which a designer may choose to make primary or not.
- Also, foreign keys model relationships implicitly in the relational model, but the 'R' in 'ER' tells you that *relationships are modelled explicitly in the ER model*.

ER Modelling: Constraints on Relationship Types

27



- A **cardinality ratio constraint** specifies the maximum number of relationship instances that an entity can participate in
- It can be **1:1, 1:N, N:1, N:M** (or, equivalently, **M:N**)

How cardinality ratio constraints are captured

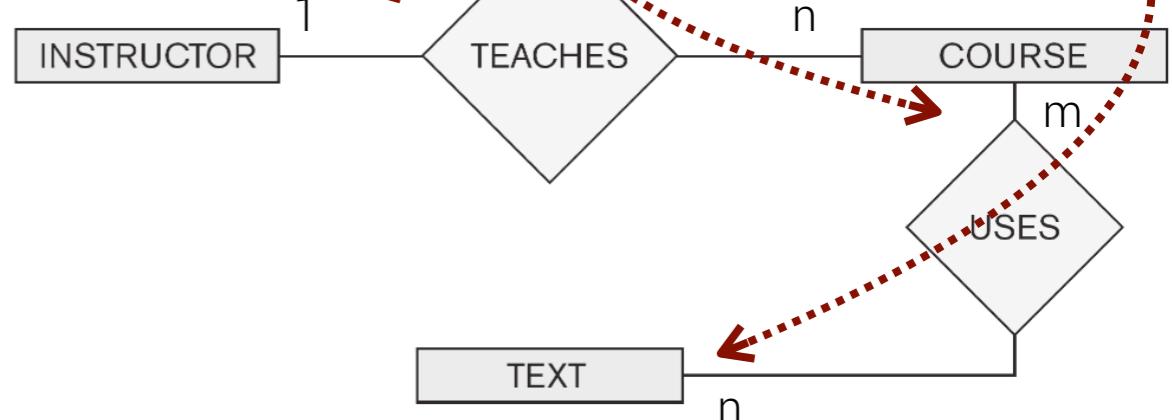
28

Instructors teach courses.
An *instructor* teaches
many courses.

A course is taught by **one** instructor only.

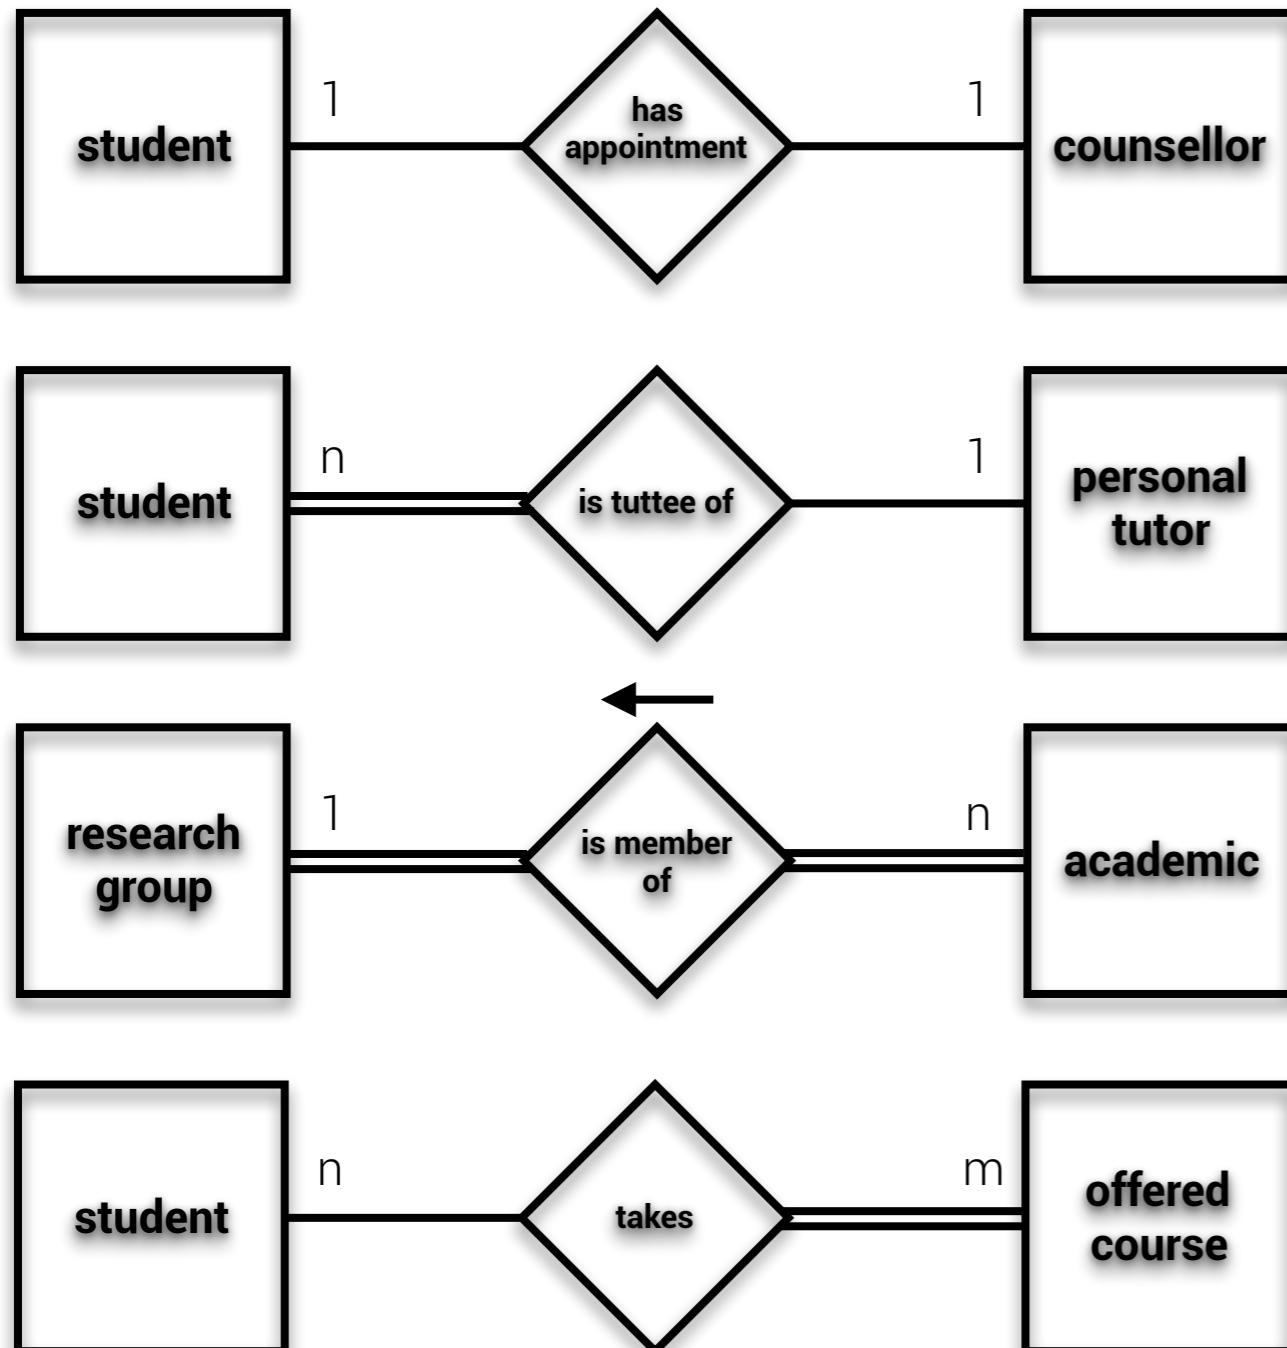
Courses use texts.

A course uses **many** texts.
A text is used by **many** courses.



ER Modelling: Constraints on Relationship Types

29



- A **participation constraint** specifies whether the existence of the entity depends on its being related to another entity
 - If participation is **total**, then every entity in the total entity set must participate in some relationship instance
 - If participation is **partial**, then some entities may not participate in any relationship instance

double line

single line

How participation constraints are captured

30

Instructors teach courses.

An instructor teaches many courses.

An *instructor* ~~may not~~ teach any course.

A course is taught by one instructor only.

A course ~~must~~ be taught by an *instructor*

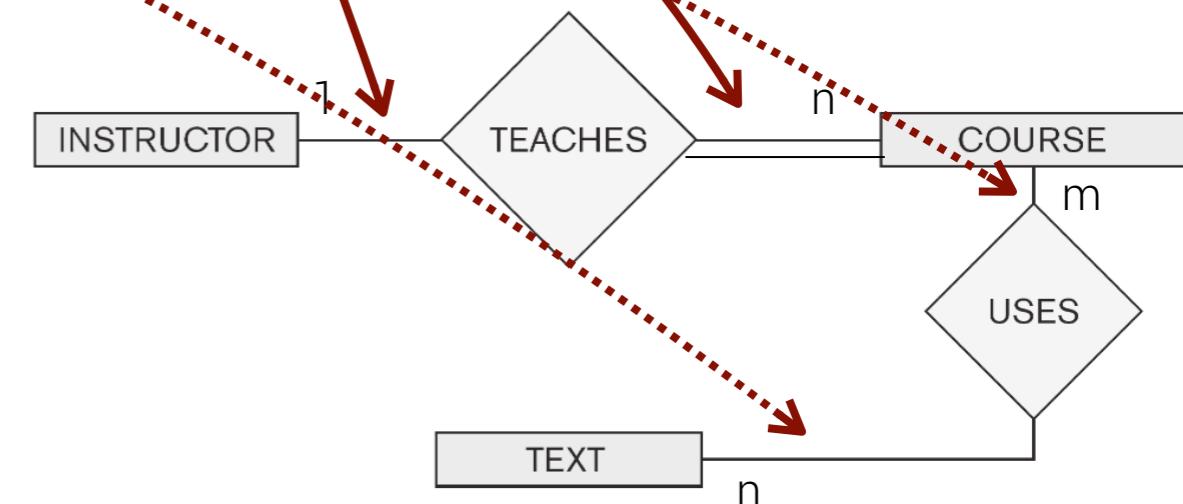
Courses use texts.

A course uses many texts.

A course ~~may not~~ use any text.

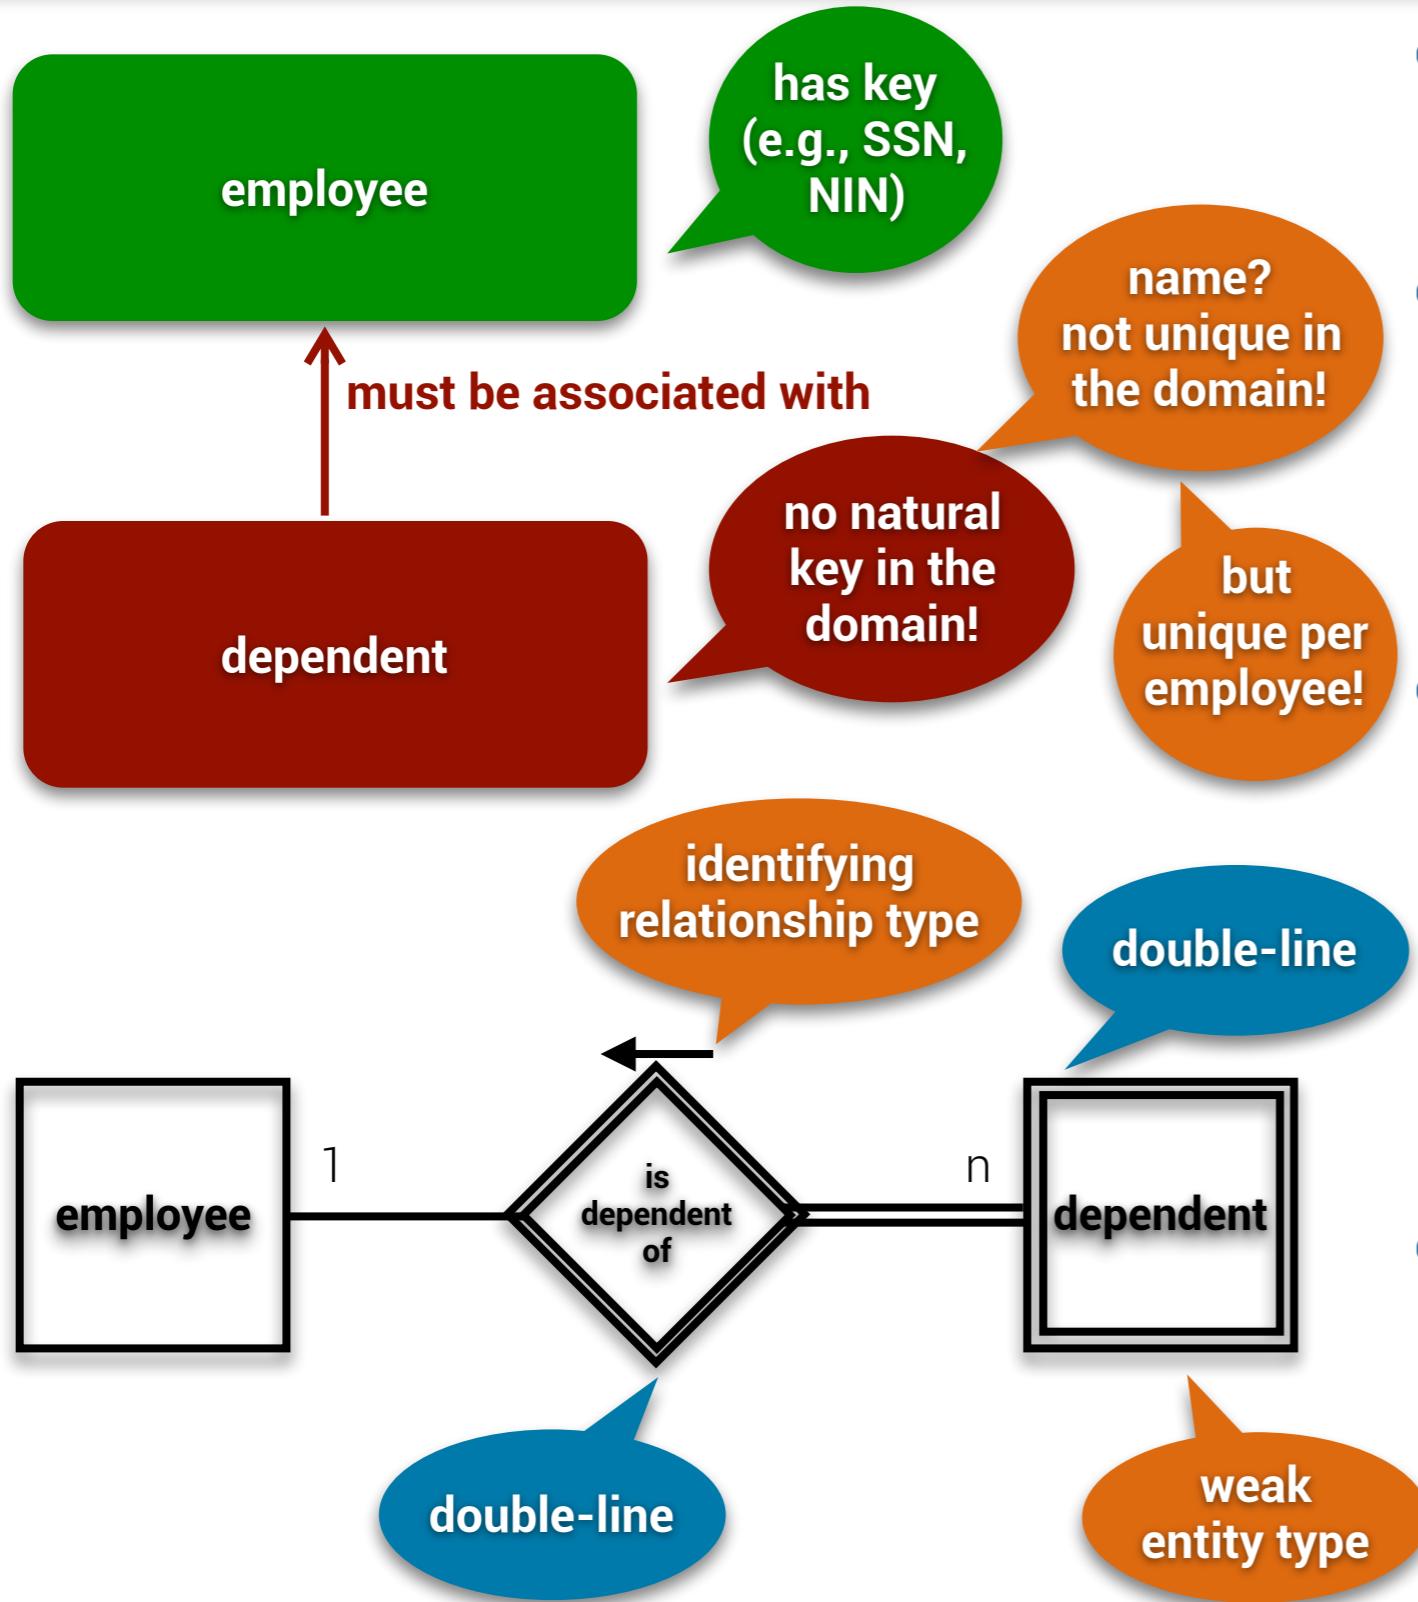
A text is used by many courses.

A text ~~may not~~ be used by any course.



ER Modelling: Weak Entity Types

31

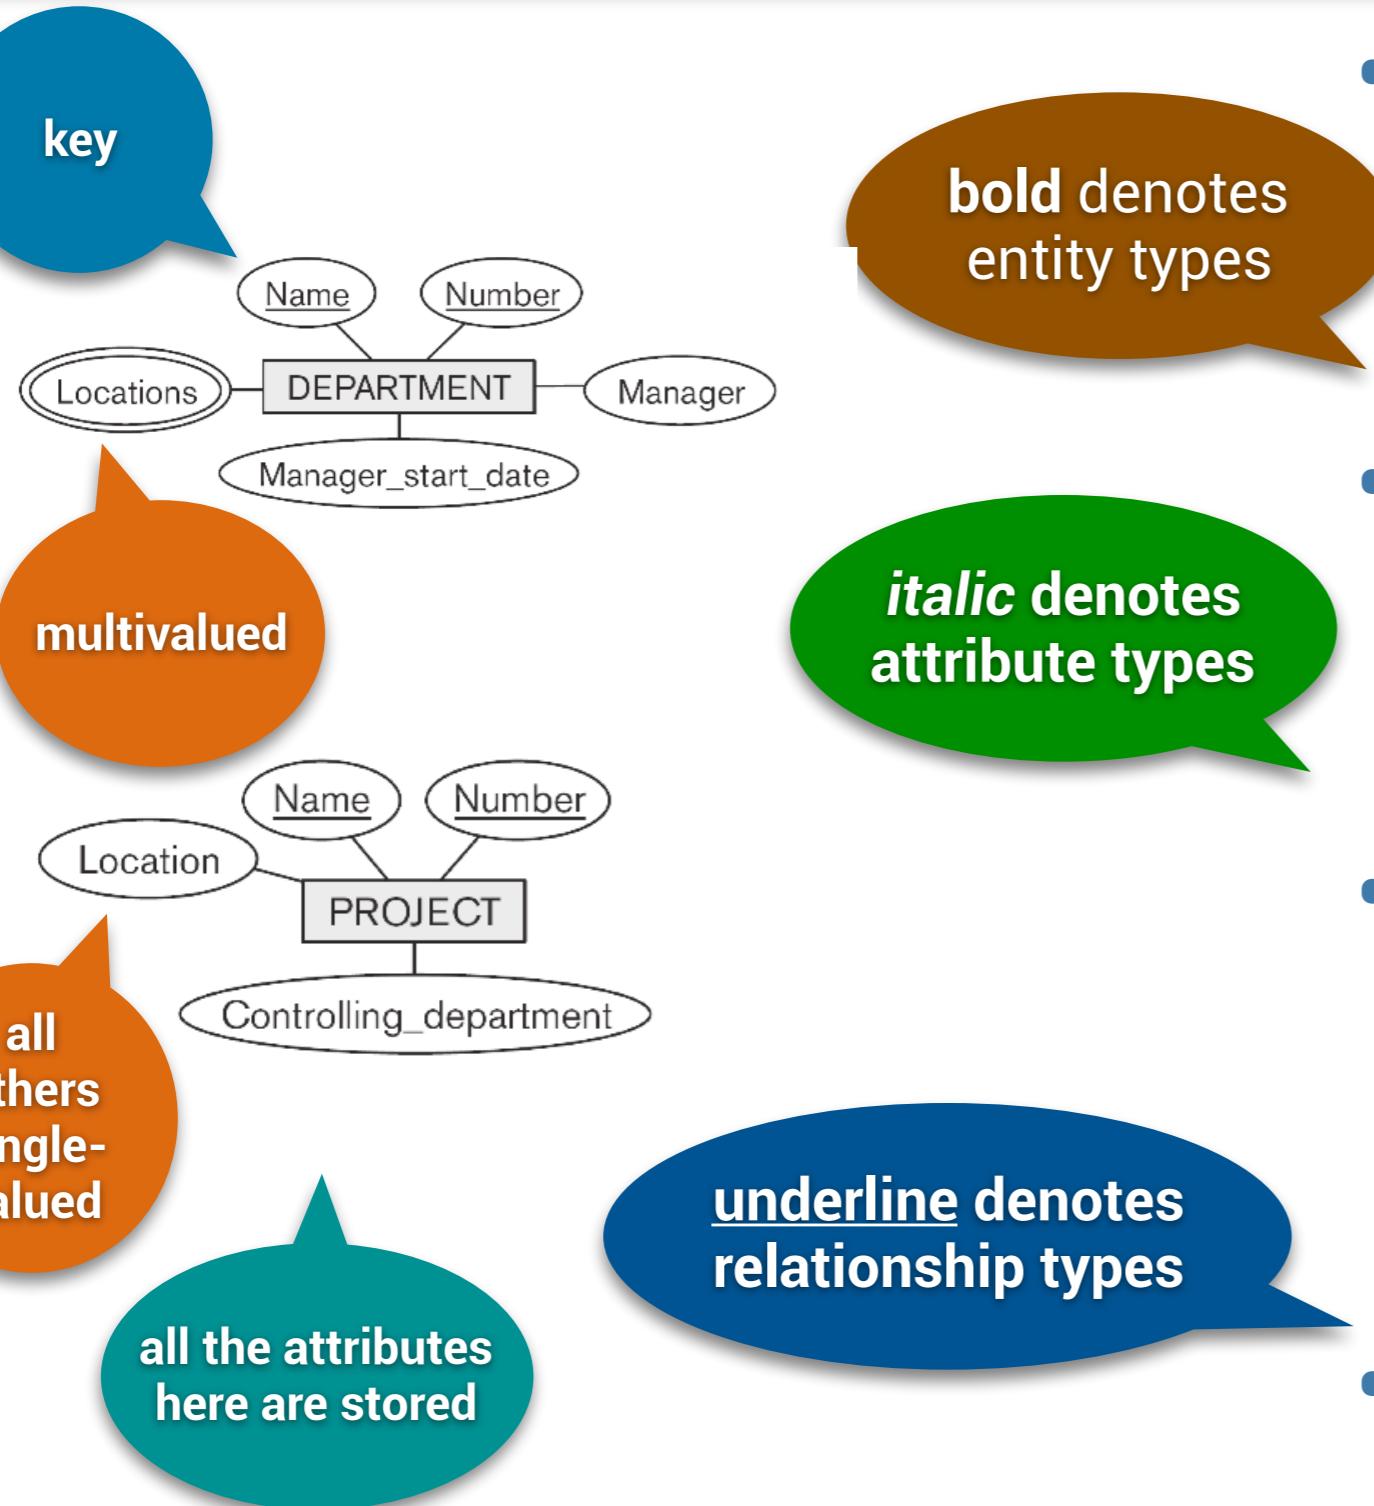


- A **weak entity type** does not have key attributes of their own
- It only has a **partial** (or **weak**) **key** that identifies each entity in the entity set in the context of another entity of another entity type
- It can only be properly identified (i.e., retain its conceptual status and have an associated entity set) by being related to specific entities of some other entity type, called the **owner**
- An **identifying relationship** is one that relates a weak entity type to its owner, and **always** has a *total* participation constraint

Deriving ER Model from Data Requirements Specification

Annotating a Requirements Specification

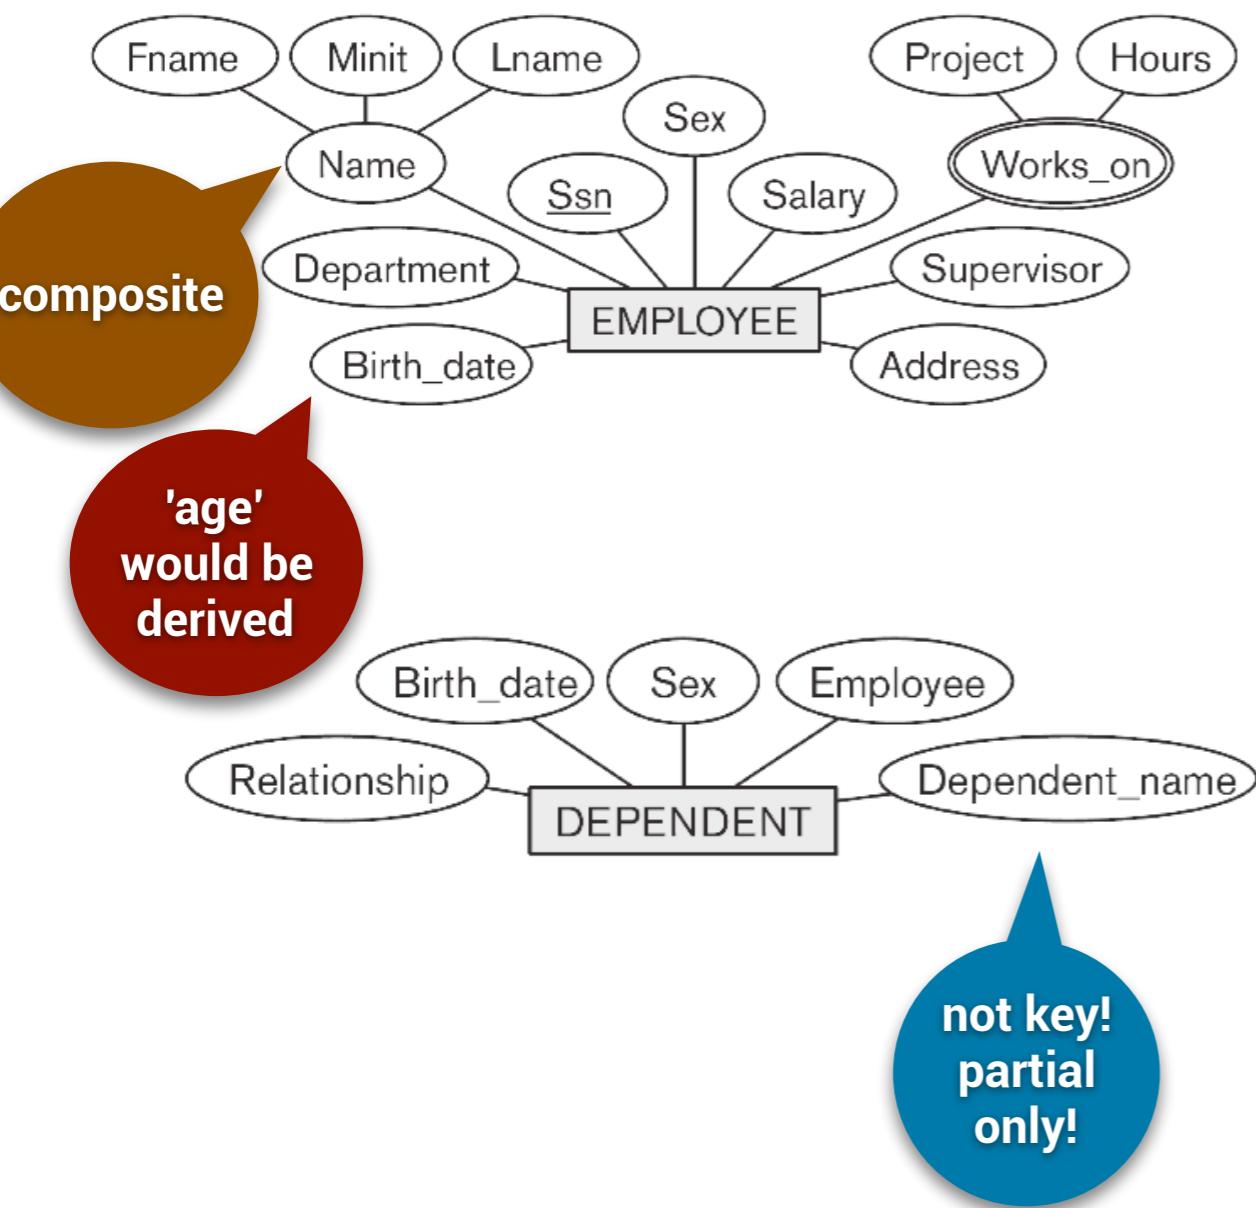
X



- Assume our miniworld encompasses the **employees**, **departments** and **projects** of a company.
- The company is organized into departments, which have a *name*, a *number* and a *location* (possibly more than one).
- Each department is managed by a manager, who is an employee, and we wish to store the *date* in which the latter started managing the department.
- Each department controls various projects.

Annotating a Requirements Specification

X

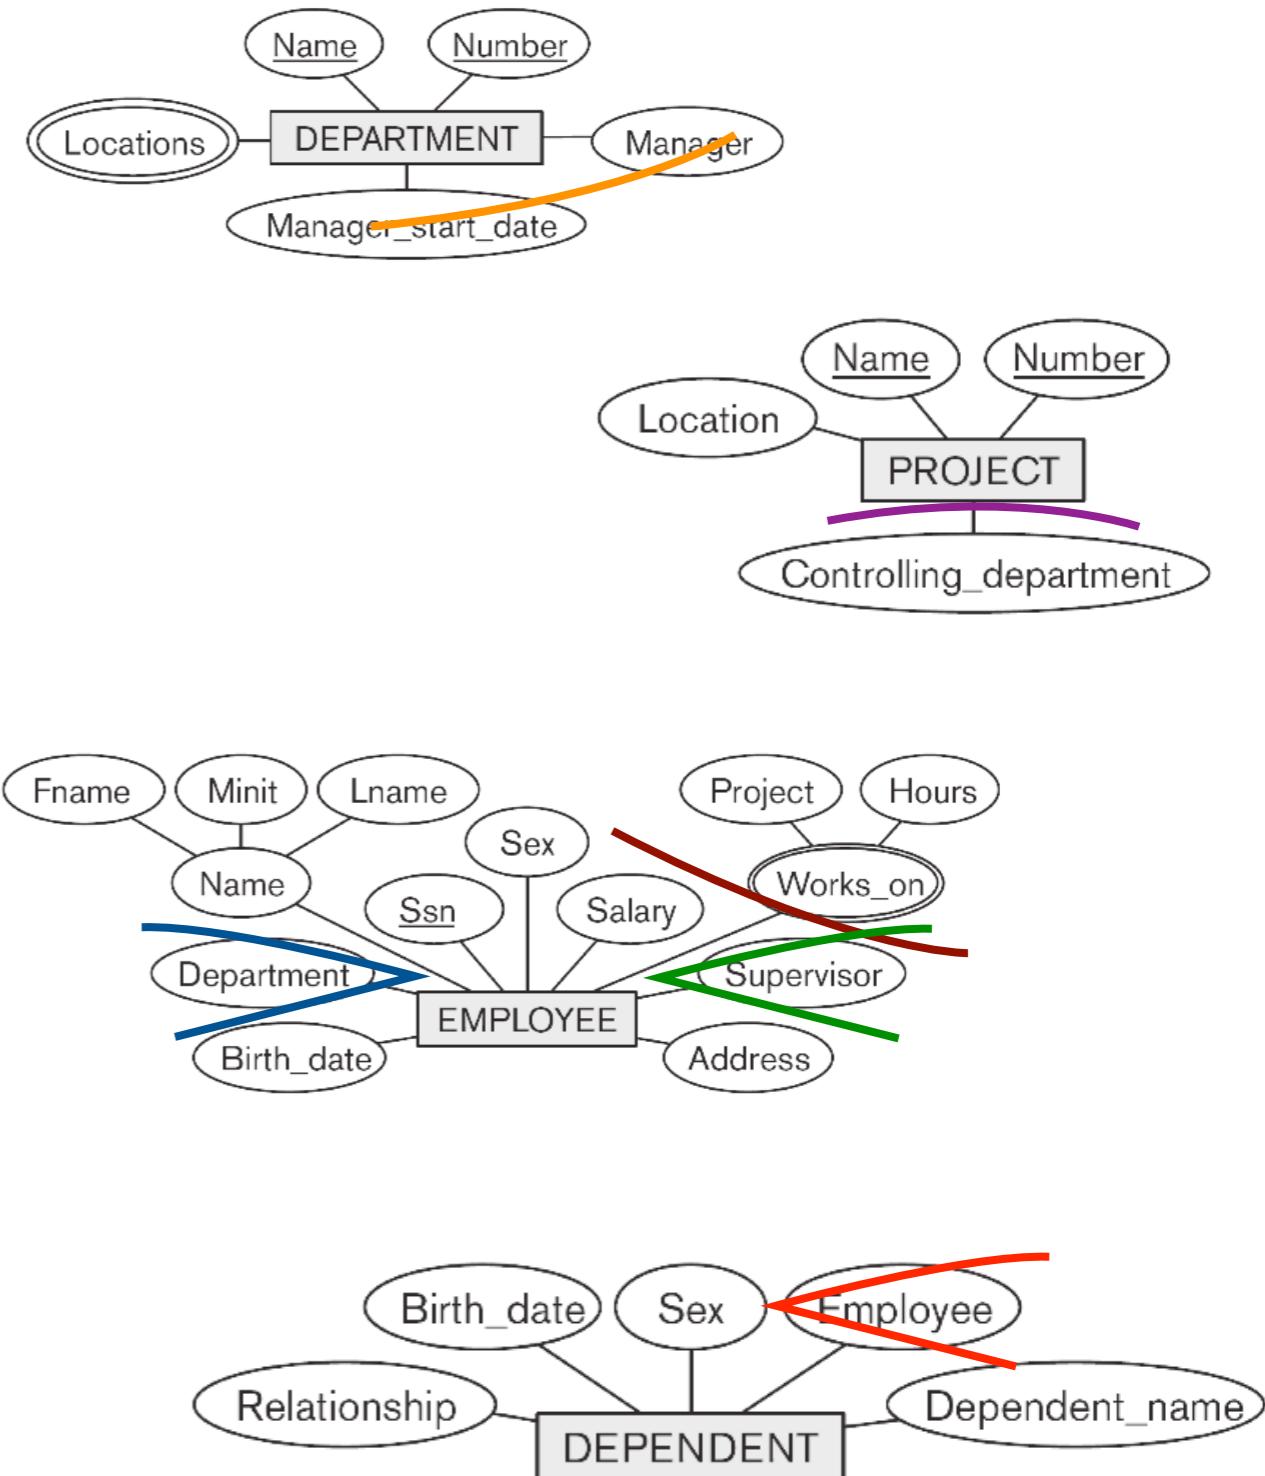


- An employee works for only one department but may work on more than one project (not necessarily controlled by the same department).
- We wish to store how many *hours* an employee works on each project and which employee supervises which employee.
- For each employee, we wish to store the employee's *name*, *social security number*, *address*, *salary*, *gender*, and *birth date*.
- We also wish to store information about the **dependents** of each employee, viz., their *name*, *gender*, *birth date*, and *relationship to the employee*.

ER Modelling: From Implicit to Explicit Relationships

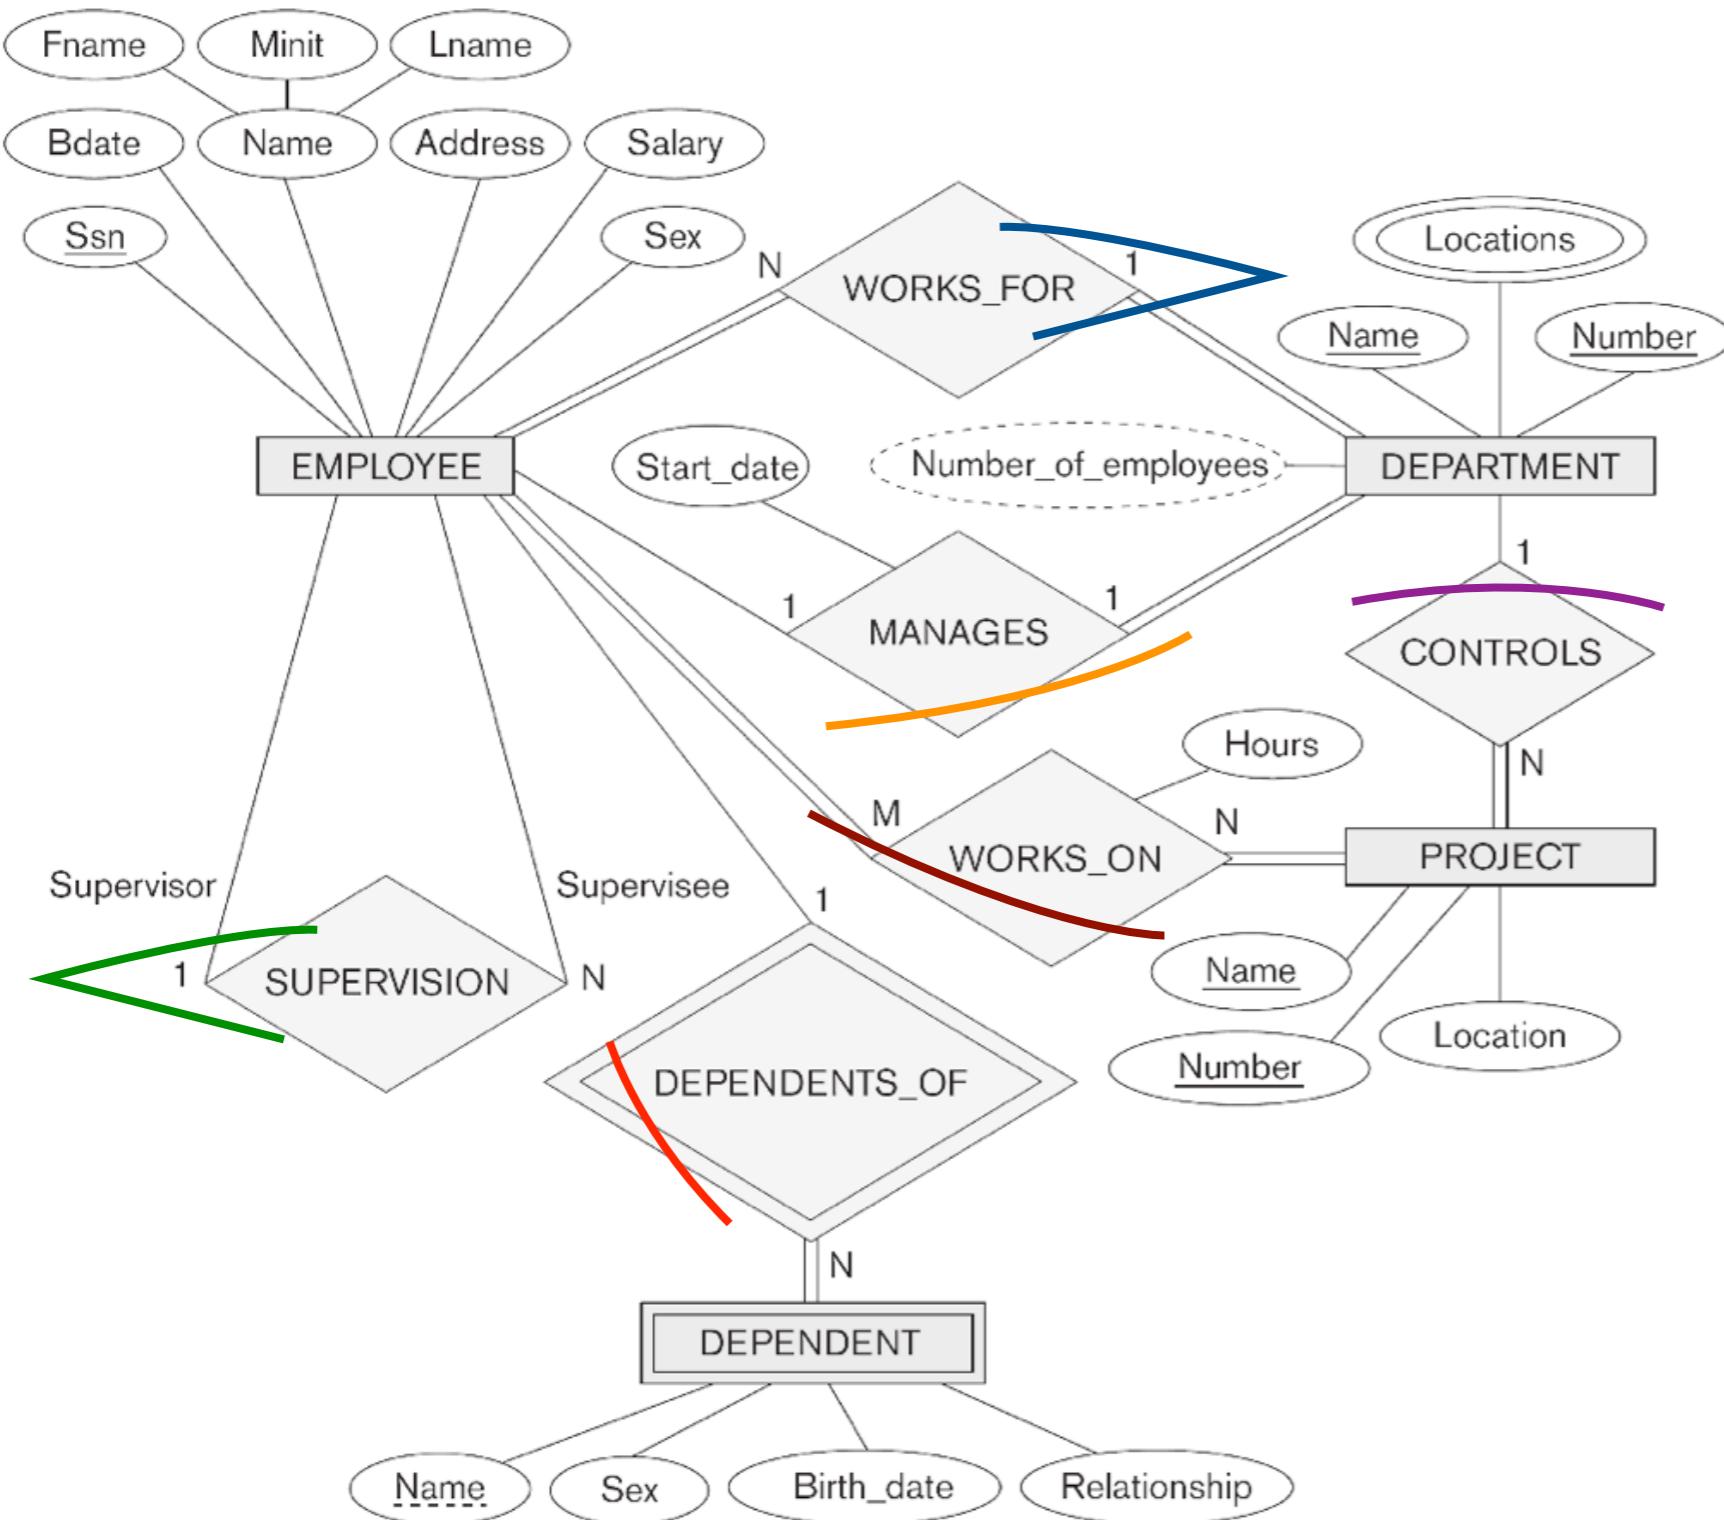
X

- To make relationships explicit we:
 - ▶ Change attributes that represent relationships into relationship types
 - ▶ Determine the cardinality ratio and participation constraints of each relationship type



Conceptual Design: COMPANY Database

X



manages = dept.manager + dept.manager_start_date

works_for = emp.department

works_on = emp.works_on

controls = proj.controlling_department

dependents_of = dep.employee

supervision = emp.supervisor

In The Next Lecture

32

- (Study the notation, read the mandatory readings!)
- We'll continue our exploration of conceptual modelling.
- We'll see that how the ER conceptual model can be enhanced to capture more application semantics.

L 3

Designing Databases: The Enhanced Entity-Relationship Approach

Fundamentals of Databases
Alvaro A A Fernandes, SCS, UoM

Readings

X

This lecture relies on your having read the assigned mandatory readings associated with this lecture:

Ramez Elmasri and Shamkant Navathe. Fundamentals of Database Systems. New International Edition. Pearson, 2013. pp. 246-286.

You can read the material online using this link:

lib.mylibrary.com/Open.aspx?id=527132

This is a shared resource. Don't hog it. When you're done, logout and close the browser to free it for others.

If you fail to study the material, you're likely to find the lecture harder to follow.

More information is available in the Learning Activities Manual.

In Previous Lectures

x

- We learned that data is an enterprise asset and that DBMSs are crucial to manage it well.
- We learned the importance of adopting different levels of abstraction in designing and implementing databases.
- We learned how data models lead to a distinction between schemas and instances that enables a logical view of the data.
- We learned of the practical benefits of performing conceptual modelling before logical design and why the ER approach serves well this purpose.

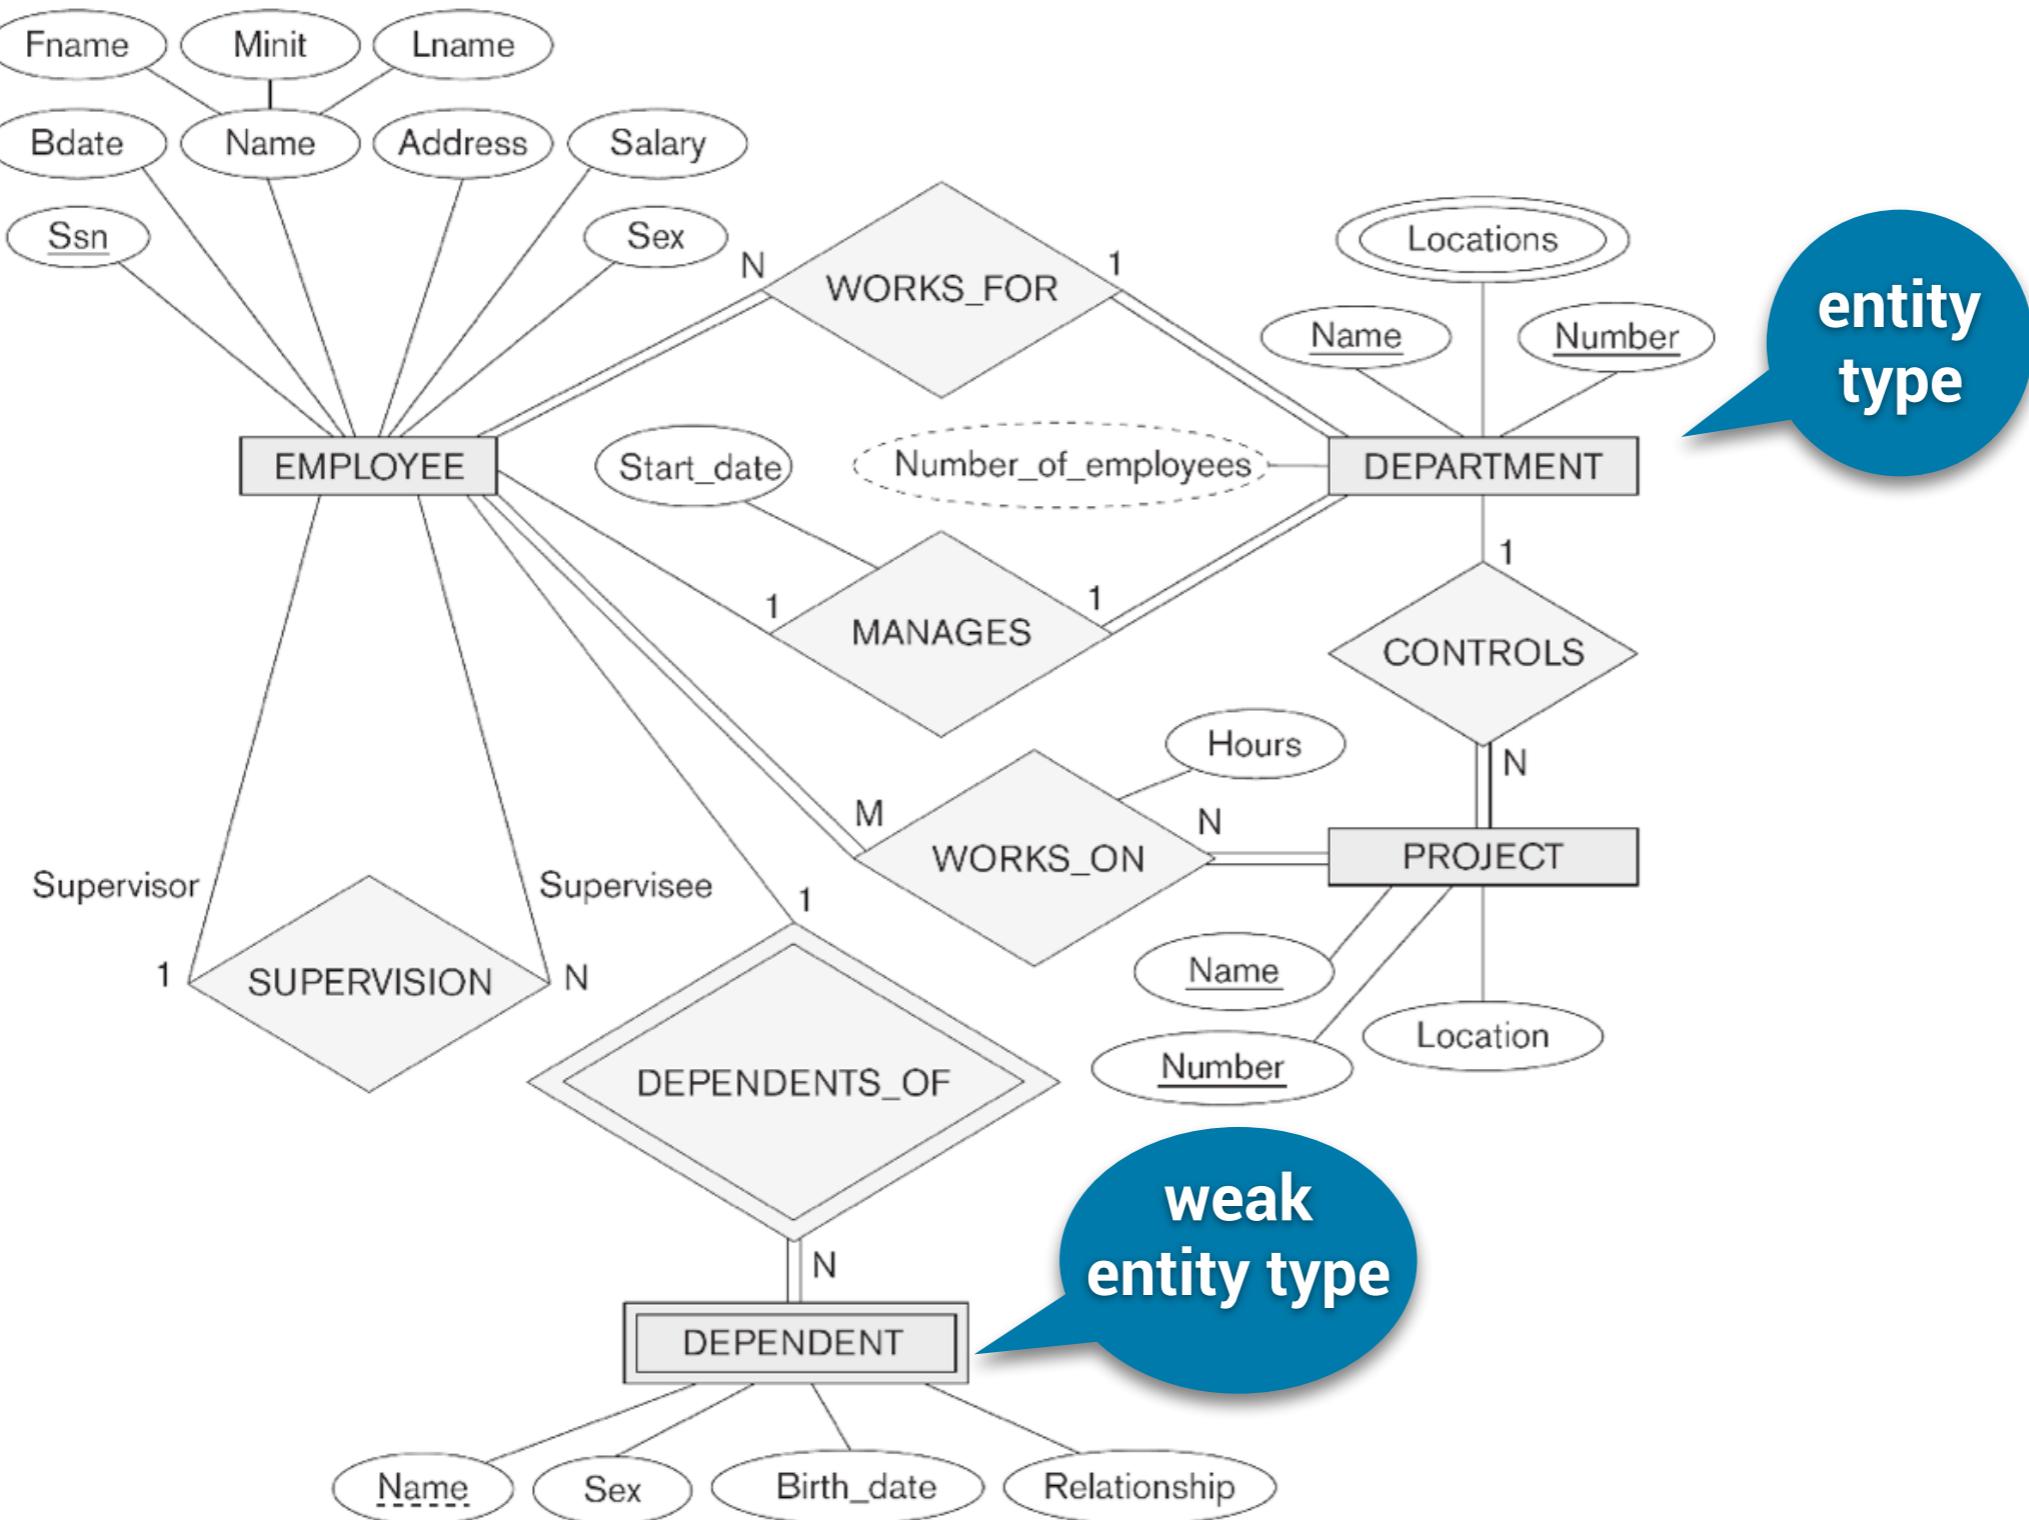
In This Lecture

x

- What is the enhanced entity-relationship model (EER)?
- What are specialization/generalization (S/G) processes in EER modelling?
- What is EER-inheritance and what does it apply to?
- What constraints apply to S/G relationships?
- How do S/G hierarchies and lattices arise?
- What are union types?
- What design guidelines are useful in EER modelling?

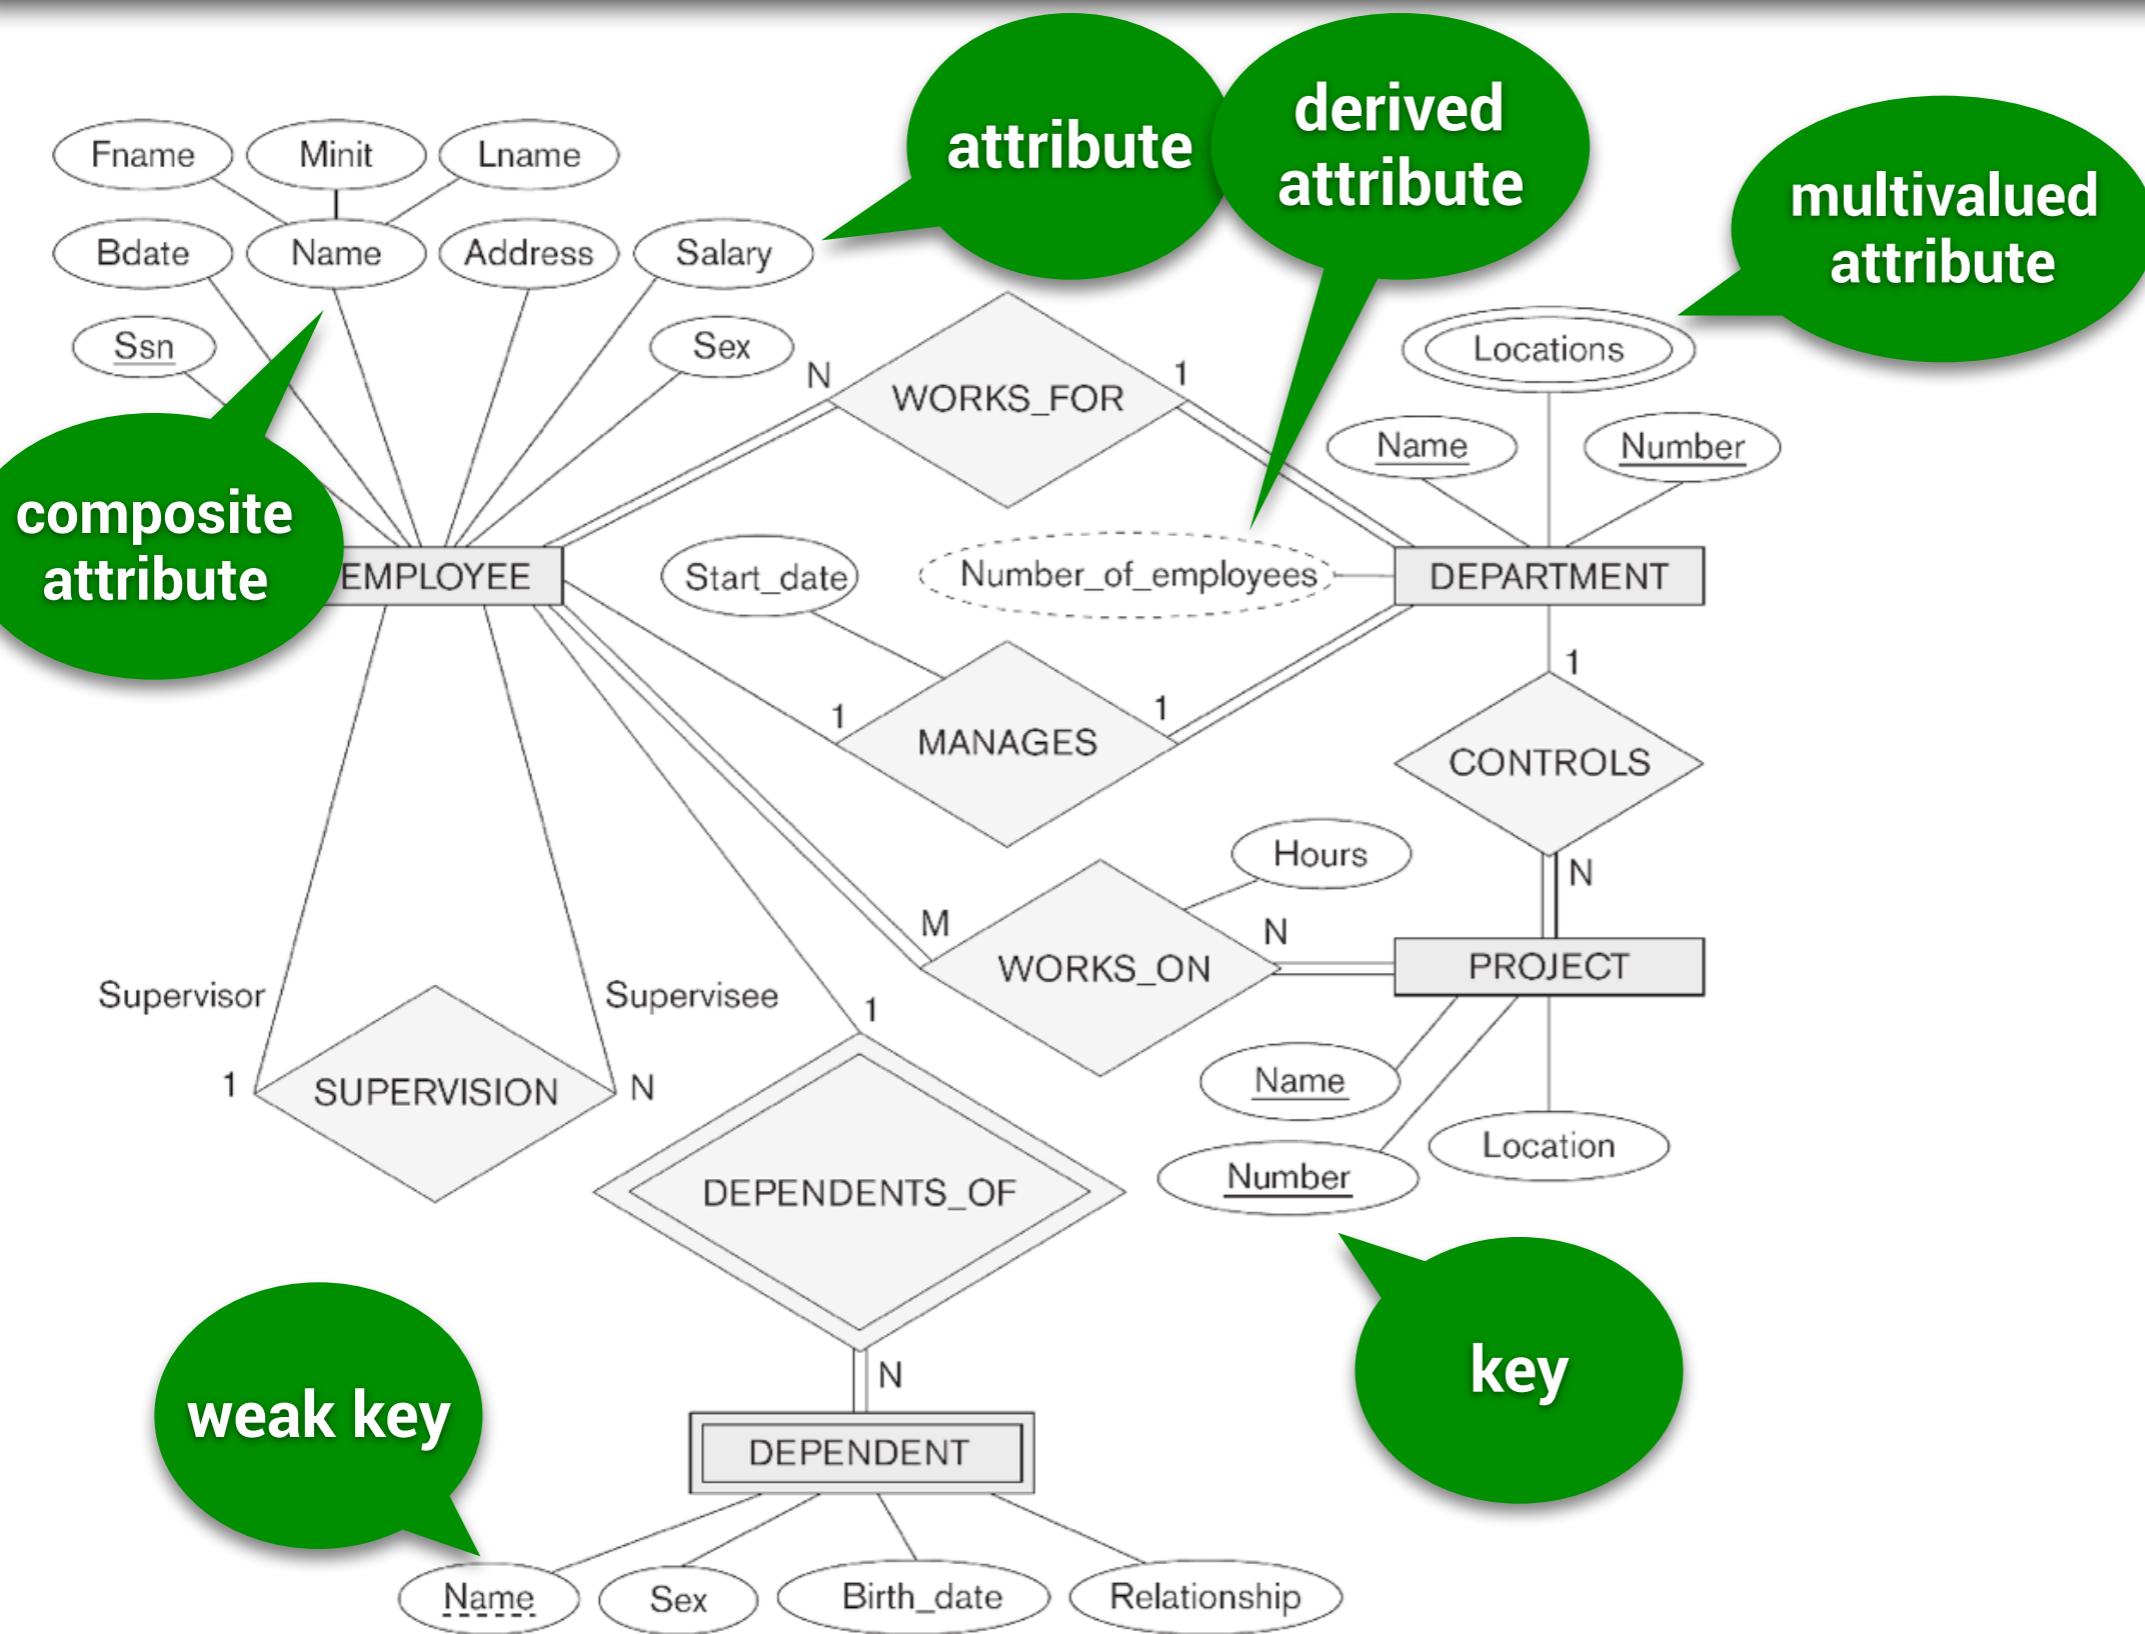
Review: ER Constructs and Notation

34



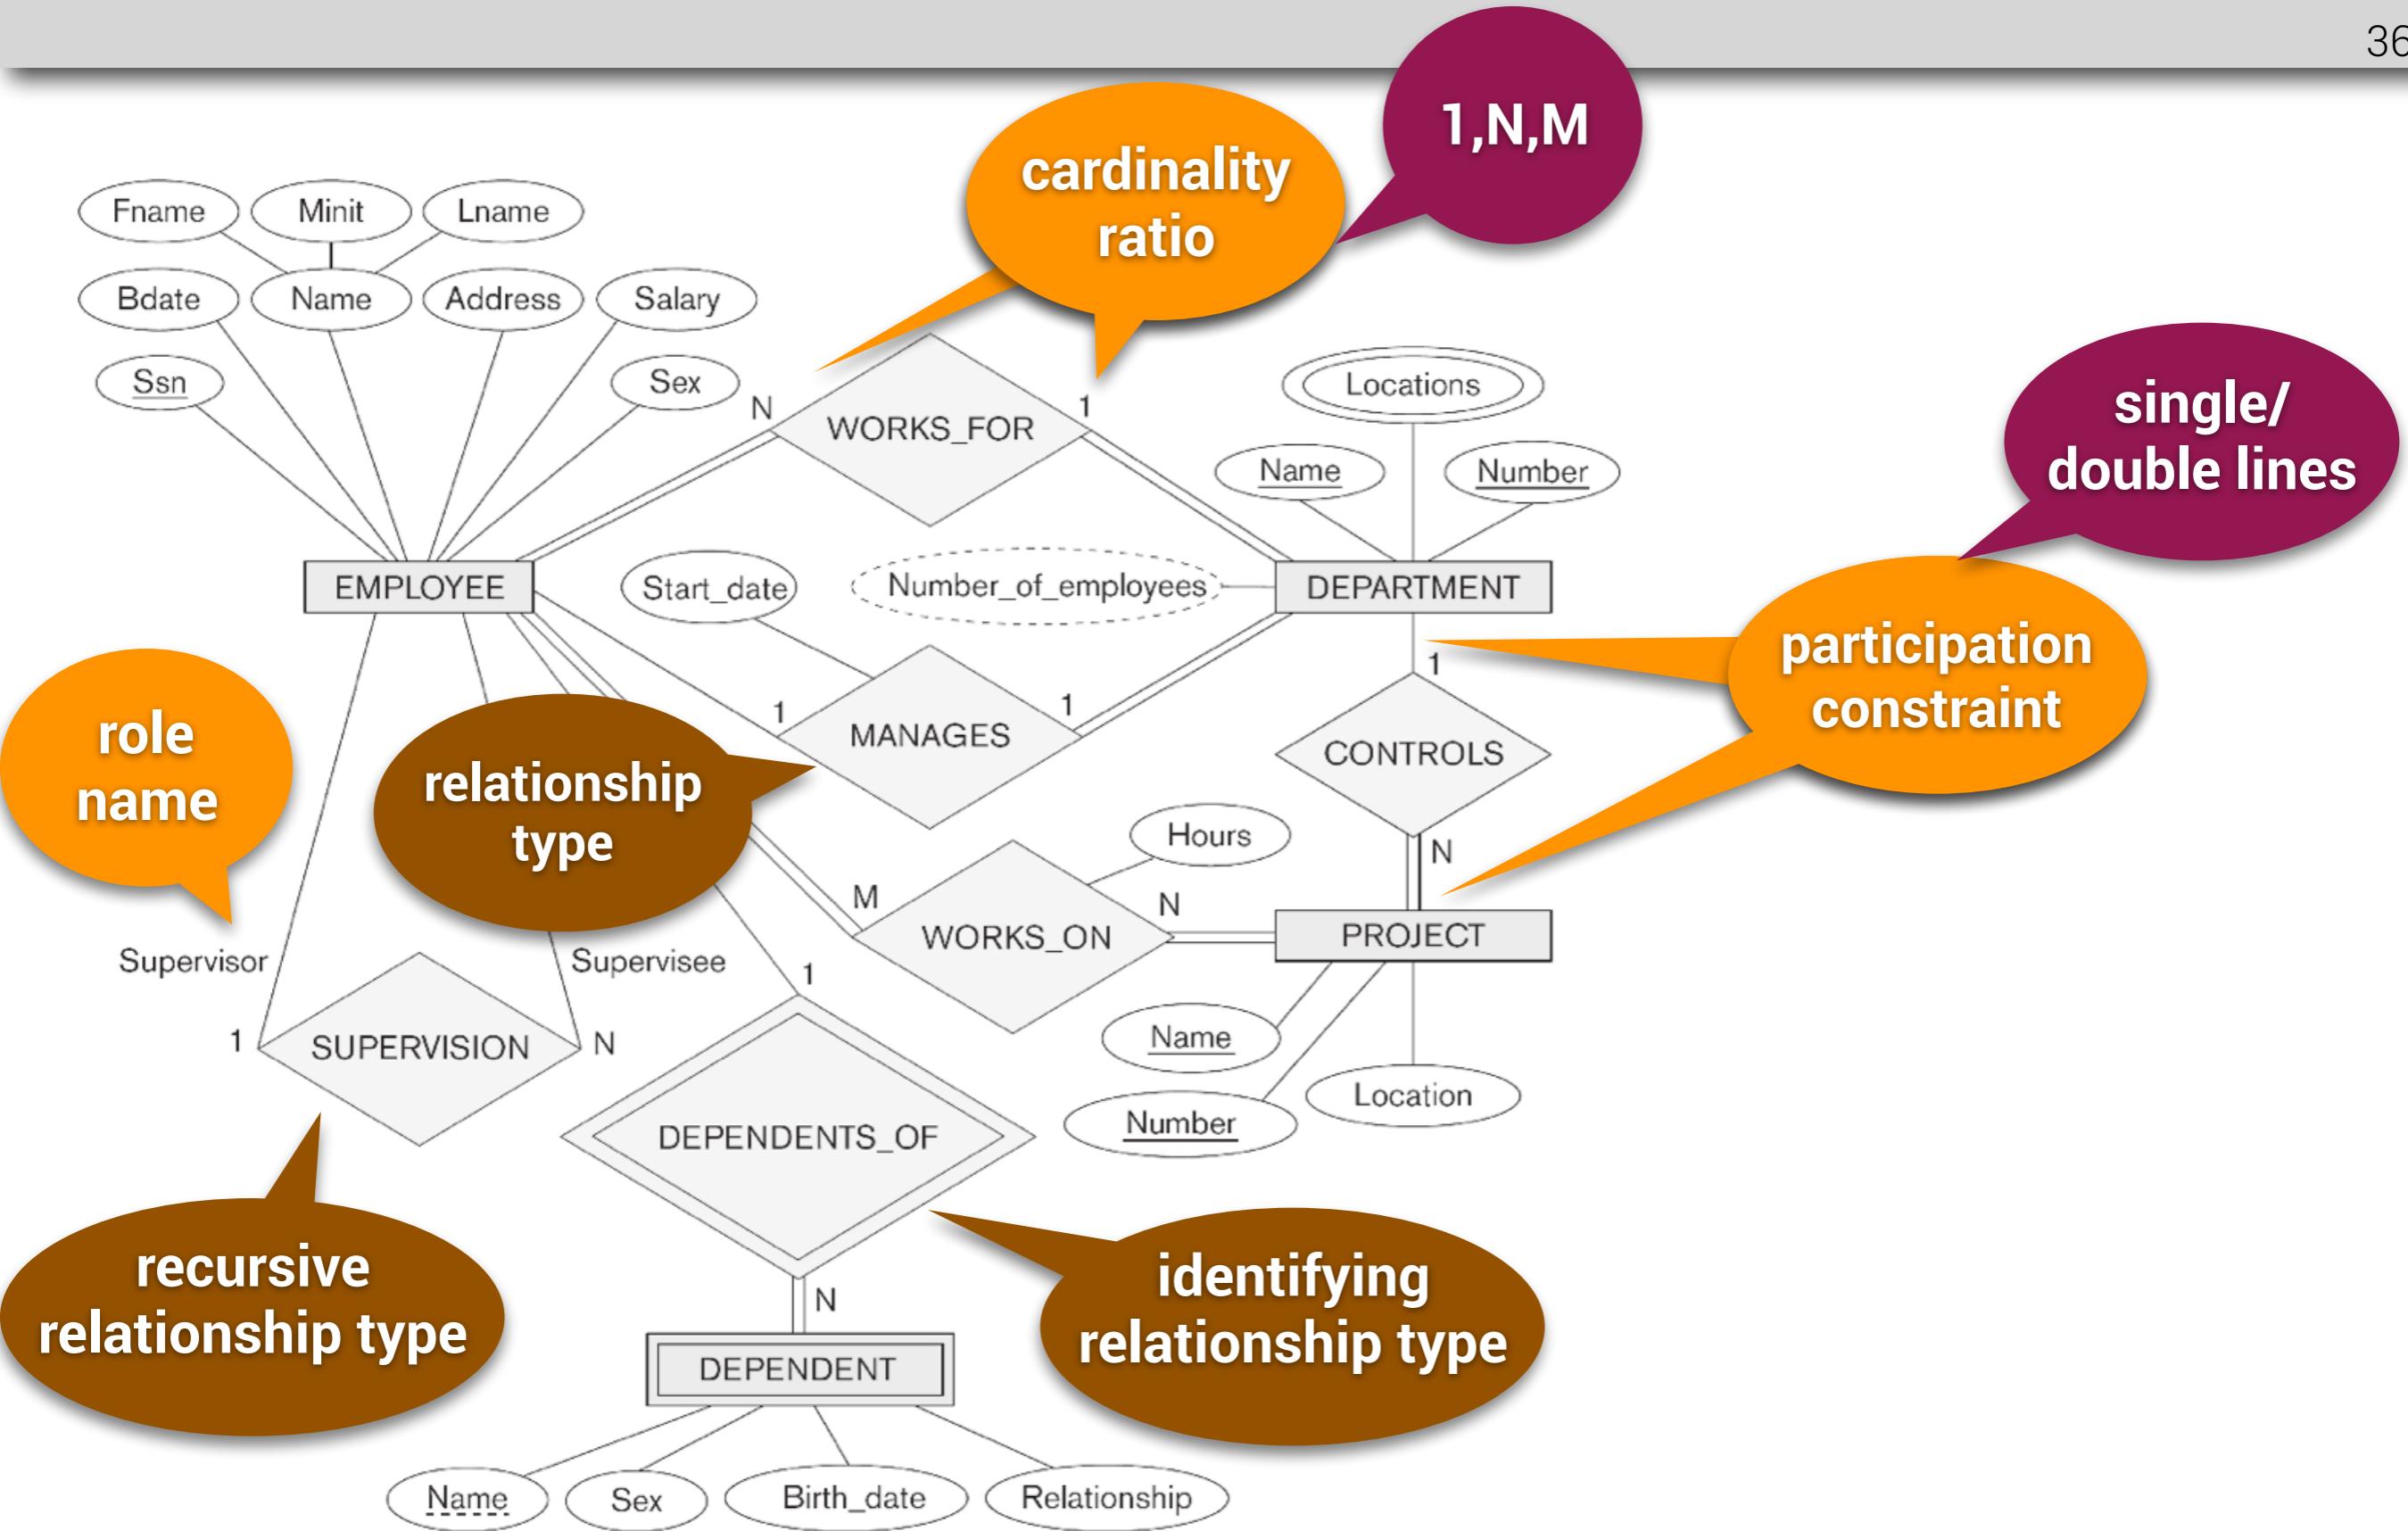
Review: ER Constructs and Notation

35



Review: ER Constructs and Notation

36



Specialization and Generalization

37

- In conceptual modelling, it is common for the designer to identify a set-subset relationship between entities.
- Some of these are useful in capturing relevant semantics.
- By 'relevant' here, we mean something that allows the final DBMS (and applications that use it) to be
 - ▶ more accurate in how they capture reality, and
 - ▶ (often) more efficient in doing so.

An Intuitive Example

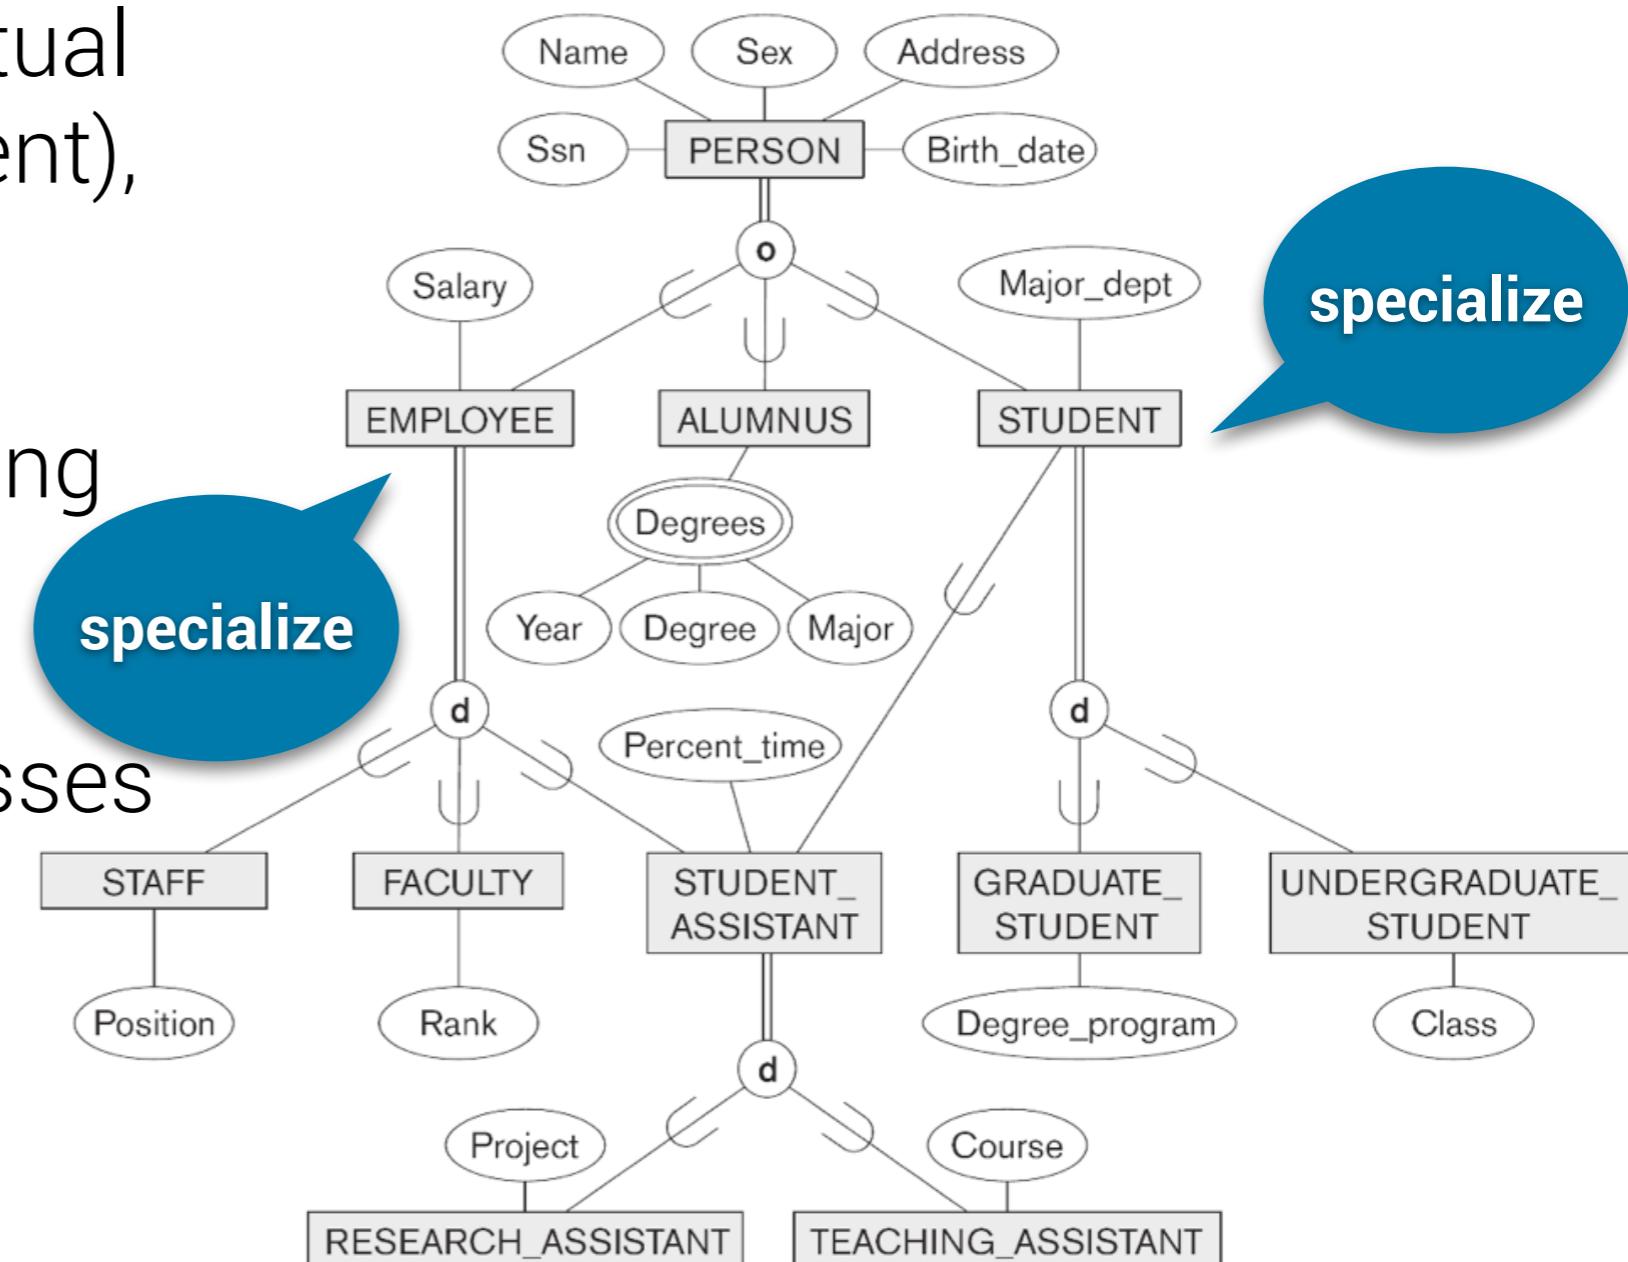
38

- Assume you have a STUDENT entity.
- You then realize that some students have different attributes that apply to them.
- For example, a PhD student may have an account number to be used when s/he is paid for teaching assistant hours, but an UG student doesn't.
- An UG student accumulates credits, but a PhD doesn't.
- This suggests that it is useful (more accurate, and possibly more efficient) to specialize STUDENT into PhD_STUDENT and UG_STUDENT.
- This is top-down (i.e., from a set to its subsets) conceptual analysis.

Refining ER Schemas with Specialization/Generalization

39

- In top-down conceptual analysis (or refinement), we:
 - ▶ start with an existing entity type
 - ▶ identify its subclasses of interests
 - ▶ then specialize



Specialization v. Generalization

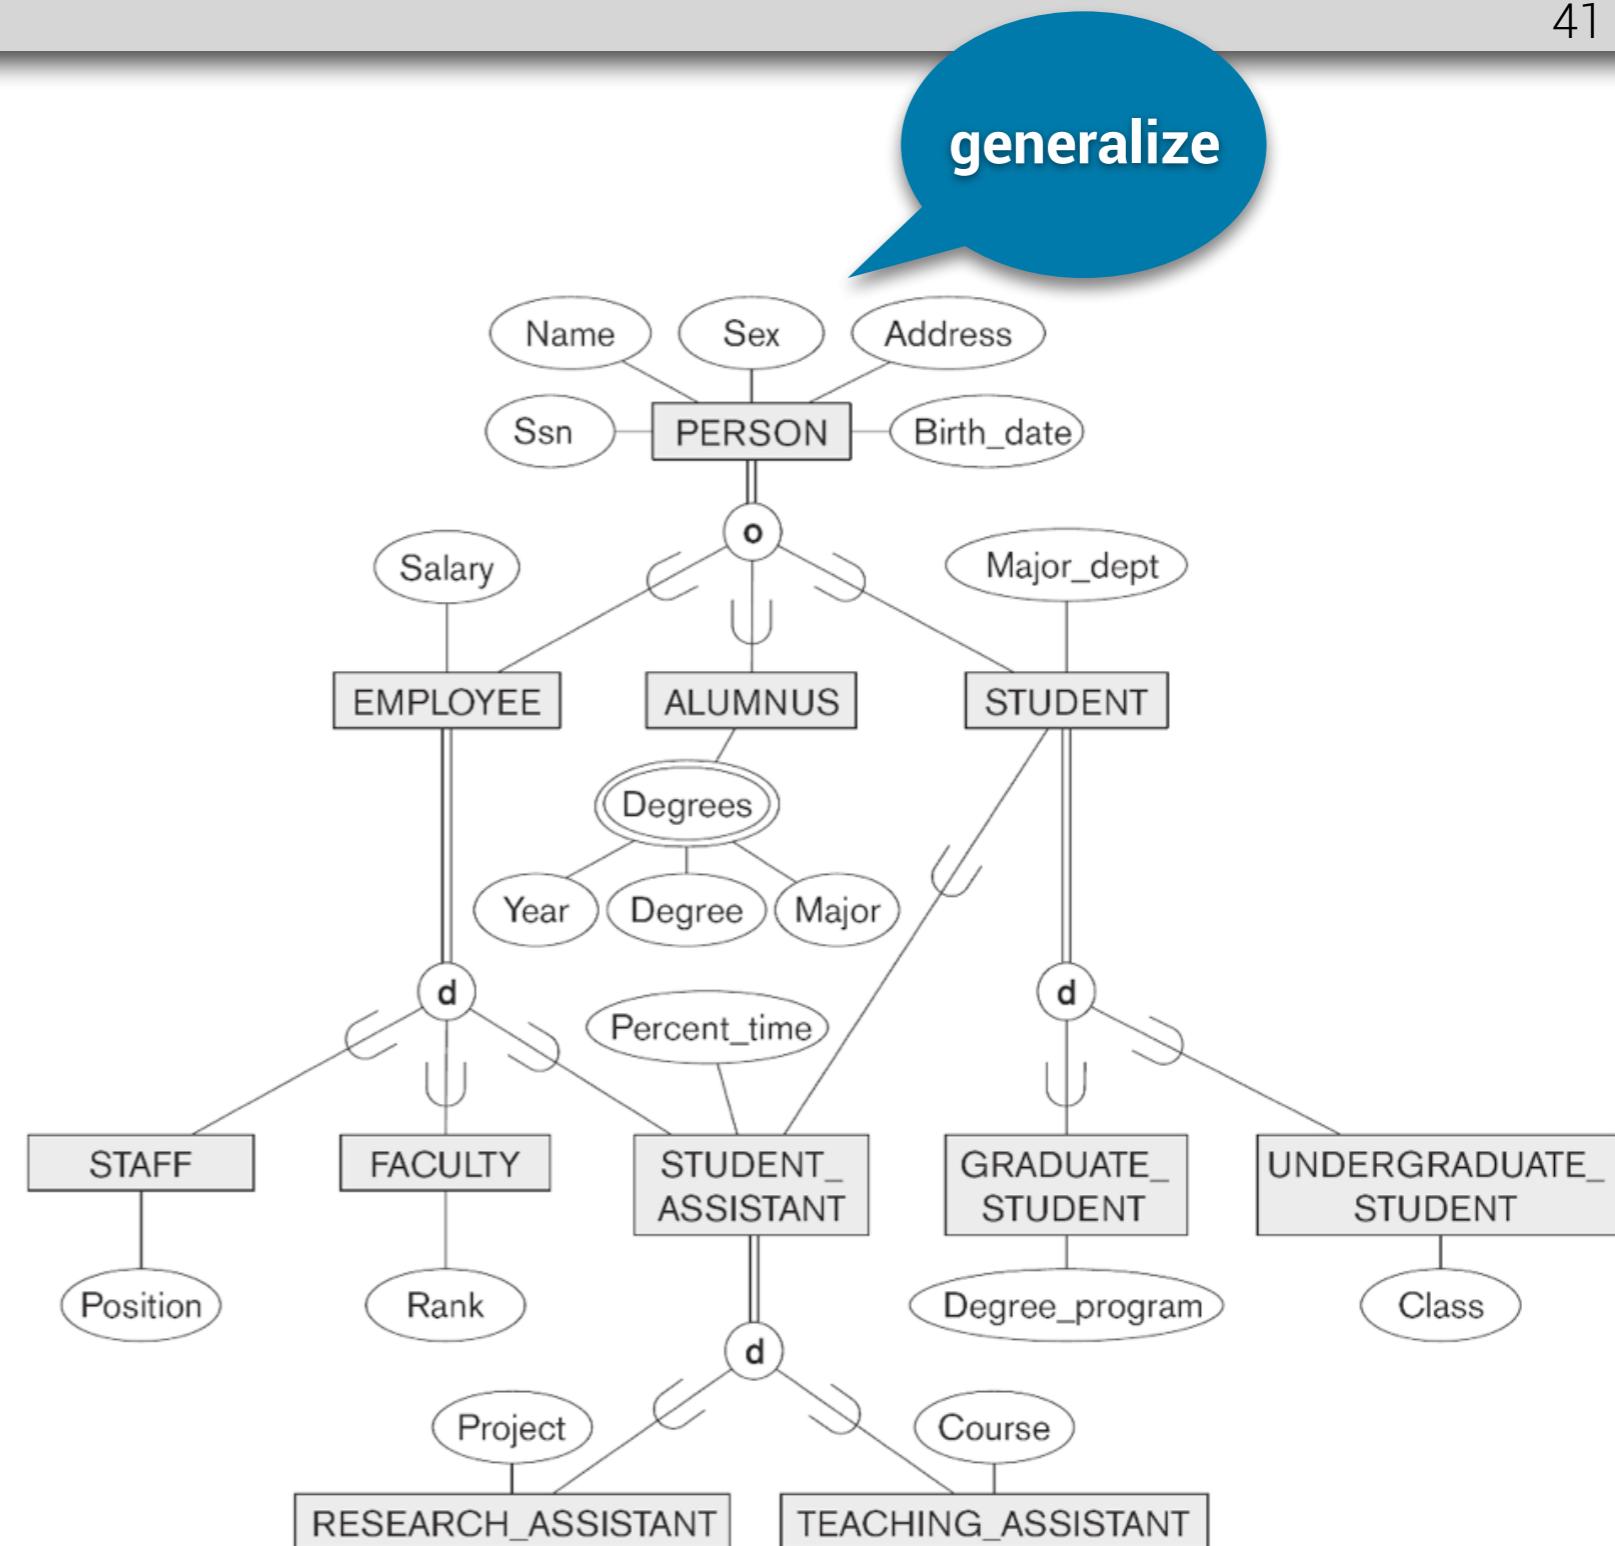
40

- The previous example is of a specialization (i.e., we start with a set and define one or more subsets of it).
- In generalization, we start with one or more sets and define a superset of it.
- For example, one might have started with entities such as BORROWER, MORTGAGE HOLDER, ACCOUNT HOLDER and then decide to generalize these into CUSTOMER.
- This is bottom-up (i.e., from several sets to their superset) conceptual synthesis.

Refining ER Schemas with Specialization/Generalization

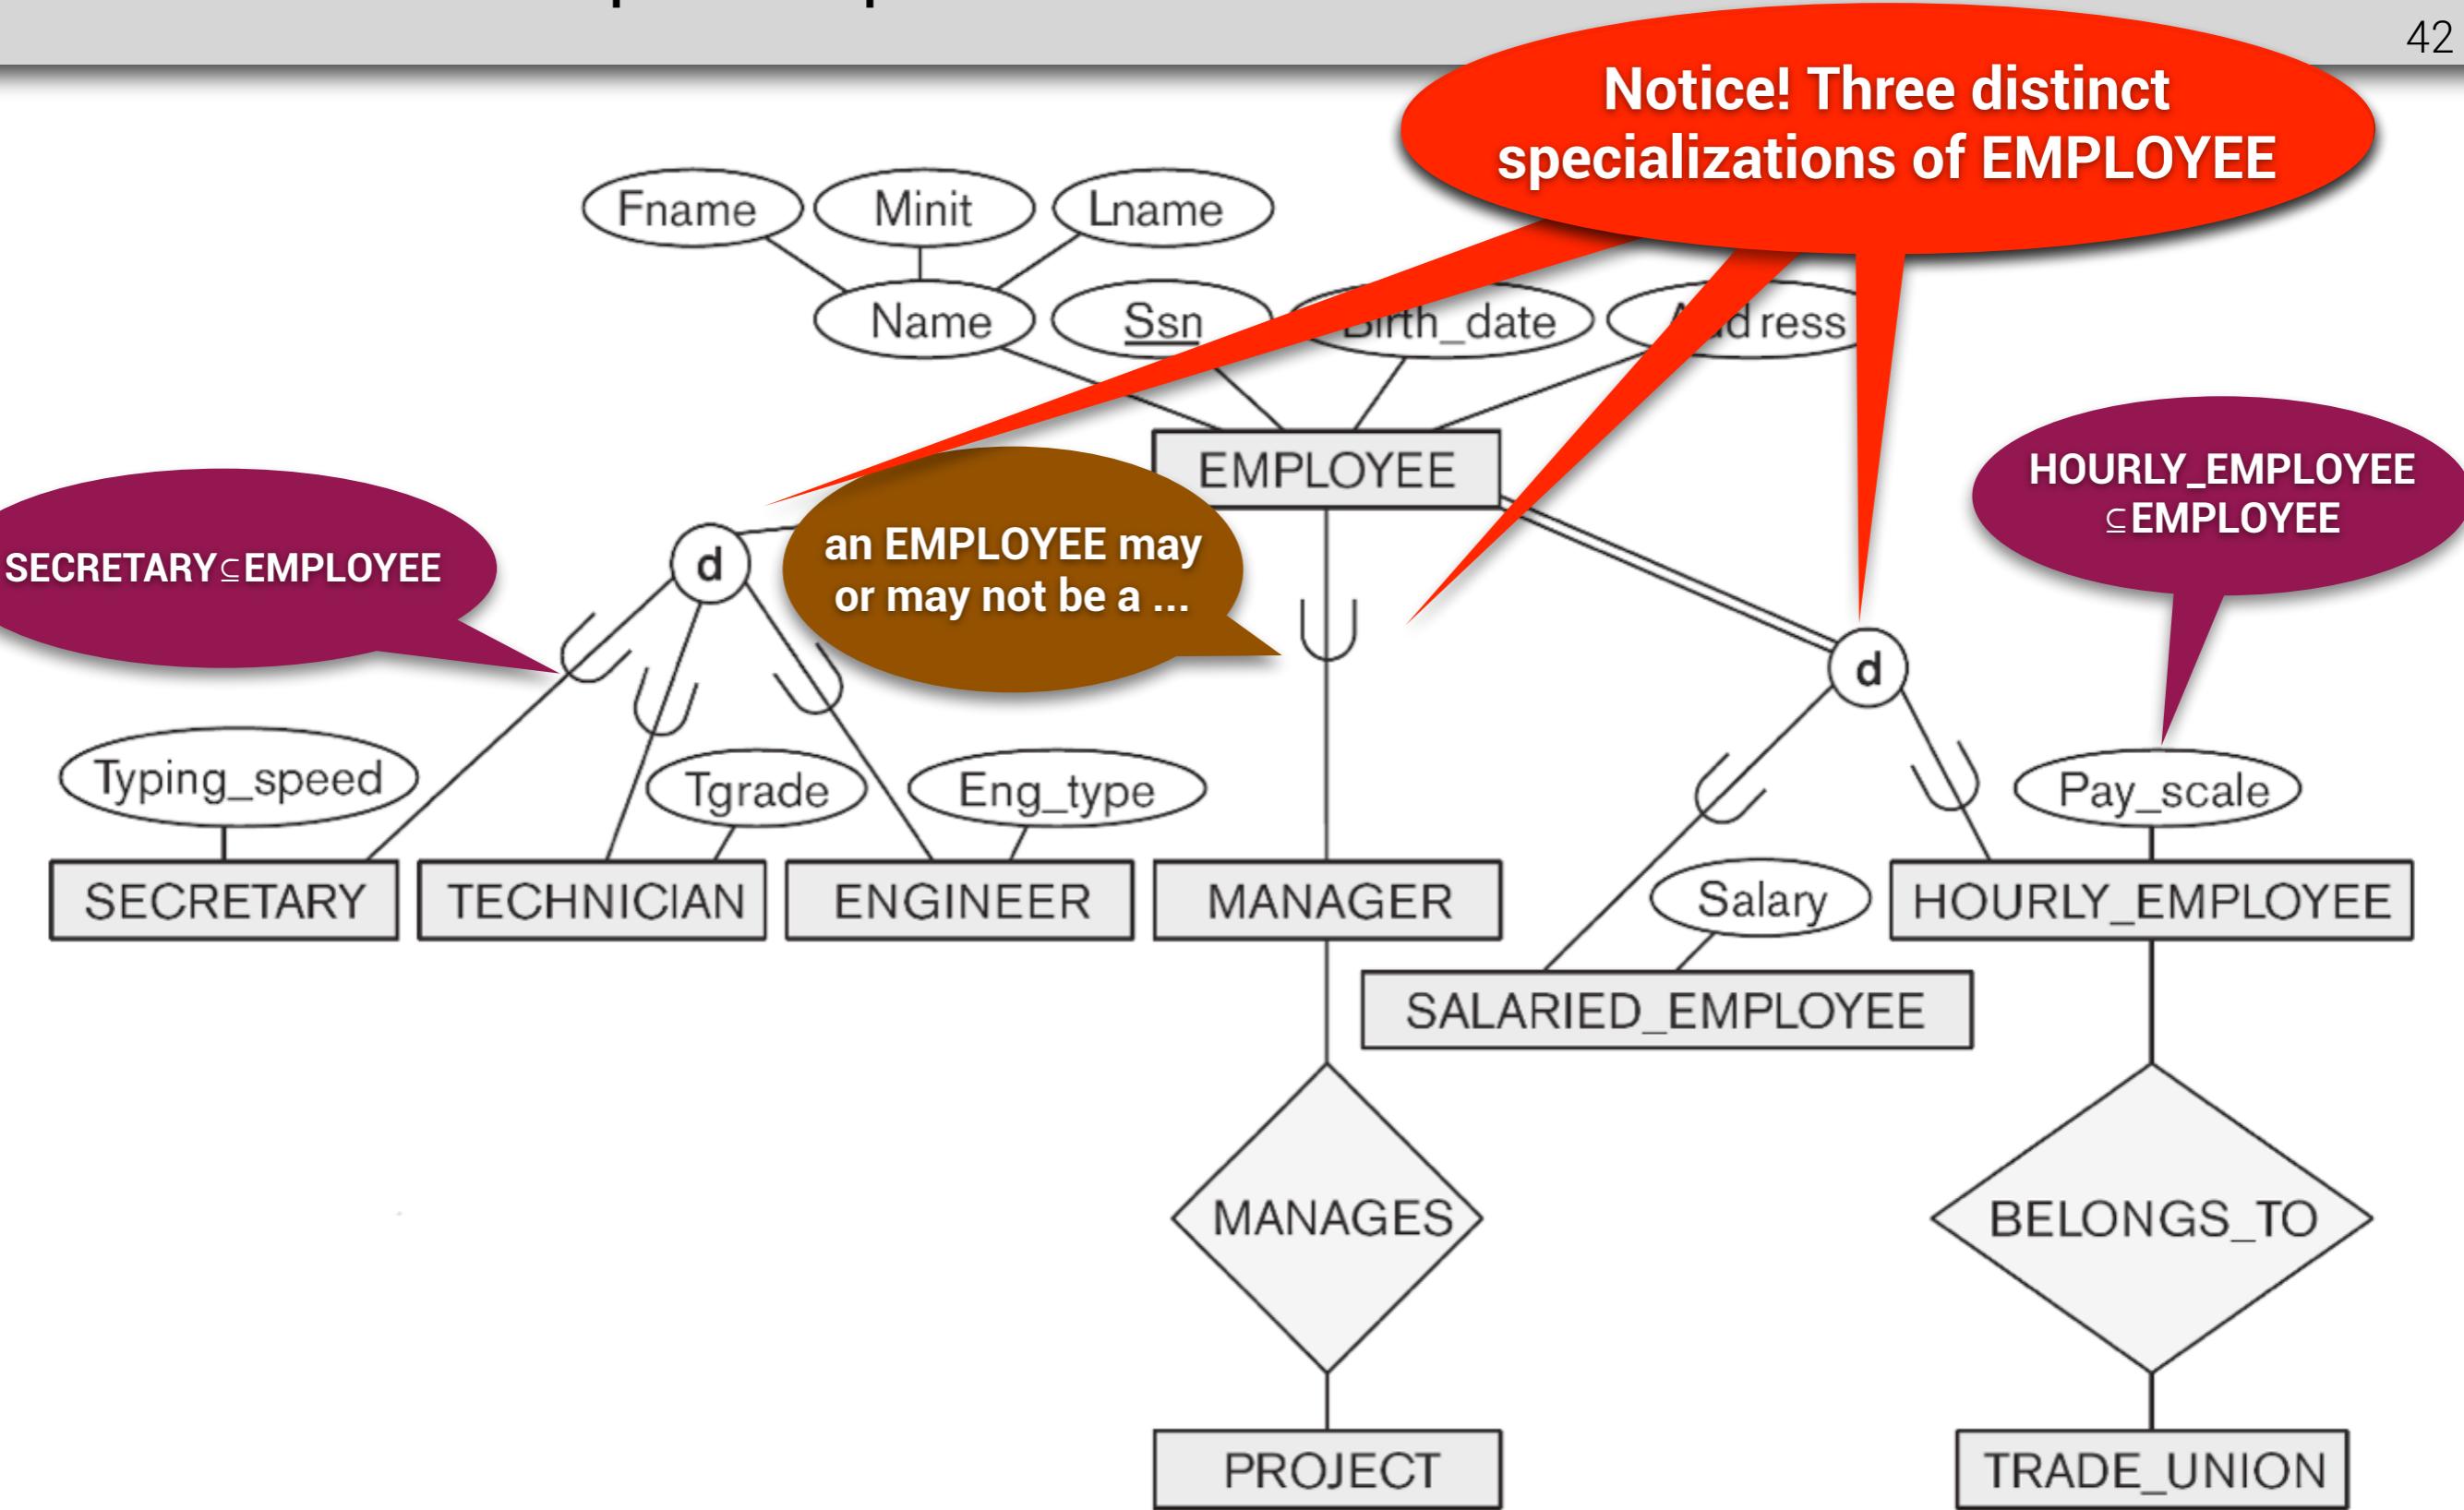
41

- In bottom-up conceptual synthesis, we:
 - ▶ start with with a set of existing entity types
 - ▶ identify their superclasses of interest
 - ▶ then generalize



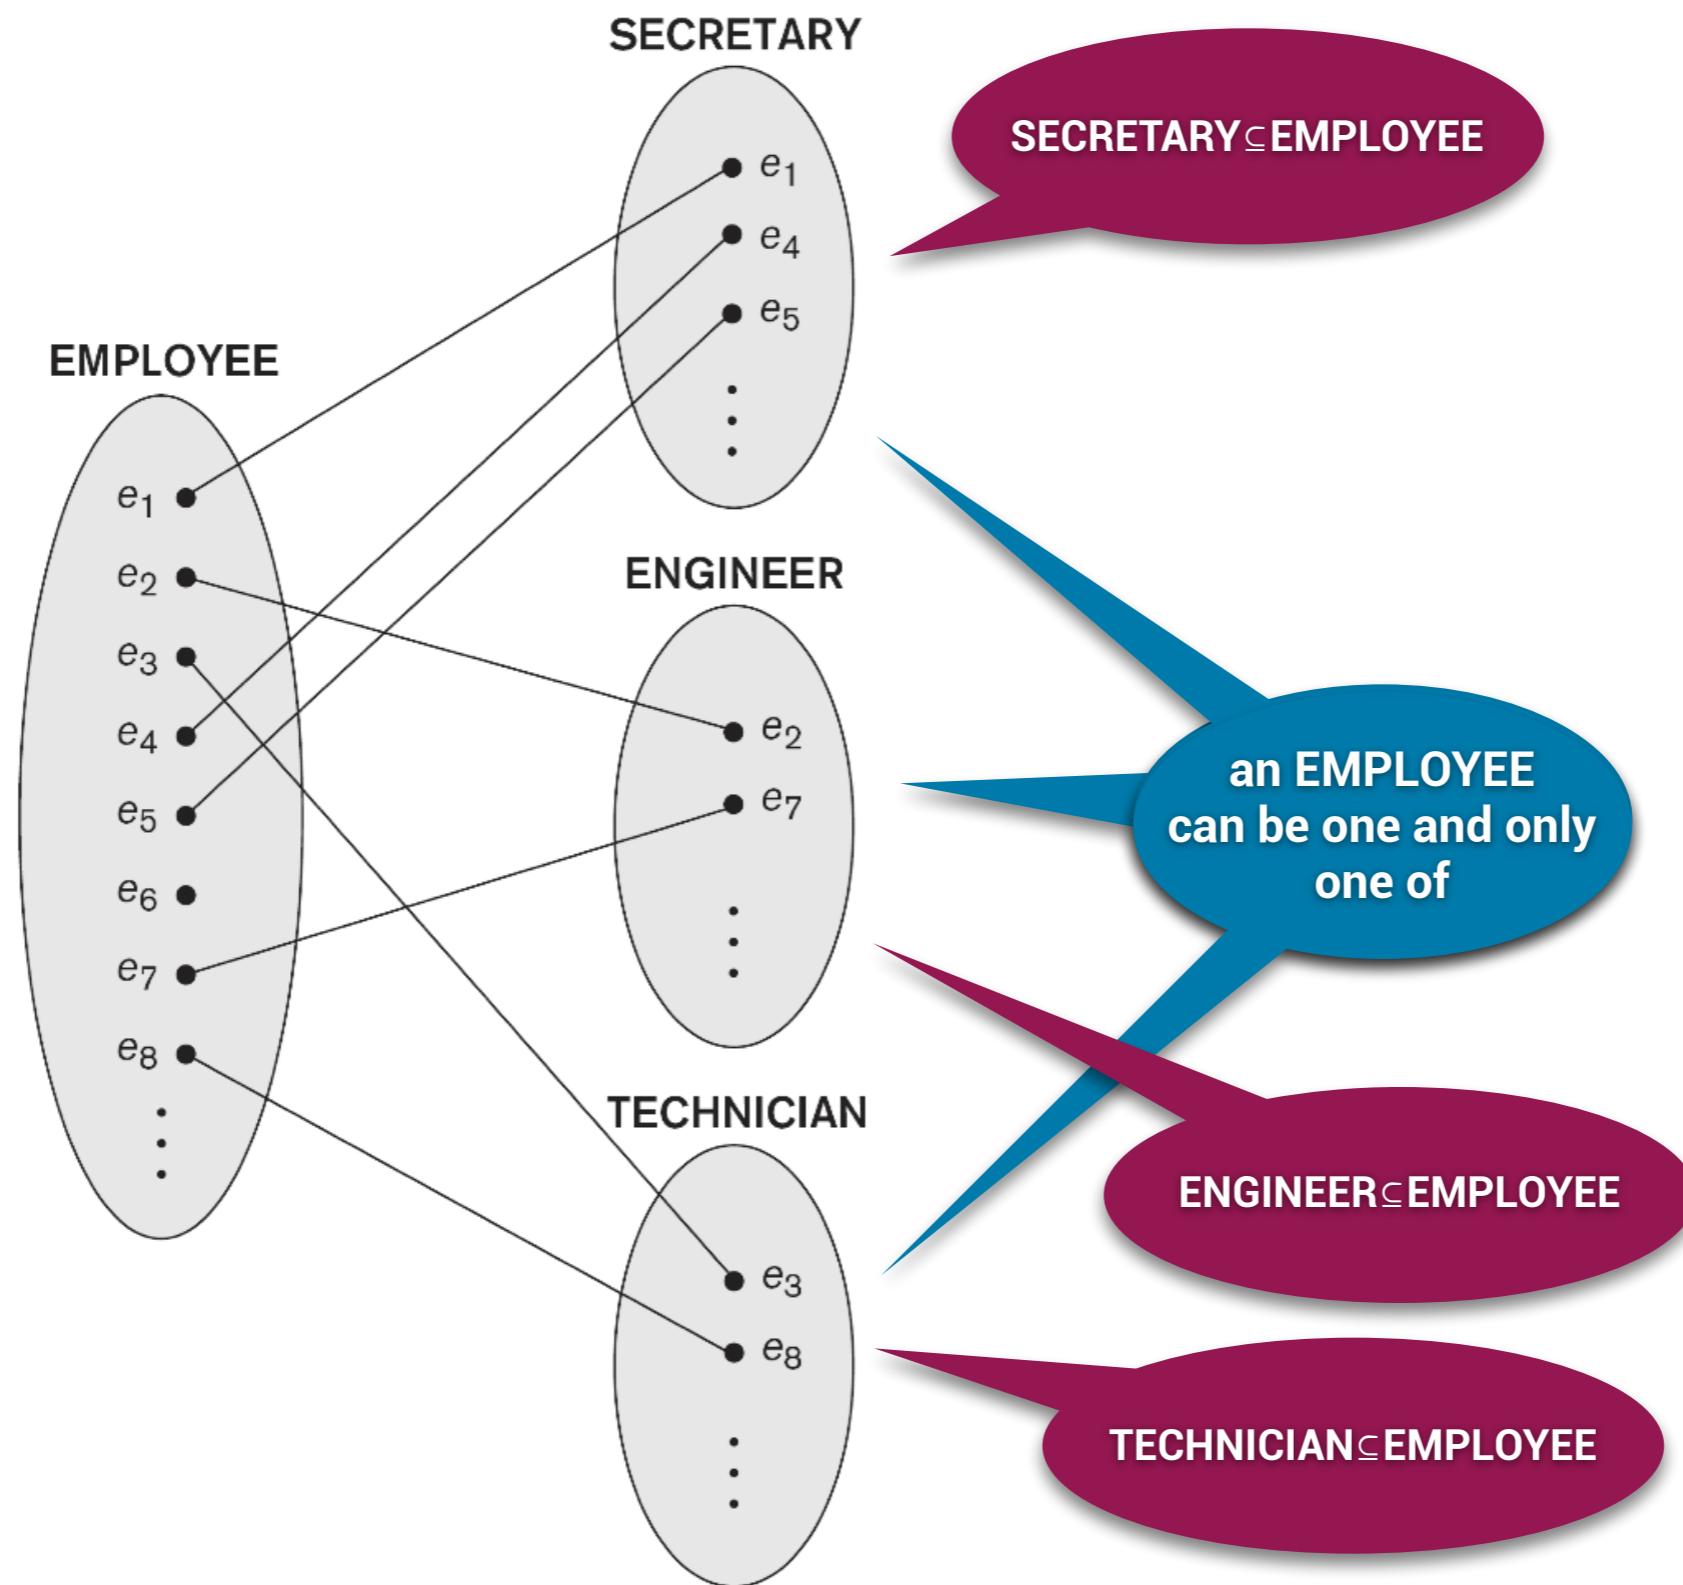
EER: Example Specializations

42



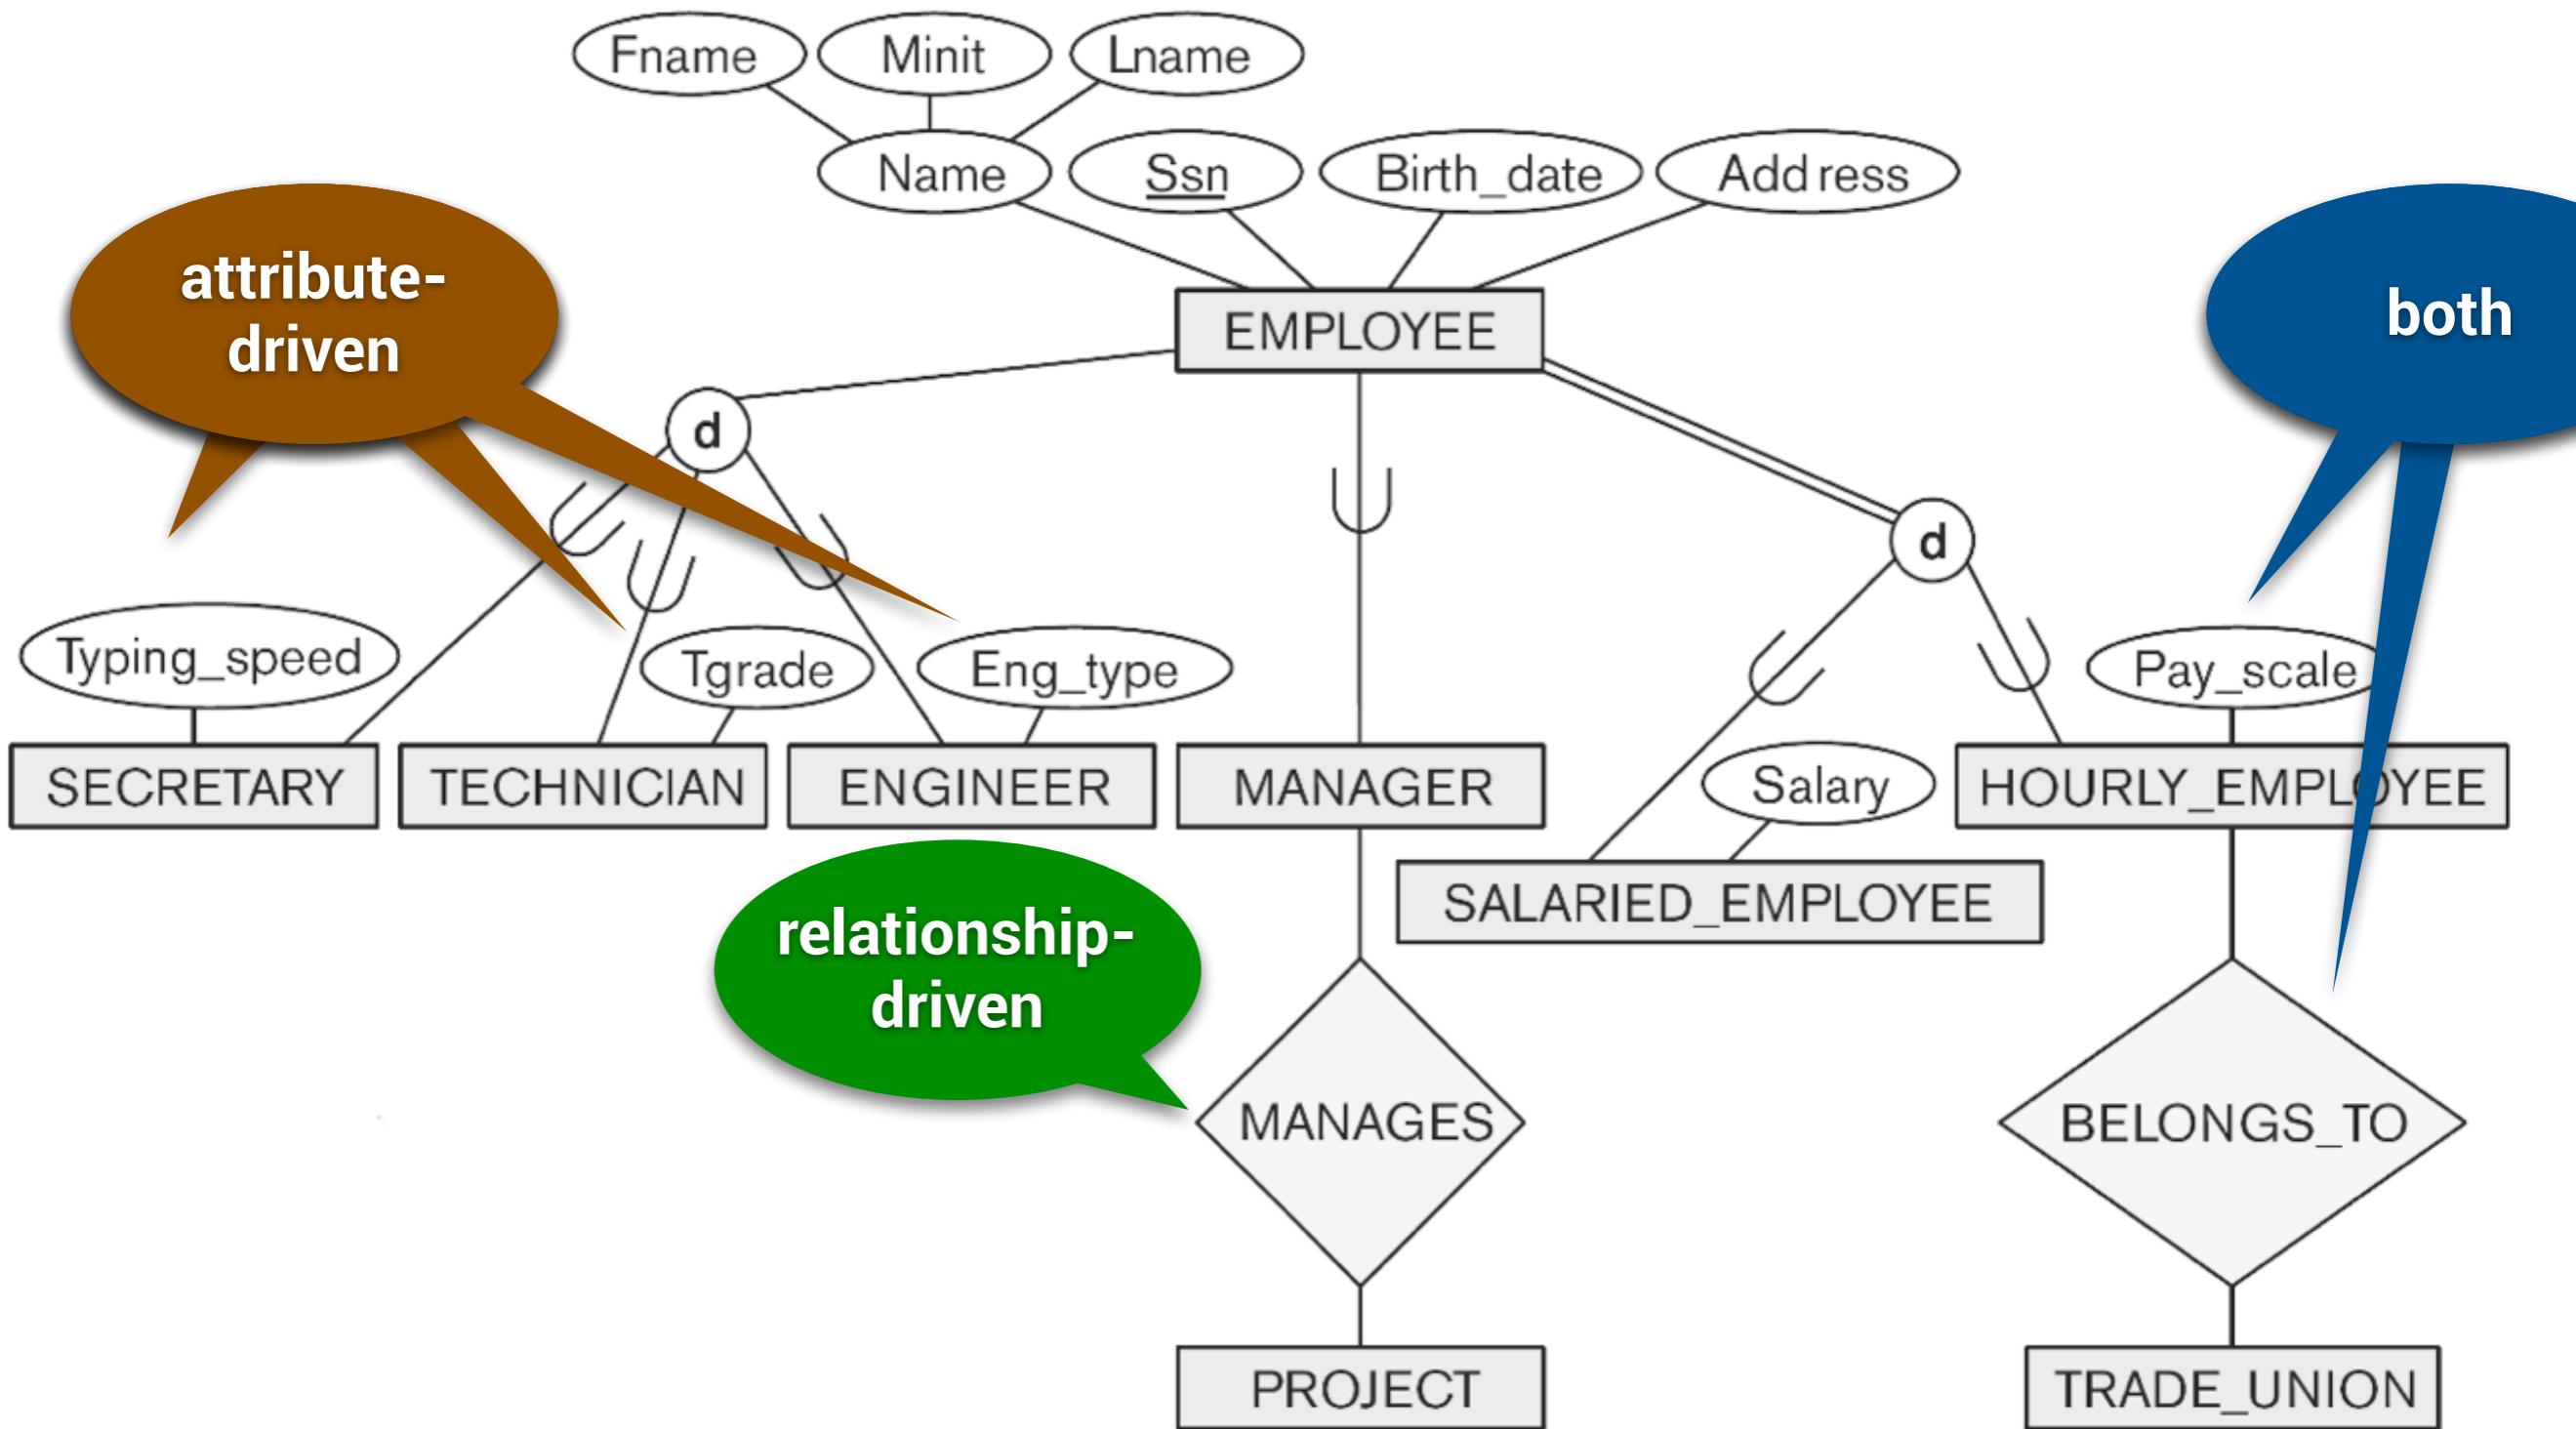
EER: Example Instances of a Specialization

43



EER: Attribute/Relationship-Driven Identification

44



EER: Subclass Characterization

45

- An entity subtype may be:

▶ **Attribute-defined**

Assuming that both salary or residency are stored attributes, then the following subtypes are **attribute-defined**:

Lower_Tax_Employee \subset Employee

\equiv

{e in Employee |
e.salary < 50K and
e.residency = 'UK'}

▶ **Predicate- (or condition-) defined**

▶ **User-defined**

Upper_Tax_Employee \subset Employee

\equiv

{e in Employee |
not e in Lower_Tax_Employee}

EER: Subclass Characterization

46

- An entity subtype may be:

▶ **Attribute-defined**

Assuming that residency is stored but EU_member is a Boolean-valued function, then the following subtypes are **predicate-defined**:

▶ **Predicate- (or condition-) defined**

Home_Student \subset Student \equiv
 $\{s \text{ in Student} \mid s.\text{residency} = \text{country} \text{ and } \text{EU_member}(\text{country})\}$

▶ **User-defined**

Overseas_Student \subset Student \equiv
 $\{s \text{ in Student} \mid \text{not } s \text{ in Home_Student}\}$

EER: Subclass Characterization

47

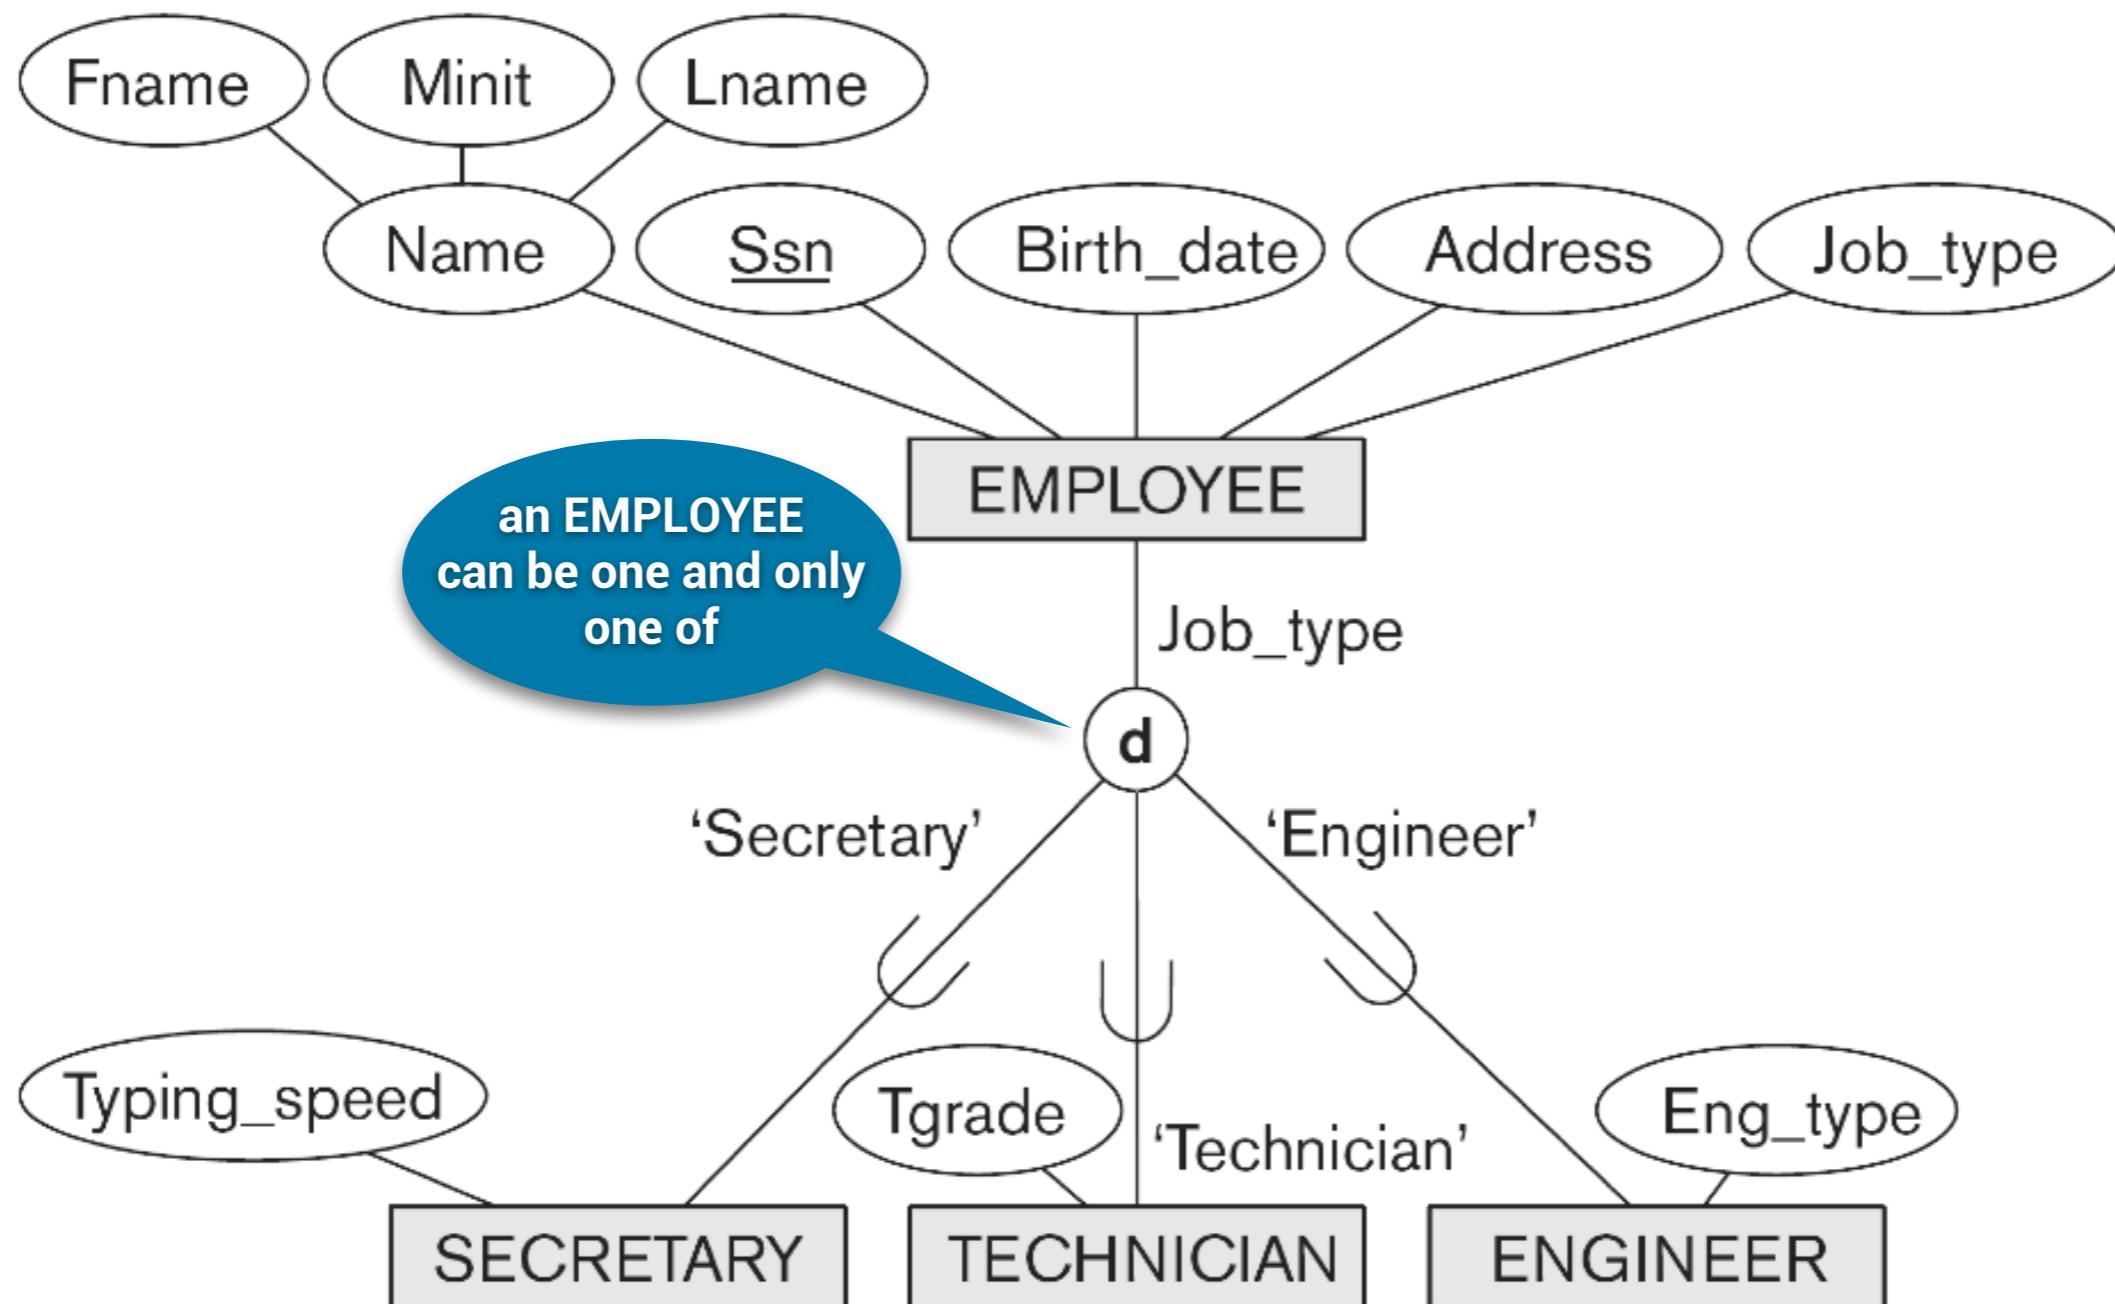
- An entity subtype may be:
 - ▶ **Attribute-defined**
 - ▶ **Predicate- (or condition-) defined**
 - ▶ **User-defined**

Assuming that prizeAwarded is a user input (as opposed to a programmed function), then the following subtype is **user-defined**:

PrizeRecipient \subset Student \equiv
 $\{s \text{ in Student} \mid$
prizeAwarded $\leftarrow\}$

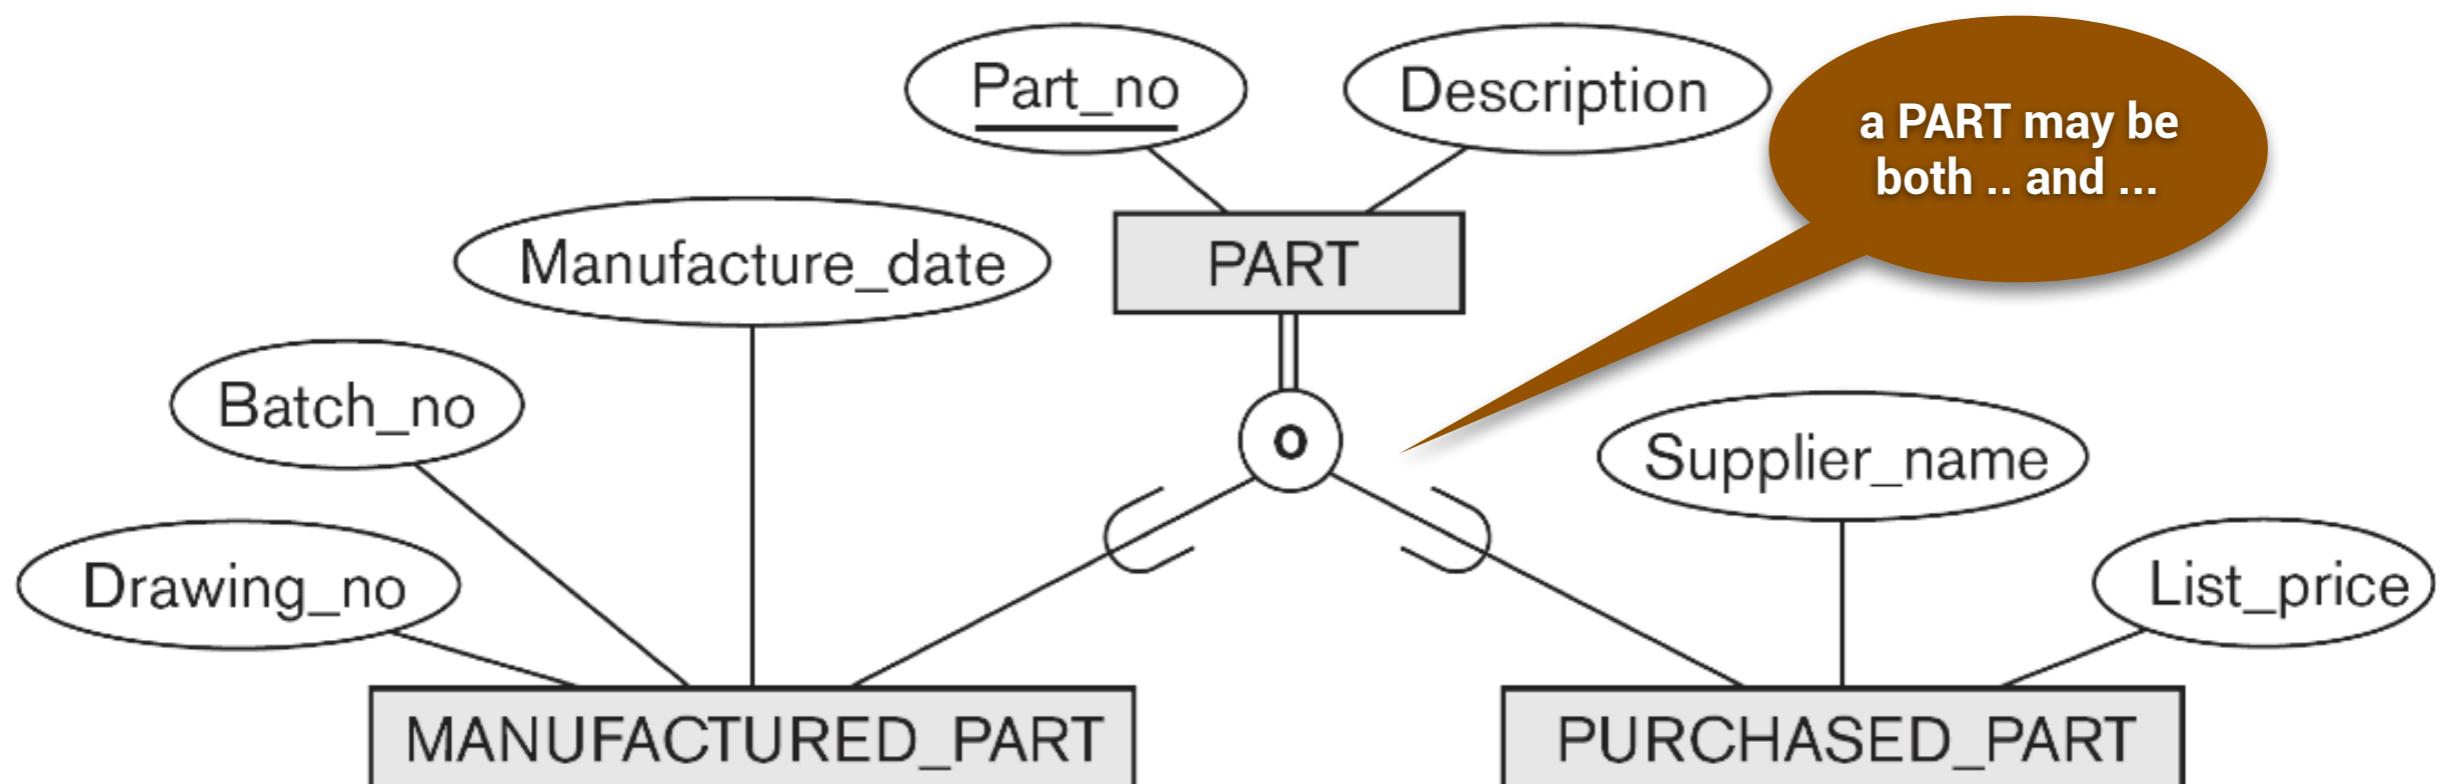
EER: Constraints on Subclass Relationships

48



EER: Constraints on Subclass Relationships

49



In The Next Lecture

x

- We'll start learning about the logical model of data known as the relational model.
- As a logical model, it is directly supported by existing DBMSs, and therefore is our target for implementation.

L 4

The Relational Model

Fundamentals of Databases
Alvaro A A Fernandes, SCS, UoM

Readings

X

This lecture relies on your having read the assigned mandatory readings associated with this lecture:

H. Garcia-Molina, J. D. Ullman, and J. Widom. Database Systems: The Complete Book. New International Edition, 2nd Edition. Pearson, 2014. pp. 13-62.

You can download the material using this link:

<https://contentstore.cla.co.uk/secure/link?id=b892c15f-3068-e611-80c6-005056af4099>

If you fail to study the material, you're likely to find the lecture harder to follow.

More information is available in the Learning Activities Manual.

In the Previous Lecture

x

- We learned about the conceptual model of data known as the entity-relationship model.
- We saw that it can be extended to capture more expressive modelling constructs such as specialization and generalization.

In This Lecture

x

- We'll start to have a deeper look at the relational model that underpins the most widely-deployed DBMSs today.
- We'll also begin to look at how SQL implements the relational languages by looking at its DDL and DML facilities.

What are the basic notions in data modelling?

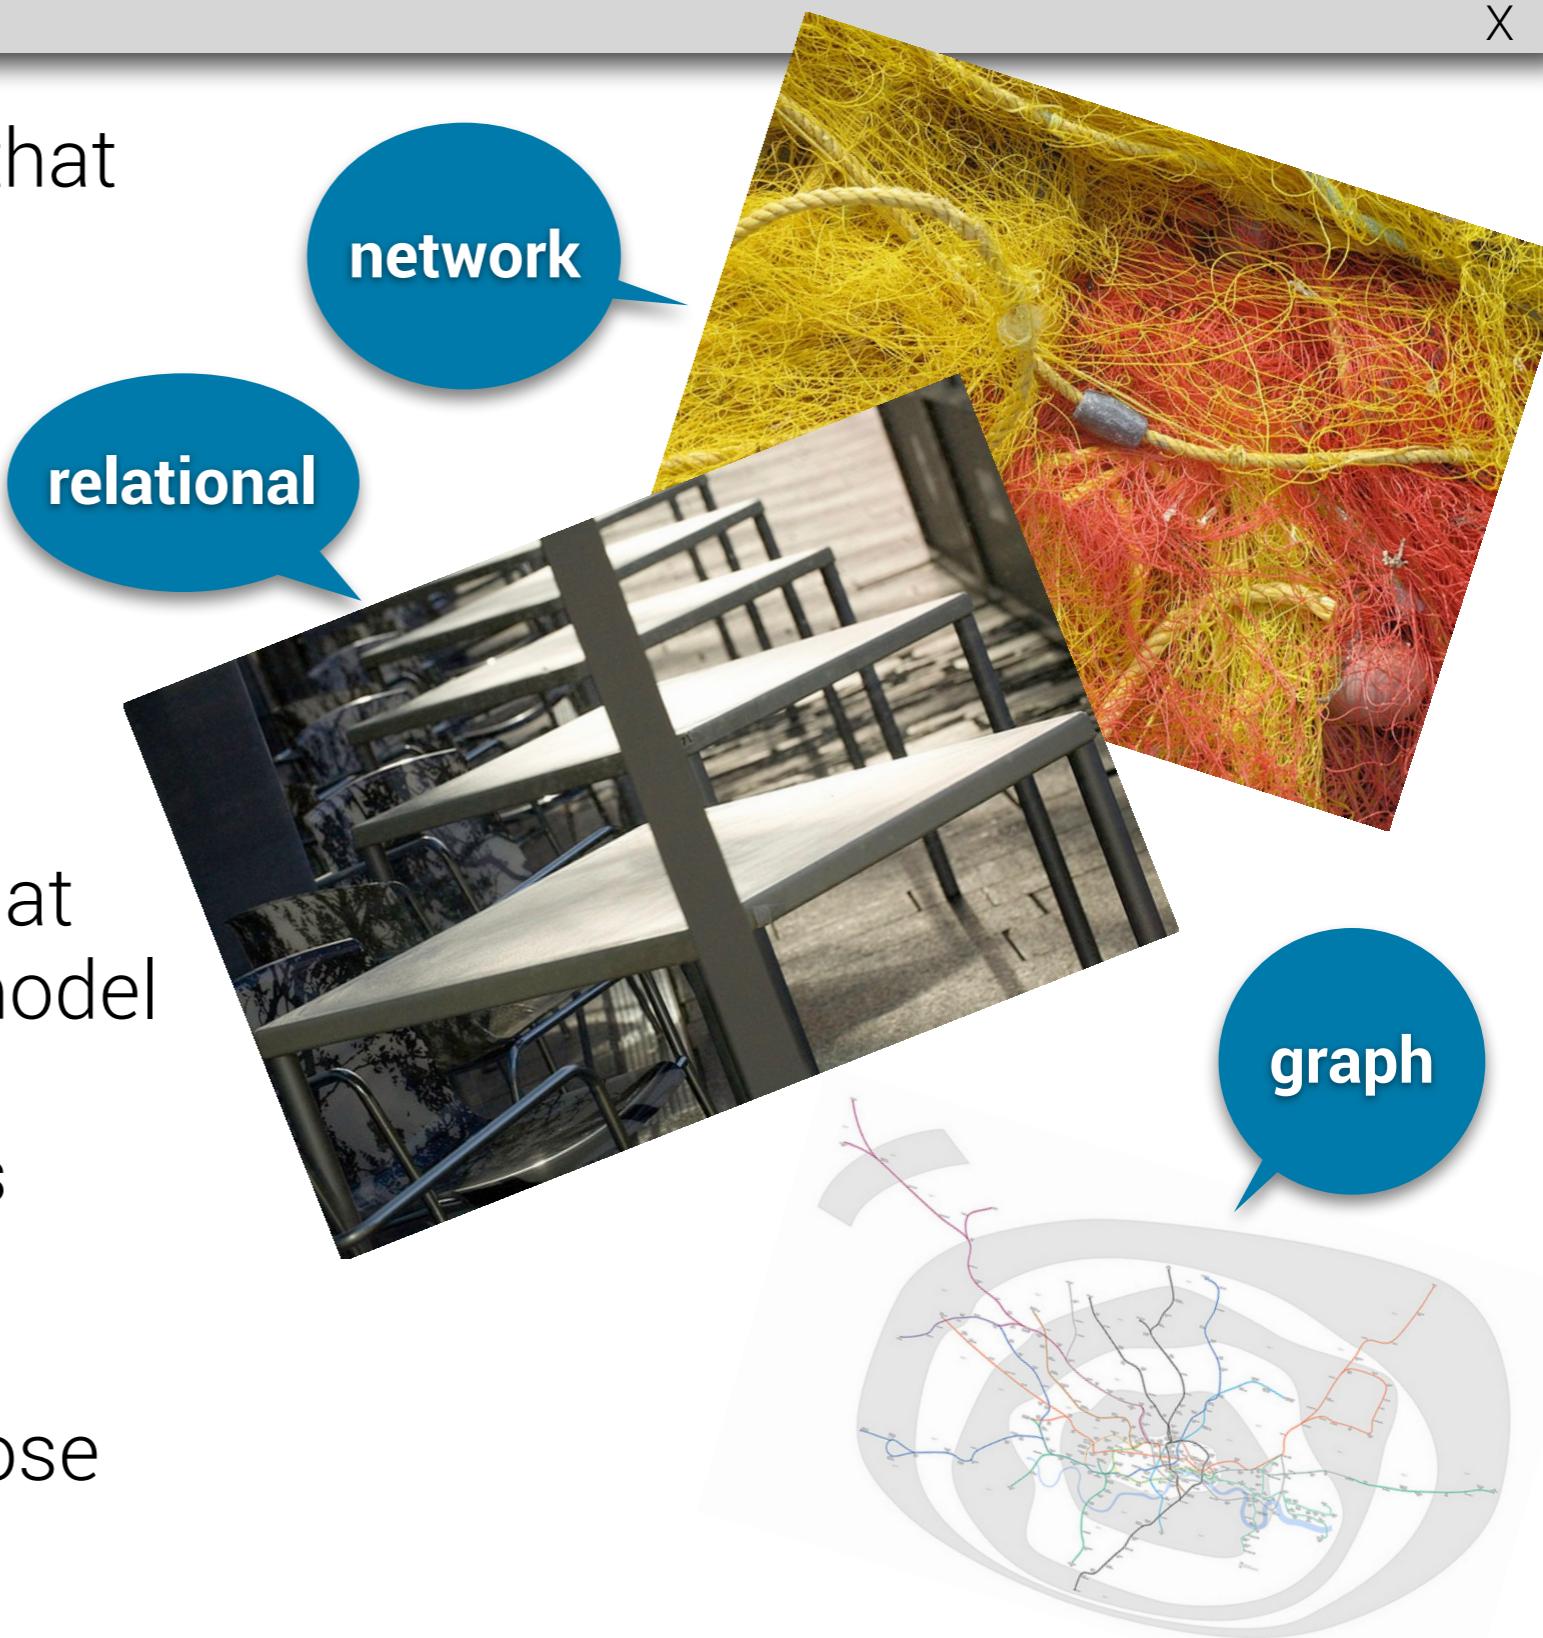
- Given an **application domain**, e.g., a university, we want to
 - ▶ identify the **information needs** that arise and
 - ▶ **formally represent** them
 - ▶ as a means to guide the development of effective and efficient **data provision**.
- We want to do so in a **wide and comprehensive manner**, keeping data and applications separate.
- We need a **DBMS** to achieve this.



What are the basic notions in data modelling?

x

- We choose a **data model** that is implemented by some DBMS.
- Then, the DBMS will be capable of:
 - ▶ making the most of the **structural properties** that characterize the data model
 - ▶ enforcing its **semantics**
 - ▶ deriving **logical consequences** from those semantics



What are the basic notions in data modelling?

X

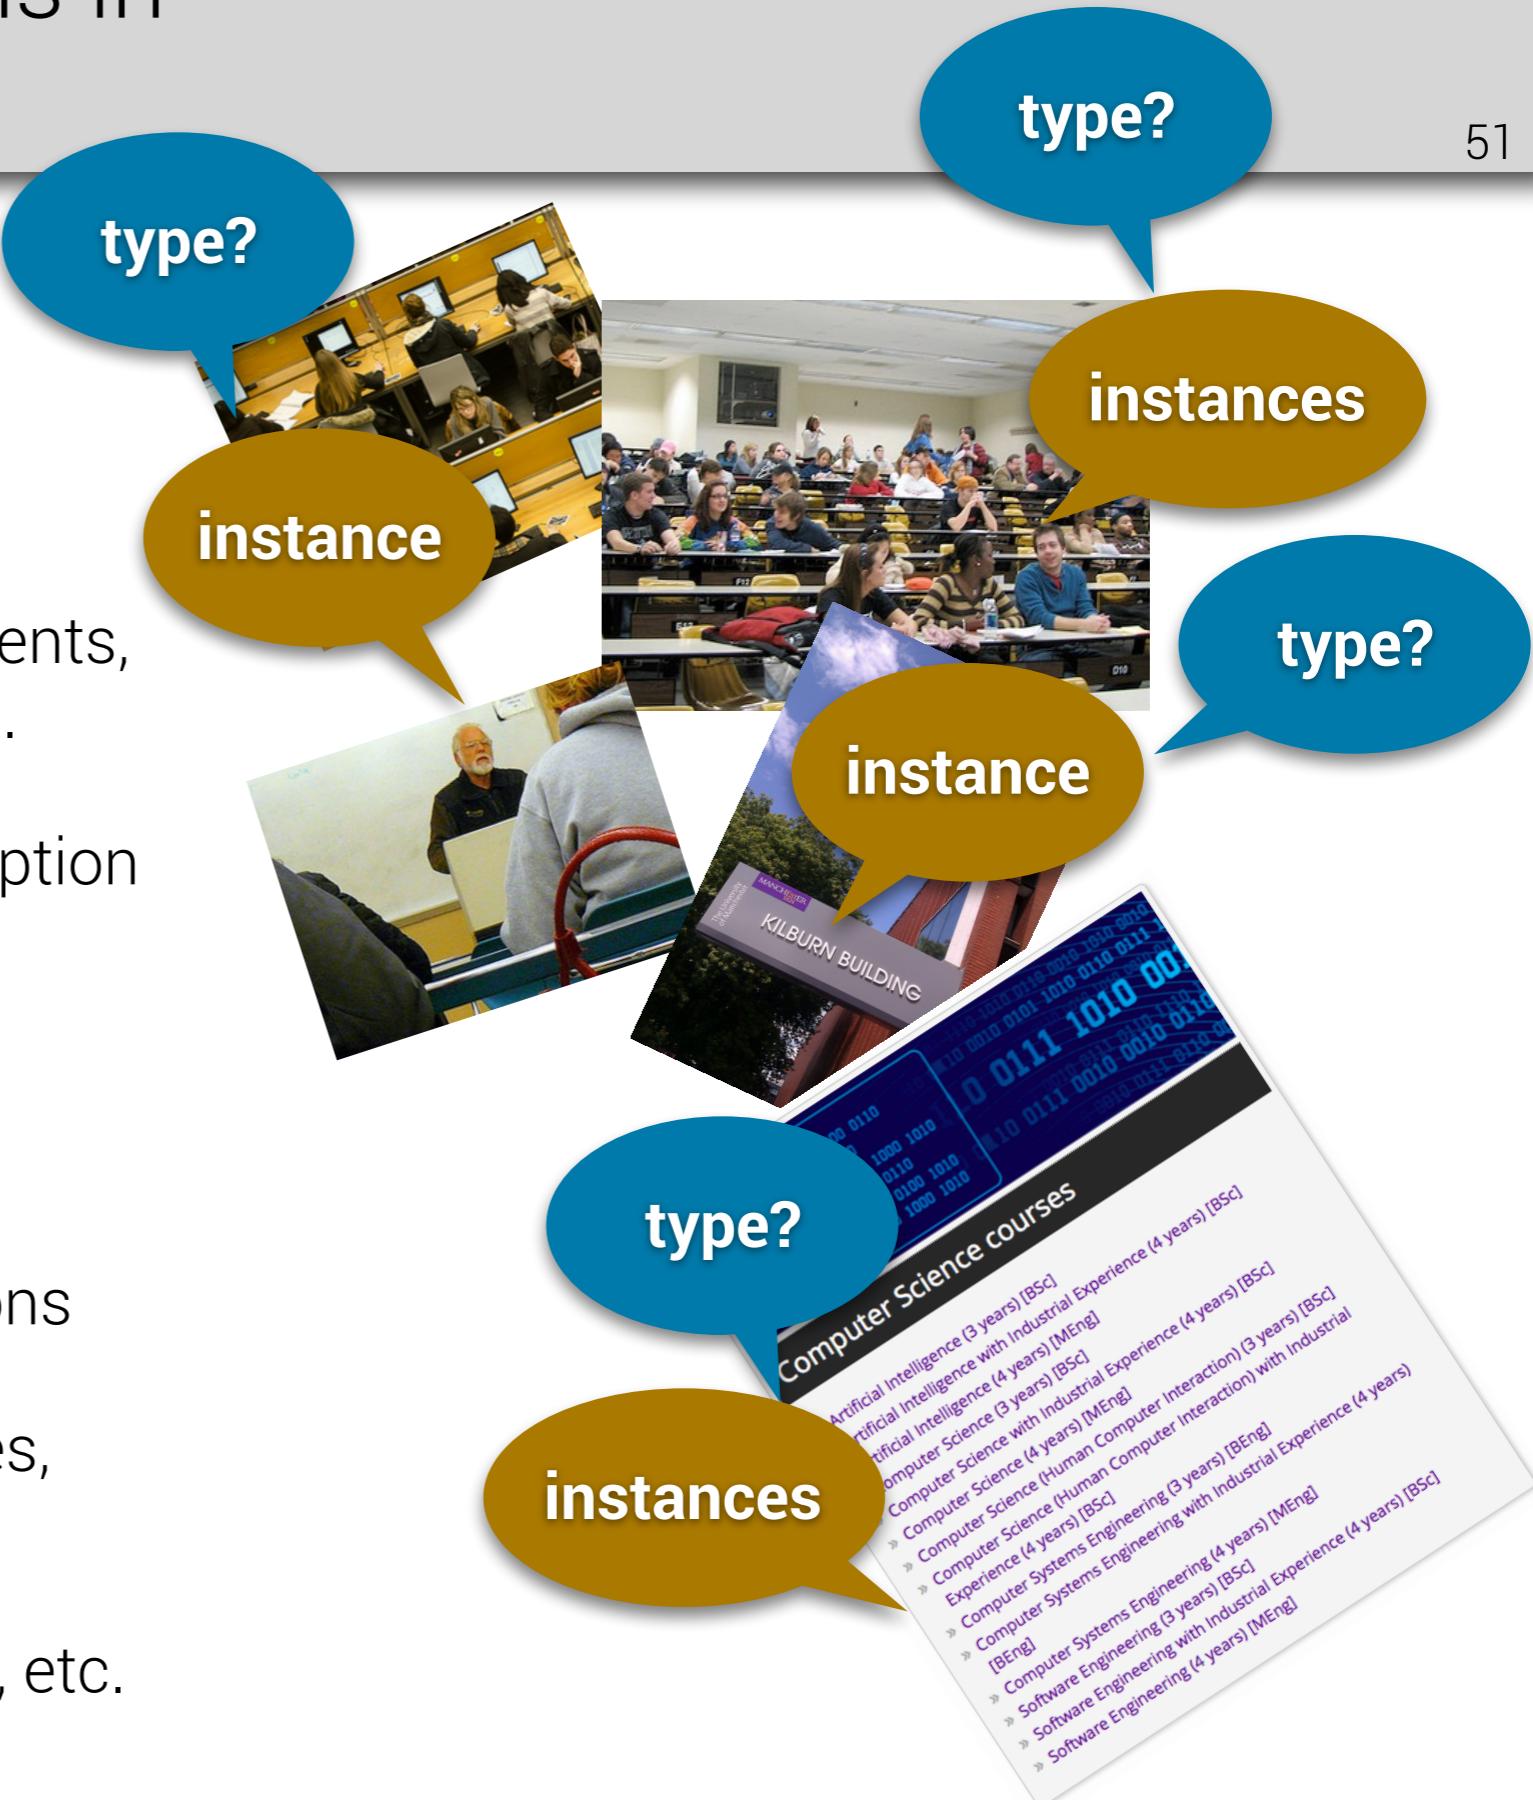
- For example, we need to store information about students, academics, departments, buildings, course units, etc.
- For each of those, we need, e.g., to know information about them
 - ▶ **students**: name, id, programme, courses taken, marks, etc.
 - ▶ **instructors**: name, id, courses taught, etc.
 - ▶ **departments**: name, building, programmes offered, etc.
 - ▶ **course units**: name, credits, programme, students, lecturer, location, etc.



What are the basic notions in data modelling?

51

- **Abstractions** for types
 - ▶ concepts, entities, objects, **relations**, nodes, etc.
 - ▶ students, instructors, departments, course units, programmes, etc.
- **Properties** for value-based description
 - ▶ properties, **attributes**, etc.
 - ▶ name, id, credits, marks, etc.
- **Relationships** between abstractions
 - ▶ relationships, **references**, edges, etc.
 - ▶ courses taught, courses taken, etc.



What are the basic notions in data modelling?

52

- **Constraints** on what can be said, e.g., **referential** ones:

▶ an instructor cannot teach a non-existent course

- Valid **calculations or operations** on what is known, e.g., **set comprehensions**

let

Jane_Eyre_Instructors

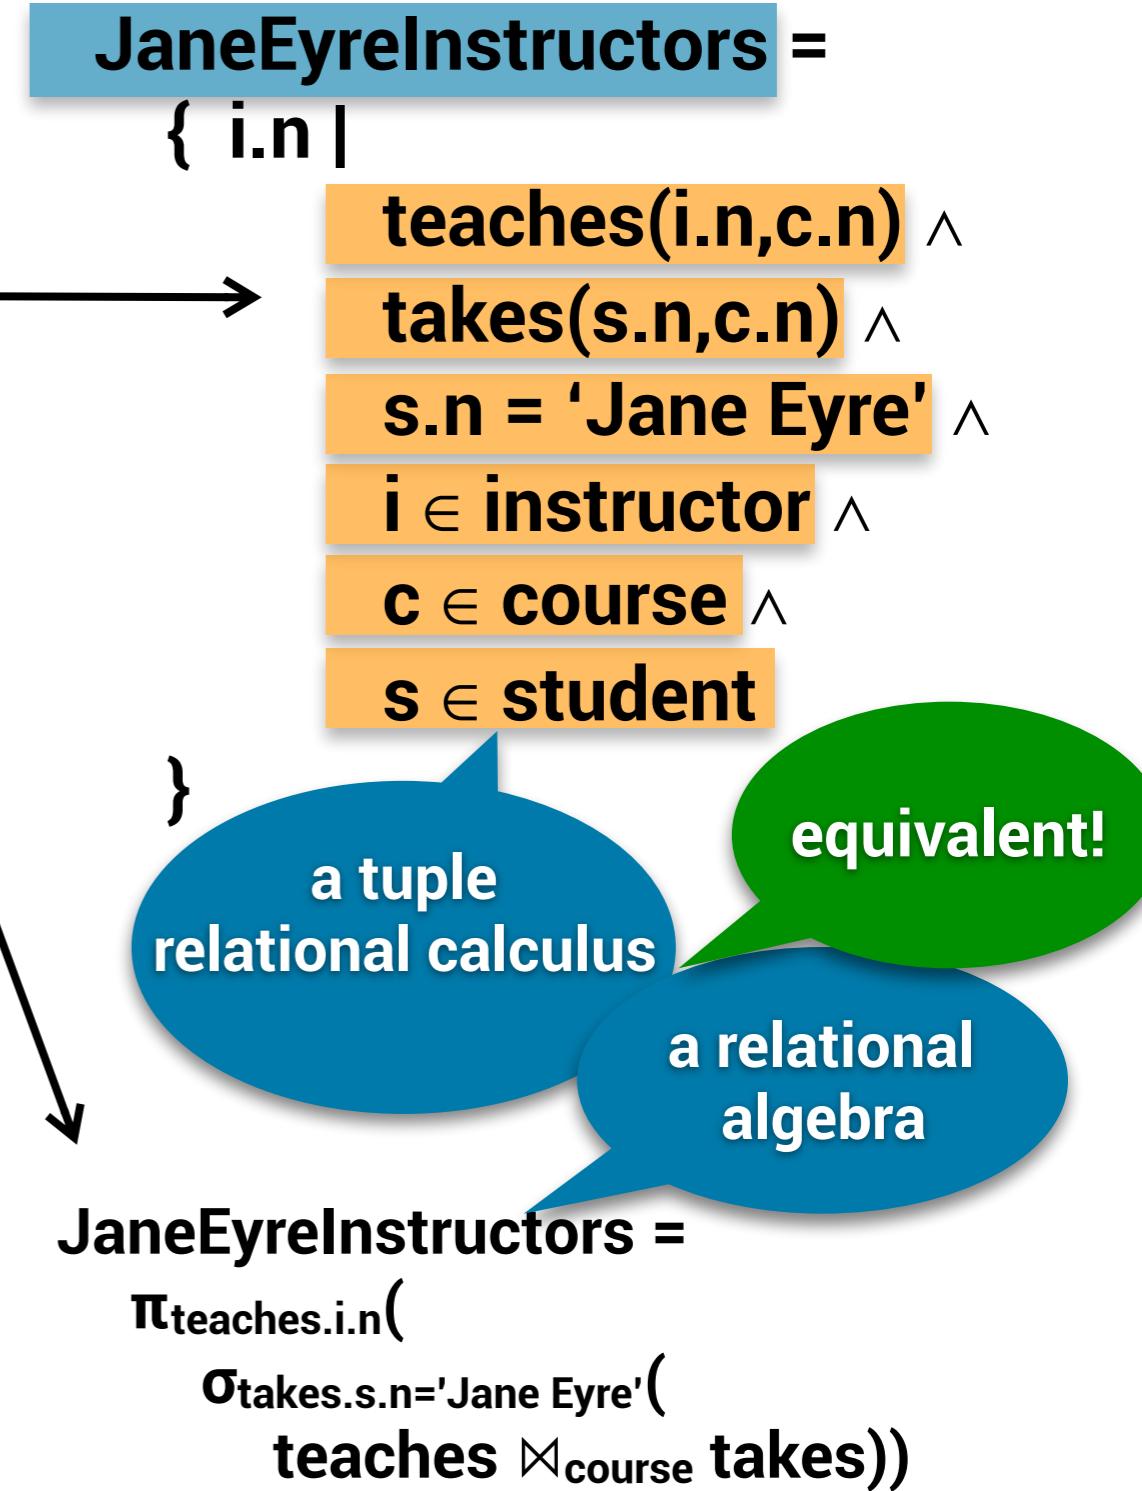
be the set of all **instructors** that teach some **course unit** that is taken by the student called Jane Eyre

even = $\{2x \mid x \in \mathbb{Z}\}$
odd = $\{2x+1 \mid x \in \mathbb{Z}\}$

The Relational Model

53

- **Single collection type**, i.e., **tables**, so a very **simple** model
- Leads to simple, but surprisingly expressive, **algebra** and **calculi**
- Effective: tables often match **how we think about data**
- Leads to **efficient implementations**: it underlies the most important commercial DBMSs today
- Basis of **SQL**: the most widely-used language for interacting with databases today



The Relational Model

54

- **Regular, rectangular tables** with **typed, atomic columns**
- Table definitions capture **relation schemas/types**, e.g., *instructors*
- **Columns** capture properties, e.g., instructor name
- **Rows** record individuals, e.g., 15151 (whose name is Mozart)
- Table states capture relation instances (i.e., type extensions)
- Formally, relation instances are subsets of the n -ary Cartesian product over attribute domains

The diagram shows a relational table named *instructor* with four columns: *ID*, *name*, *dept_name*, and *salary*. The table has 13 rows, each representing an instructor. The rows are color-coded: rows 1 through 10 are light blue, row 11 is pink, row 12 is light green, and row 13 is light blue. Four blue speech bubbles provide annotations:

- domain integrity**: Points to the *dept_name* column, which contains values from a limited set of department names.
- homogeneous, type-conformant columns**: Points to the *name* column, which contains all string values.
- entity integrity**: Points to the first two columns (*ID* and *name*), indicating that each row must be uniquely identifiable by its *ID*.
- uniquely identifiable rows**: Points to the last two columns (*dept_name* and *salary*), indicating that each row must have a unique combination of these two attributes.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
44443	Singh	Finance	80000

(a) The *instructor* table

The Relational Model

values in the
referring column must
occur in the referred
column

55

- Relationships are captured by references from a column in one table to a column in another
- For example, to which department an *instructor* belongs
- If we have operations on tables (e.g., selecting rows, projecting and aggregating column values, combining tables, etc.) we can answer a question such as
 - Is Computer Science broke??!

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

referential
integrity

(a) The *instructor* table

dept_name	building	budget
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

What is SQL?

56

- SQL is the most widely-used relational language
- It is based on a (declarative) relational calculus (i.e., set comprehensions)
- The most important SQL queries can be translated into an equivalent (procedural) algebraic expression for evaluation
- This approach of translating from a declarative expression to an equivalent procedural one
 - ▶ frees application designers from having to write efficient query execution plans
 - ▶ the DBMS is delegated the task of searching to find one such plan

What is SQL?

57

- SQL has an expressive QL subset centred on the well-known

SELECT

FROM

WHERE

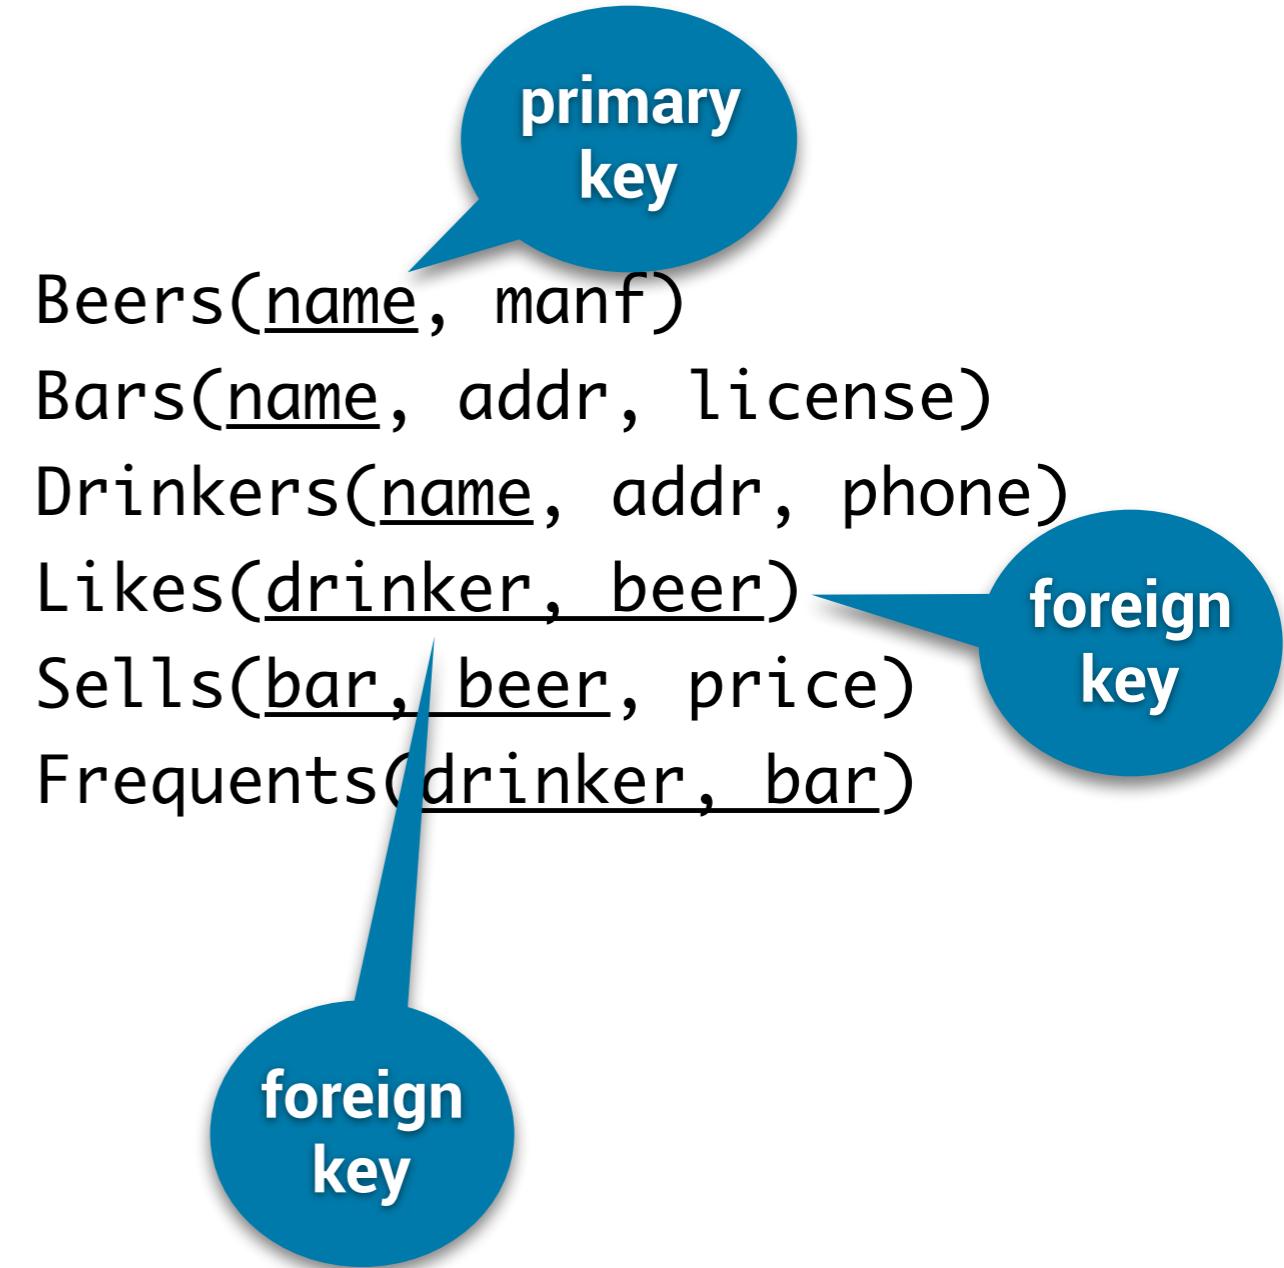
GROUP BY idiom

- SQL has expressive DML capabilities:
 - Reactive triggers
 - Procedural constructs (e.g, loops, assignments, etc.)

BBD DB :: A Running Example

58

- Assume a database about beers, bars and drinkers.
- There must exist a choice of a set of attribute values that uniquely identifies a tuple: we call it the **primary key** of the relation.
- One relation may use the primary key of another relation to model a relationship: we call it a **foreign key**.



SQL as a DDL: Schemas

59

CREATE
ALTER
DROP

INDEX
VIEW
...

INTEGER
FLOAT
REAL
CHAR(n)
VARCHAR(n)
DATE
TIME
...

```
CREATE TABLE Bars (
    name    CHAR(20),
    manf   VARCHAR2(20),
    CONSTRAINT ba_pk PRIMARY KEY (name)
);
```

```
CREATE TABLE Beers (
    name    VARCHAR2(20),
    addr   VARCHAR2(20),
    license CHAR(3),
    CONSTRAINT be_pk PRIMARY KEY (name)
);
```

```
CREATE TABLE Sells (
    bar      CHAR(20),
    beer    VARCHAR2(20),
    price   REAL,
    CONSTRAINT s_pk PRIMARY KEY (bar, beer)
    CONSTRAINT s_fk_bar FOREIGN KEY (bar)
        REFERENCES Bars(name),
    CONSTRAINT s_fk_beer FOREIGN KEY (beer)
        REFERENCES Beers(name)
);
```

SQL as a DDL: Constraints

60

specifies what the DBMS must do to rows in this table when events take place in the Bars or the Beers table

the events are either deletions or updates in the primary key

the default is to reject (i.e., prevent Bars from being changed)

```
CREATE TABLE Sells (
    bar      CHAR(20),
    beer     VARCHAR2(20),
    price    REAL,
    CONSTRAINT s_pk
        PRIMARY KEY (bar, beer)
    CONSTRAINT s_fk_bar
        FOREIGN KEY (bar)
            REFERENCES Bars(name)
            ON DELETE SET NULL
            ON UPDATE CASCADE,
    CONSTRAINT s_fk_beer
        FOREIGN KEY (beer)
            REFERENCES Beers(name)
            ON DELETE SET NULL
            ON UPDATE CASCADE,
);
```

we can set the foreign key to NULL (e.g., on delete)

we can CASCADE, i.e., update the foreign key here to the new value of the primary key in Bars

SQL as a DML: Insert Single

61

```
INSERT INTO Likes  
VALUES ('Sally', 'Bud');
```

implicit
order

```
INSERT INTO  
Likes(beer, drinker)  
VALUES ('Bud', 'Sally');
```

explicit
order

```
CREATE TABLE Drinkers (  
    name CHAR(30) PRIMARY KEY,  
    addr CHAR(50) DEFAULT '123 Sesame St.',  
    phone CHAR(16));
```

implicit
default

```
INSERT INTO Drinkers(name)  
VALUES ('Sally');
```

this input

```
('Sally', '123 Sesame St', NULL)
```

this
result

SQL as a DML: Insert Set

62

```
INSERT INTO SallyPossibleBuddies  
(SELECT d2.drinker  
FROM Frequent d1, Frequent d2  
WHERE d1.drinker = 'Sally' AND  
d2.drinker <> 'Sally' AND  
d1.bar = d2.bar  
);
```

d1 is made to
point to Sally

d2 is made to
point to all the
other drinkers

This ensures
that Sally
(d1) and the
other drinker
(d2) frequent
the same bar.

SQL as a DML: Delete/Update

63

```
DELETE FROM Likes;
```

unconditional:
deletes everything

delete from Beers(name, manf)
all beers for which there is
another beer by the same
manufacturer.

only deletes the rows that
satisfy the condition

```
DELETE FROM Beers b1  
WHERE EXISTS (SELECT name  
FROM Beers b2  
WHERE b2.manf=b1.manf  
AND b2.name<>b1.name);
```

```
UPDATE Drinkers  
SET phone = '555-1212'  
WHERE name = 'Fred';
```

updates one
row
(name is PK)

```
UPDATE Sells  
SET price = 4.00  
WHERE price > 4.00;
```

updates many rows
(potentially)

In The Next Lecture

x

- We'll delve into the relational algebra.
- We'll look into how SQL can be mapped into the relational algebra.
- We'll explore how this declarative-to-procedural transformation is a key one in DBMSs.

L_5

The Relational Languages

Fundamentals of Databases
Alvaro A A Fernandes, SCS, UoM

Readings

X

This lecture relies on your having read the assigned mandatory readings associated with this lecture:

Karen Morton et al. Pro Oracle SQL. 2nd Edition. Apress, 2013. pp. 1-24.

You can download the material using this link:

https://link.springer.com/chapter/10.1007/978-1-4302-6221-3_1

If you fail to study the material, you're likely to find the lecture harder to follow.

More information is available in the Learning Activities Manual.

In Previous Lectures

x

- We learned that data is an enterprise asset and that DBMSs are crucial to manage it well.
- We learned the importance of adopting different levels of abstraction in designing and implementing databases.
- We learned how data models lead to schemas and instances that enable a logical view of the data.
- We learned about the main components of DBMSs and the various architectures used to deploy them.
- We learned about conceptual modelling.
- We started learning about the relational approach to logical data modelling.
- We started learning SQL, focussing on its DDL and DML capabilities.

In This Lecture

x

- What is a relational algebra?
- What are the basic operations of relational algebra?
- How do we relate SQL to relational algebra?
- What are the main constructs and query types in SQL?

A Familiar Example: Elementary Algebra

65

operator
semantics leads to
evaluation

e.g., 6

closed:
inputs and outputs
from the same set

operator
semantics leads to
equivalence laws

operators:
 $+, -, *, /$

$$2+2$$

$$4*(3/2)$$

$$(4*3)/2$$

$$x+1$$

operands:
values from the
Reals

operands:
variables

$$2/(4*x)$$

$$x+y = y+x$$

$$(x*y)*z = x*(y*z)$$

e.g.,
commutativity of $+$

e.g.,
associativity of $*$

A Relational Algebra

66

Which operator semantics?
How to evaluate?

Which operands?

What do variables denote?

Which operators?

Closed? Inputs
and outputs one and the
same type?

relations!

$\sigma, \pi, \bowtie, \rho, \delta!$
 $\cup, \cap, \setminus!$

e.g., which
commutative?

Which equivalence laws
does operator semantics lead to?

yes!

e.g., \bowtie, \dots

e.g., which
associative?

$$\sigma_\theta(R \bowtie S) \equiv \sigma_\theta(R) \bowtie \sigma_\theta(S)$$

How do we want to operate on relations?

67

R(a:int, b:int, c:str, d:str, e:flt)

rows

single table

combining
tables

columns

S(b:int, g:int, x:str, z:flt)

T(v:int, w:int, x:str, y:str, z:flt)

How do we want to operate on relations?

The unary case

68

$R(a:\text{int}, b:\text{int}, c:\text{str}, d:\text{str}, e:\text{flt})$

grow/
shrink left-to-right

projection:
 π

$R'(a:\text{int}, e:\text{flt})$

add/keep
some columns

How do we want to operate on relations?

The unary case

69

$R(a:\text{int}, b:\text{int}, c:\text{str}, d:\text{str}, e:\text{flt})$

shrink top-to-bottom

$R'(a:\text{int}, b:\text{int}, c:\text{str}, d:\text{str}, e:\text{flt})$

selection: σ

keep only some rows

How do we want to operate on relations?

The binary case

70

$R(a:\text{int}, b:\text{int}, c:\text{str}, d:\text{str}, e:\text{flt})$

grow/shrink
top-to-bottom

keep some rows
and add some more

$T'(v:\text{int}, w:\text{int}, x:\text{str}, y:\text{str}, z:\text{flt})$

$T(v:\text{int}, w:\text{int}, x:\text{str}, y:\text{str}, z:\text{flt})$

union: \cup

except:
\\

intersect:
 \cap

How do we want to operate on relations?

The binary case

71

$T(v:\text{int}, w:\text{int}, x:\text{str}, y:\text{str}, z:\text{flt})$

$S(b:\text{int}, g:\text{int}, x:\text{str}, z:\text{flt})$

grow left-to-right

keep/add some columns and drop some columns

$S'(v:\text{int}, w:\text{int}, x:\text{str}, y:\text{str}, z:\text{flt}, b:\text{int}, g:\text{int})$

product:
x

join: \bowtie

Relational-Algebraic Operations: The (prefix) unary case

x

$R(a:\text{int}, b:\text{int}, c:\text{str}, d:\text{str}, e:\text{flt})$
(*<operator>*)

<parameters>

π : projection
σ : selection
ρ : rename

π : attribute list
σ : predicate
ρ : attribute list

relation-
denoting
variable (or
expression)

$\pi_{(a,e)}(R)$

$\sigma_{a>0}(R)$

$\rho_{a \rightarrow z}(R)$

$\rho_{R(z,b,c,d,e)}(R)$

Relational-Algebraic Operations: The (infix) binary case

x

$T(v:\text{int}, w:\text{int}, x:\text{str}, y:\text{str}, z:\text{flt})$

$S(b:\text{int}, g:\text{int}, x:\text{str}, z:\text{flt})$

((left-operand)<operator>

(right-operand))

<parameters>

relation-denoting variable (or expression)

\cup : union
 \setminus : except
 \cap : intersect
 \times : product
 \bowtie : join

$(T \cup S)$

$(T \setminus S)$

$(T \cap S)$

$(T \times S)$

\cup : <none>
 \setminus : <none>
 \cap : <none>
 \times : <none>
 \bowtie : predicate

relation-denoting variable (or expression)

$(T \bowtie_{T.z = S.z} S)$

Selection: Notation and Example

72

Sells		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00



JoeMenu		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75

$\text{JoeMenu} := \sigma_{\text{bar} = 'Joe's'}(\text{Sells})$

Projection: Notation and Example

73

bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

$\text{BeerPrices} := \pi_{\text{beer}, \text{price}}(\text{Sells})$

BeerPrices	
beer	price
Bud	2.50
Miller	2.75
Miller	3.00

(Cartesian) Product: Notation and Example

74

R_1	
A	B
1	2
3	4

$$R_3 := R_1 \times R_2$$

R_2	
C	D
5	6
7	8
9	10

R_3			
A	B	C	D
1	2	5	6
1	2	7	8
1	2	9	10
3	4	5	6
3	4	7	8
3	4	9	10

T'1' R1 + T'5' in R2!

T'1' R1 + T'7' in R2!

T'1' R1 + T'9' in R2!

T'3' R1 + T'5' in R2!

⋮
⋮
⋮

Renaming: Notation and Example

75

before!

Bars	
name	addr
Joe's	Maple St.
Sue's	River Rd.

Barz := $\rho_{\text{Bars}}(\text{bar}, \text{addr})$ (Bars)

after!

bar	addr
Joe's	Maple St.
Sue's	River Rd.

θ -Join: Notation and Example

76

one
from Sells!

Sells		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Coors	3.00

one
from Bars!

Bars	
name	addr
Joe's	Maple St.
Sue's	River Rd.

keep those
concats for which θ is
true!

Is this table a
bit odd?!

BarInfo := Sells $\bowtie_{\text{Sells.bar} = \text{Bars.name}}$ Bars

BarInfo				
bar	beer	price	name	addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Joe's	Maple St.
Sue's	Bud	2.50	Sue's	River Rd.
Sue's	Coors	3.00	Sue's	River Rd.

(Natural) Join: Notation and Example

X

Sells			Barz
bar	beer	price	
Joe's	Bud	2.50	
Joe's	Miller	2.75	
Sue's	Bud	2.50	
Sue's	Coors	3.00	

Same name?

BarInfo := Sells \bowtie Barz

Better!

Keep only one of
the columns!

bar	beer	price	addr
Joe's	Bud	2.50	Maple St.
Joe's	Miller	2.75	Maple St.
Sue's	Bud	2.50	River Rd.
Sue's	Coors	3.00	River Rd.

Extended Projection: Notation and Example

X

R	
A	B
1	2
3	4
5	6
7	8

$$S := \pi_{A+B \rightarrow C, A, A \rightarrow A2}(R)$$

simple arithmetic

replication

A+B! A A again!

S

C	A	A2
3	1	1
7	3	3
11	5	5
15	7	7

Complex Expression: Alternative Forms

X

expression

linear, textual form

$\Pi_{teaches.i_n}($

$\sigma_{takes.s_n='Jane Eyre'}($
teaches \bowtie_{course} takes))

named intermediate results

assignment sequence

operator tree

dataflow form

teaches

takes

$\Pi_{teaches.i_n}$

and this one?

$\sigma_{takes.s_n='Jane Eyre'}$

name for this node?

\bowtie_{course}

$Tmp1 := \text{teaches} \bowtie_{course} \text{Takes}$

$Tmp2 := \sigma_{takes.s_n='Jane Eyre'}(Tmp1)$

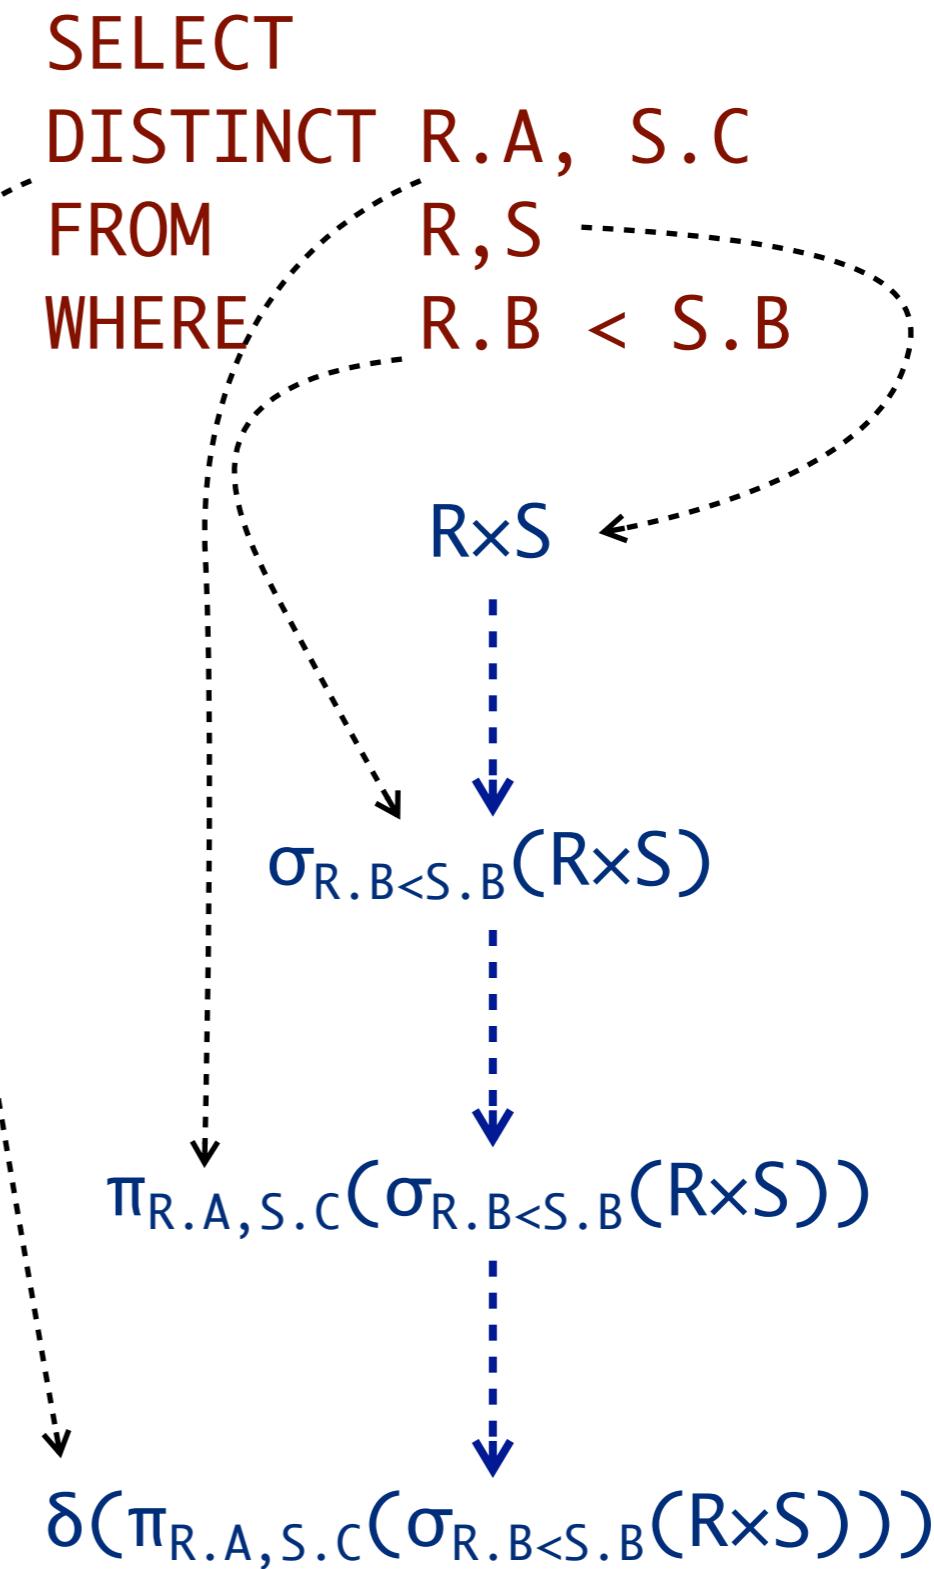
$\text{JaneEyreInstructors} := \Pi_{teaches.i_n}(Tmp2)$

Direct Translation of Core SQL to RA

77

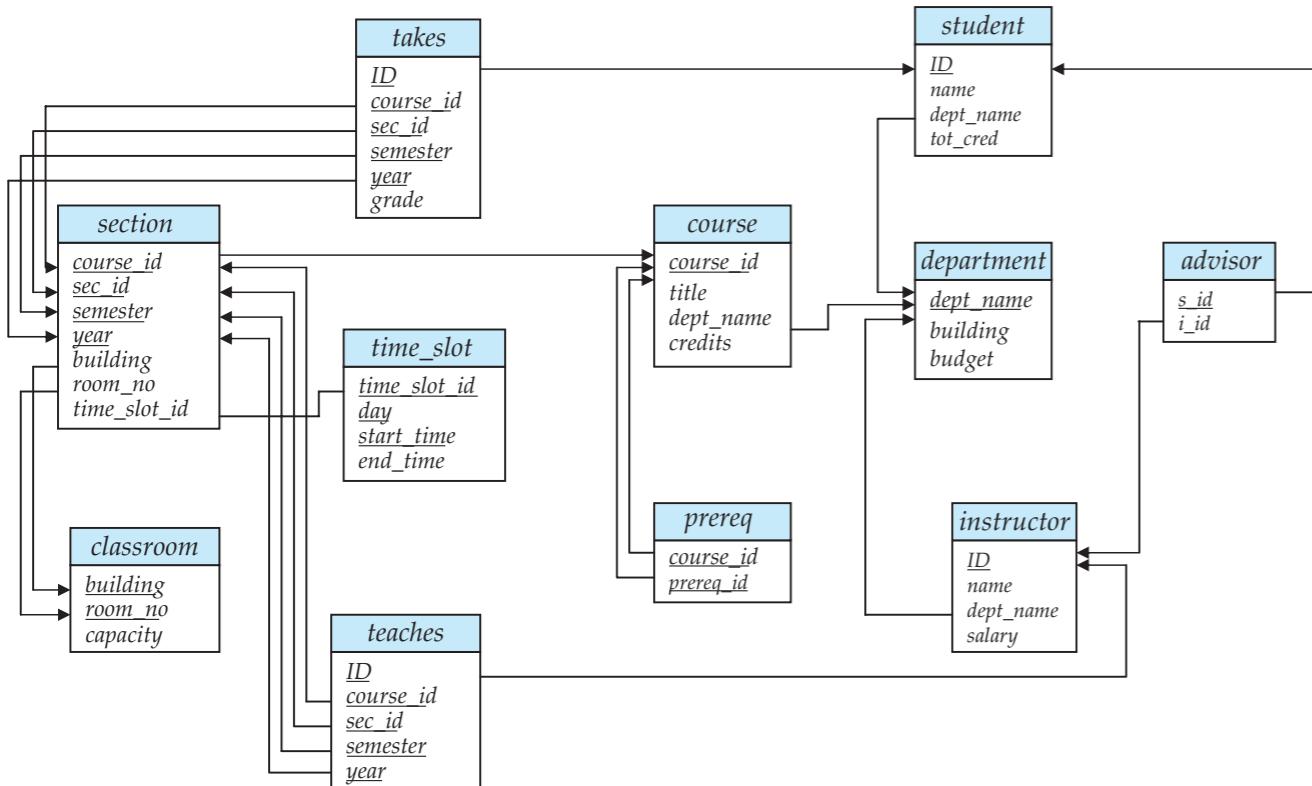
- The basic SQL query blocks map to RA as follows:

1. Form the (cascaded binary) **product** of the relations in the **FROM** clause
2. To this result, apply a **selection** operator whose predicate is the one in the **WHERE** clause
3. To this result, apply a **projection** whose attribute list is that in the **SELECT** clause
4. If **DISTINCT** is requested, to this result, apply a **duplicate removal** operator



SQL: The SELECT Clause

78



`SELECT name
FROM instructor;`

just this column!

`SELECT DISTINCT name
FROM instructor;`

remove duplicates!

`SELECT ALL name
FROM instructor;`

keep duplicates!

`SELECT *
FROM instructor;`

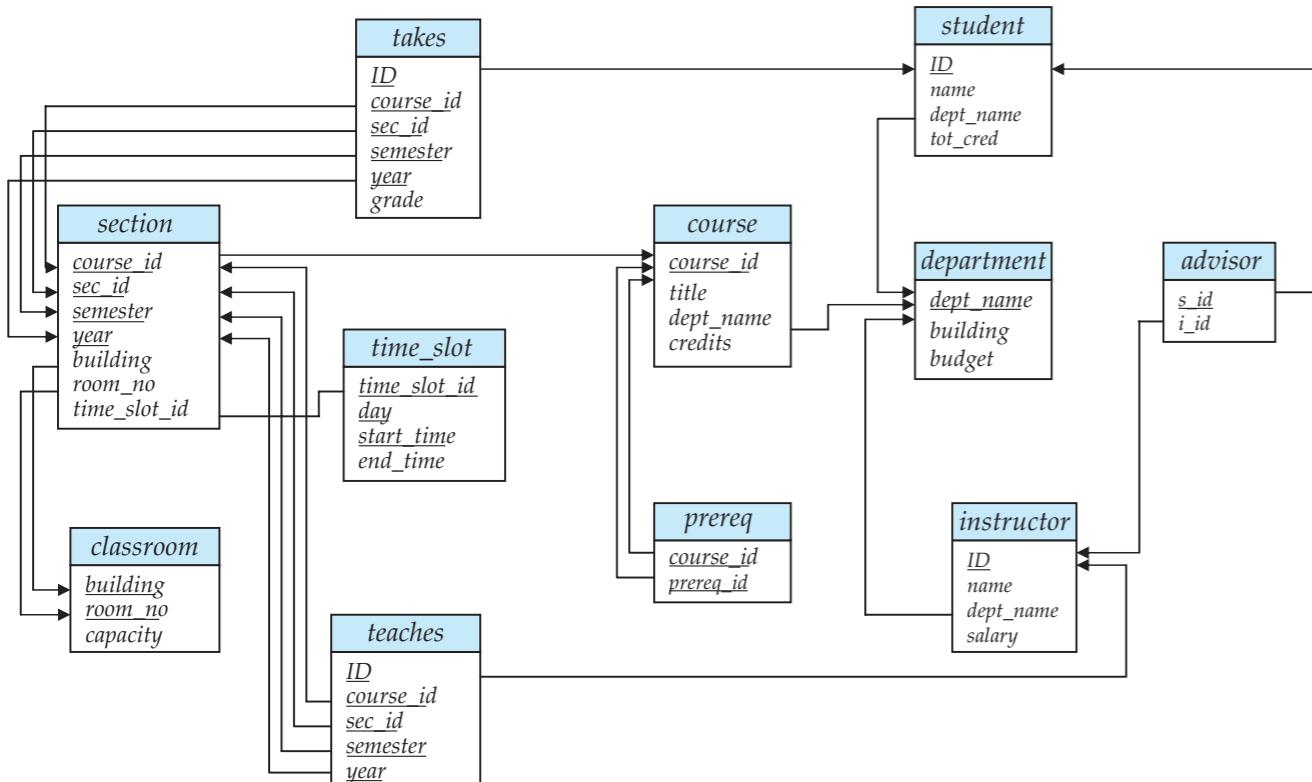
all columns!

`SELECT ID, name, salary/12
FROM instructor;`

derive this column!

SQL: The WHERE and FROM Clauses

79



SELECT name
FROM instructor
WHERE dept_name = 'Comp. Sci.'
AND salary > 80000;

complex predicates!

SELECT *
FROM instructor, teaches;

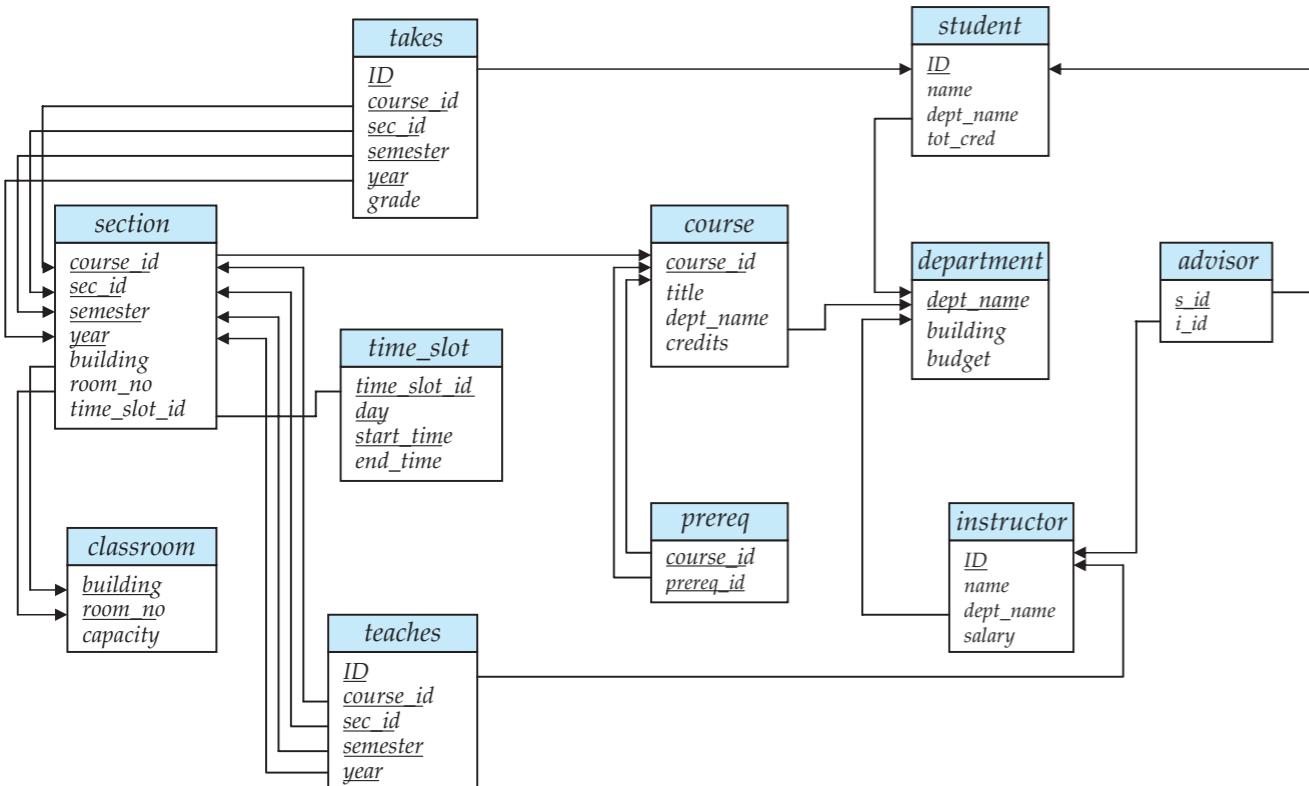
more than one? product!

SELECT name
FROM instructor, department
WHERE instructor.dept_name =
department.dept_name
AND salary > budget;

unless there is a predicate: join!

SQL: Set Operations

X



(SELECT course_id
FROM section
WHERE sem = 'Fall' and year = 2009)
UNION ALL
(SELECT course_id
FROM section
WHERE sem = 'Spring' and year = 2010);

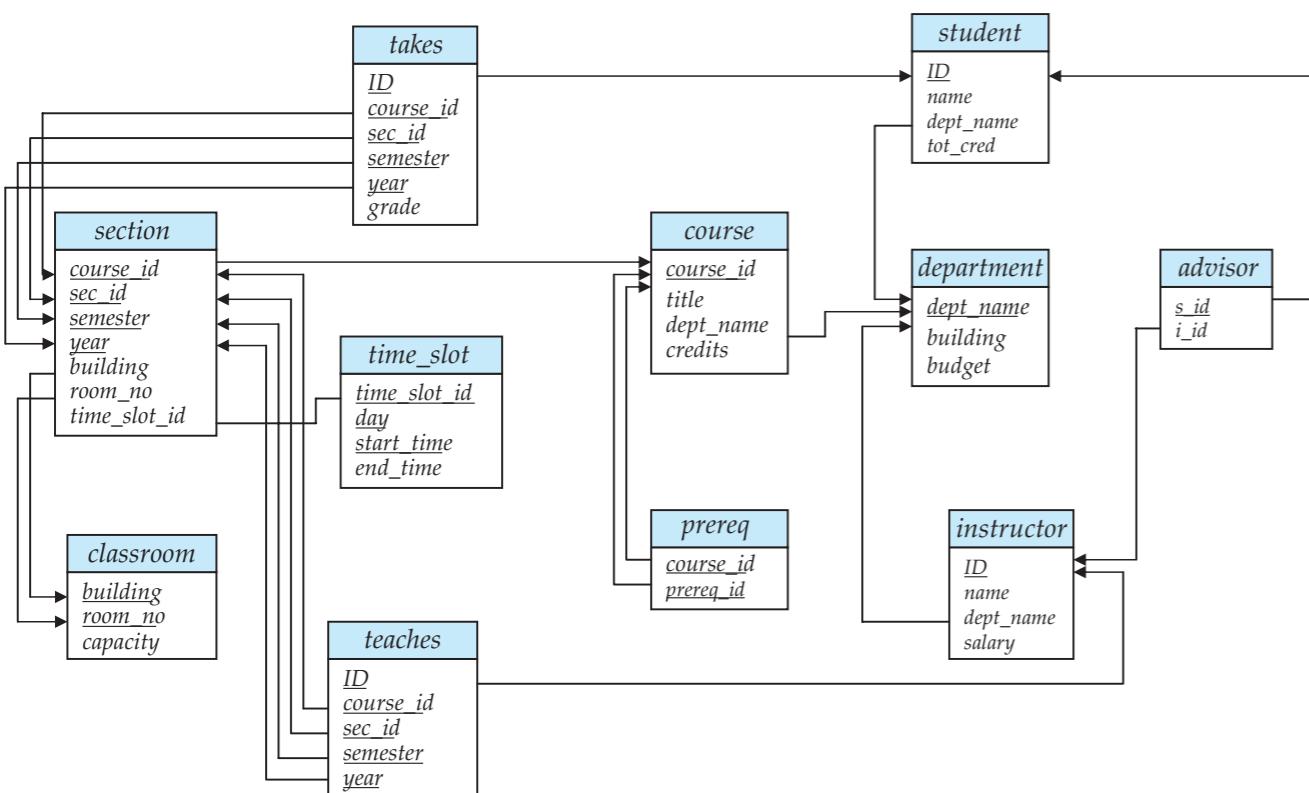
(SELECT course_id
FROM section
WHERE sem = 'Fall' and year = 2009)
EXCEPT ALL
(SELECT course_id
FROM section
WHERE sem = 'Spring' and year = 2010);

(SELECT course_id
FROM section
WHERE sem = 'Fall' and year = 2009)
INTERSECT ALL
(SELECT course_id
FROM section
WHERE sem = 'Spring' and year = 2010);

SQL: Joins

commonest
form of join!

X



```
SELECT instructor.name, teaches.course_id  
FROM instructor, teaches  
WHERE instructor.ID = teaches.ID;
```

SELECT name, course_id
FROM instructor JOIN teaches
USING (ID);

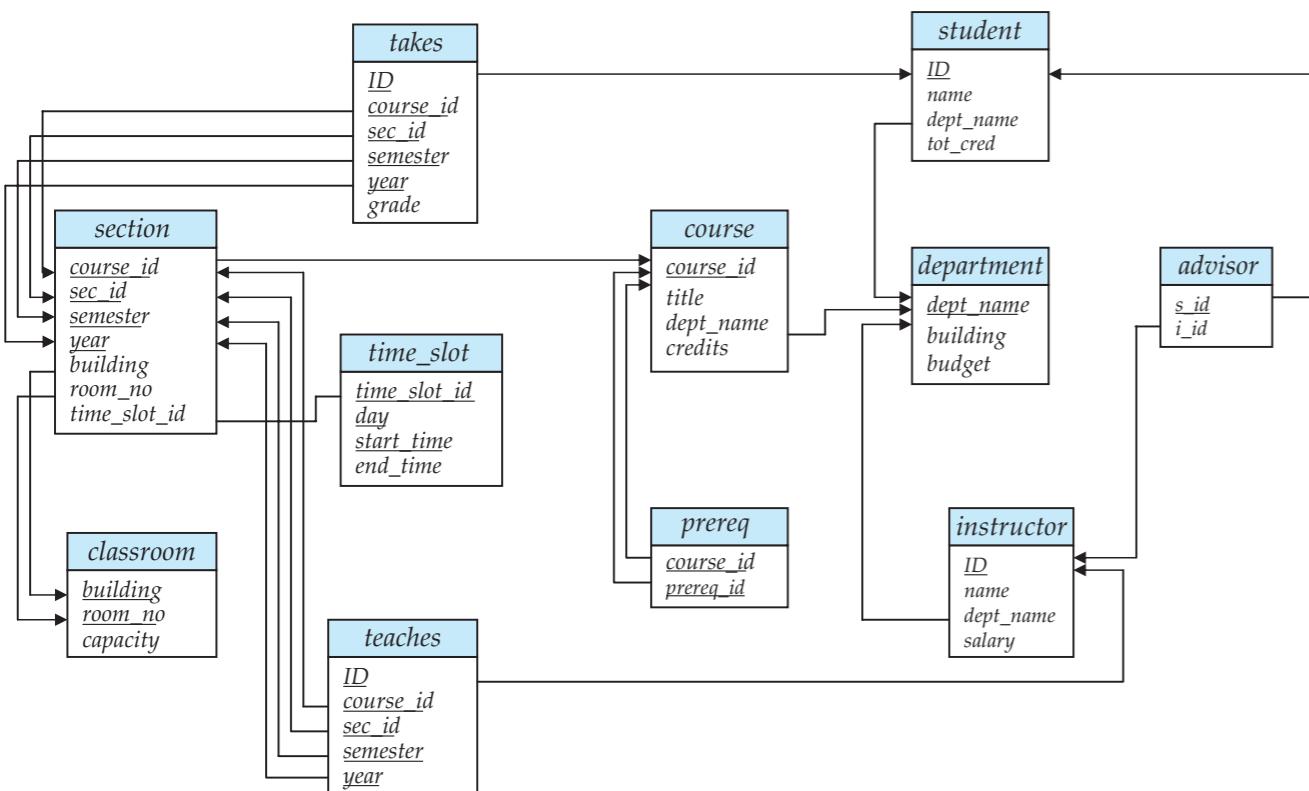
SELECT name, course_id
FROM instructor NATURAL JOIN teaches;

alternatively!

if possible!

SQL: Renaming

X



```
SELECT ID, name, salary/12 AS monthly_salary  
FROM instructor;
```

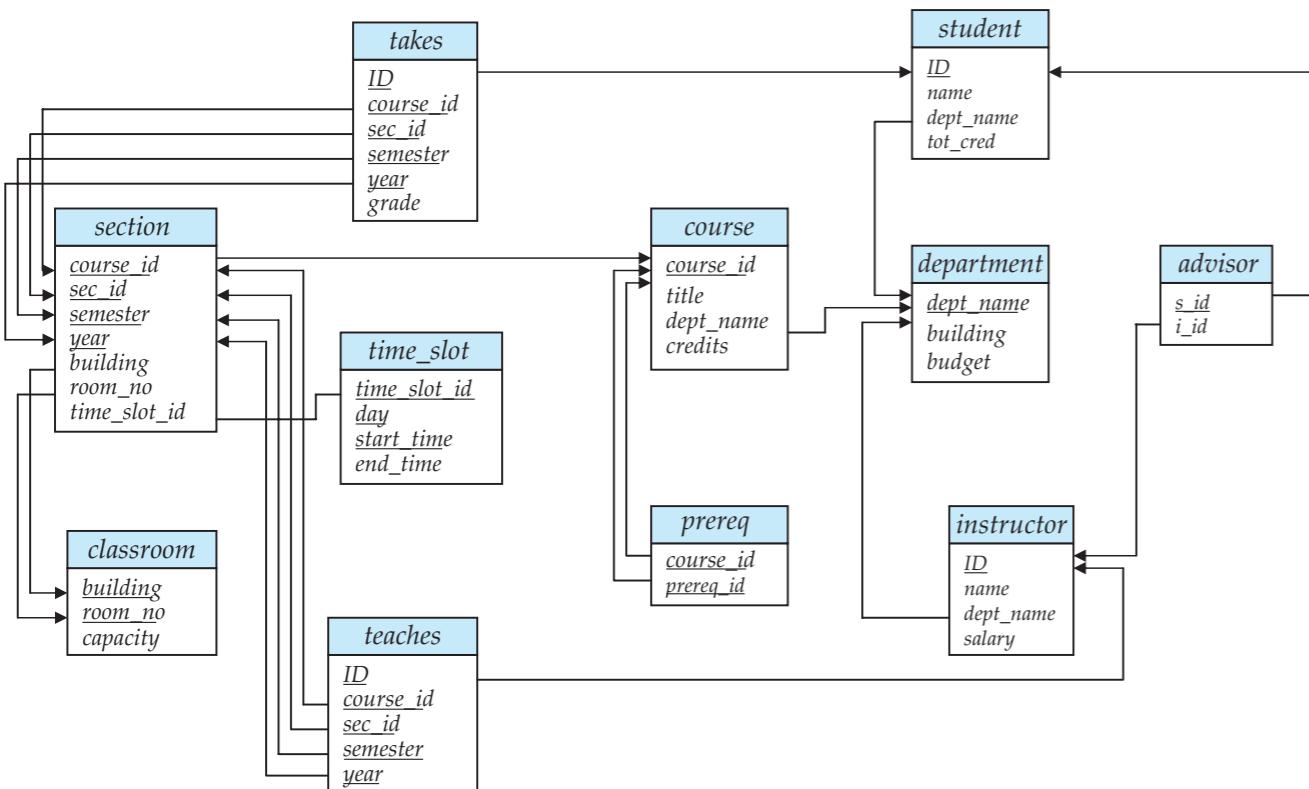
optional but recommended!

```
SELECT DISTINCT t.name  
FROM instructor AS t, instructor AS s  
WHERE t.salary > s.salary  
AND s.dept_name = 'Comp. Sci.';
```

must omit in Oracle!

SQL: More Predicates

X



string pattern matching

```
SELECT name  
FROM instructor  
WHERE name LIKE '%DAR%';
```

interval comparison

```
SELECT name  
FROM instructor  
WHERE salary BETWEEN 80000 AND 90000;
```

tuple

```
SELECT name, course_id  
FROM instructor i, teaches t  
WHERE (i.ID, t.dept_name) = (t.ID, 'Biology');
```

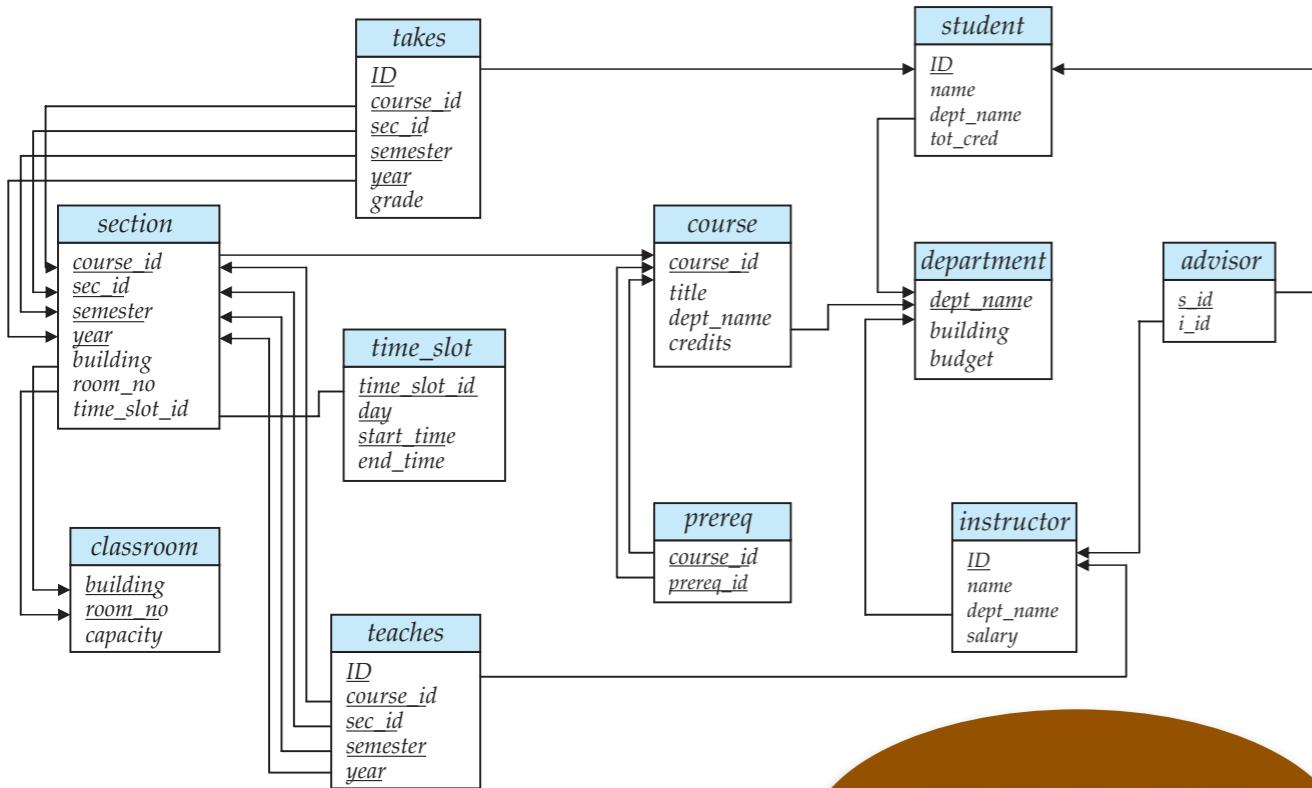
testing for NULL

```
SELECT name  
FROM instructor  
WHERE salary IS NULL;
```

SQL: Ordering and Aggregation

simplest
form!

80



`SELECT AVG(salary)
FROM instructor;`

remove
duplicates!

`SELECT COUNT(DISTINCT name)
FROM instructor;`

count all!

`SELECT COUNT(*)
FROM instructor;`

average per
department

`SELECT dept_name, AVG(salary)
FROM instructor
GROUP BY dept_name;`

but not all!

`SELECT dept_name, AVG(salary)
FROM instructor
GROUP BY dept_name
HAVING AVG(salary) > 42000;`

In The Next Lecture

x

- We'll begin our exploration of conceptual modelling.
- We'll see that conceptual modelling is aimed at providing an intuitive approach to identifying information needs.
- We'll observe that a good conceptual model facilitates the generation of well-formed logical models.
- In particular, we'll learn about the formalism for conceptual modelling known as Entity-Relationship (ER) Modelling.

L6

Designing Databases: Mapping a Conceptual into a Logical Design

Fundamentals of Databases
Alvaro A A Fernandes, SCS, UoM

Readings

X

This lecture relies on your having read the assigned mandatory readings associated with this lecture:

Ramez Elmasri and Shamkant Navathe. Fundamentals of Database Systems. New International Edition. Pearson, 2013. pp. 287-308.

You can read the material online using this link:

lib.mylibrary.com/Open.aspx?id=527132

This is a shared resource. Don't hog it. When you're done, logout and close the browser to free it for others.

If you fail to study the material, you're likely to find the lecture harder to follow.

More information is available in the Learning Activities Manual.

In Previous Lectures

x

- We learned how data models lead to a distinction between schemas and instances that enables a logical view of the data.
- We learned about the relational approach to logical data modelling.
- We learned about the relational algebra and SQL, both its DDL and DML capabilities and its querying constructs.
- We learned of the practical benefits of performing conceptual modelling before logical design and why the EER approach serves well this purpose.

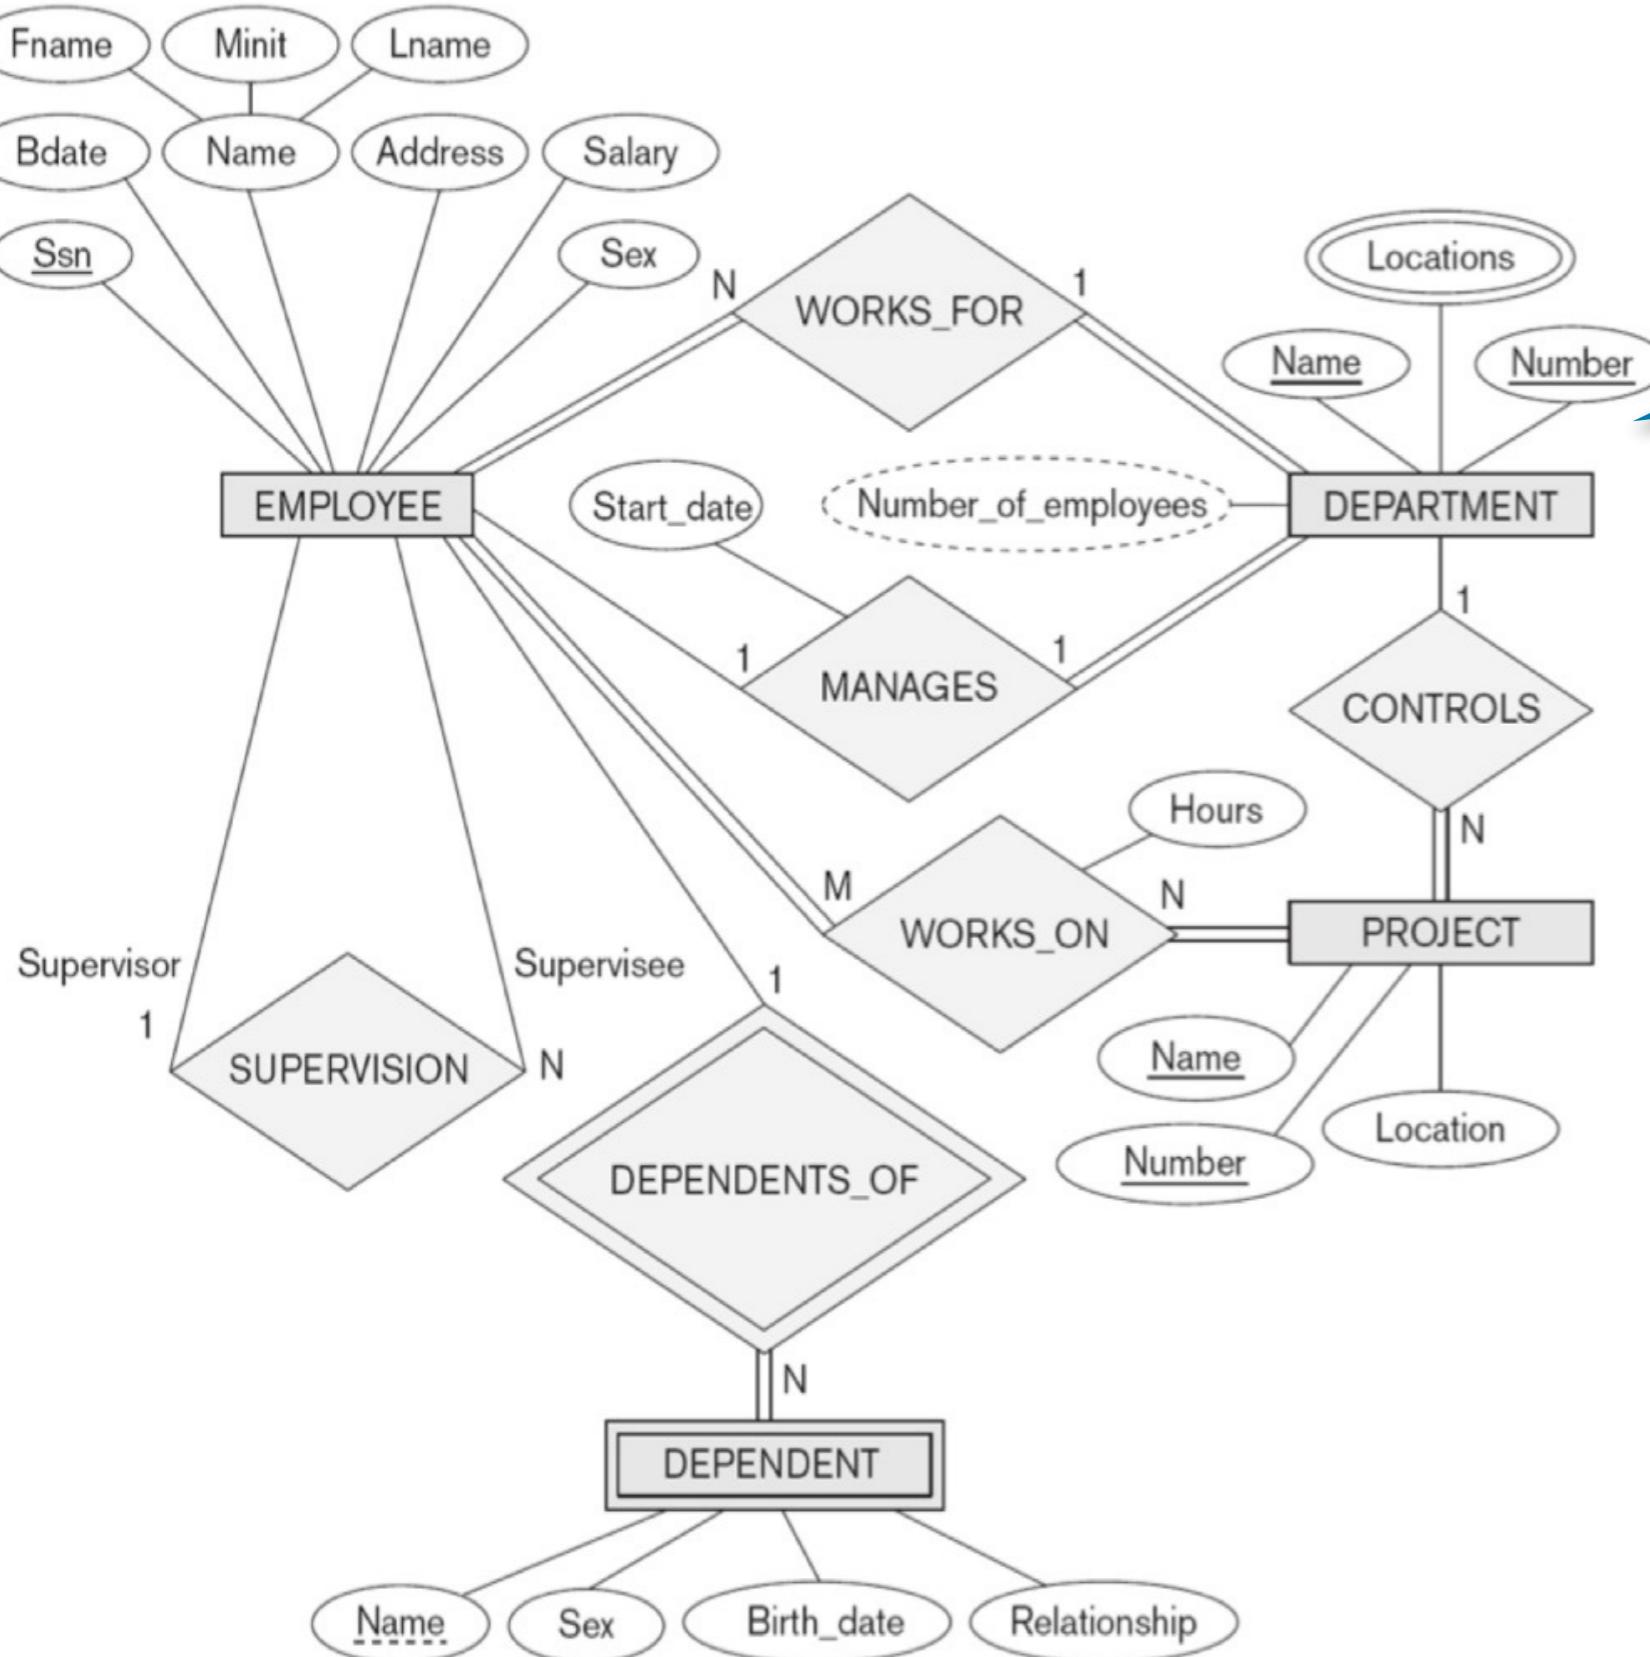
In This Lecture

x

- We'll learn how a conceptual model can be mapped into a logical model that is capable of being implemented in a relational DBMS.

An Example ER Diagram for a COMPANY DB

82



A
model like this one is the
starting point.

An Example Relational Schema for a COMPANY DB

83

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

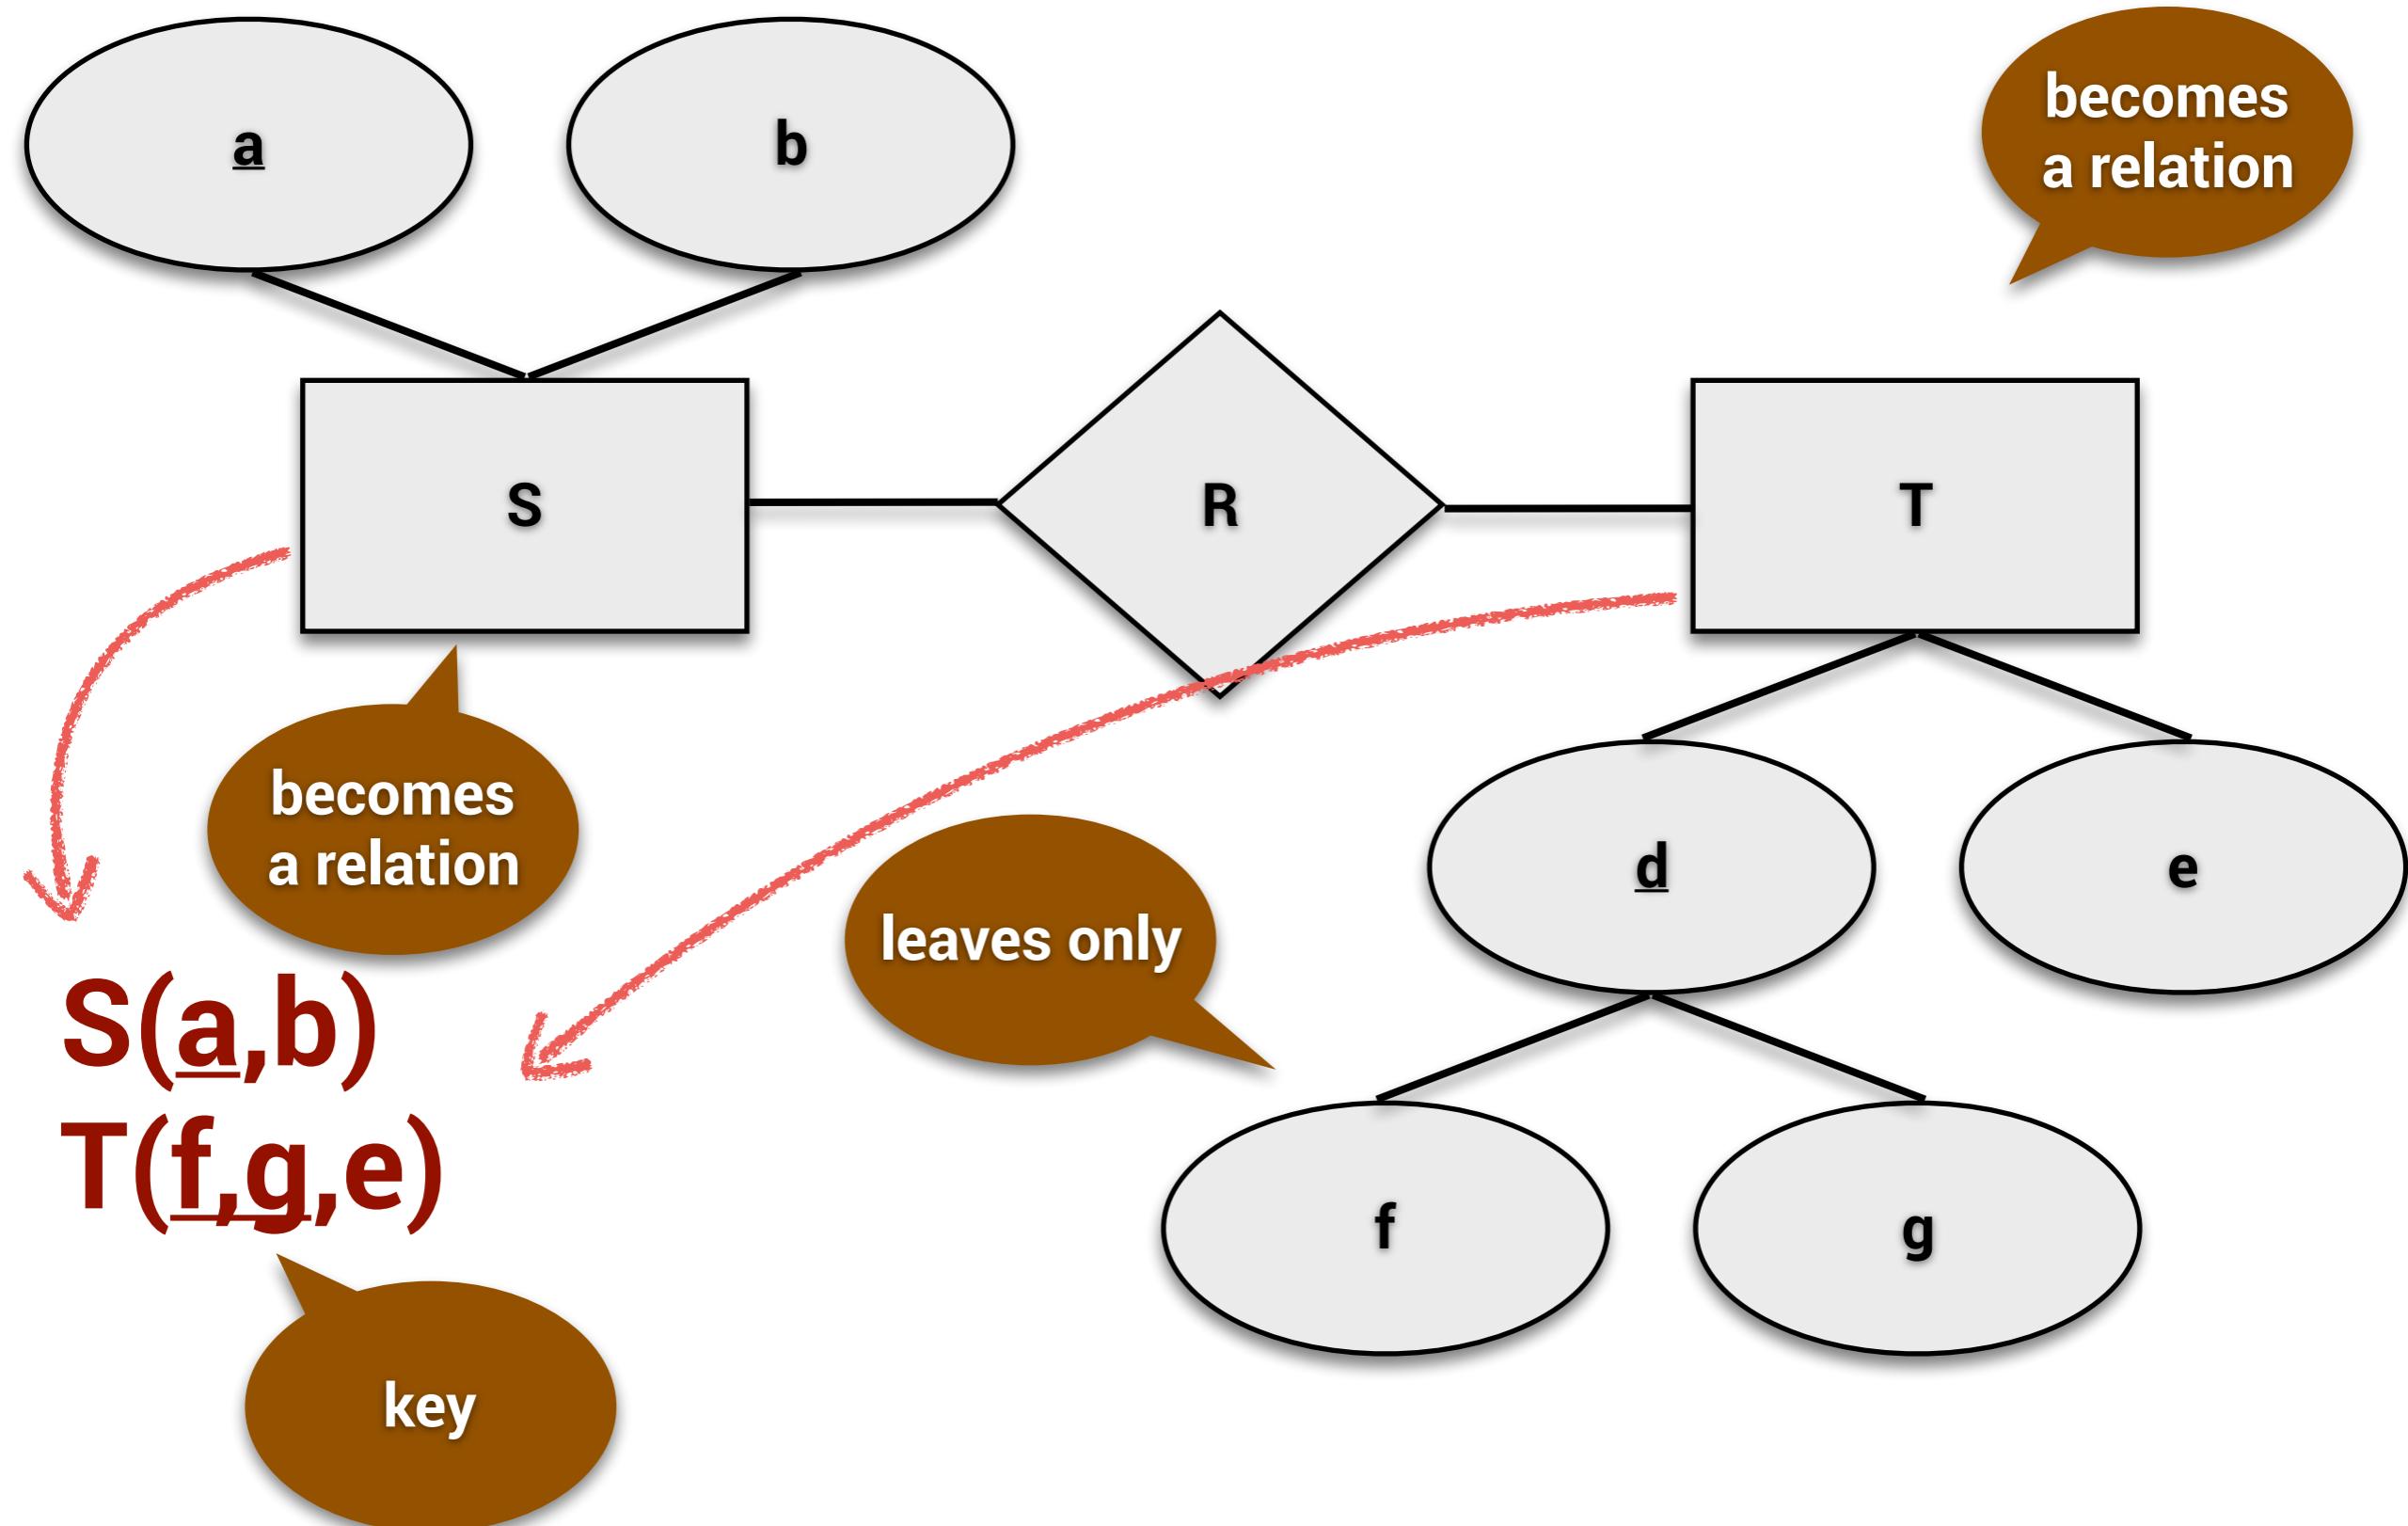
DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

A schema like this
one is the end point.

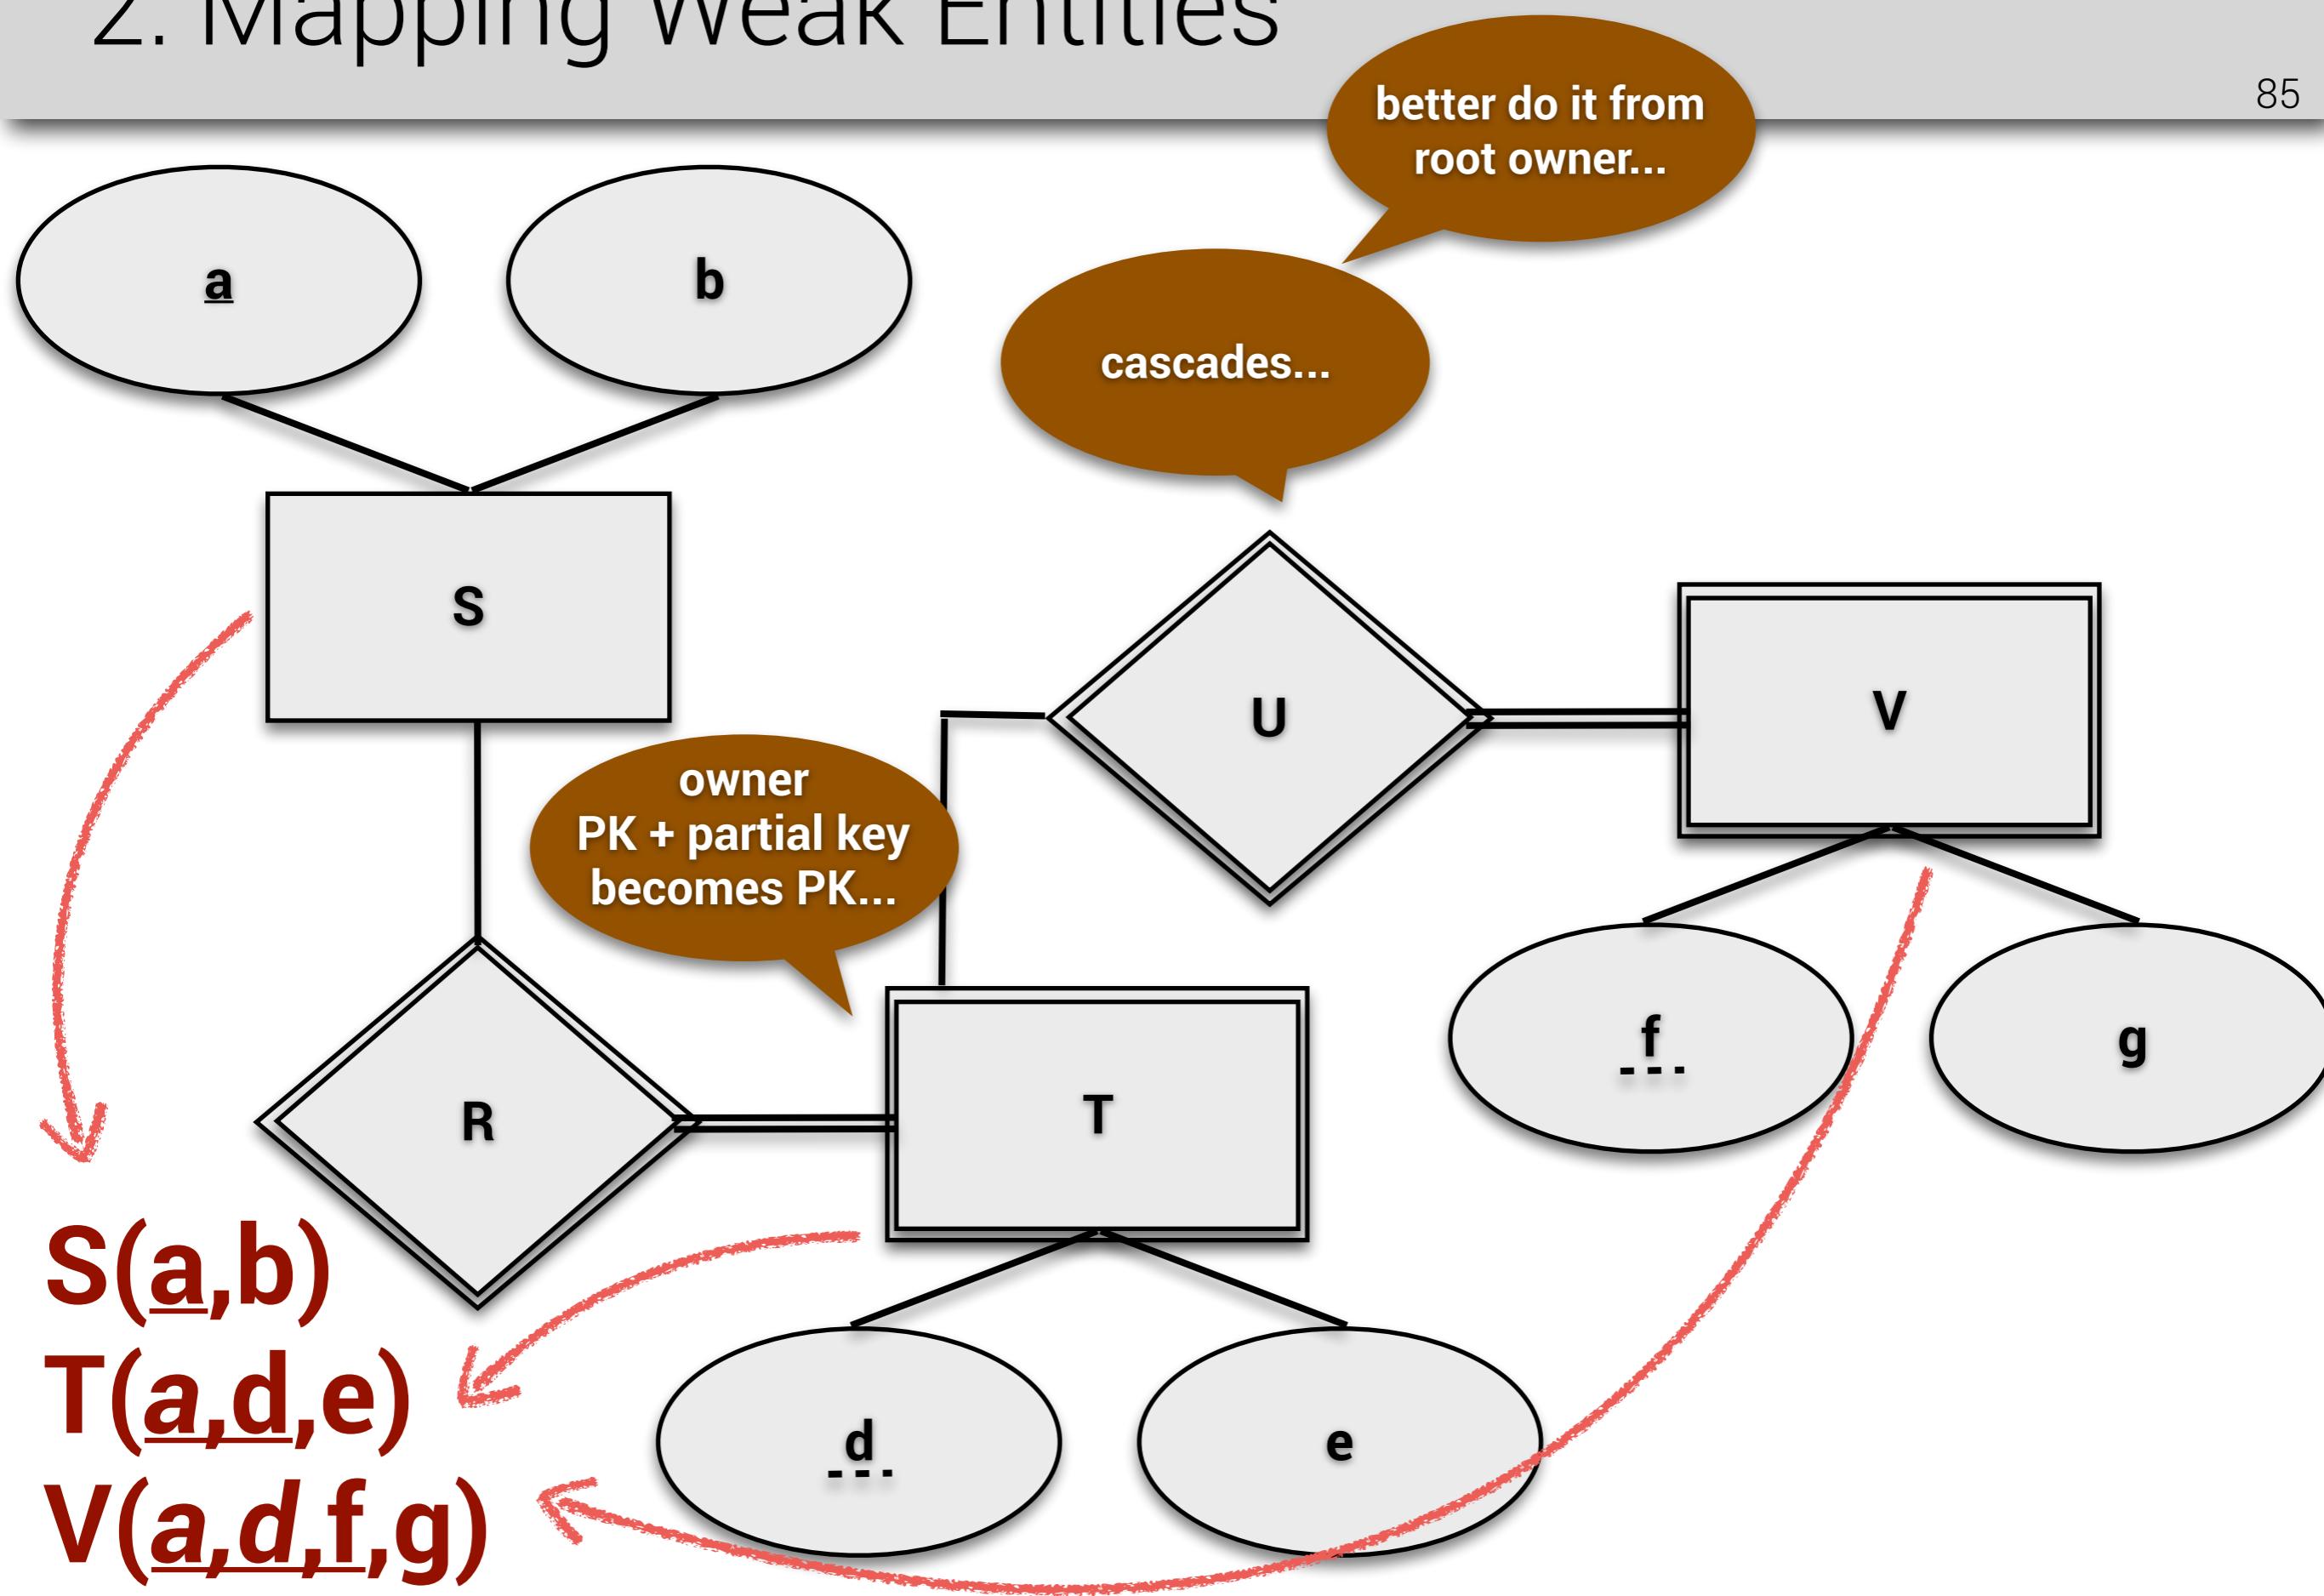
1: Mapping (Normal) Entities

84



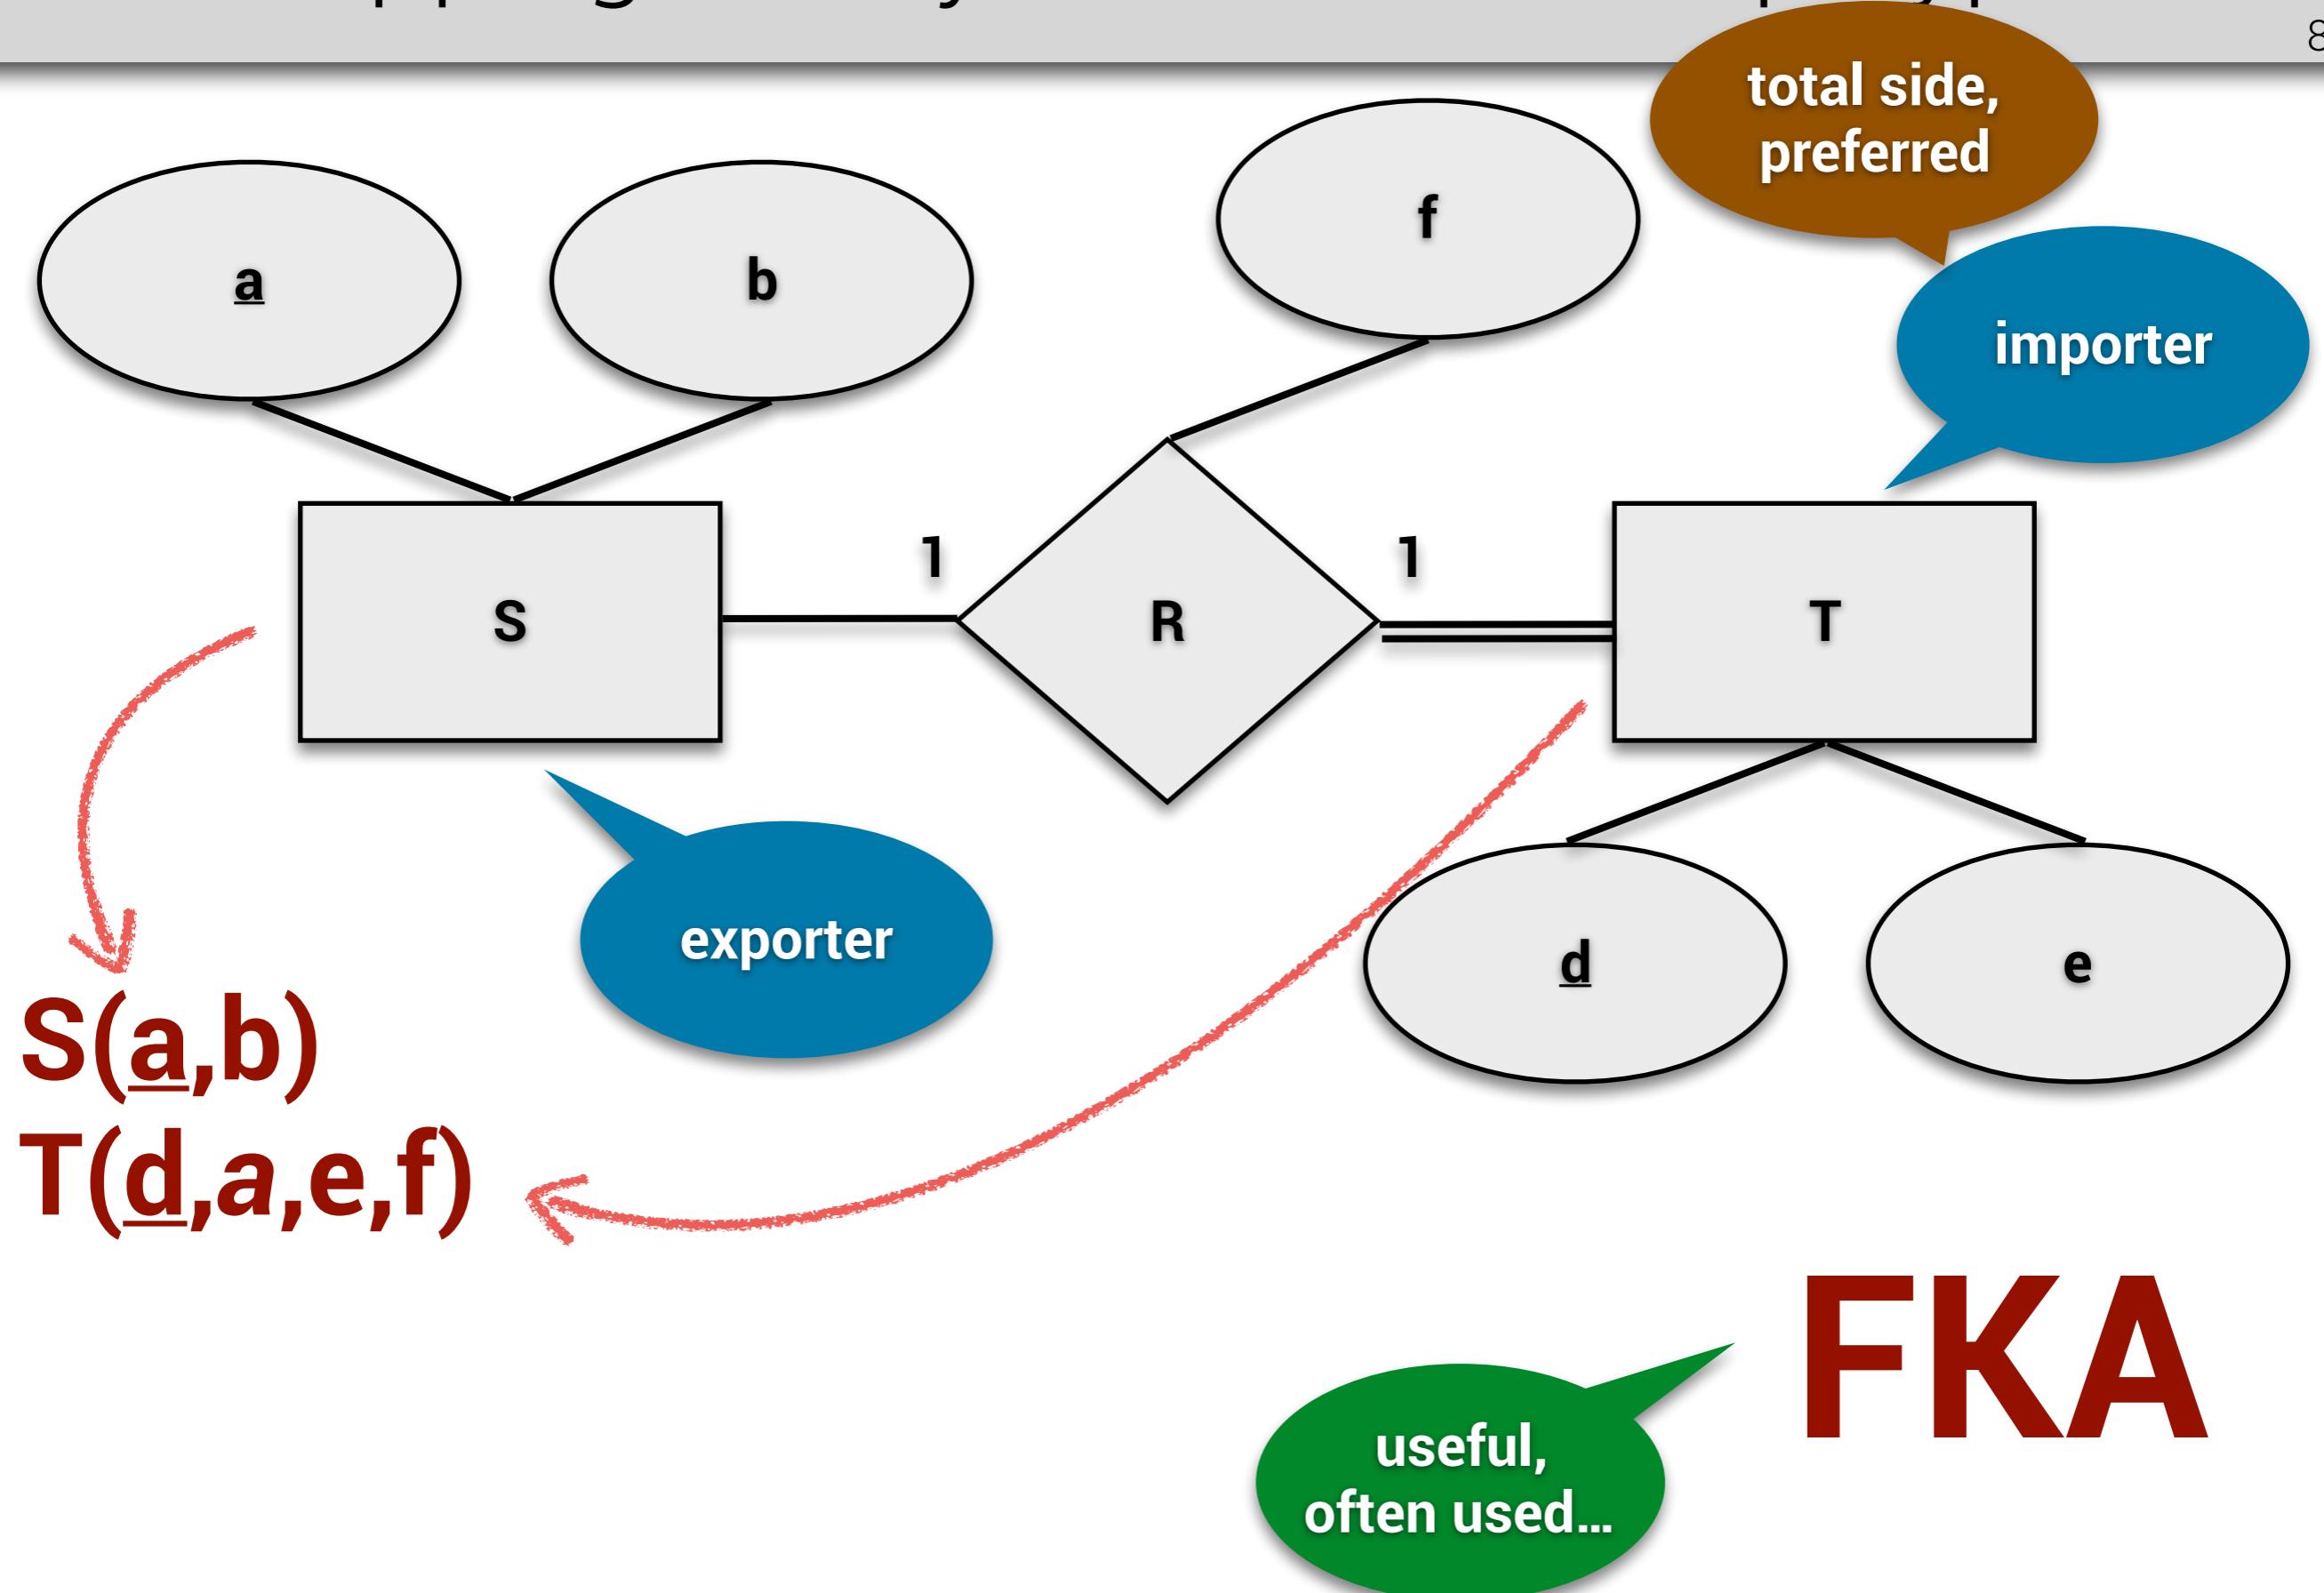
2: Mapping Weak Entities

85



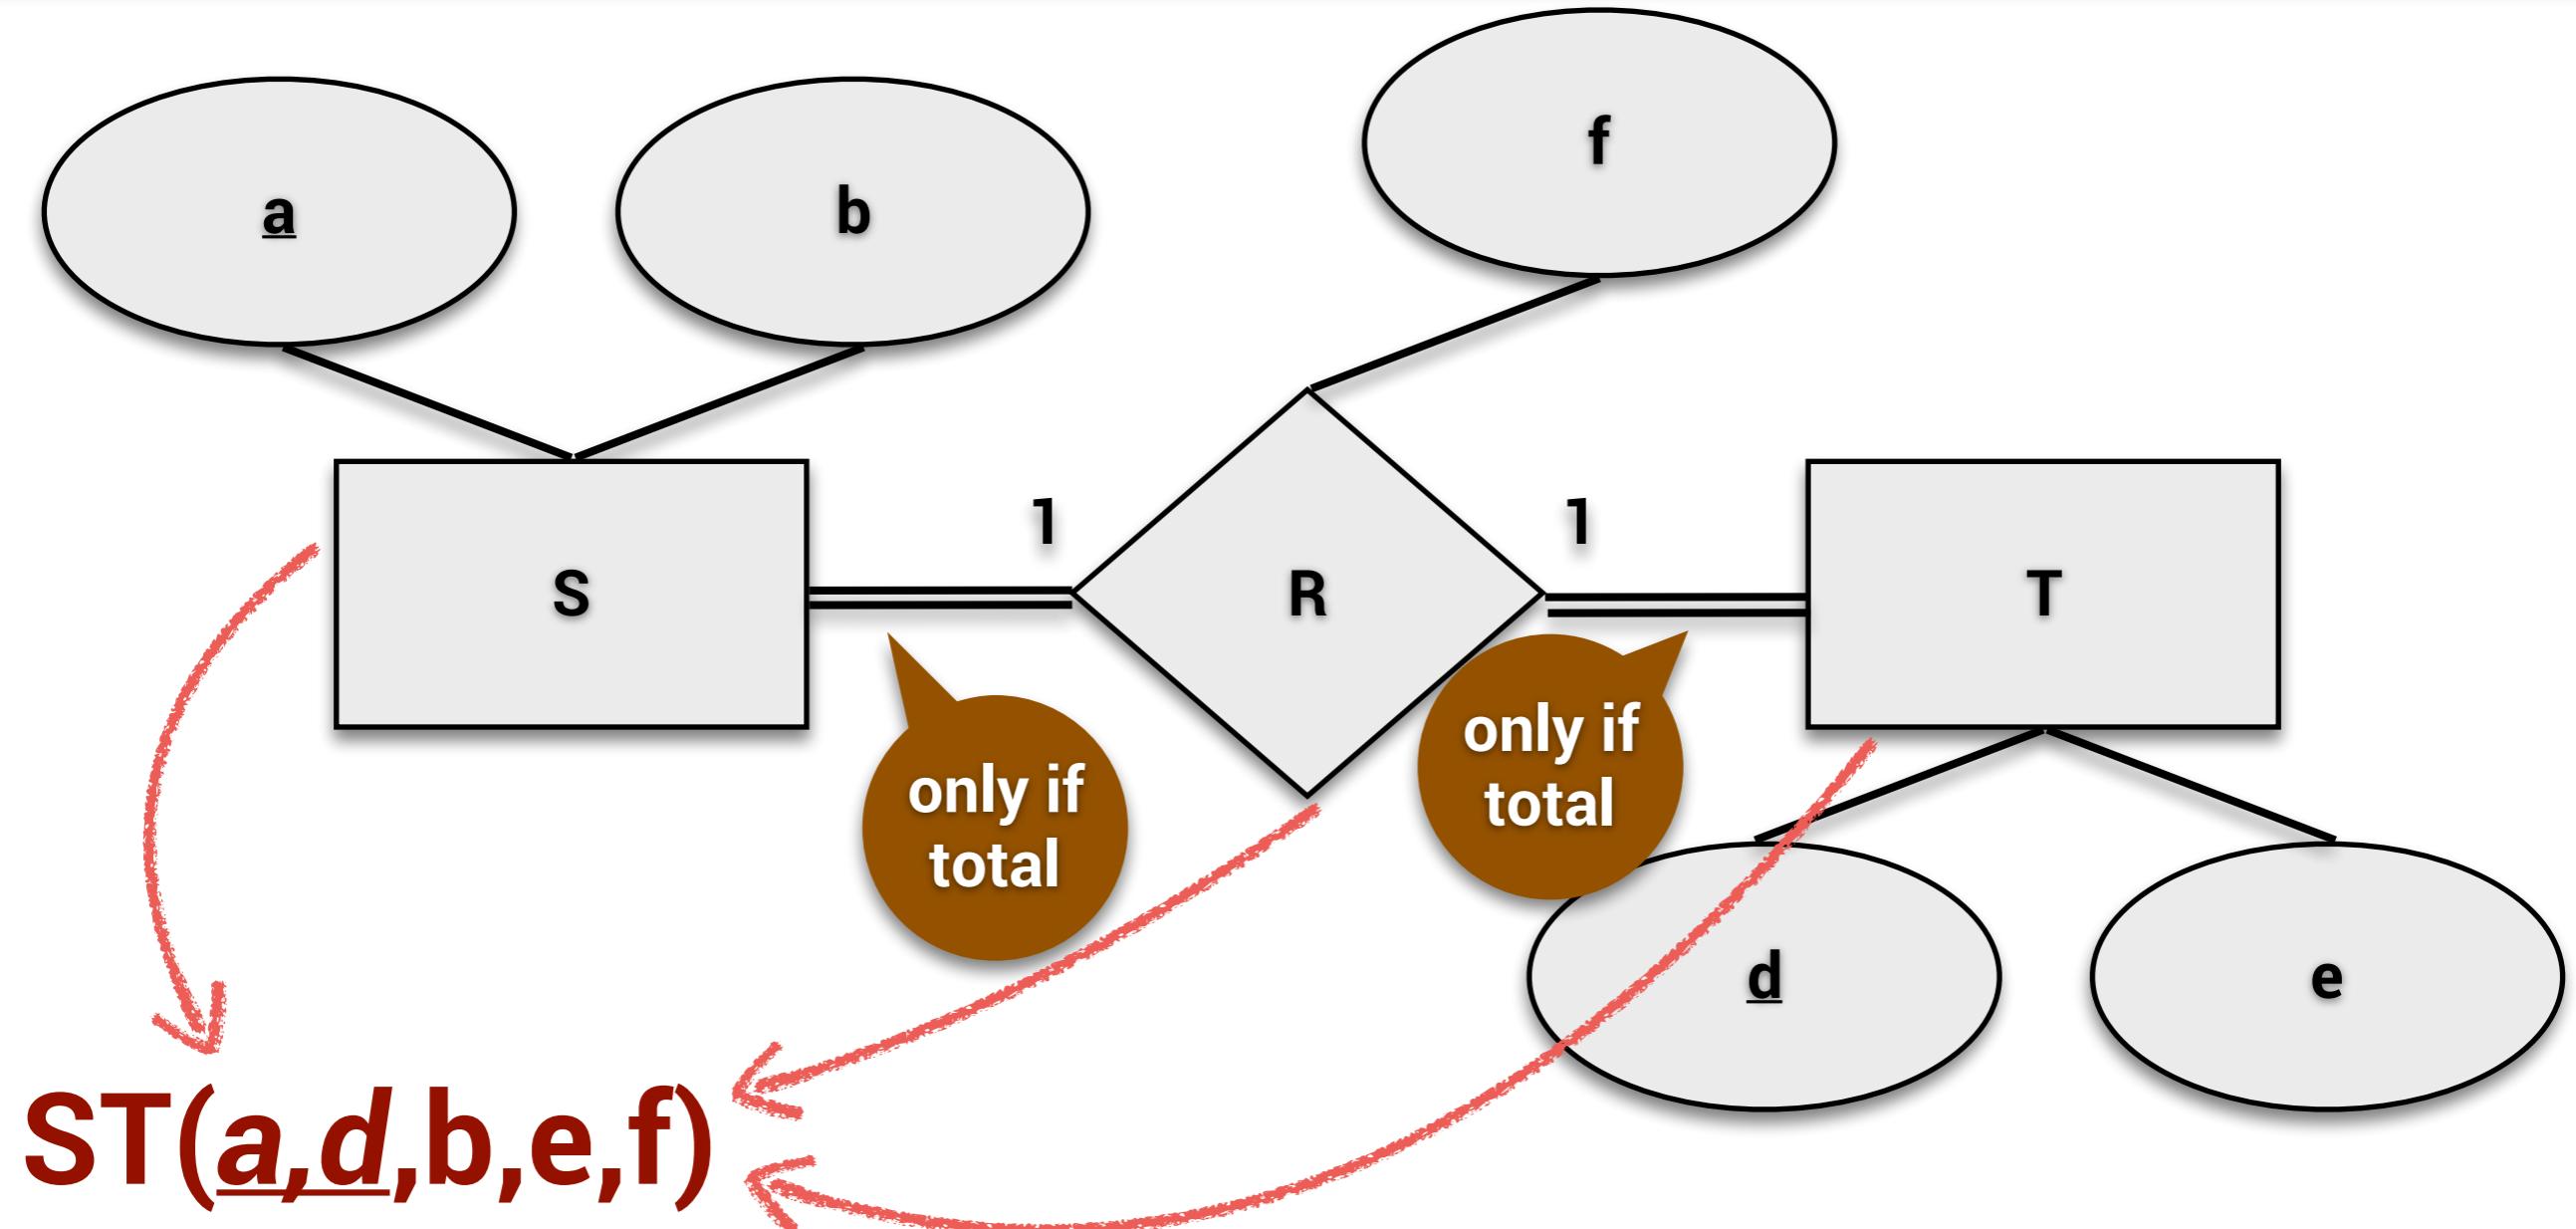
3.i: Mapping Binary 1:1 Relationship Types

86



3.ii: Mapping Binary 1:1 Relationship Types

87

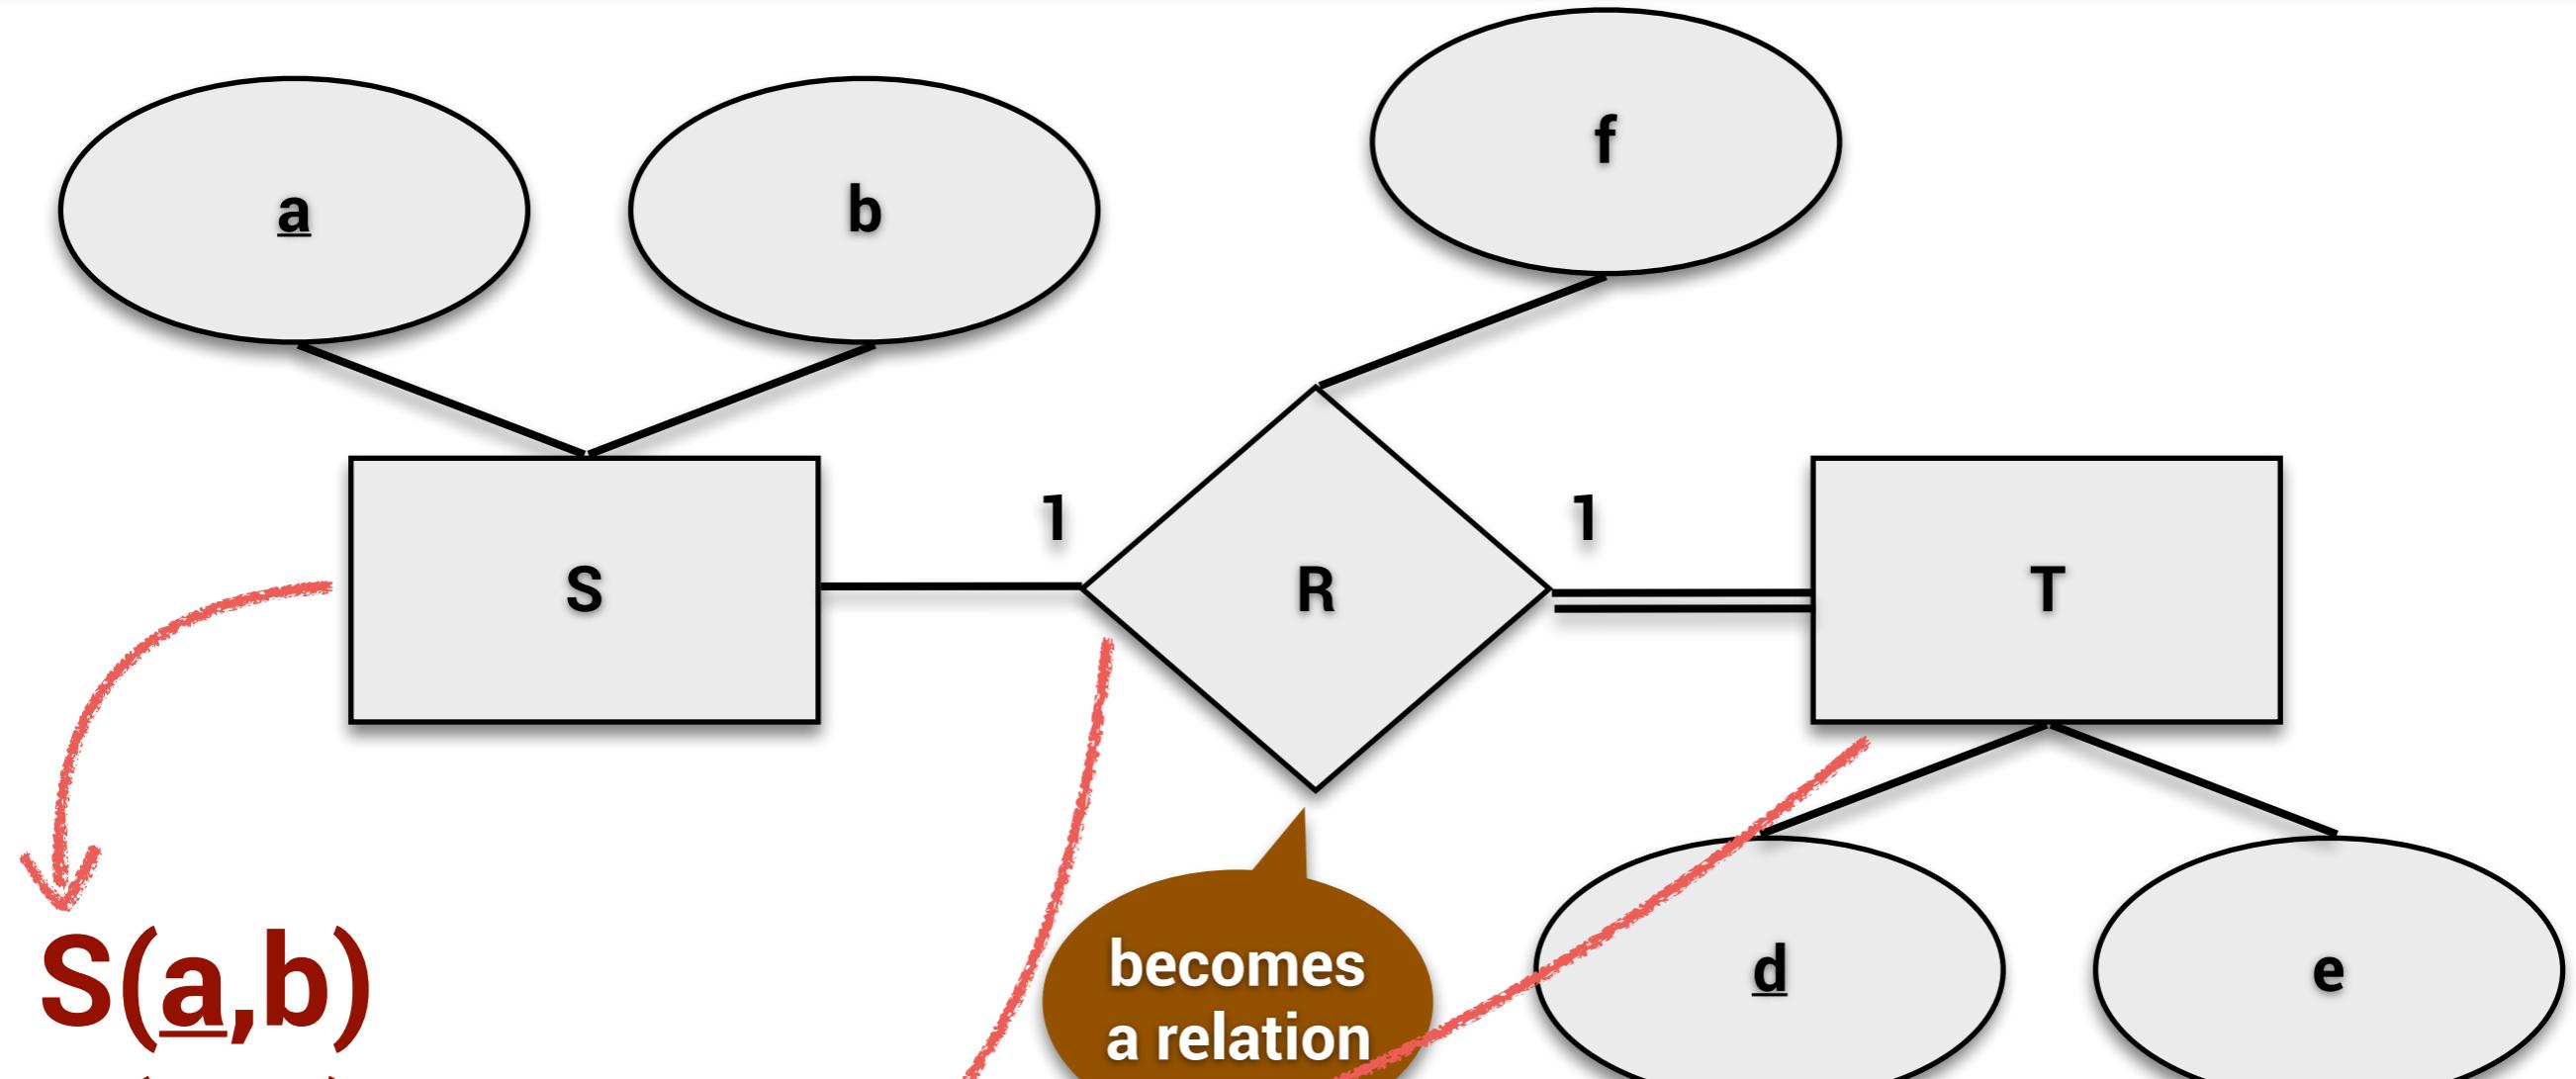


not often
used...

MRA

3.iii: Mapping Binary 1:1 Relationship Types

88



S(a,b)

T(d,e)

R(a,d,f)

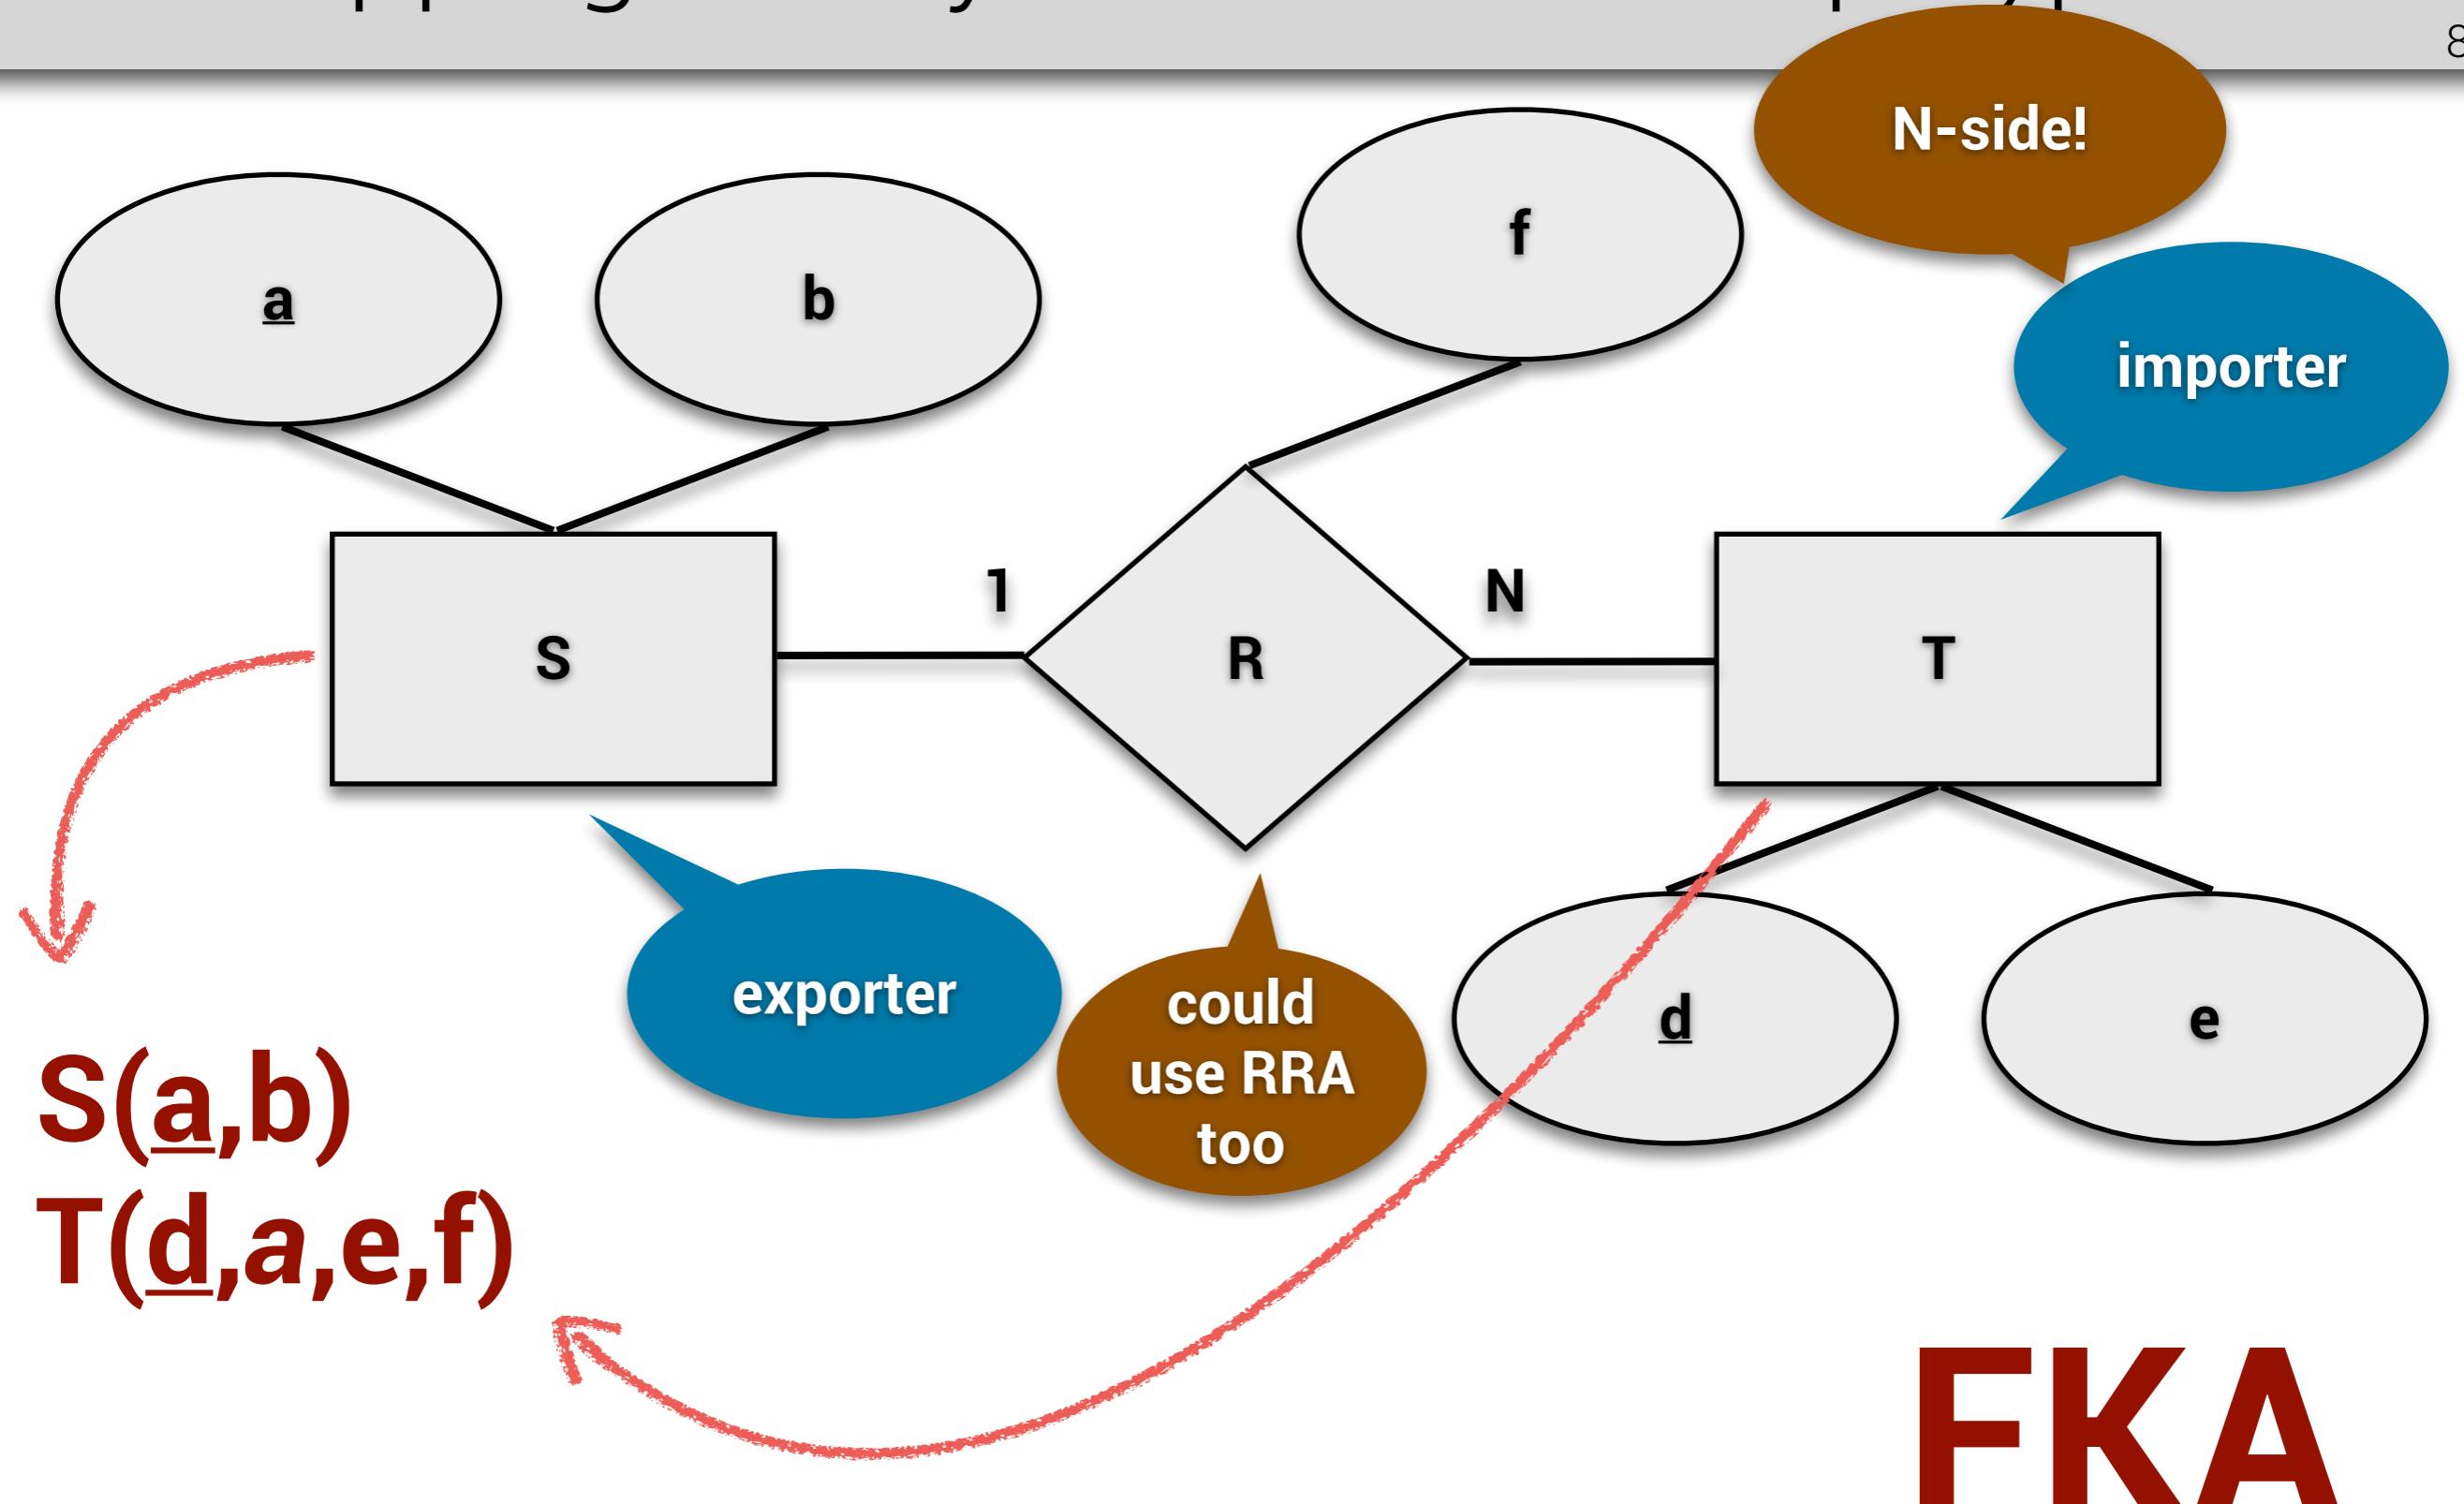
becomes
a relation

useful,
often used...

RRA

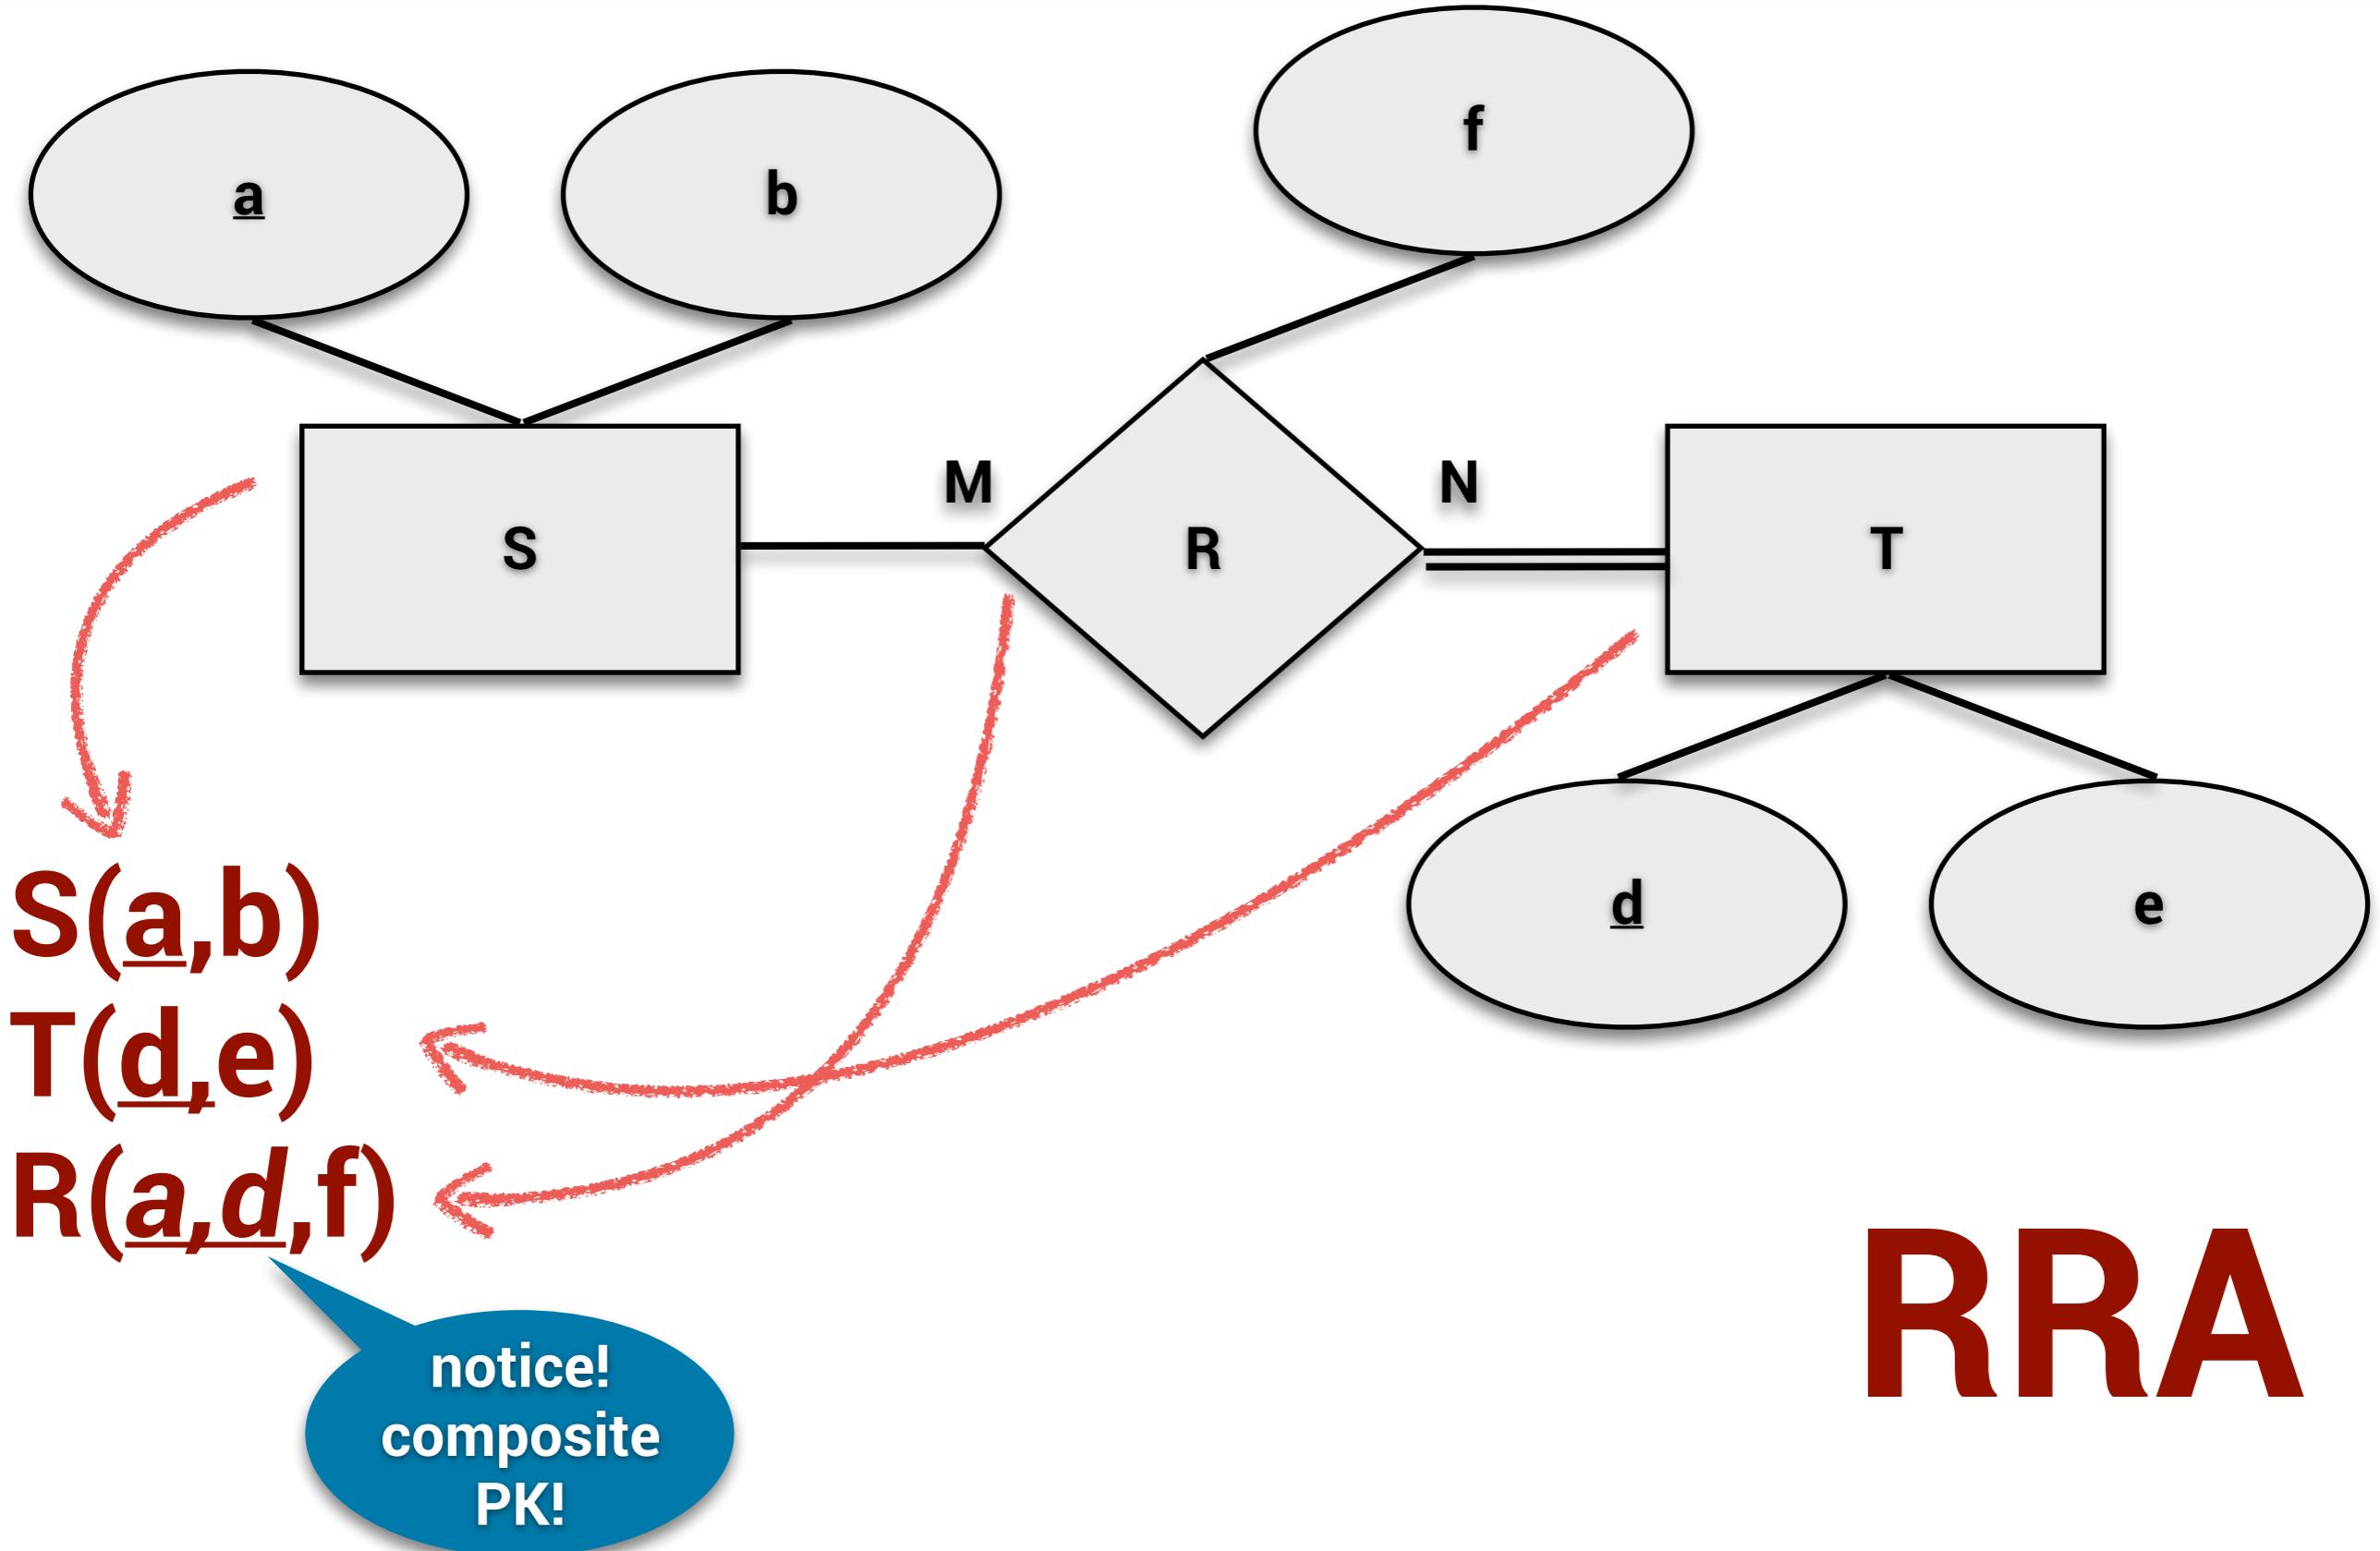
4: Mapping Binary 1:N Relationship Types

89



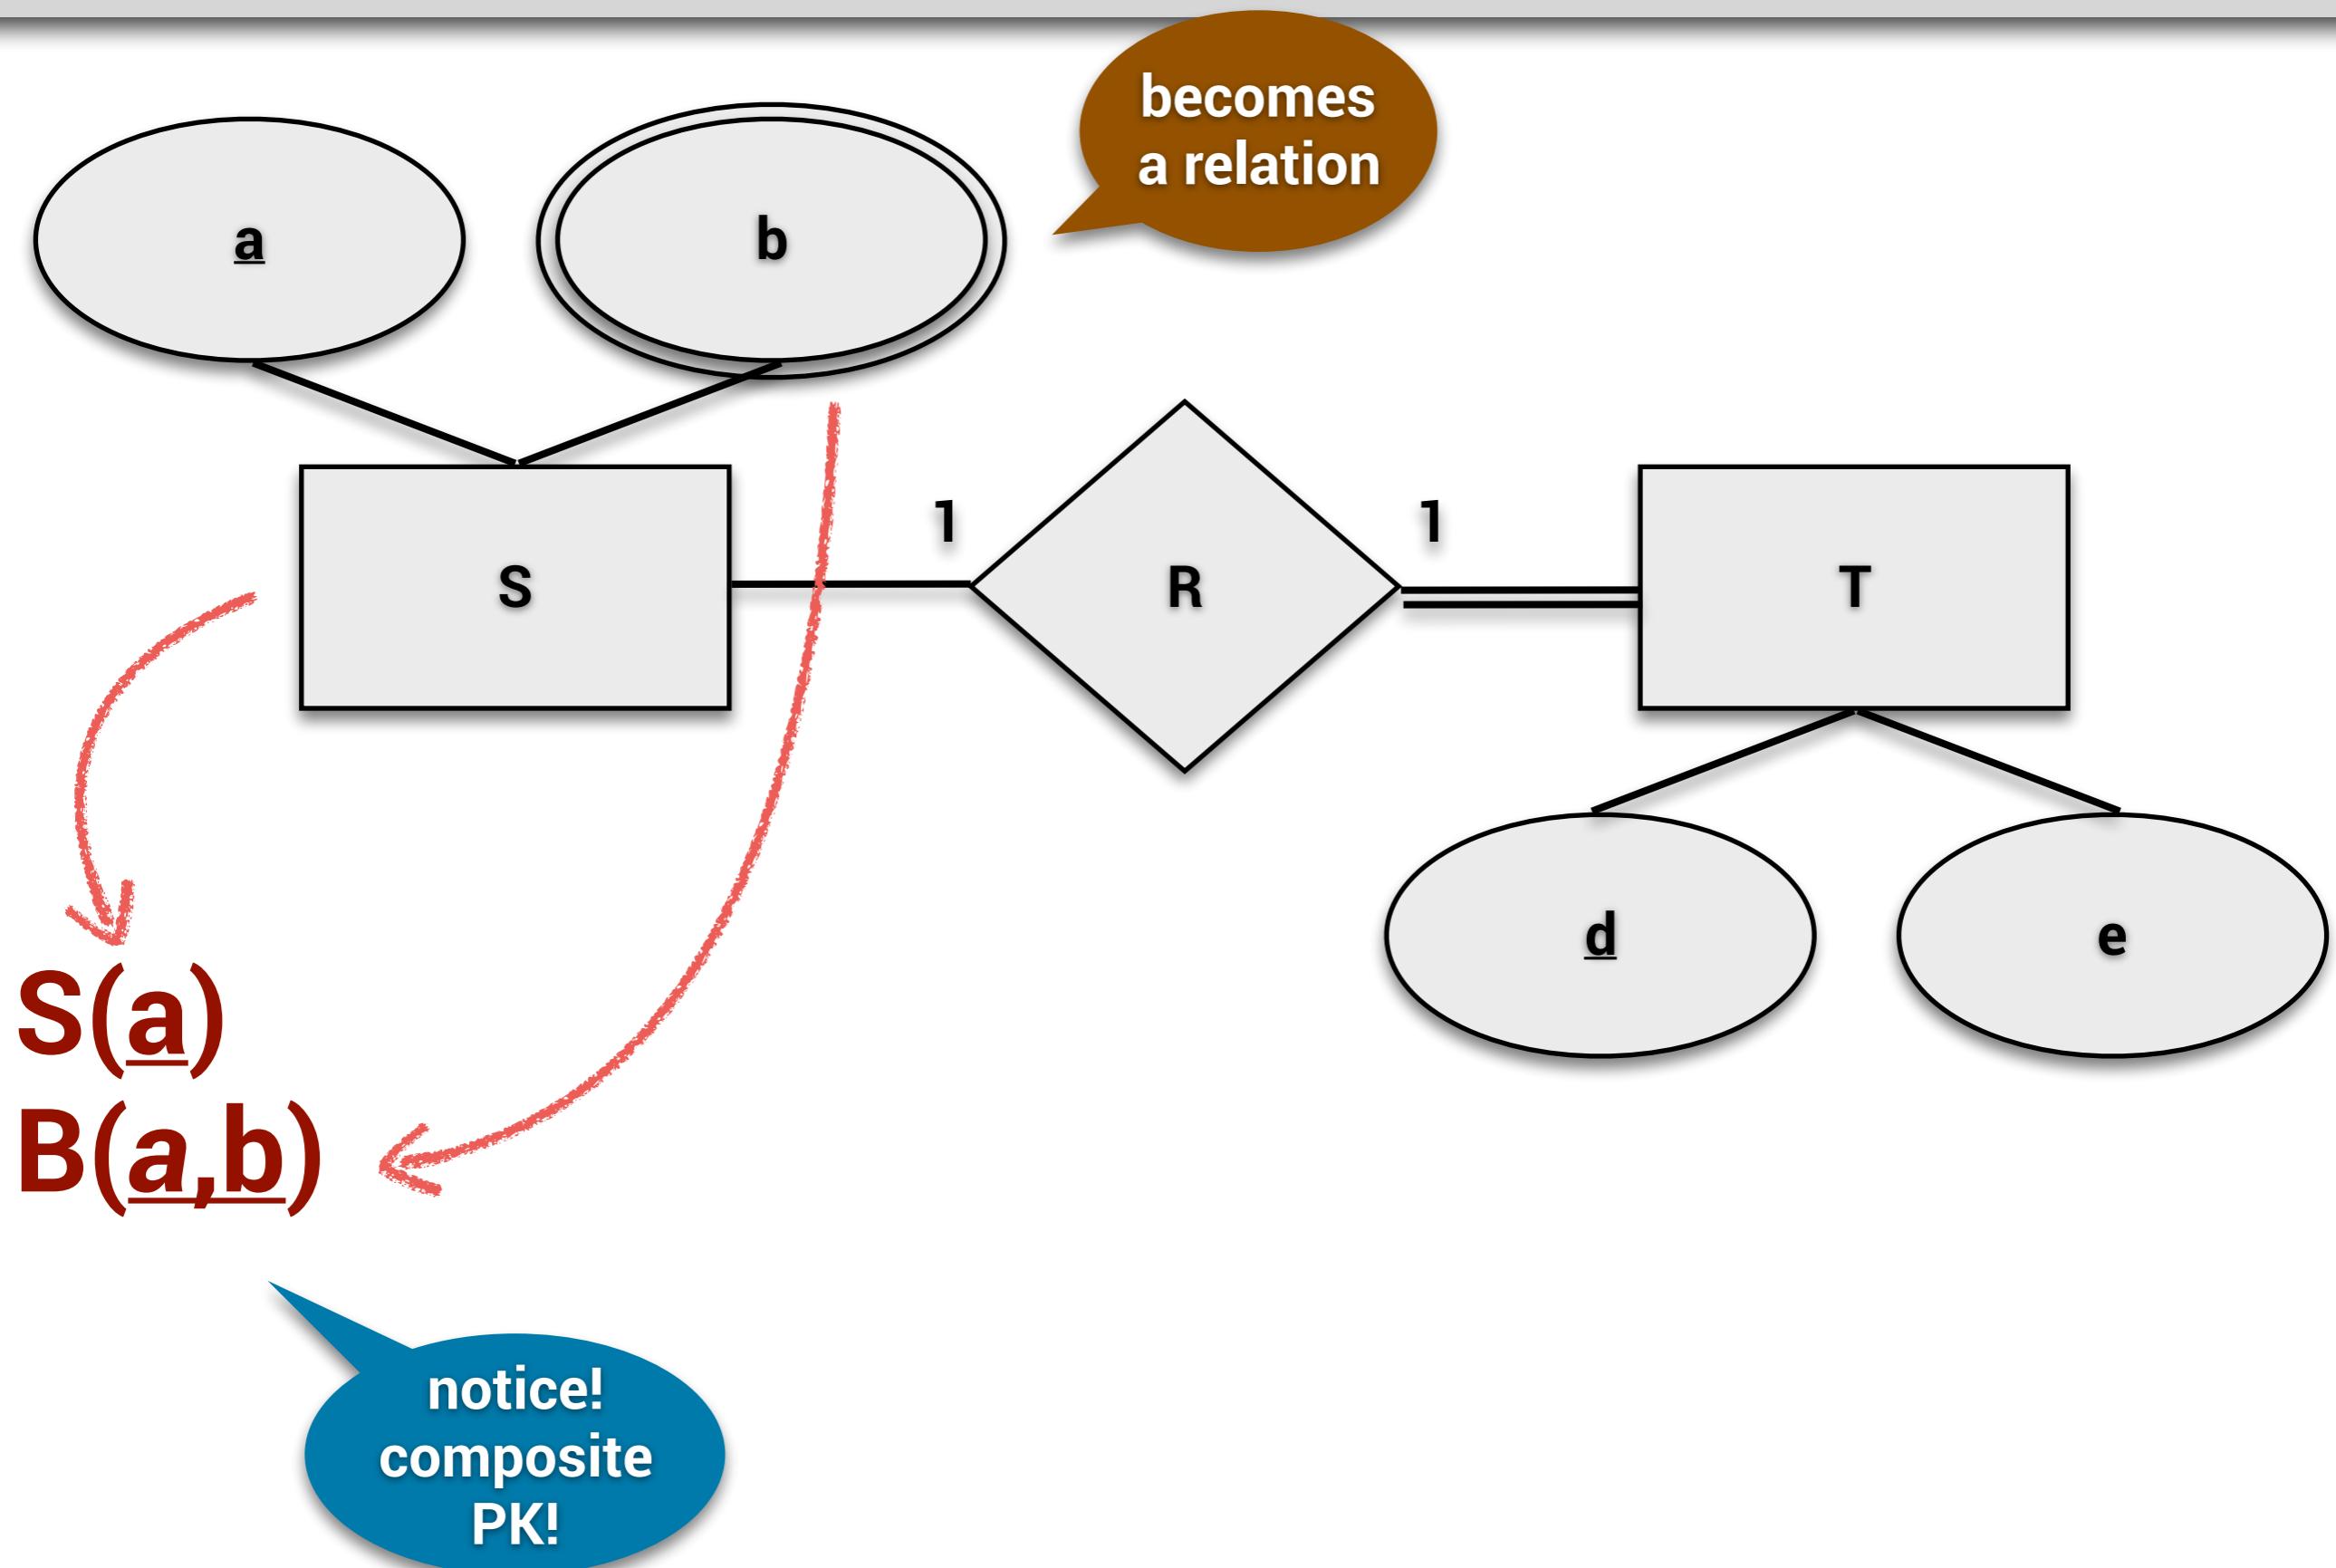
5: Mapping Binary N:M Relationship Types

90



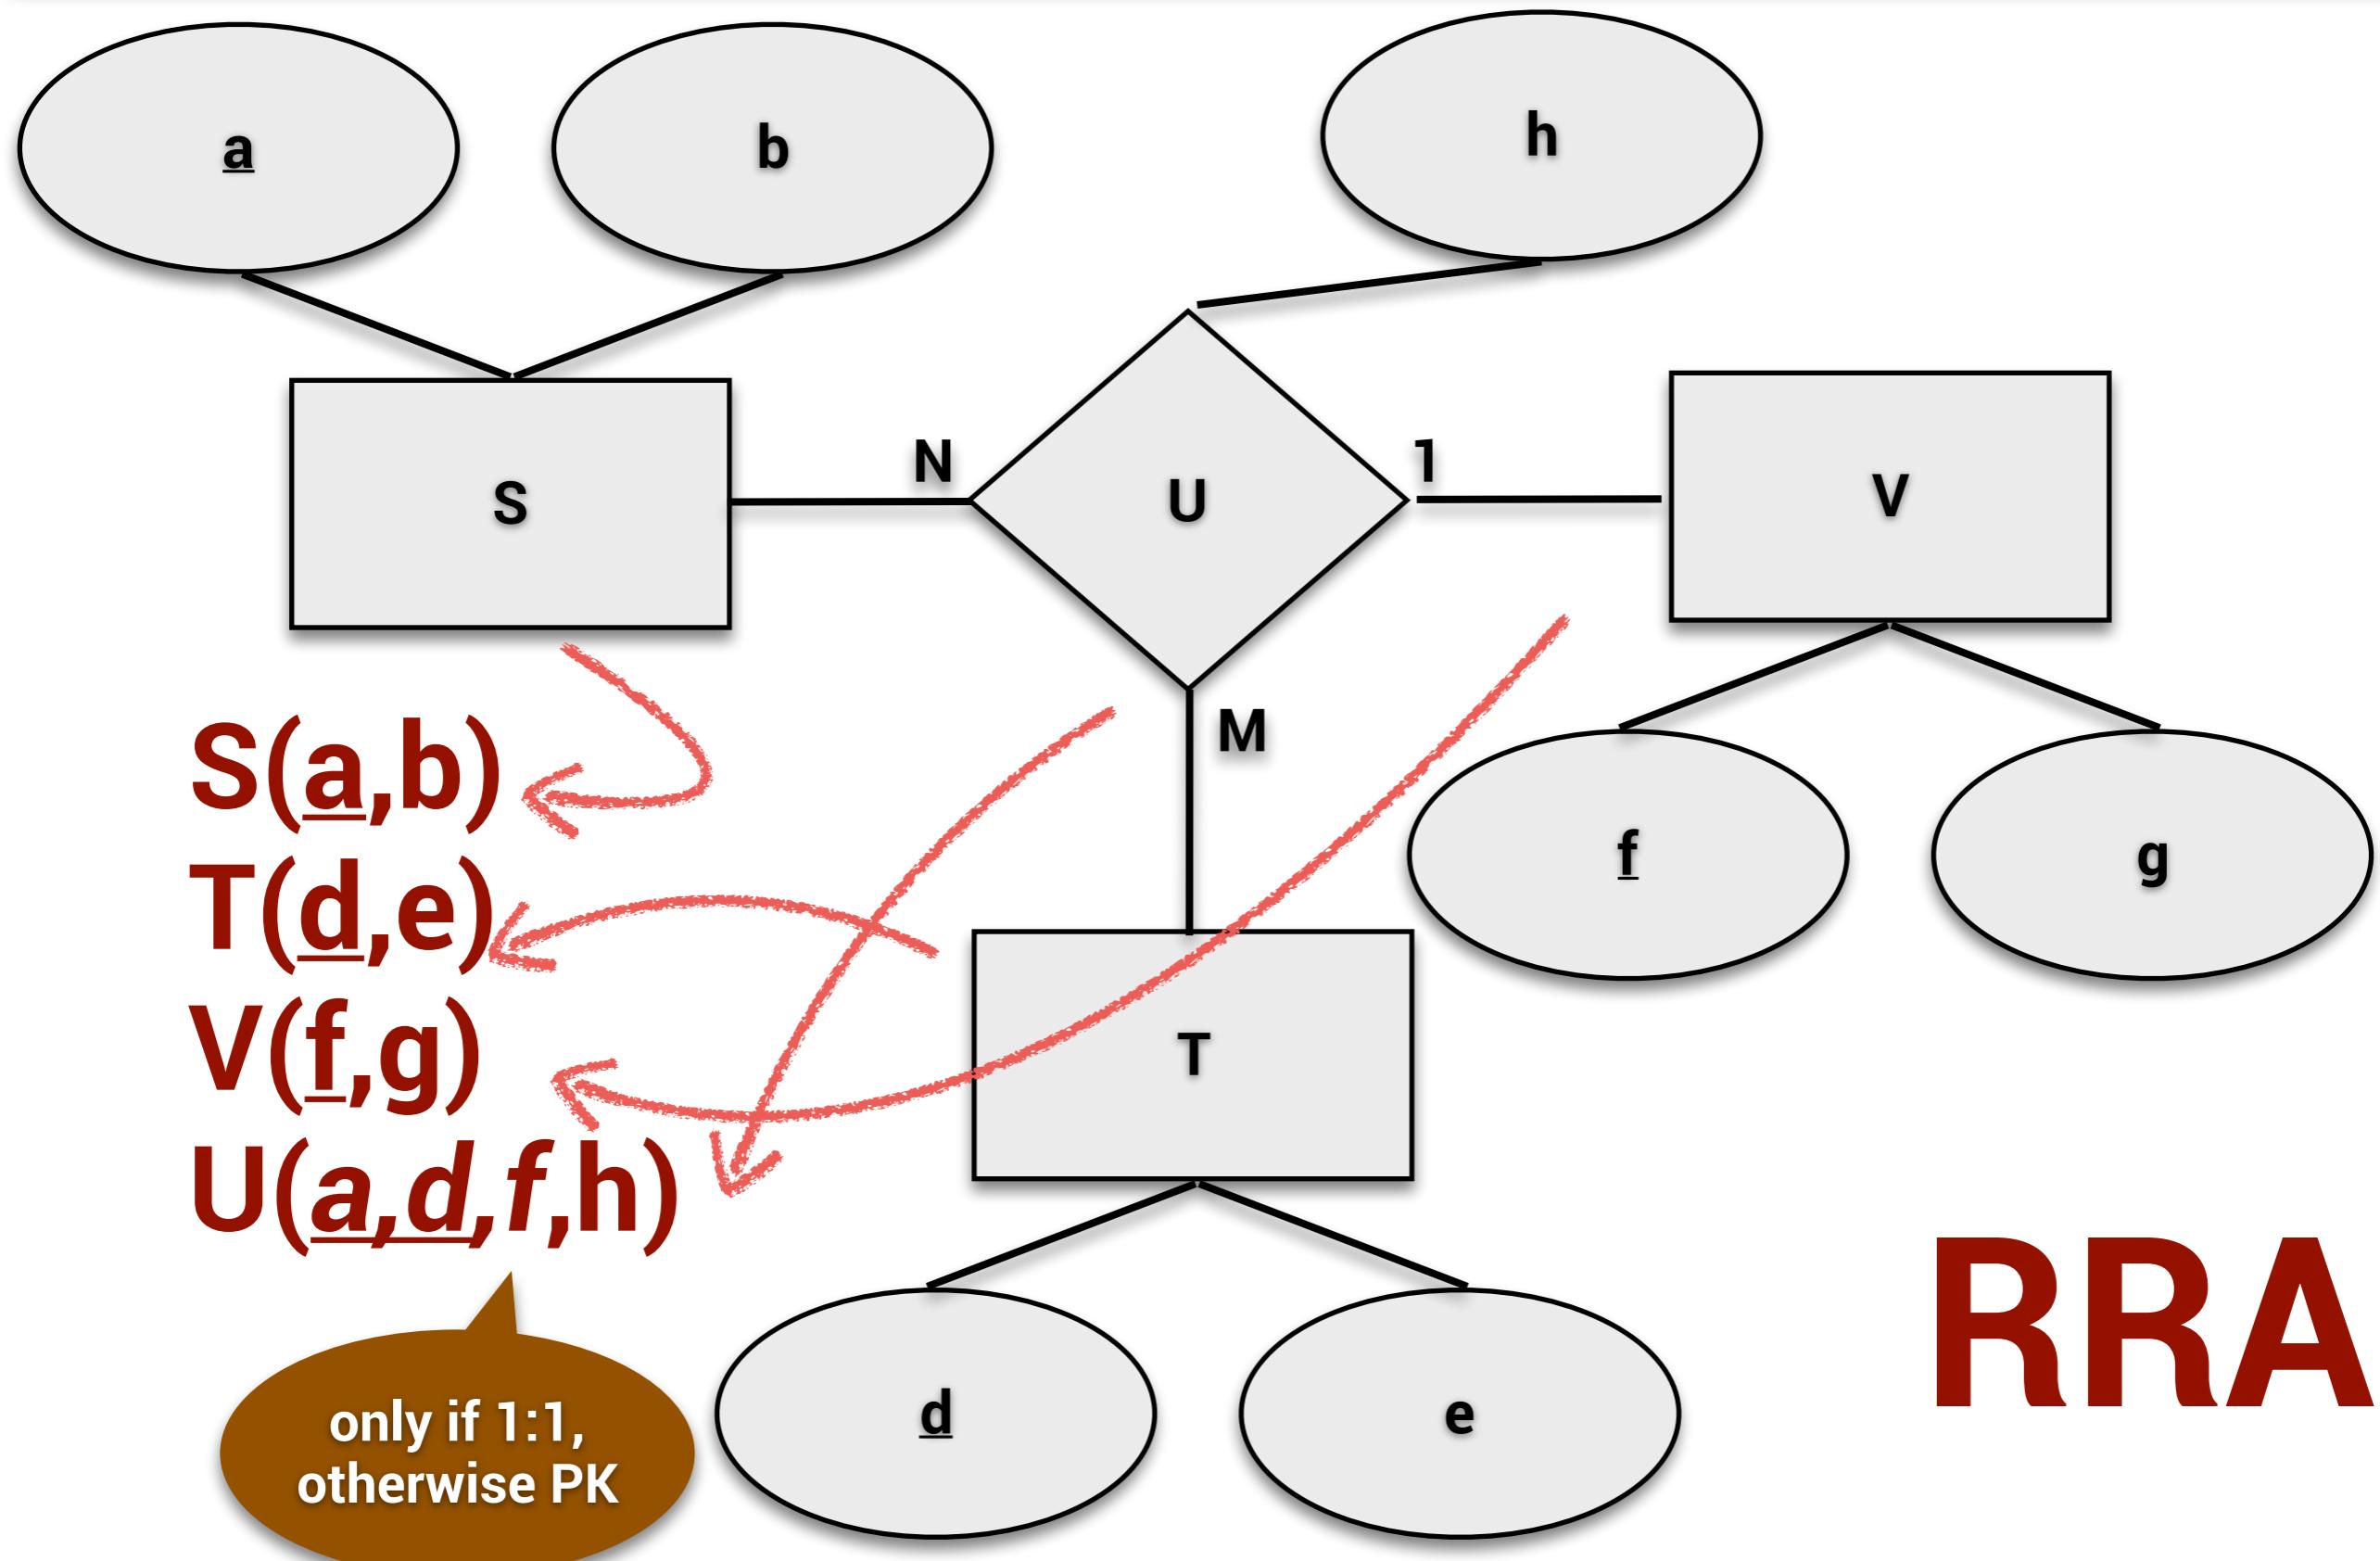
6: Mapping Multivalued Attributes

91



7: Mapping N-ary Relationship Types

92



Spec/Gen Case-by-Case [8A]

93

notice!
PK propagates
"down"!

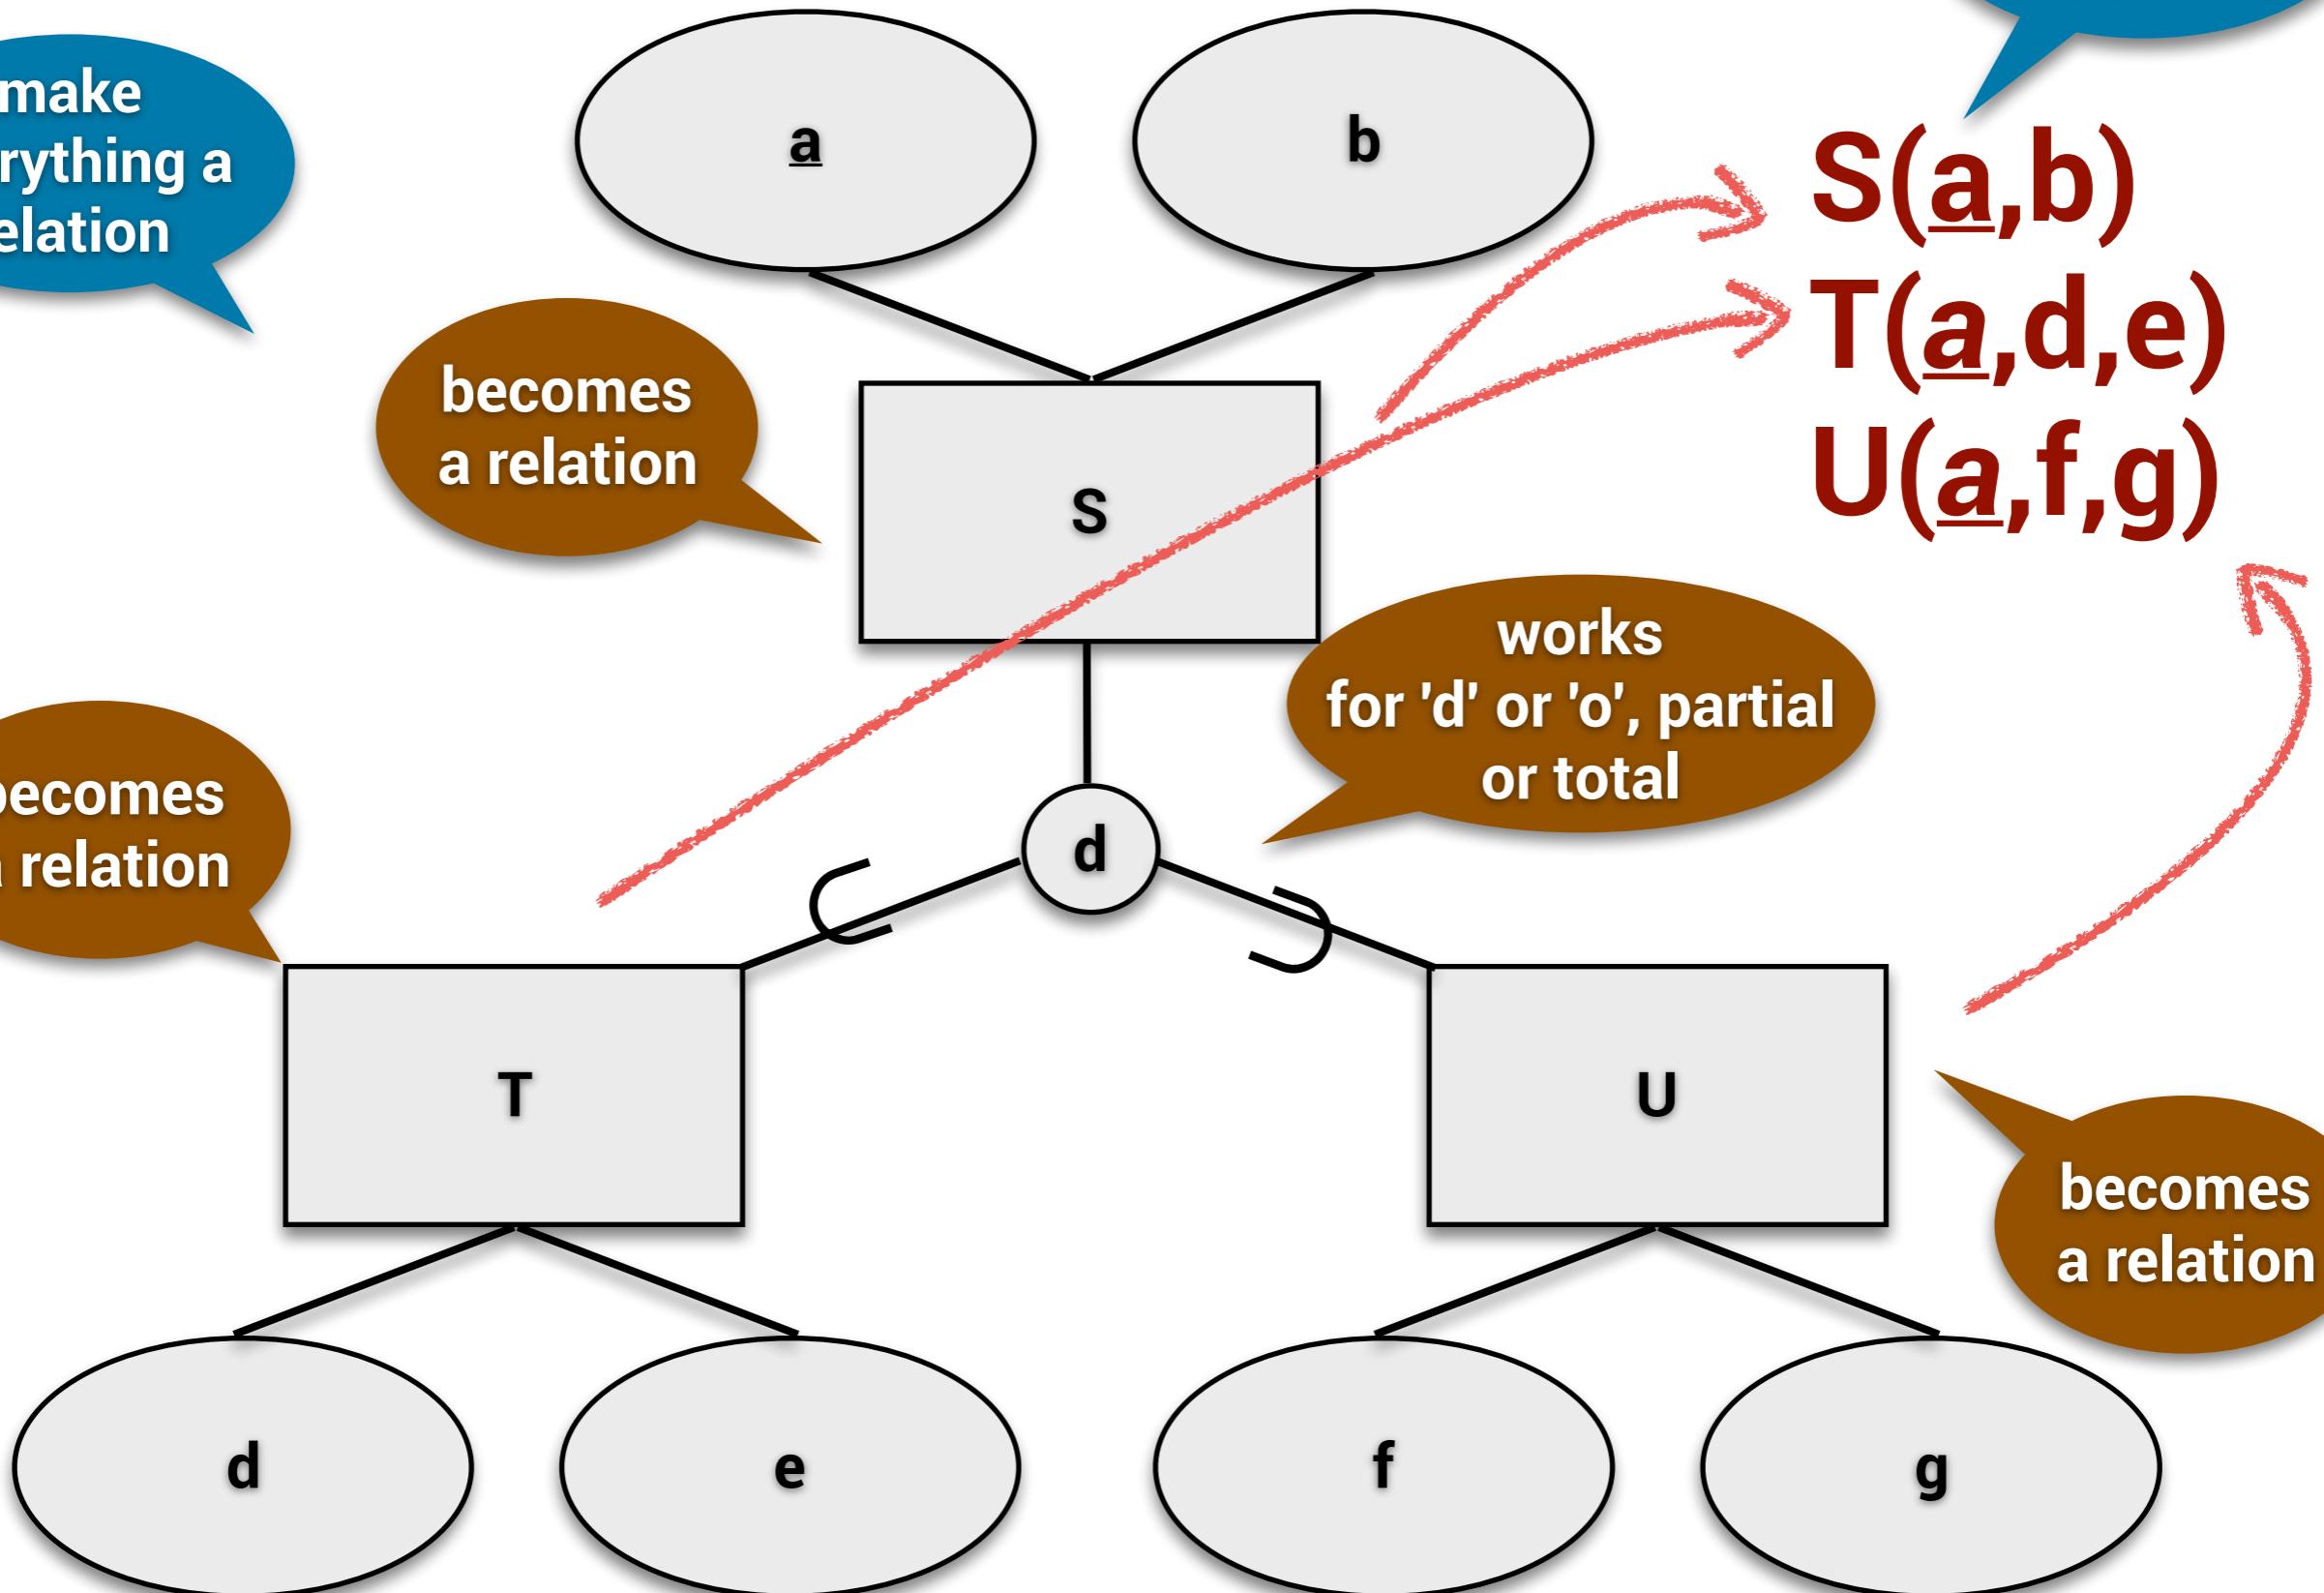
make
everything a
relation

becomes
a relation

becomes
a relation

works
for 'd' or 'o', partial
or total

becomes
a relation



Spec/Gen Case-by-Case [8B]

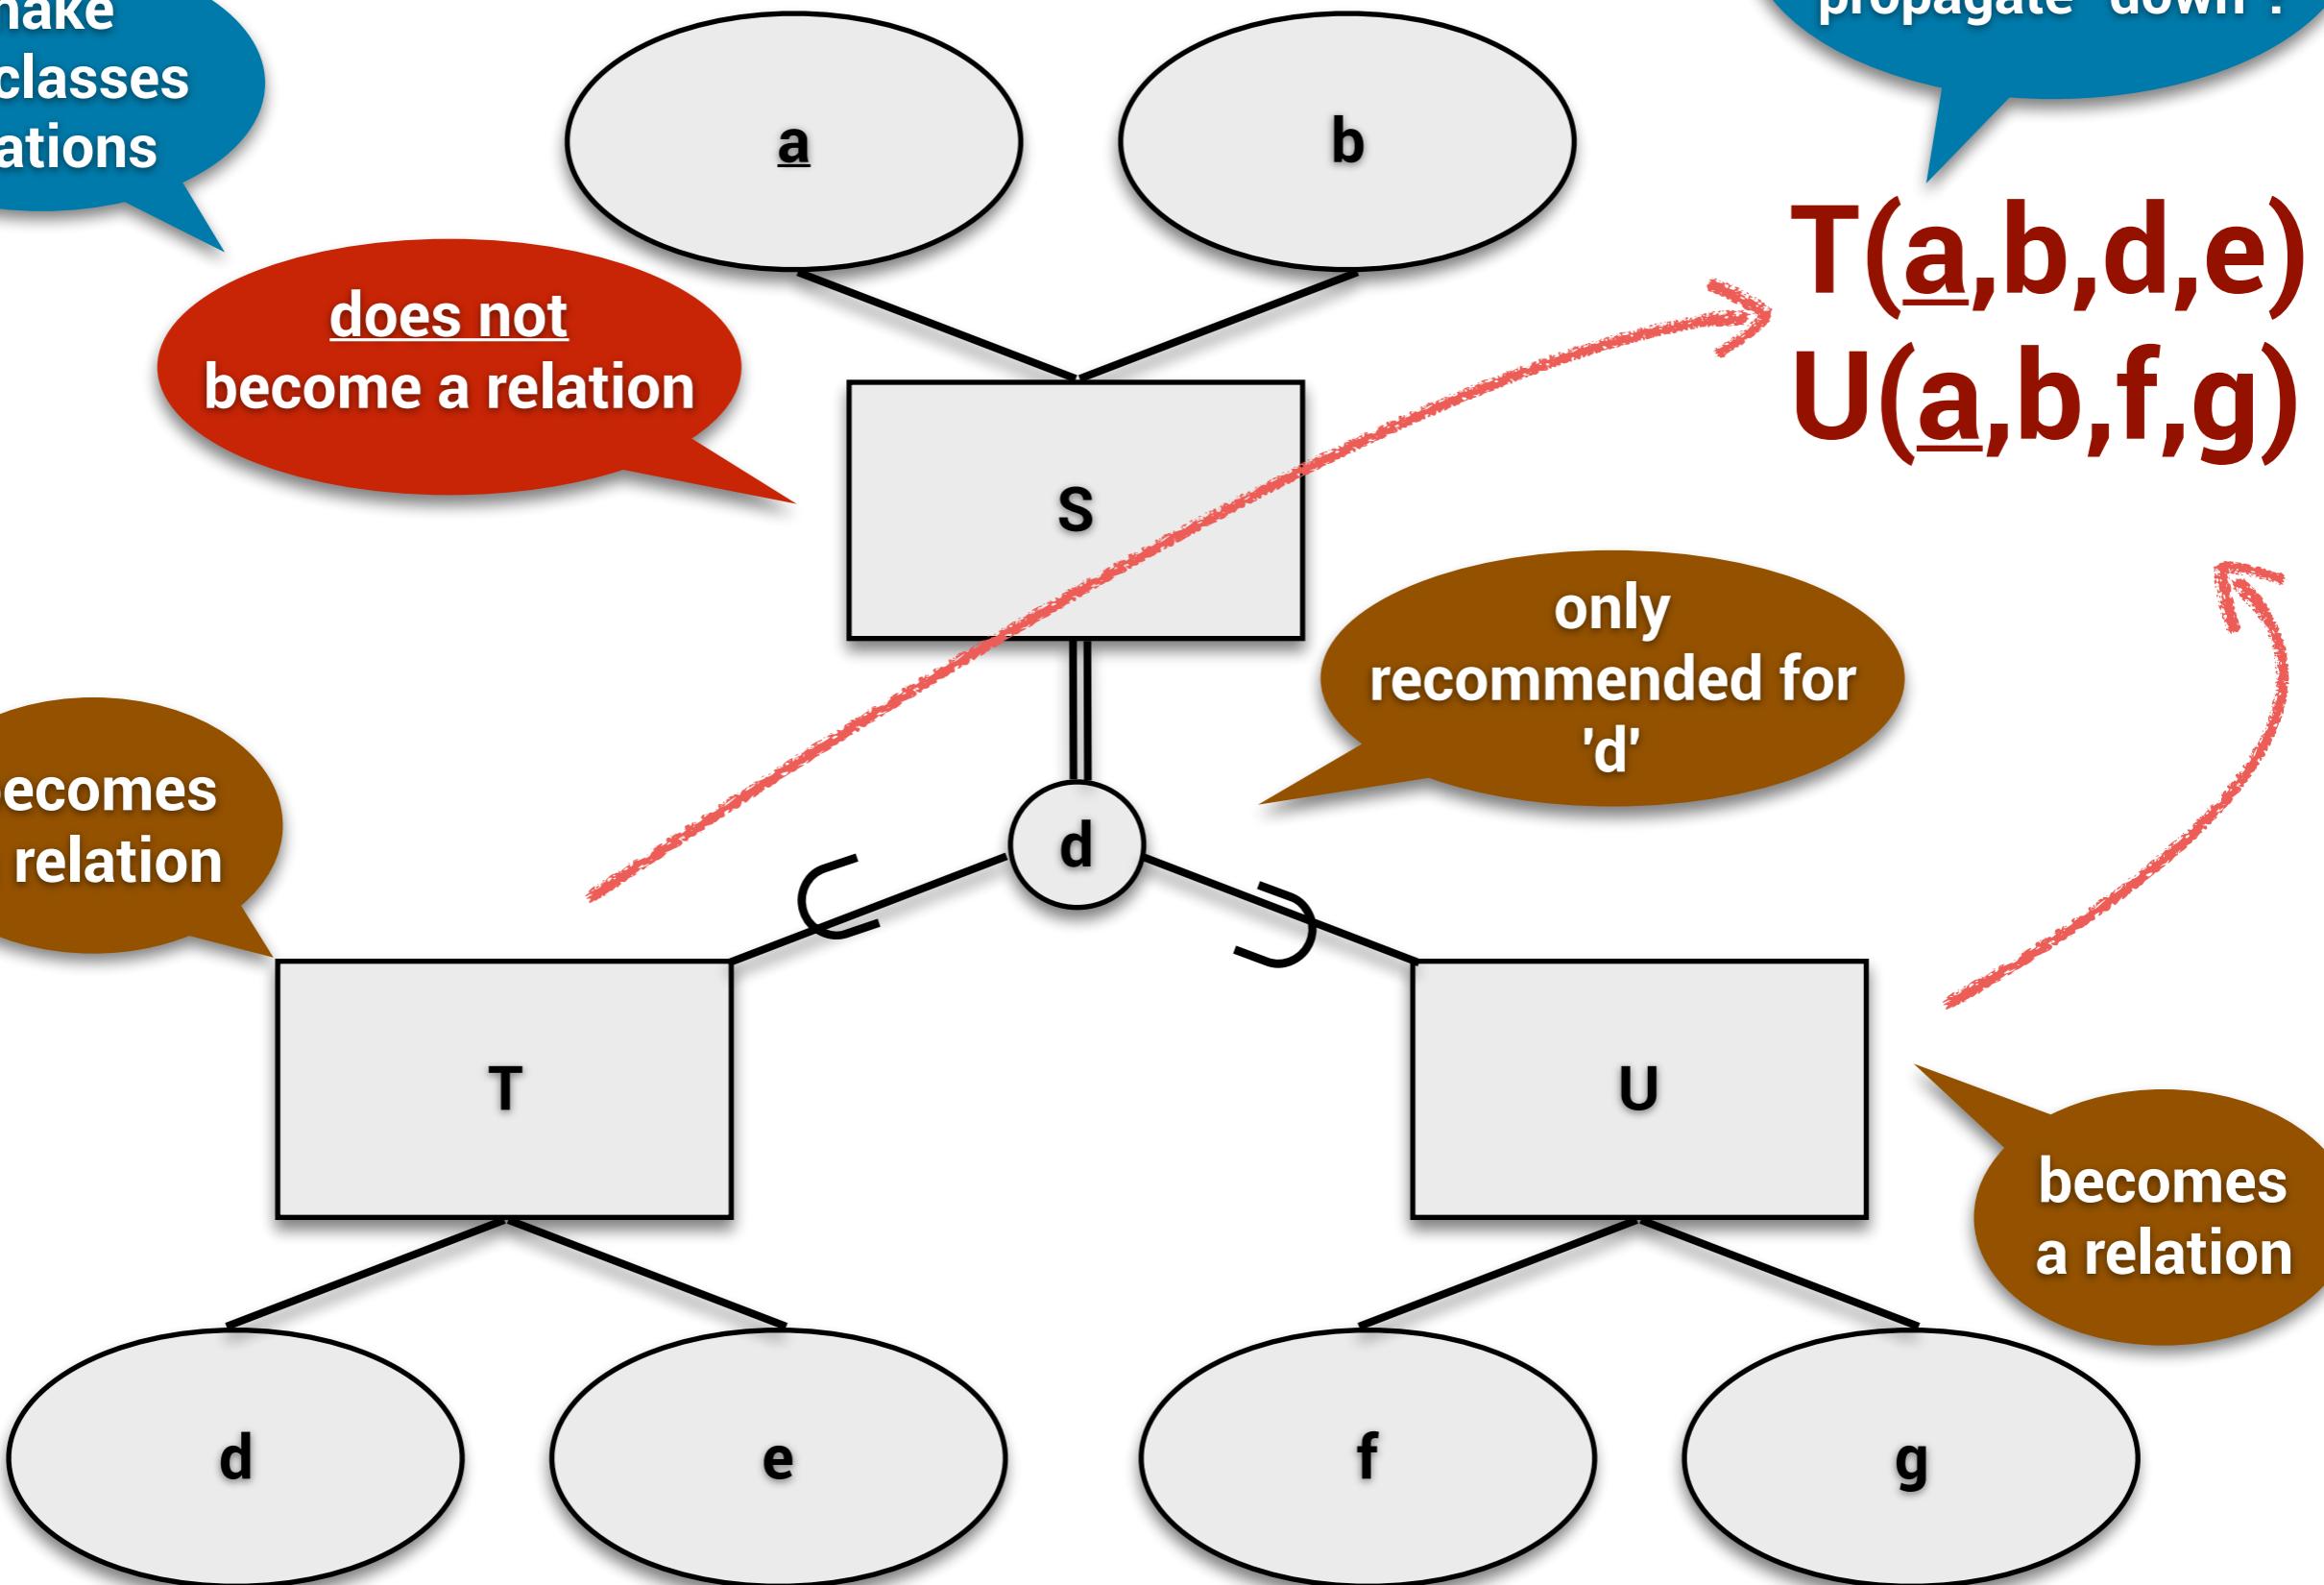
94

make
subclasses
relations

does not
become a relation

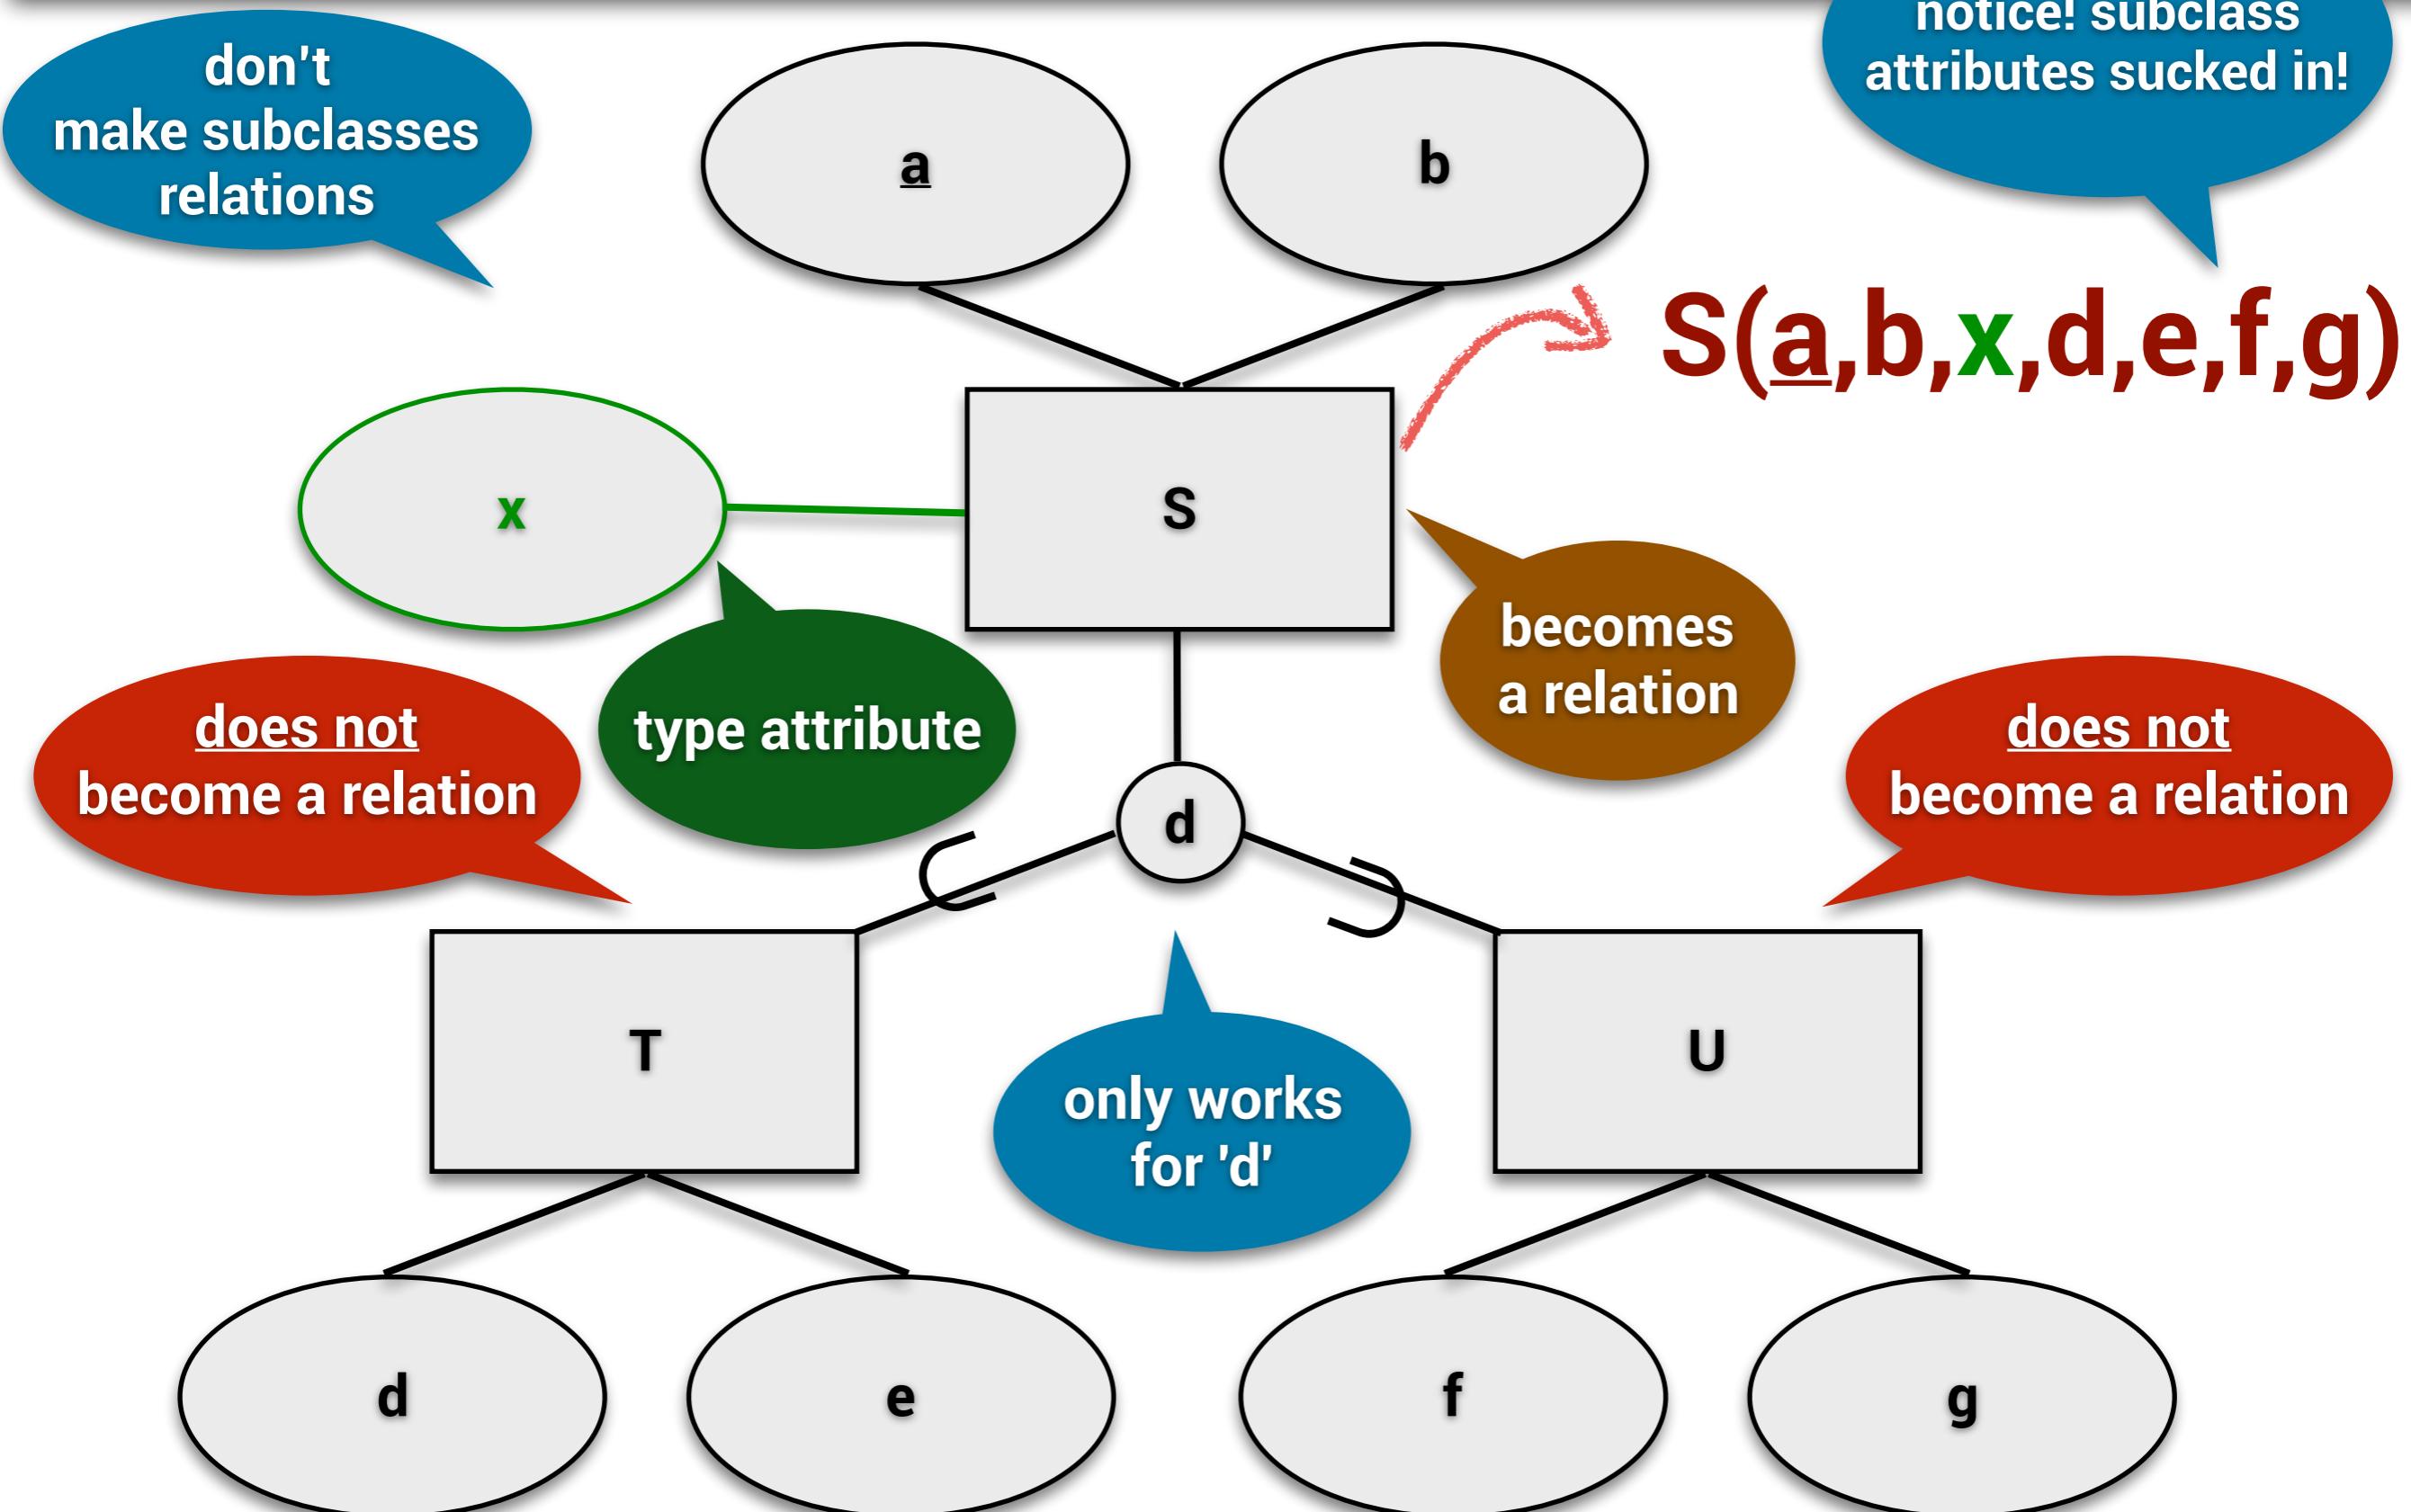
becomes
a relation

notice! PK and
non-key attributes
propagate "down"!



Spec/Gen Case-by-Case [8C]

95



Spec/Gen Case-by-Case [8D]

96

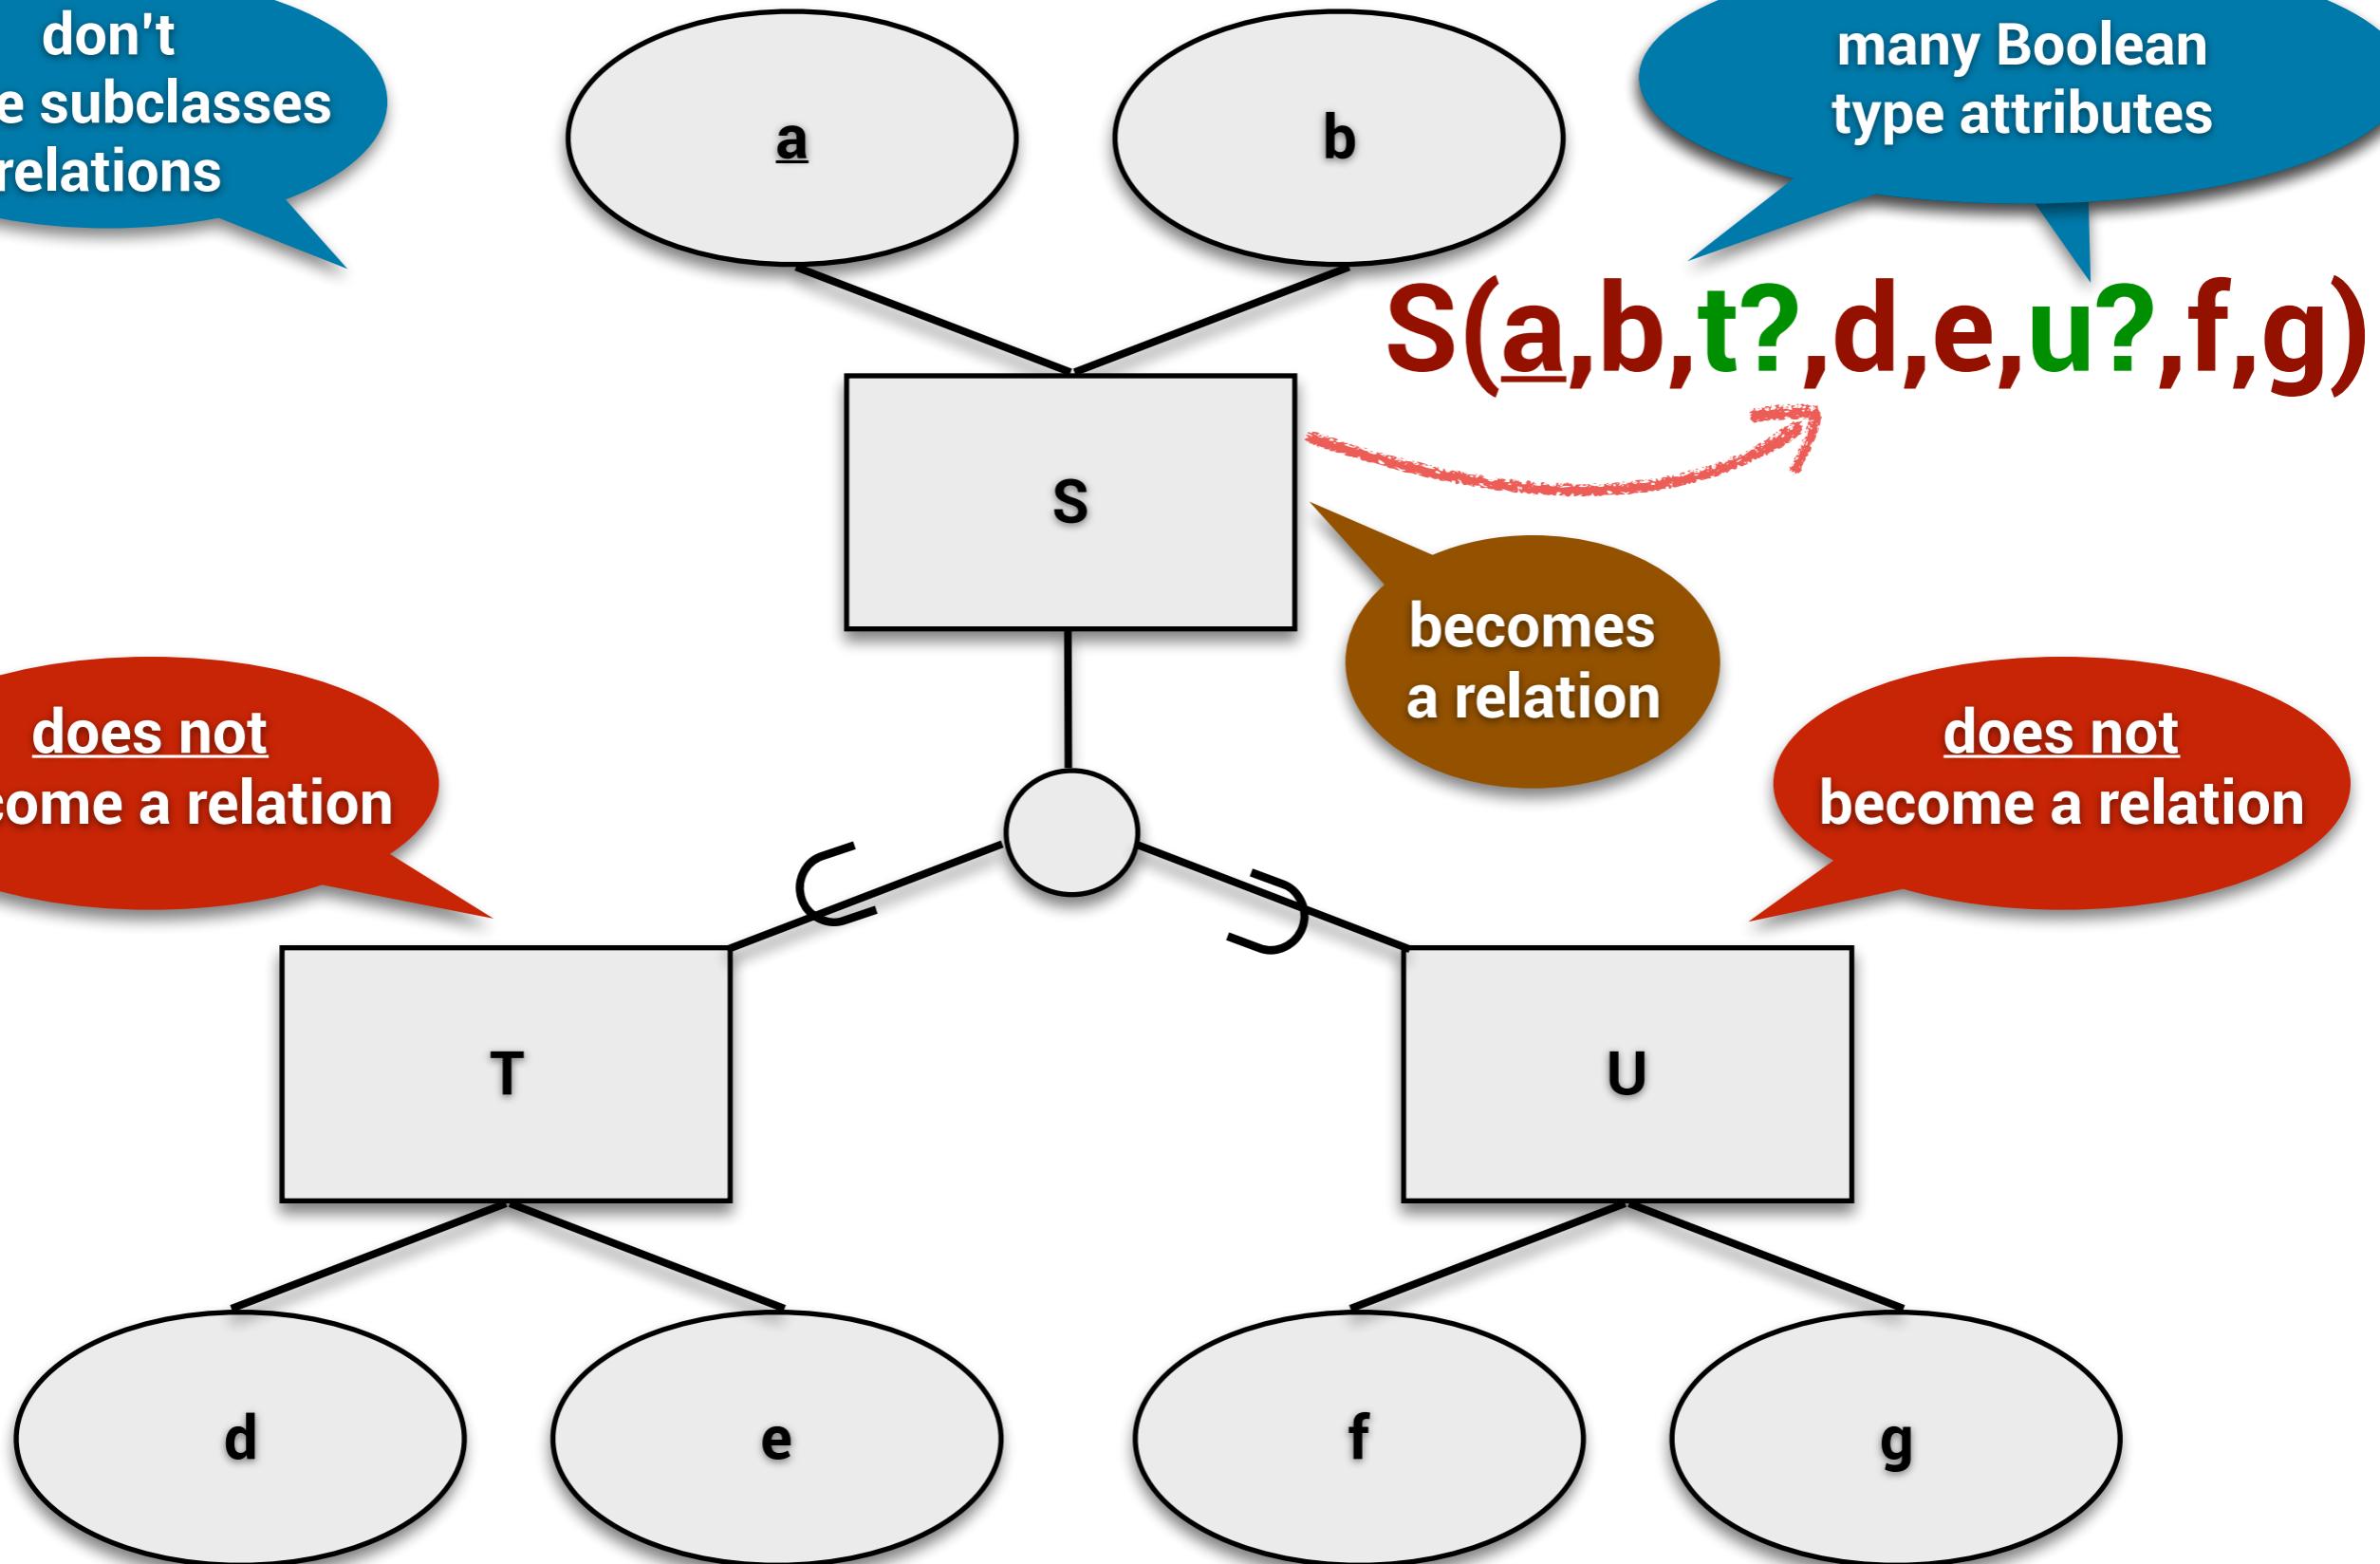
messy?...

don't
make subclasses
relations

many Boolean
type attributes

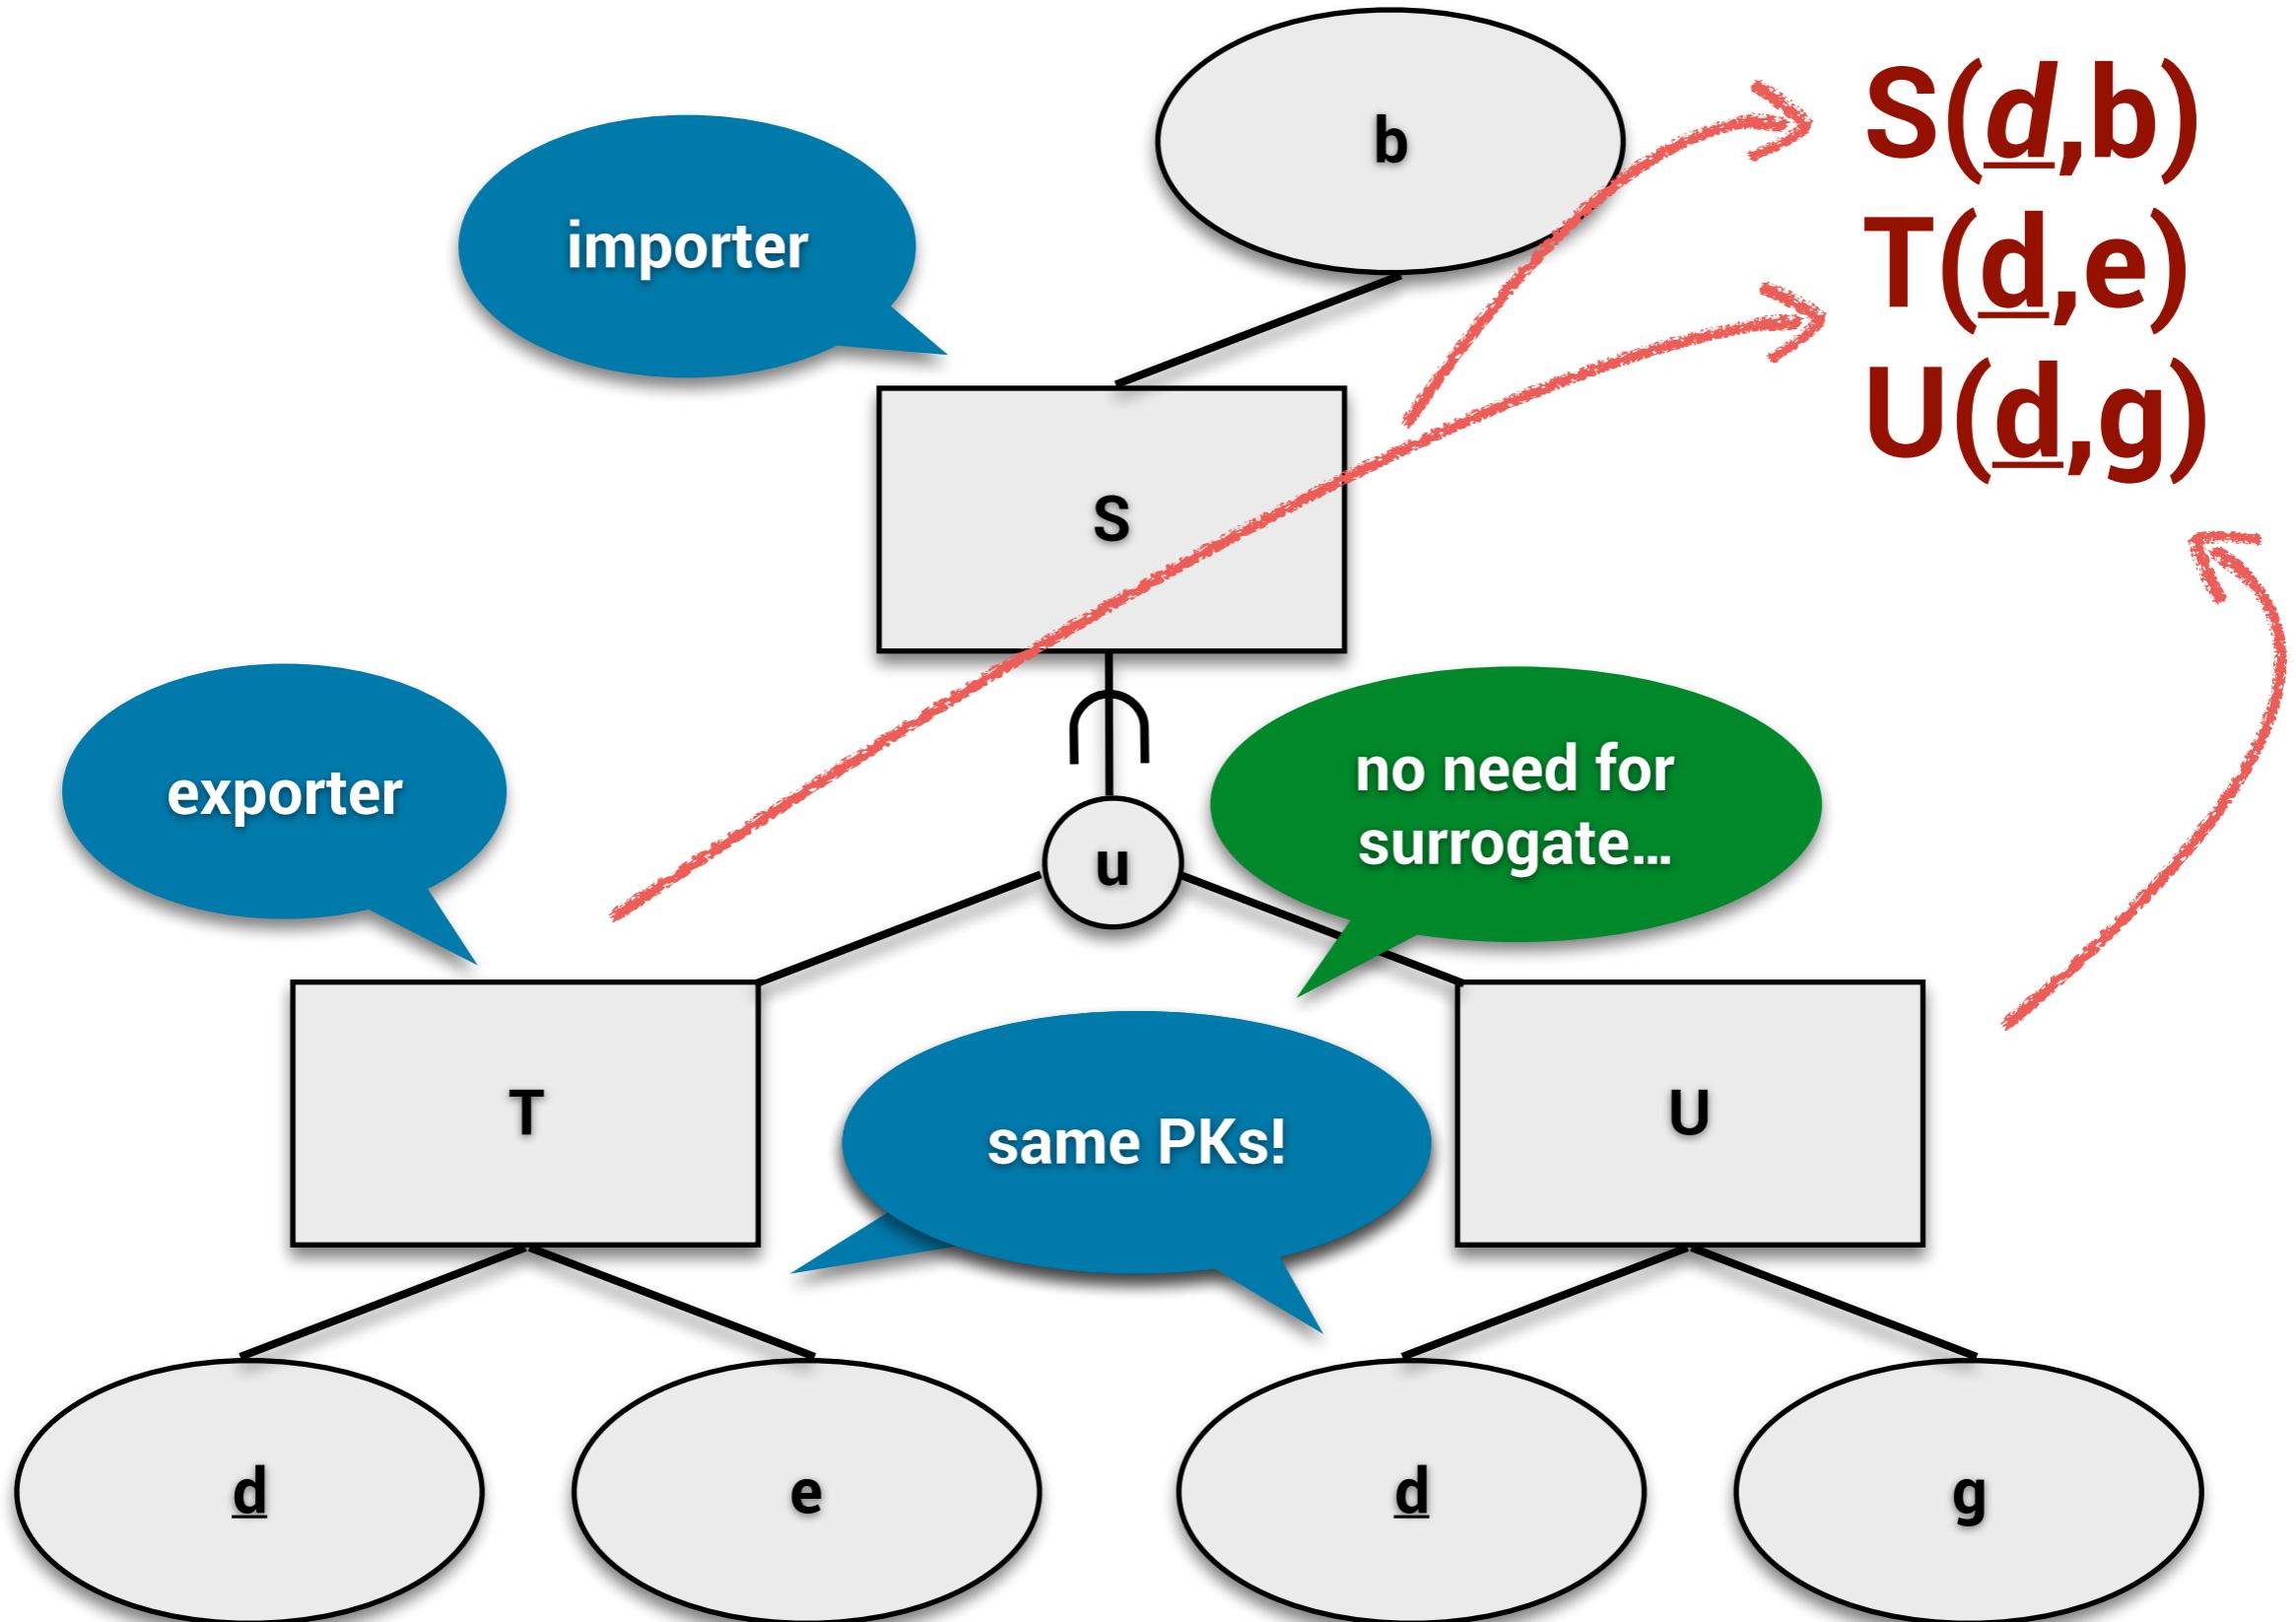
does not
become a relation

does not
become a relation



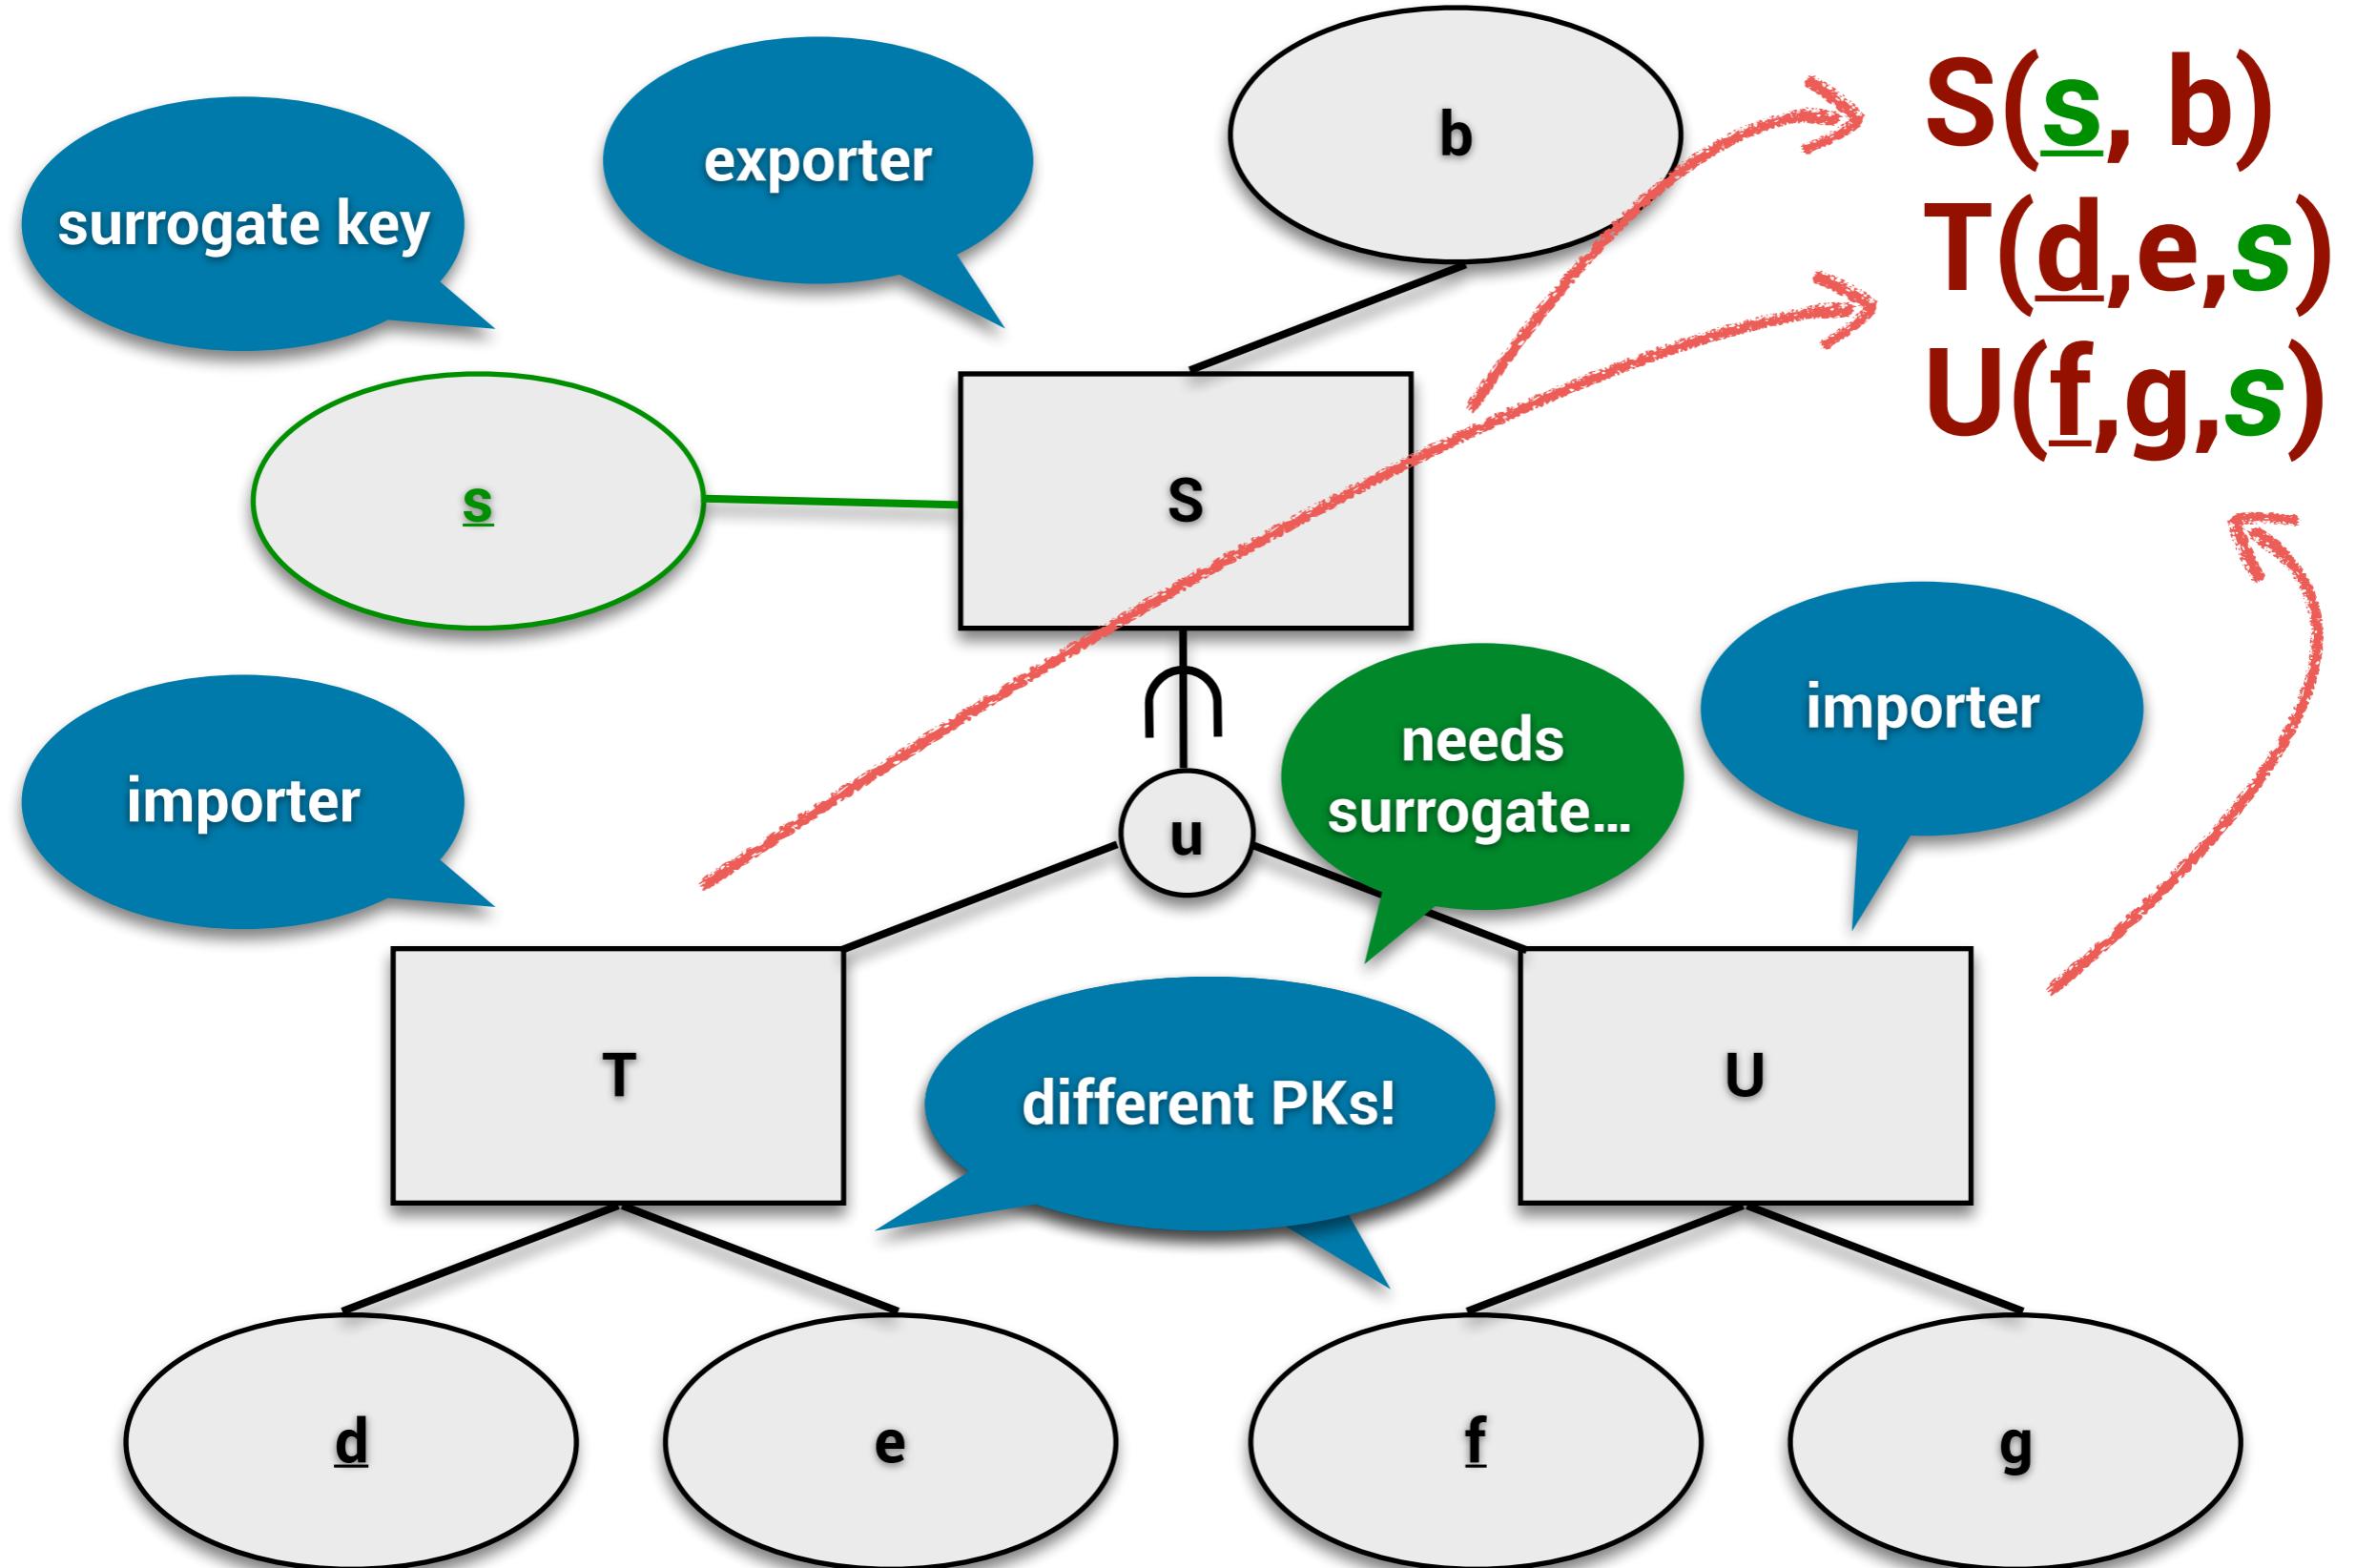
Categories (Union Types) Case-by-Case [9a]

97



Categories (Union Types) Case-by-Case [9b]

98



What Next?

99

- Sandra Sampaio takes over.
- She'll cover:
 - Functional dependencies and normal forms
 - Normalization
 - Advanced SQL
 - Transactions
 - File Organizations and Indices
- A final lecture, in the last week of term by Alvaro A A Fernandes will discuss DBMS architectures.