

COMP24111 lecture 3

The Linear Classifier,
also known as the “Perceptron”



LAST WEEK: our first “machine learning” algorithm



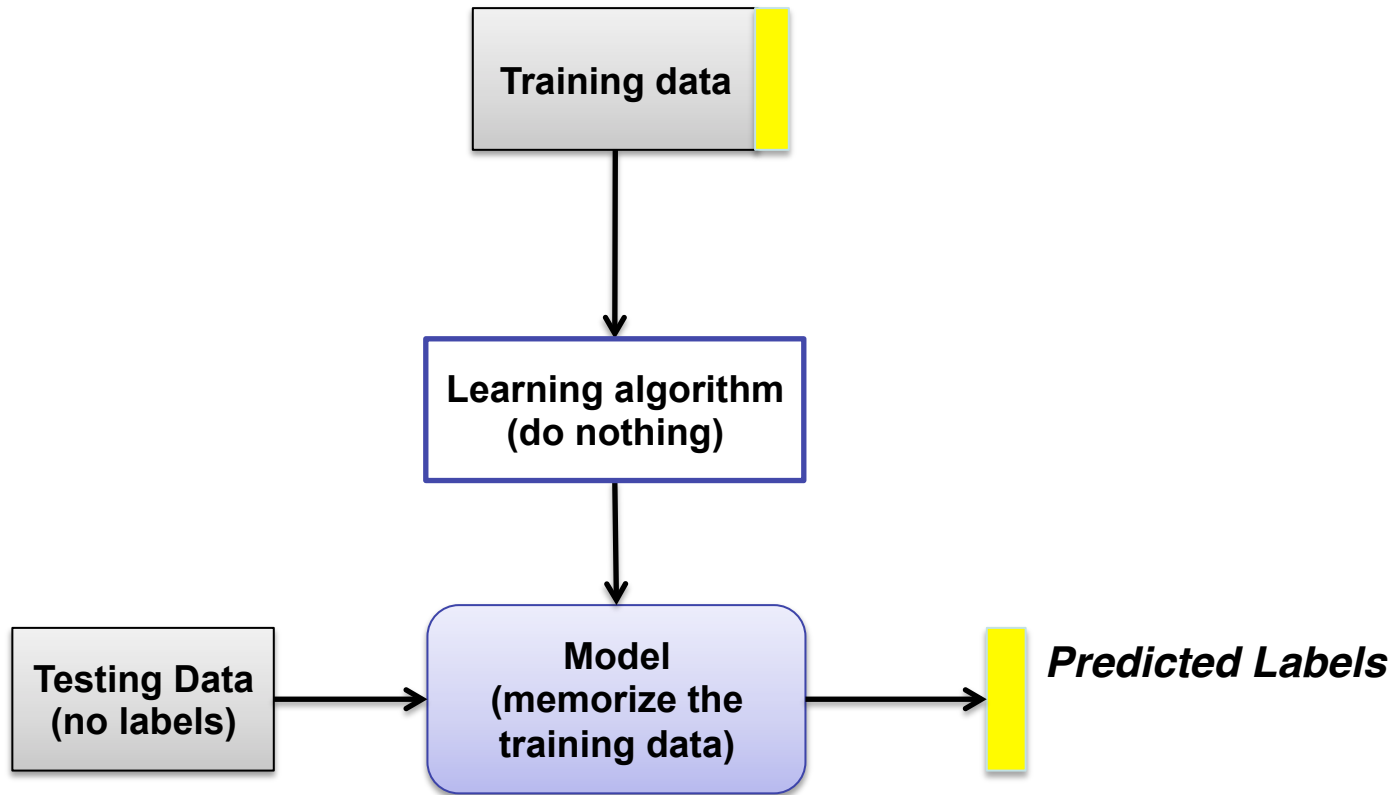
The K-Nearest Neighbour Classifier

```
Testing point  $x$   
For each training datapoint  $x'$   
    measure distance( $x, x'$ )  
End  
Sort distances  
Select  $K$  nearest  
Assign most common class
```

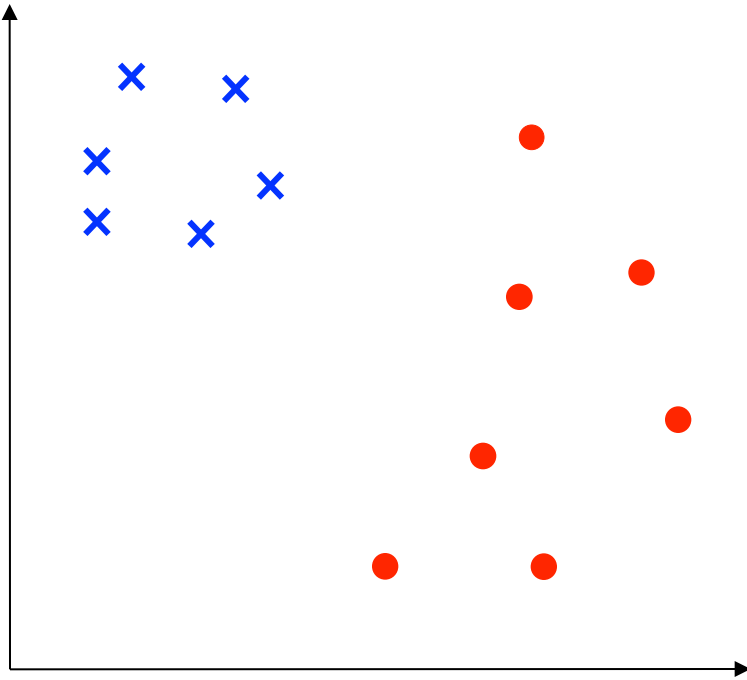
Make your own notes on its advantages / disadvantages.

I'll ask for volunteers next time we meet.....

Supervised Learning Pipeline for Nearest Neighbour

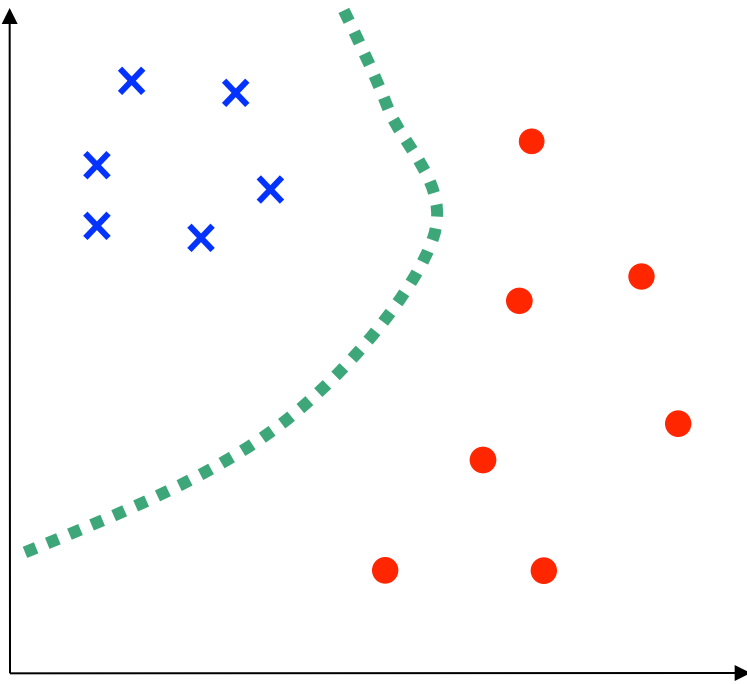


The *most important* concept in Machine Learning



The *most important* concept in Machine Learning

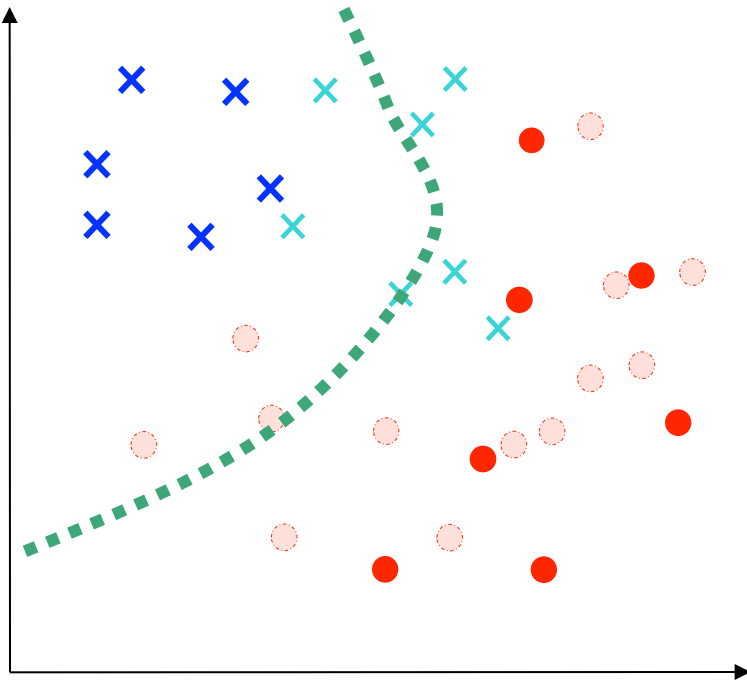
Looks good so far...



The *most important* concept in Machine Learning

Looks good so far...

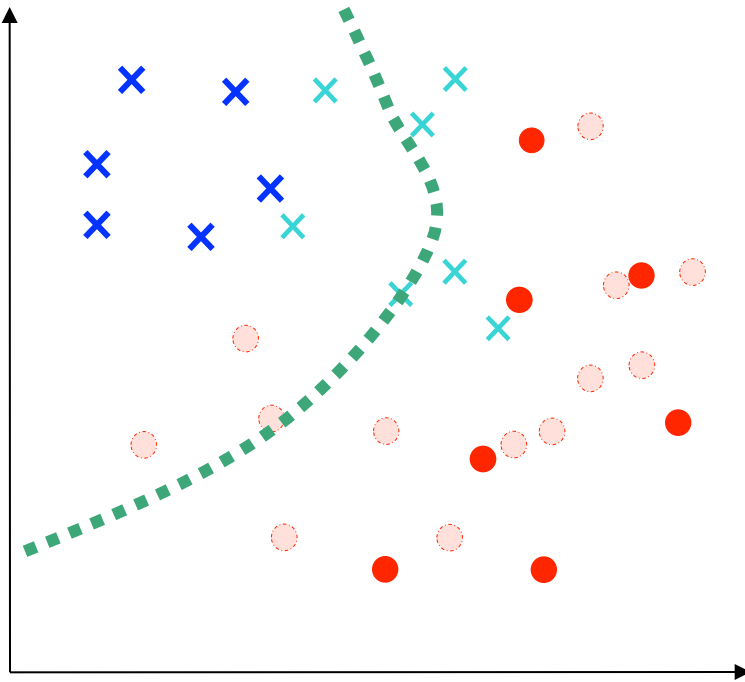
*Oh no! Mistakes!
What happened?*



The *most important* concept in Machine Learning

Looks good so far...

Oh no! Mistakes!
What happened?



We didn't have all the data.

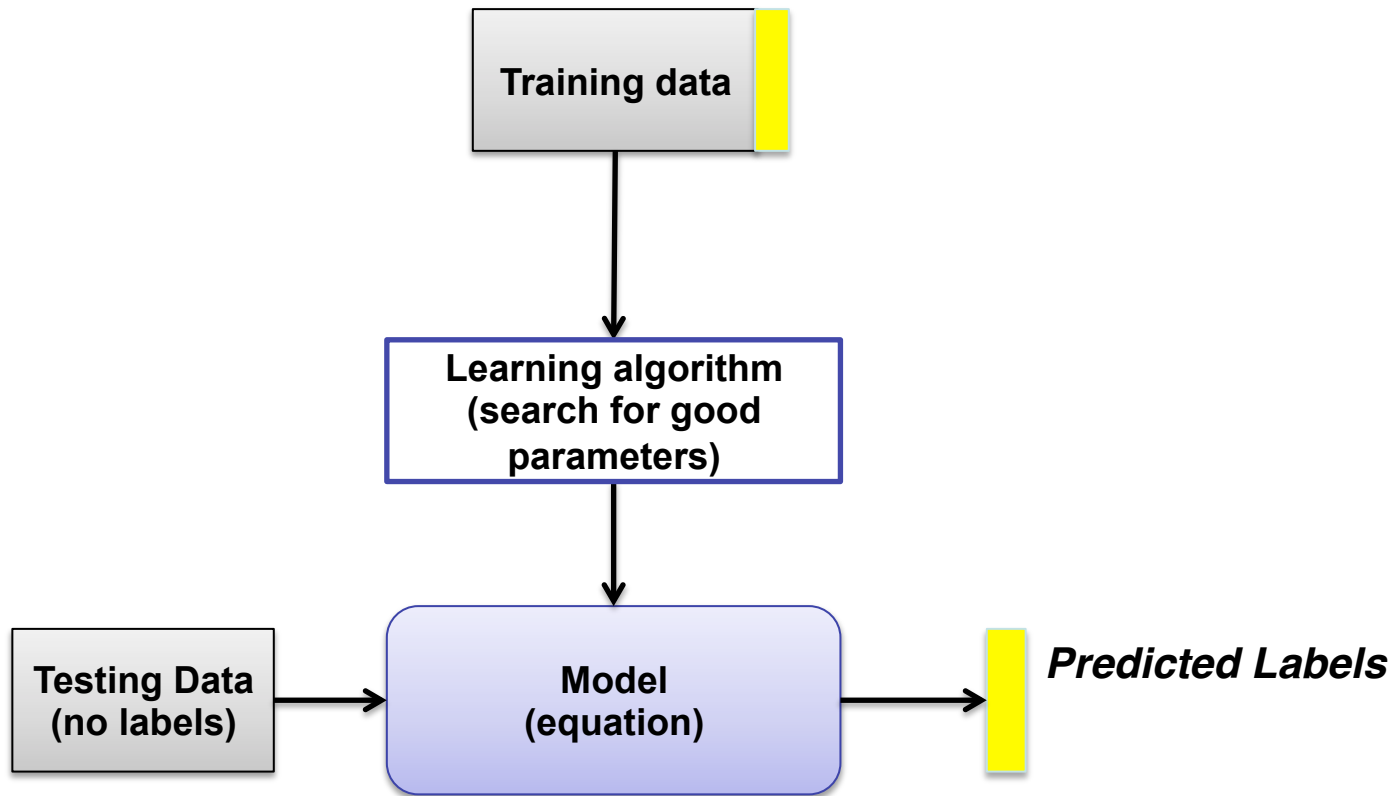
We can never assume that we do.

This is called “OVER-FITTING”
to the small dataset.

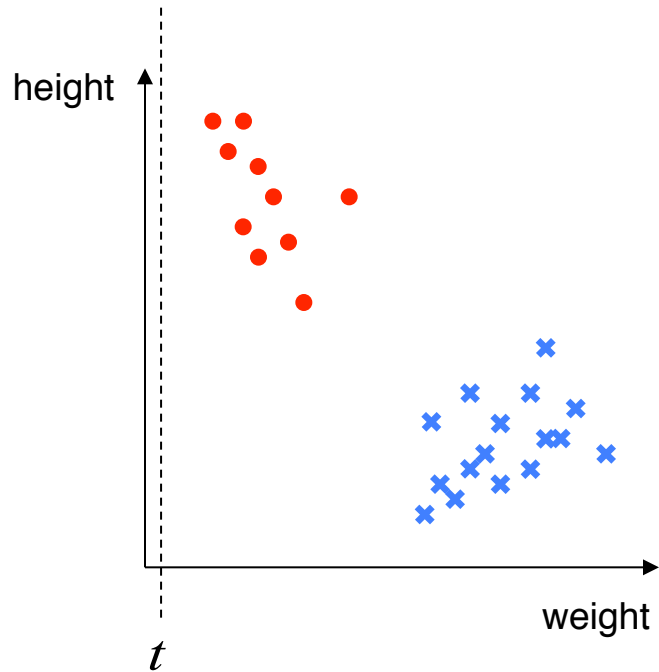
COMP24111 lecture 3

The Linear Classifier

Supervised Learning Pipeline for Linear Classifiers




A more simple, compact model?



if ($weight > t$) then "player" else "dancer"

What's an algorithm to find a good threshold?

```
t=40
while ( numMistakes != 0 )
{
    t = t + 1
    numMistakes = testRule(t)
}
```



if (*weight* > *t*) then "player" else "dancer"

We have our second Machine Learning procedure.

The threshold classifier (also known as a “Decision Stump”)

if (*weight* > *t*) then "player" else "dancer"

```
t=40
while ( numMistakes != 0 )
{
    t = t + 1
    numMistakes = testRule(t)
}
```



Three “ingredients” of a Machine Learning procedure

“Model”

The final product, the thing you have to package up and send to a customer. A piece of code with some parameters that need to be set.

“Error function”

The performance criterion: the function you use to judge how well the parameters of the model are set.

“Learning algorithm”

The algorithm that optimises the model parameters, using the error function to judge how well it is doing.

Three “ingredients” of a Threshold Classifier

Error function

```
t=40
while ( numMistakes != 0 )
{
    t = t + 1
    numMistakes = testRule(t)
}
```

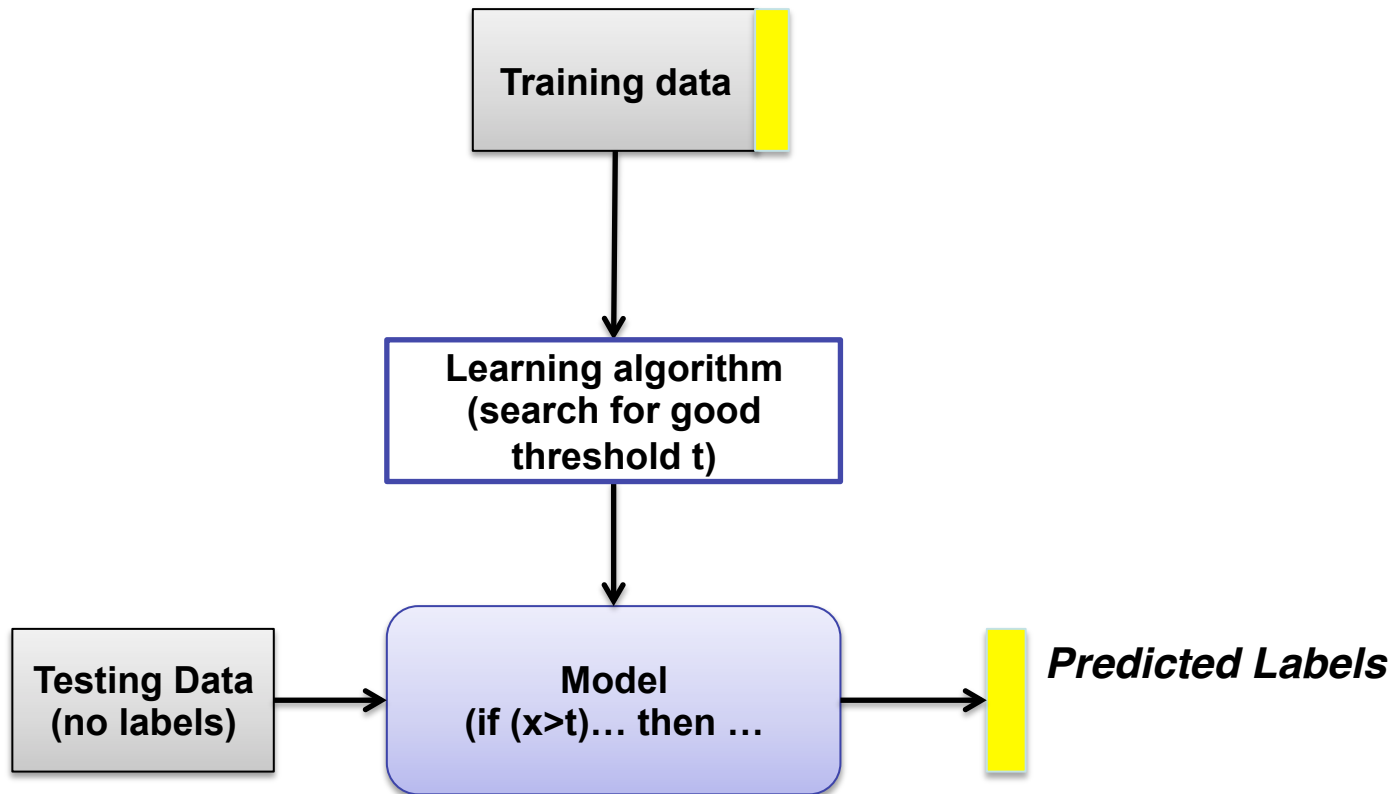
Learning
algorithm

if ($x > t$) then "player" else "dancer"

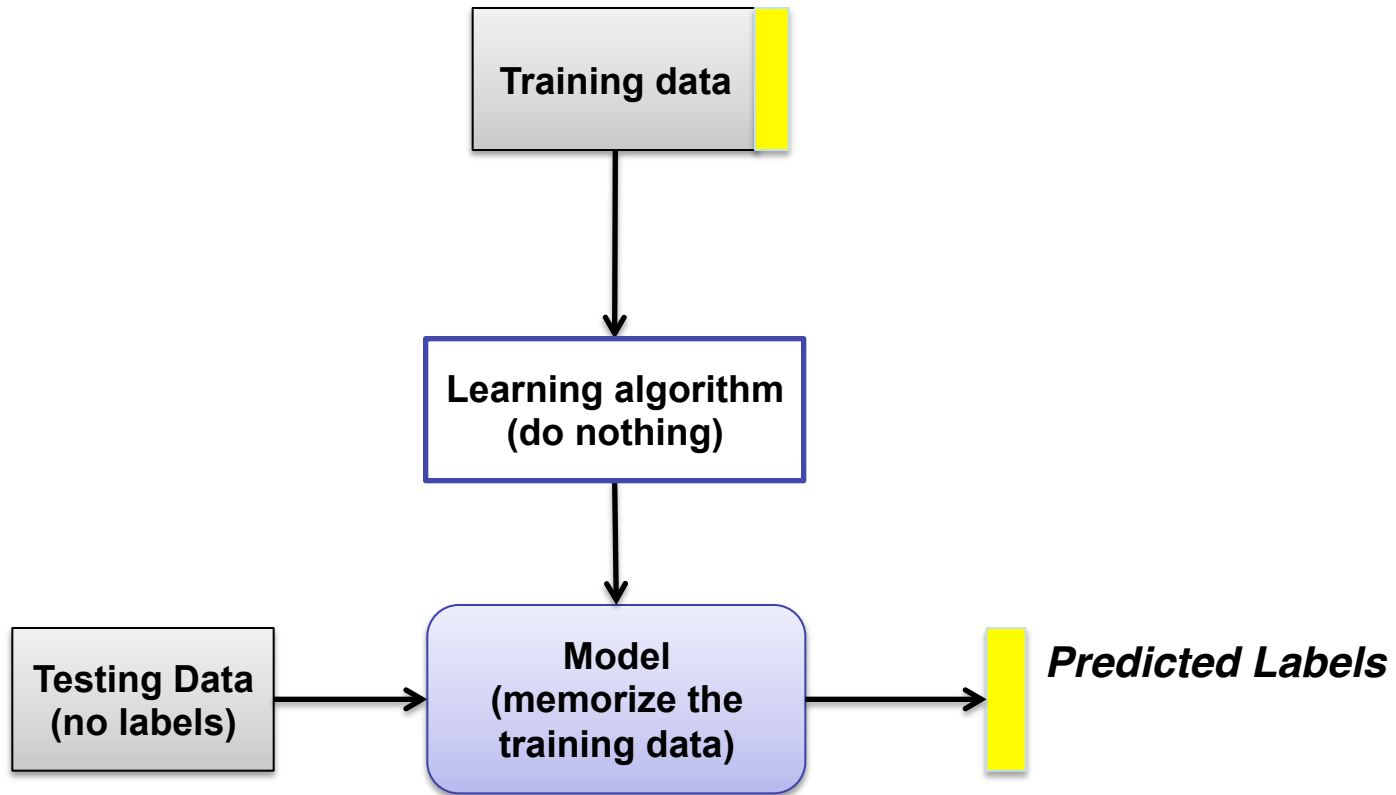
Model

General case – we’re not just talking about the weight of a rugby player – a threshold can be put on any feature ‘x’.

Supervised Learning Pipeline for Threshold Classifier



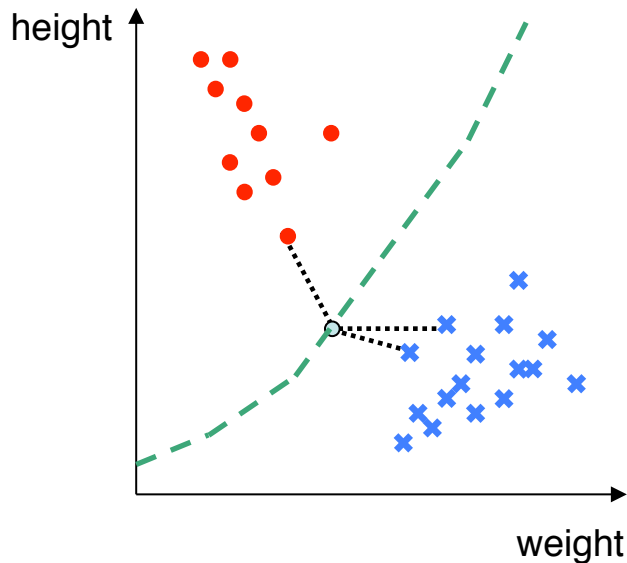
Supervised Learning Pipeline for Nearest Neighbour



What's the “model” for the Nearest Neighbour classifier?

For the k-nn, the model is the *training data* itself !

- *very good accuracy* 😊
- *very computationally intensive!* ☹️



Testing point x

For each training datapoint x'

measure $\text{distance}(x, x')$


End

Sort distances

Select K nearest

Assign most common class

New data: what's an algorithm to find a good threshold?

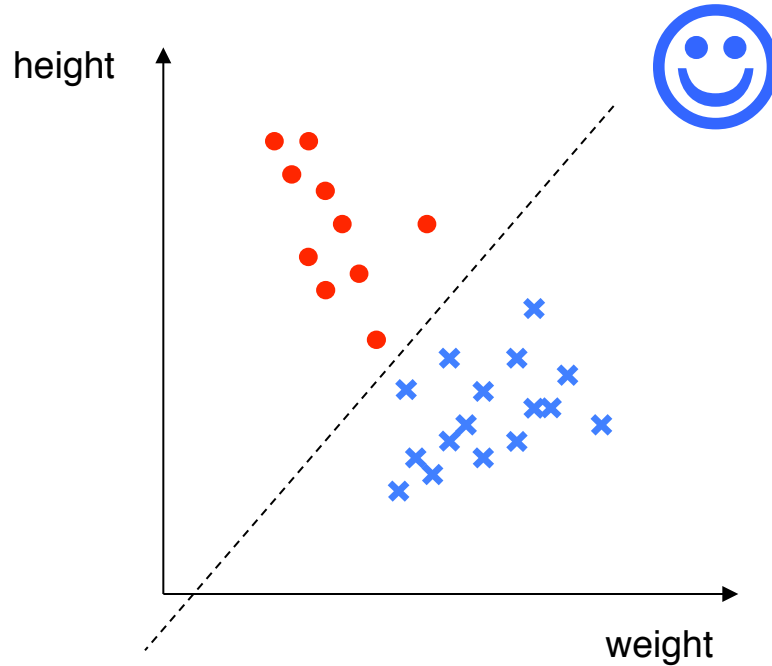


Our model does not match the problem!

if ($weight > t$) then "player" else "dancer"

1 mistake...

New data: what's an algorithm to find a good threshold?



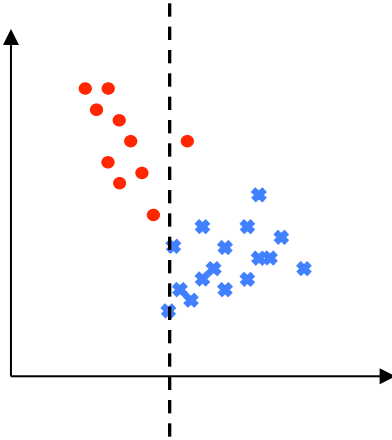
But our current model
cannot represent this...

if (*weight* > *t*) then "player" else "dancer"




We need a more sophisticated model...

if $(x > t)$ then "player" else "dancer"





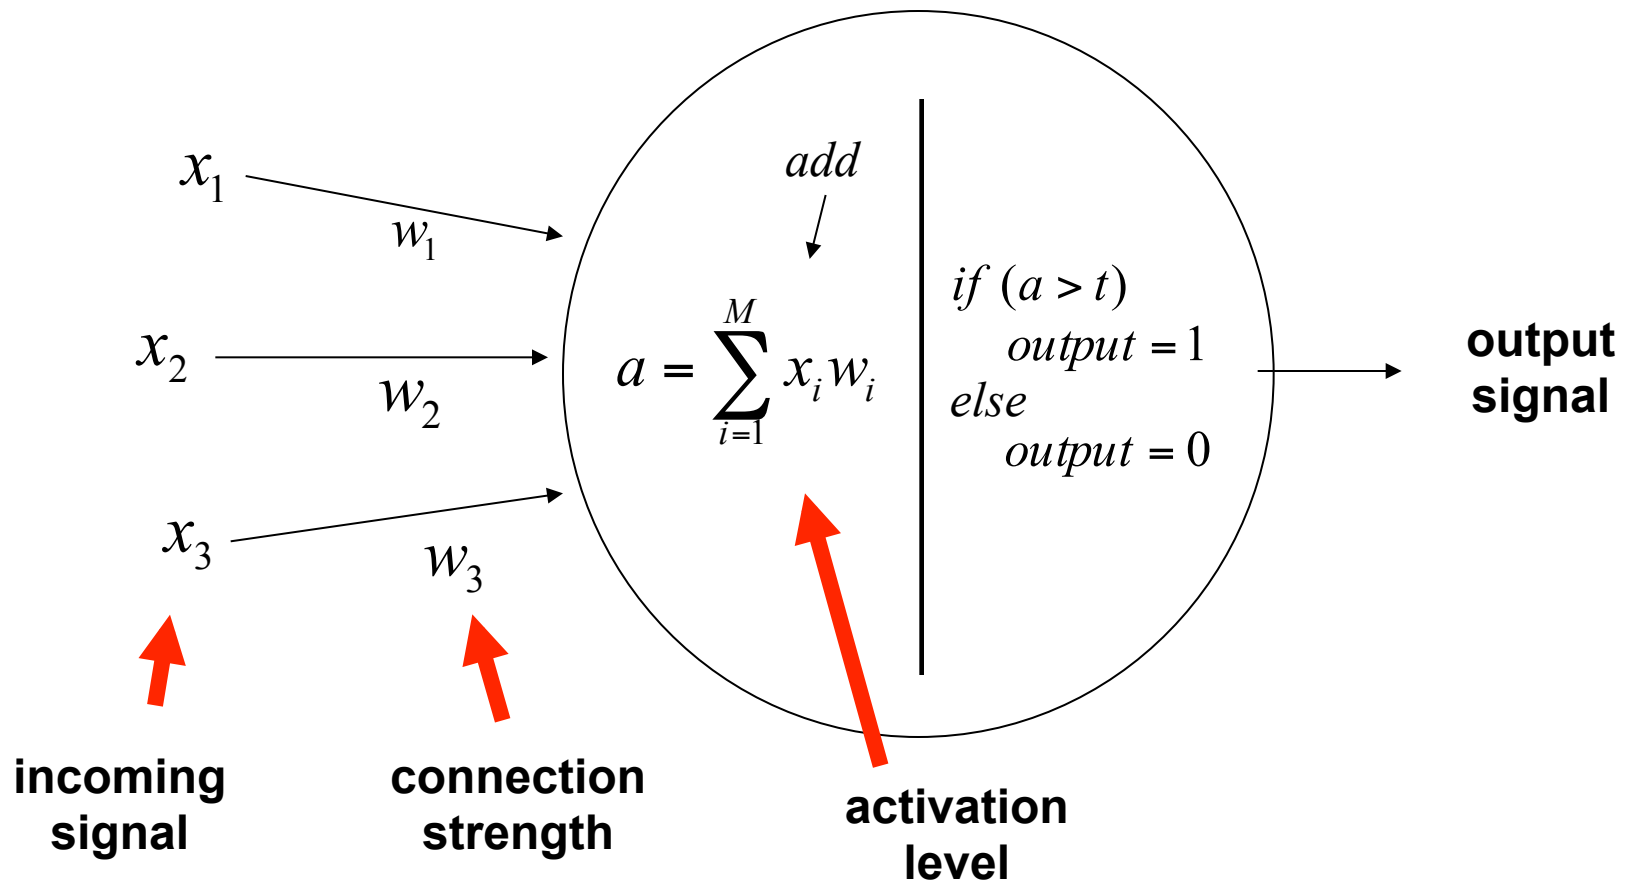
**Input signals sent
from other neurons**



**If enough sufficient
signals accumulate, the
neuron fires a signal.**

**Connection strengths determine
how the signals are accumulated**

- input signals 'x' and coefficients 'w' are multiplied
- weights correspond to connection strengths
- signals are added up – if they are enough, FIRE!

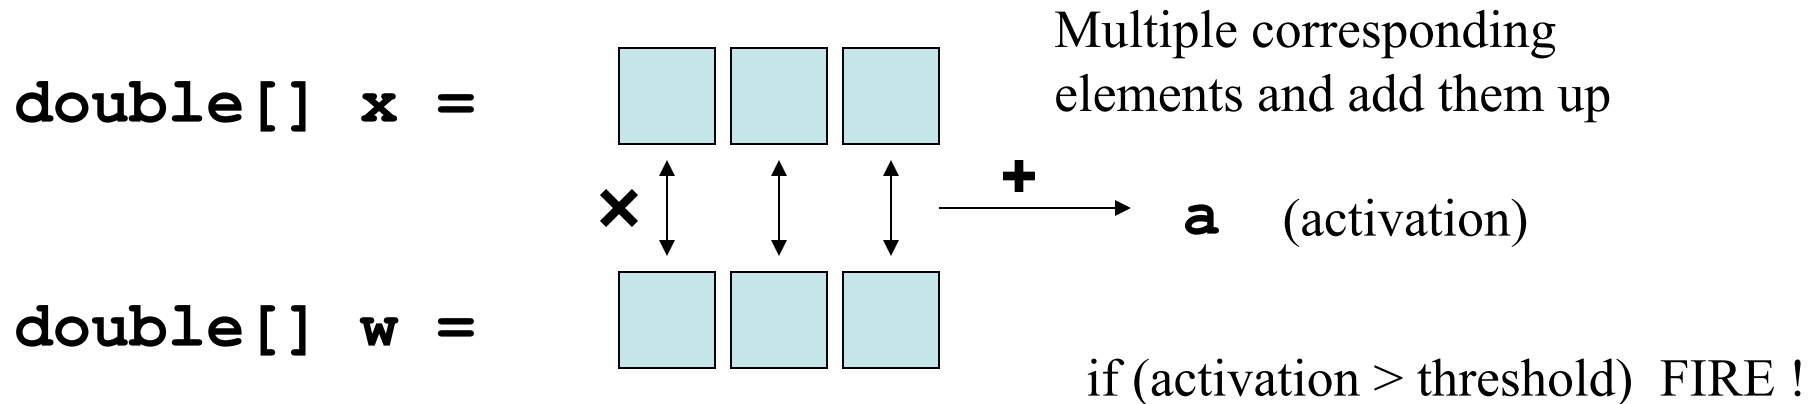


Calculation...

$$a = \sum_{i=1}^M x_i w_i$$

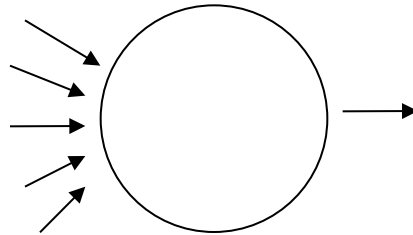
Sum notation

(just like a loop from 1 to M)

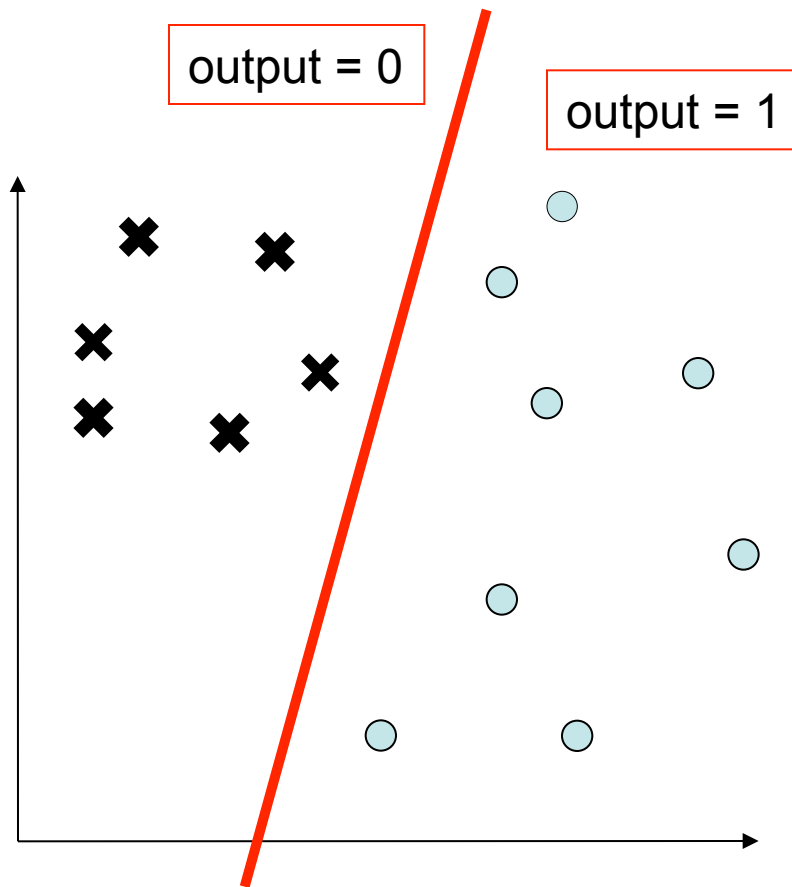


The Perceptron Decision Rule

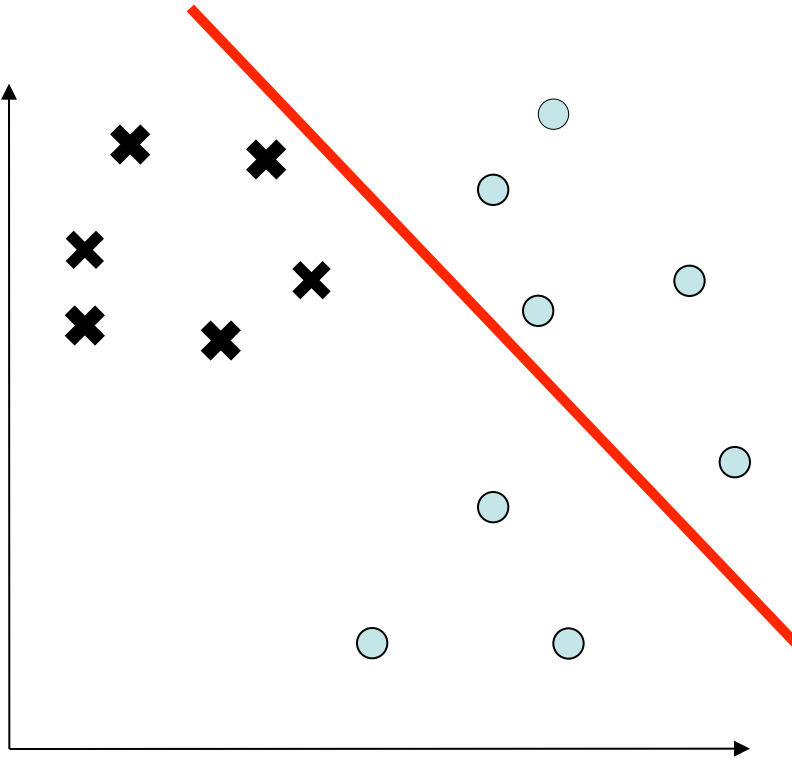
$$\text{if } \left(\sum_{i=1}^M x_i w_i \right) > t \quad \text{then } output = 1, \text{ else } output = 0$$



$$\text{if } \left(\sum_{i=1}^M x_i w_i \right) > t \quad \text{then } \textit{output} = 1, \text{ else } \textit{output} = 0$$

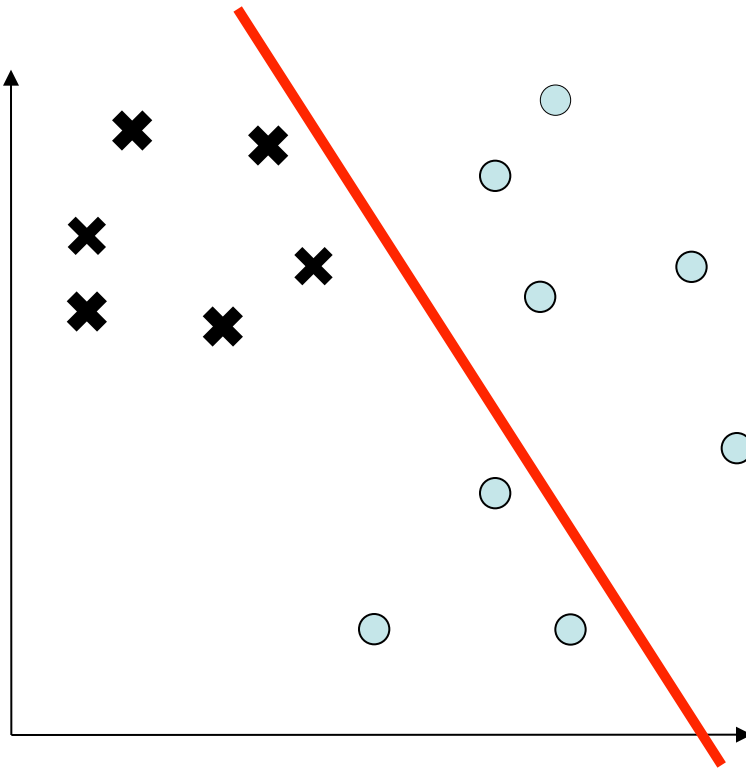


Rugby player = 1
Ballet dancer = 0



Is this a good decision boundary?

$$\text{if } \left(\sum_{i=1}^M x_i w_i \right) > t \quad \text{then } output = 1, \text{ else } output = 0$$

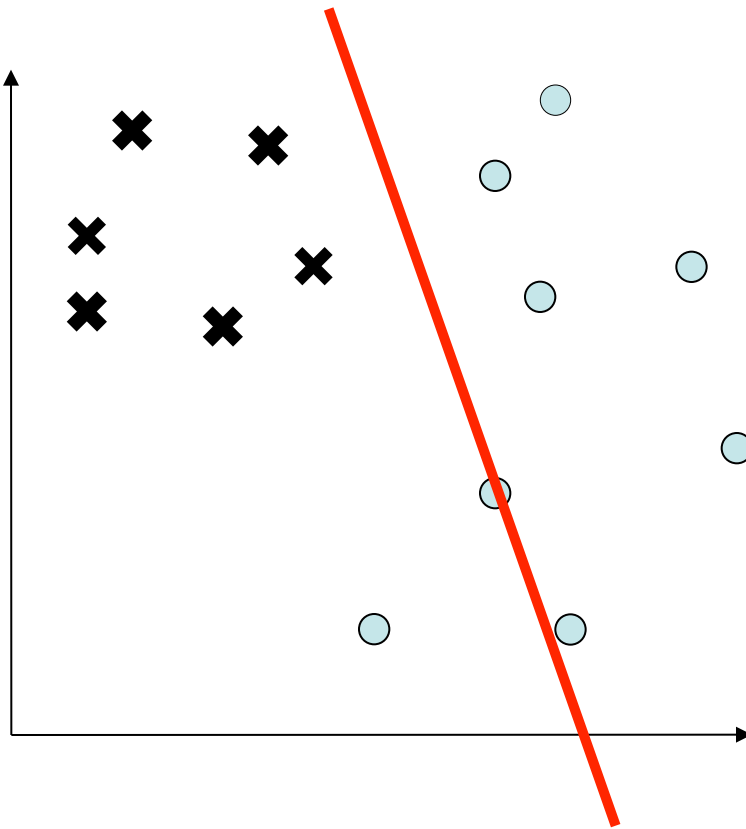


$$w_1 = 1.0$$

$$w_2 = 0.2$$

$$t = 0.05$$

$$\text{if } \left(\sum_{i=1}^M x_i w_i \right) > t \quad \text{then } output = 1, \text{ else } output = 0$$

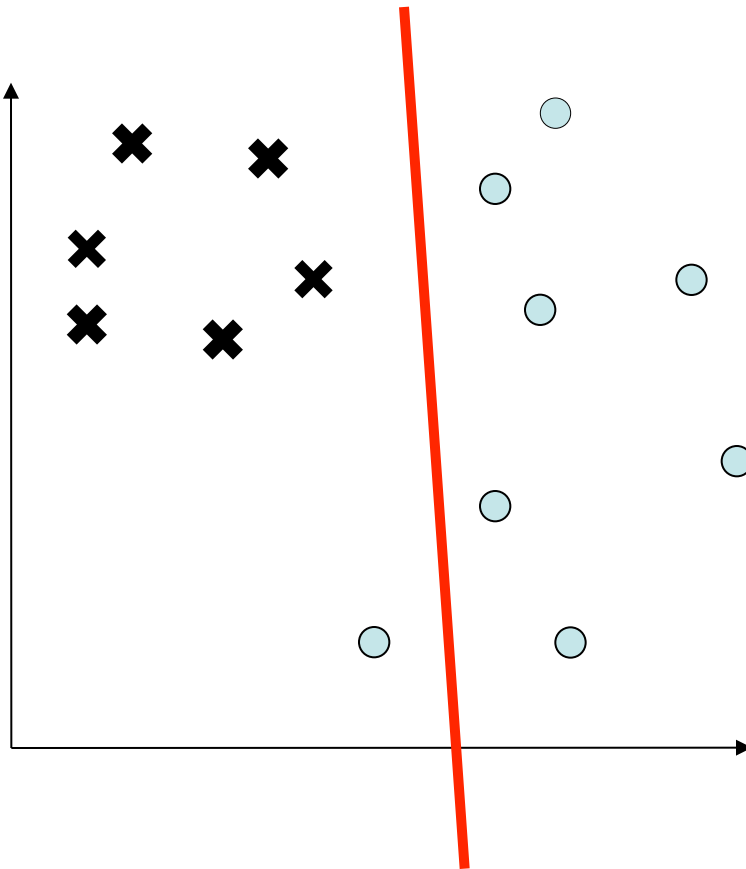


$$w_1 = 2.1$$

$$w_2 = 0.2$$

$$t = 0.05$$

$$\text{if } \left(\sum_{i=1}^M x_i w_i \right) > t \quad \text{then } output = 1, \text{ else } output = 0$$

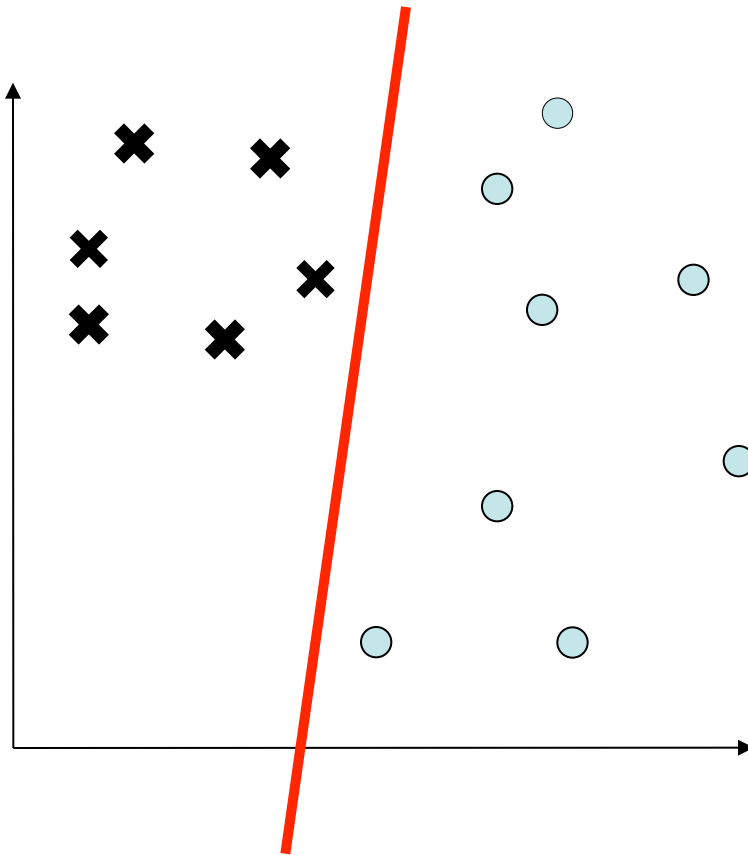


$$w_1 = 1.9$$

$$w_2 = 0.02$$

$$t = 0.05$$

$$\text{if } \left(\sum_{i=1}^M x_i w_i \right) > t \quad \text{then } output = 1, \text{ else } output = 0$$



$$w_1 = -0.8$$

$$w_2 = 0.03$$

$$t = 0.05$$

Changing the weights/threshold makes the decision boundary move.

Pointless / impossible to do it by hand – only ok for simple 2-D case.


We need an algorithm....

$$x = [1.0, 0.5, 2.0]$$

$$w = [0.2, 0.5, 0.5]$$

$$t = 1.0$$

$$a = \sum_{i=1}^M x_i w_i$$



Q1. What is the activation, a , of the neuron?

Q2. Does the neuron fire?

Q3. What if we set threshold at 0.5 and weight #3 to zero?

**Take a 20 minute break
and think about this.**


20 minute break

$$x = [1.0, 0.5, 2.0]$$

$$w = [0.2, 0.5, 0.5]$$

$$t = 1.0$$

$$a = \sum_{i=1}^M x_i w_i$$



Q1. What is the activation, a , of the neuron?

$$a = \sum_{i=1}^M x_i w_i = (1.0 \times 0.2) + (0.5 \times 0.5) + (2.0 \times 0.5) = 1.45$$

Q2. Does the neuron fire?

if (activation > threshold) output=1 else output=0


.... So yes, it fires.

$$x = [1.0, 0.5, 2.0]$$

$$w = [0.2, 0.5, 0.5]$$

$$t = 1.0$$

$$a = \sum_{i=1}^M x_i w_i$$



Q3. What if we set threshold at 0.5 and weight #3 to zero?

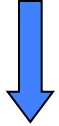
$$a = \sum_{i=1}^M x_i w_i = (1.0 \times 0.2) + (0.5 \times 0.5) + (2.0 \times 0.0) = 0.45$$

if (activation > threshold) output=1 else output=0

.... So no, it does not fire..

We need a more sophisticated model...


if (*weight* > *t*) then "player" else "dancer"



if ($f(\vec{x}) > t$) then "player" else "dancer"

$x_1 = \text{height (cm)}$

$x_2 = \text{weight (kg)}$



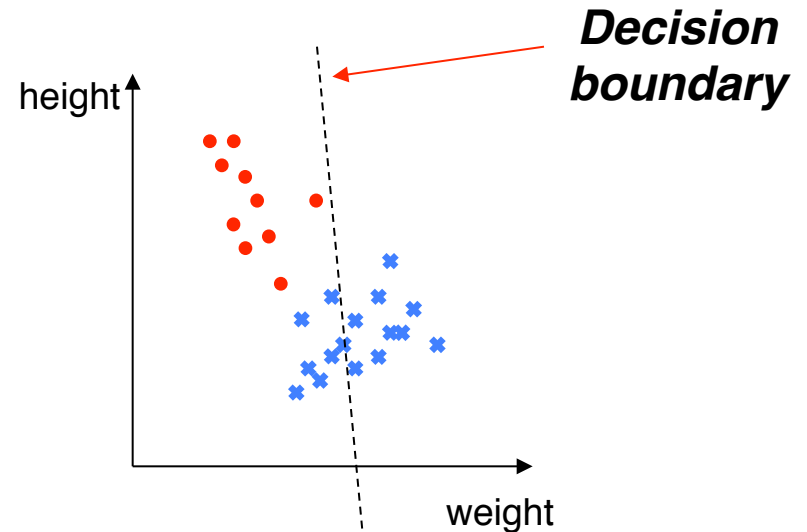
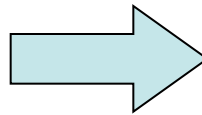
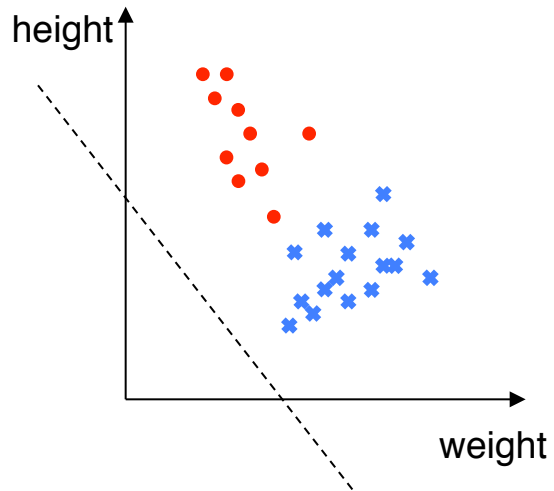
The Perceptron

$$\begin{aligned} f(\vec{x}) &= (w_1 * x_1) + (w_2 * x_2) \\ &= \sum_{i=1}^d w_i x_i \end{aligned}$$

The Perceptron

if $f(\vec{x}) > t$ then "player" else "dancer"

$$f(\vec{x}) = (w_1 * x_1) + (w_2 * x_2)$$
$$= \sum_{i=1}^d w_i x_i$$



w_1 , w_2 and t change the position of the DECISION BOUNDARY

The Perceptron

Model

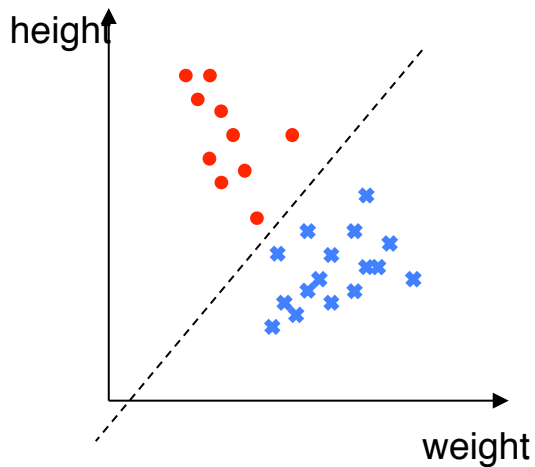
$$\text{if } \sum_{i=1}^d w_i x_i > t \text{ then } \hat{y} = 1 \text{ else } \hat{y} = 0 \quad \begin{cases} \text{"player"} = 1 \\ \text{"dancer"} = 0 \end{cases}$$

Error function

Number of mistakes (a.k.a. classification error)

Learning algo.

??? need to optimise the w and t values...



Perceptron Learning Rule

$$\text{new weight} = \text{old weight} + \underbrace{0.1 \times (\text{trueLabel} - \text{output}) \times \text{input}}_{\text{update}}$$

What weight updates do these cases produce?

if... (**target** = 0, **output** = 0) then update = ?

if... (**target** = 0, **output** = 1) then update = ?

if... (**target** = 1, **output** = 0) then update = ?

if... (**target** = 1, **output** = 1) then update = ?

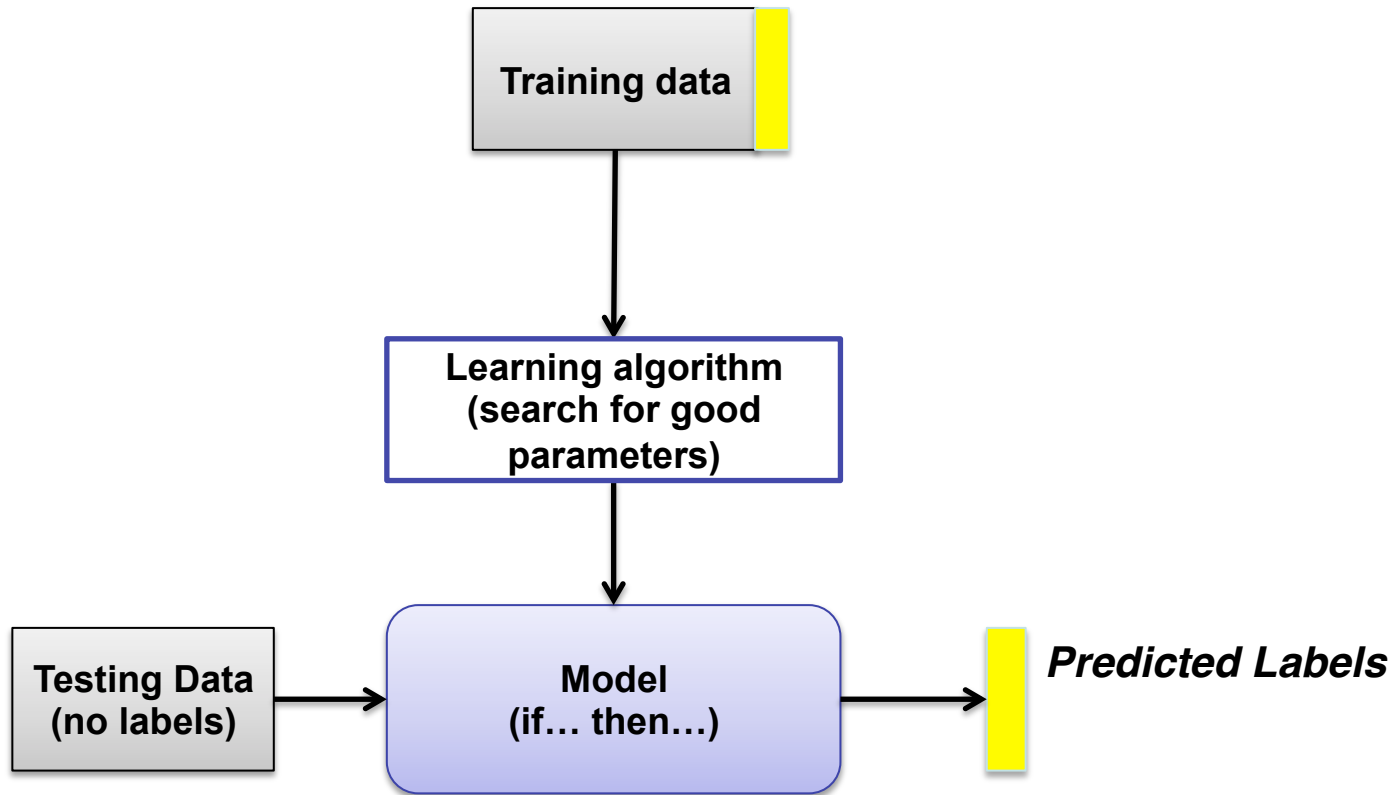
Learning algorithm for the Perceptron

```
initialise weights to random numbers in range -1 to +1
for n = 1 to NUM_ITERATIONS
    for each training example (x,y)
        calculate activation
        for each weight
            update weight by learning rule
        end
    end
end
end
```

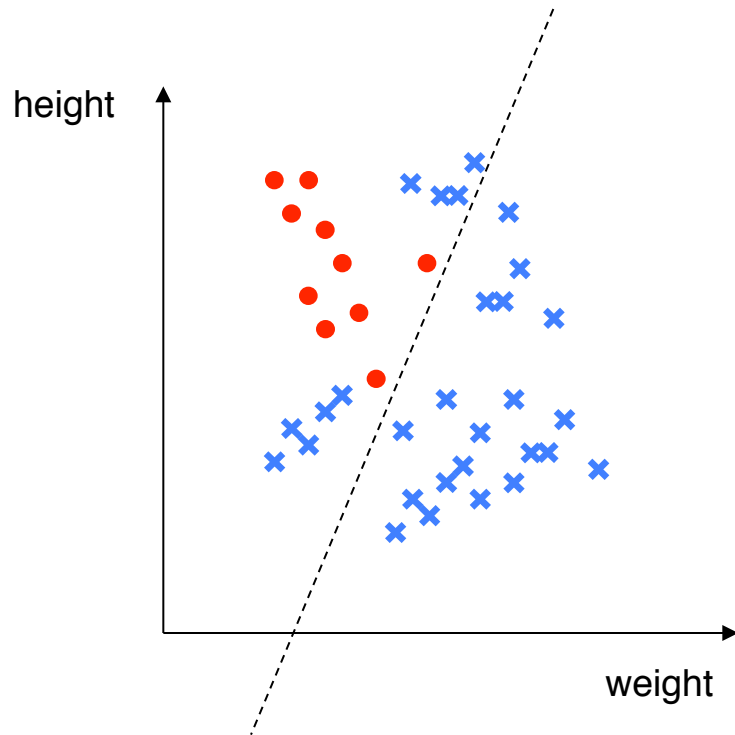
Perceptron convergence theorem:

If the data is linearly separable, then application of the Perceptron learning rule will find a separating decision boundary, within a finite number of iterations

Supervised Learning Pipeline for Perceptron



New data.... “non-linearly separable”



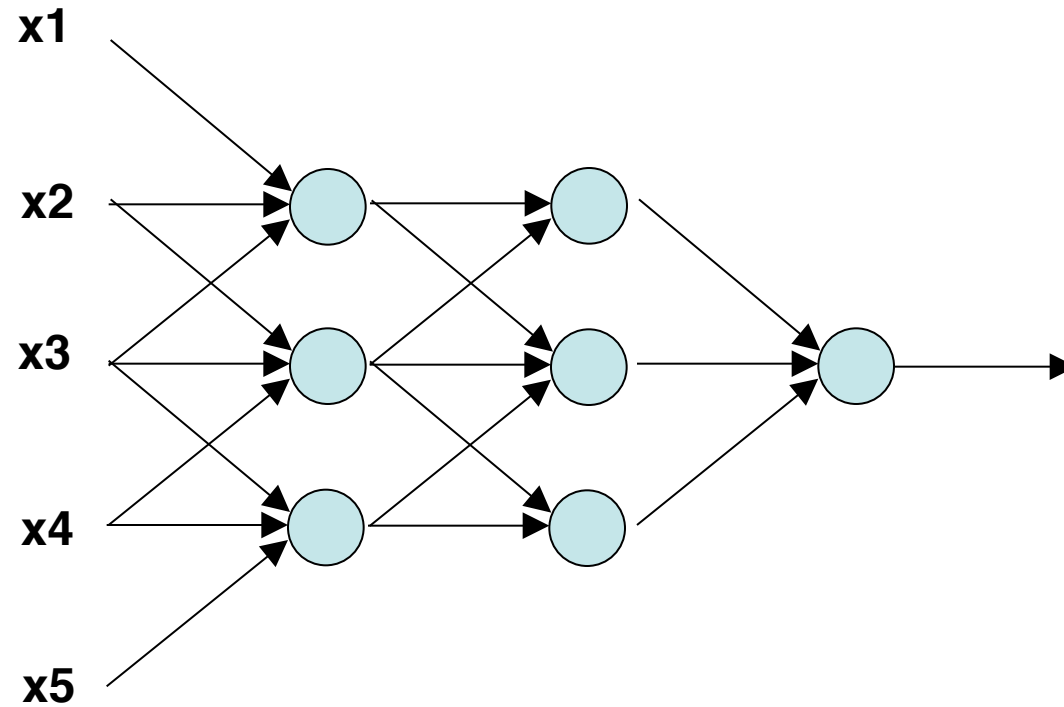
Our model does not
match the problem!

(AGAIN!)

if $\sum_{i=1}^d w_i x_i > t$ then "player" else "dancer"

Many mistakes!

Multilayer Perceptron

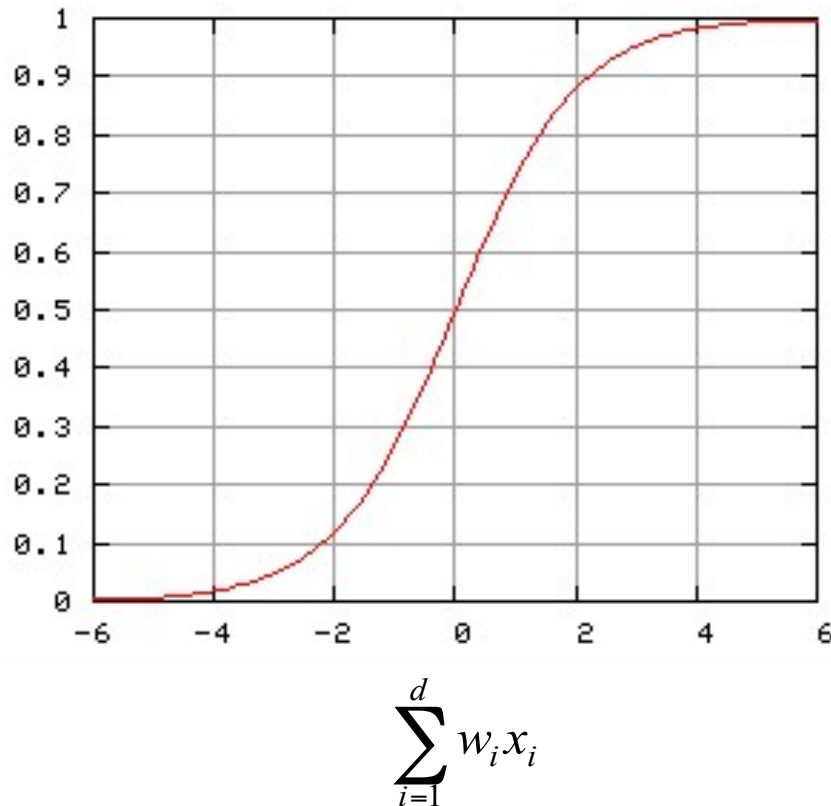


Sigmoid activation – no more thresholds needed 😊


~~if $\sum_{i=1}^d w_i x_i > t$ then $\hat{y} = 1$ else $\hat{y} = 0$~~

$$a = \frac{1}{1 + \exp\left(-\sum_{i=1}^d w_i x_i\right)}$$

activation level



MLP decision boundary – nonlinear problems, solved!



Neural Networks - summary

Perceptrons are a (simple) emulation of a neuron.

Layering perceptrons gives you... a multilayer perceptron.
An MLP is one type of neural network – there are others.

An MLP with sigmoid activation functions can solve highly nonlinear problems.

Downside – we cannot use the simple perceptron learning algorithm.

Instead we have the “backpropagation” algorithm.

This is outside the scope of this introductory course.