

From last time



A file system uses inodes which contain 8 block-numbers. These are for the first 7 blocks of the file and an indirect block, which just contains block-numbers for the remaining blocks in the file.
A block-number occupies 2 bytes. Each block is 4k bytes.

What is the maximum size of a file in bytes?

What is the maximum total size of directories and files in a single disk partition?

COMP25111: Operating Systems
Lecture 18: Windows XP Case Study

Oscar Palomar
School of Computer Science, University of Manchester
Autumn 2016

COMP25111 Lecture 18
Overview & Learning Outcomes

- 1 Background
- 2 Components: layers & managers
 - Scheduling
 - Virtual Memory
 - Input/Output
 - File system

COMP25111 Lecture 18
Goals

- Portability: written in C and C++
(>29M lines – Linux <2M + 1.5M for X windows)
– HAL (Hardware Abstraction Layer) = processor-dependent
- Extensibility: layered architecture; executive in kernel mode.
- Compatibility (Win16, Win32, POSIX, OS/2)
– DOS emulation in “Virtual DOS machine”
- Performance, Scalability, Multiprocessor support
- Reliability
- International support (Unicode etc.)

1/27COMP25111 Lecture 18
History & Motivation

- Early Windows systems (3.1, 95, 98)
 - demonstrated the usefulness of a GUI-based OS
 - designed for “home use” (not secure or multiprocessing)
 - initially 16-bit
- Windows NT:
 - designed as a commercial strength 32-bit system
 - led to Windows 2000, XP, Vista, 7, ...
- Modern Windows:
 - mobile systems (power efficiency, touch vs. mouse, ...)
 - Windows 8, 10

3/27COMP25111 Lecture 18
Unifying Themes

- For non-programmers
- Users (& programmers) see Windows as a single entity
 - GUI built into OS
 - Win32 API hides library/system split
- “Object-Oriented”
 - Handles
 - methods
- Registry

Structure (OSC/J)

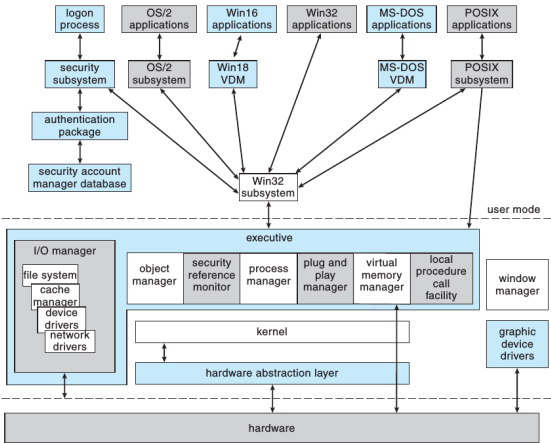


Figure 22.1 Windows XP block diagram.

Architecture – layered system of modules

- Protected mode:
- HAL: e.g. read/write device registers, interrupt handling so independent of bus; timers/clock
 - “kernel”: thread scheduling, interrupt handling, low-level processor synchronisation, recovery after power failure, etc.
 - executive: managers; “system services” distribute system calls to managers
- User mode:
- service processes
 - environment subsystems (Win32, Posix, OS/2)

System services

- System call interface not publicly available
- Well defined Win32 API
- very large library of functions
 - which hide system calls
 - includes GUI
- (Unix: minimal set of system calls, all public)

Object manager

- “object-oriented” – uses objects (“handles”) for all services & entities
- Objects: process, thread, section of memory, open file, comms port, semaphore, timer, registry key, device, driver, ...
- Searchable directory
- Named objects can be shared,
Unnamed objects are private to creating process
- Each object can have Access Control List (ACL)

Configuration manager – Registry

- Hive = name of a registry database
e.g. HKEY_CURRENT_USER (HKCU)
- Subkey/Key = sub/directory
e.g. HKEY_CURRENT_USER\Control Panel\desktop
- entry = (name, type, data) i.e. value
e.g. “ScreenSaveTimeout”=“60”
“ScreenSaveUsePassword”=dword:00000000
- + replaces many scattered configuration files
- hard to copy, share, edit, back-up
 - heavily used, single point of failure

Integrity

Kernel – scheduling

- Job – Process – Thread – Fiber
- process owns resources
 - threads scheduled by OS (system calls)
 - fibers scheduled by library (outside OS)
- Also deals with
- pipes, sockets, rpc ...
 - semaphores ...
 - multi-core CPUs

Scheduling

Pre-emptive

Priorities (0:zero-pages, 1-15:user, 16-31:system)

Selectable time quantum (e.g. 20ms)

Dynamically change user priorities (initial-15 only):

- time quantum expired: -1
- boost priority of threads released from I/O (e.g.+6 keyboard)
- extend quantum of foreground application
- starvation prevention (temporarily boost priority)
 - find threads that haven't run in a while

Cannot guarantee real-time deadlines

Memory Management

assumes H/W support; based on Intel PC memory architecture

2^{32} = 4GB address space; top half=OS, bottom=user pages (e.g. 4kB); no segments, but 64kB boundaries
top & bottom 64kB unused

processes can share pages

- "copy on write"
- "position independent code"

version of LRU

- keeps some pages free, so can always load new page
- dynamically calculated max working set size per process
- lots of heuristics and kludges

can memory-map files

Input/Output

Manager interrogates each slot of each bus

- at boot time (mostly)
 - at any time (e.g. USB) – Plug&Play
- device info → where to find driver

Drivers:

- I/O Request Packet (IRP)
- Object-based (list of methods):
 - init, add device, interrupt, fast I/O, DMA, abandon ...
- unplug
- configurable
- multi-processor safe

Monolithic or “stacked”

NTFS

“clusters” (blocks) e.g. 4KB

File = set of attributes (byte streams – max 2^{64} B)

Master File Table (MFT) = set of 1kb records

0-26: metadata (MFT, log, boot, rootdir, bitmap, ACL, ...)

MFT record = attributes, name, list of (block start & count)
(if needed: list of overflow records)

directory: data = list of (file-name, MFT record, ...)
(large directories use B+ trees)

Integrity: updates within transactions, logging

Compression & Encryption

MFT (MOS4)




Figure 11-39. The NTFS master file table.




Figure 11-41. An MFT record for a three-run, nine-block stream.

Summary of key points

1 Background

2 Components: layers & managers

- Scheduling
- Virtual Memory
- Input/Output
- File system

Your Questions

Glossary

- Hardware Abstraction Layer (HAL)
- DLL
- Win32 API
- Handle
- Registry
- Hive, Key
- Fiber
- NTFS
- Master File Table (MFT)

Reading

- newer OSC/J: Ch 22
- older OSC/J: Ch 21
- MOS: Ch 11