

Overview & Learning Outcomes

- 1 Background
- 2 Shell
- 3 Components: layers & managers
 - Scheduling
 - Memory
 - Files
 - Input/Output

COMP25111: Operating Systems

Lecture 19: Linux Case Study

Oscar Palomar

School of Computer Science, University of Manchester

Autumn 2016

COMP25111 Lecture 19

UNIX History & Motivation

1/31COMP25111 Lecture 19

Unifying Themes

2/31

(weak pun on “MULTICS”)
Initially (1969) single-user, soon (1973) multi-user timesharing system
Written in C (developed at Bell Labs to support Unix)
Fundamental Architectural Design finished in 1978
POSIX standards 1988, 1992,3,5, 2001,4,8

by programmers, for programmers
small, modular, clean design
UI consistency, brevity
source distribution

Everything is a file

Composition & re-usability:
– shell, I/O redirection, pipes, filters

Simplicity & minimality:
– 1 reusable best tool for each purpose

COMP25111 Lecture 19

Background

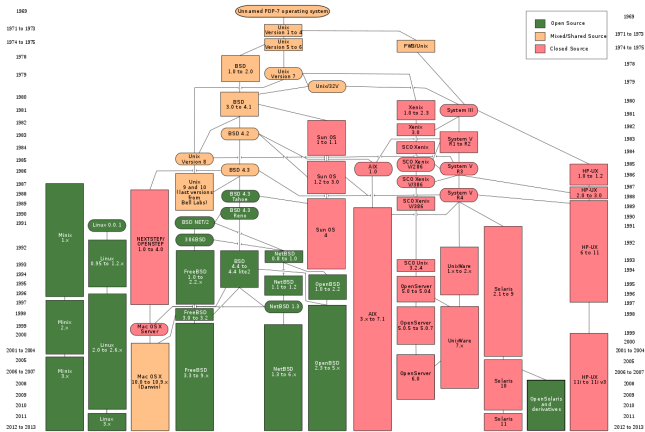
4/31COMP25111 Lecture 19

Background

5/31

Development (wikimedia: Unix_history-simple.png)

Shell



user-level process, executes programs
(command interpreter, CLI)

– reads next user command
– searches path for program
– forks child process & execs program
– waits for termination of child

COMP25111 Lecture 19

Background

6/31COMP25111 Lecture 19

Shell

8/31

fork & exec (MOS figs 10-4, 10-7)

```
pid = fork();
if (pid < 0)      { /* error */ }
else if (pid > 0) { /* parent */ }
else /* pid==0 */ { /* child */ }

while (1) {
    type_prompt();
    read_command(&command, &params);
    pid = fork();
    if (pid < 0)      { /* error */ }
    else if (pid > 0) waitpid(-1, &status, 0);
    else /* pid==0 */ execve(command, params, 0);
}
```

I/O redirection

processes start with 3 open files: standard input, output, error

Can redirect to/from files:
e.g. compile <input >output 2>errors
or compile <input &>combined-output

or another process in a pipeline:
e.g. ls -la | grep Nov | grep 23
pipe() system call

Interfaces (MOS fig. 10-1)

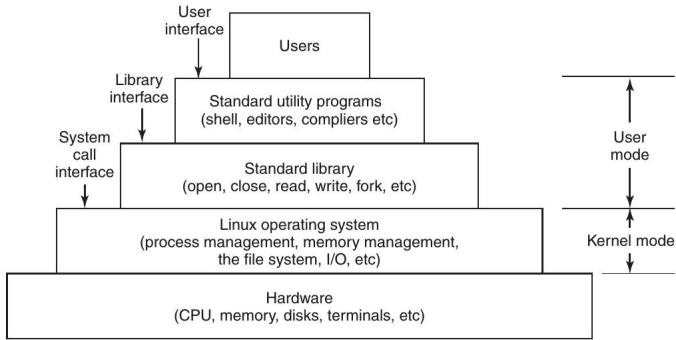


Figure 10-1. The layers in a Linux system.

Architecture – Overview

User level (non-privileged):
user processes = application, library etc.

Programmer Interface:
– System programs (e.g., mkdir, rm, cp, ...)
– System calls (file manipulation; process control; information)

Kernel level (privileged):
managers (file, process, memory, ...) & device drivers

Kernel layers (MOS fig. 10-3)

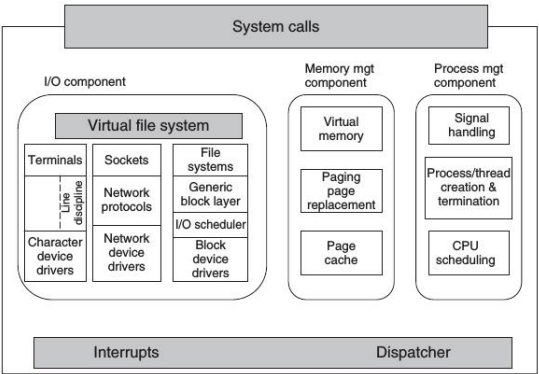



Figure 10-3. Structure of the Linux kernel

Process

“process descriptor” (see handout 5 – PCB)
– task_struct
– unique PID
– address space
– UFIDs (handout 16)
– signal handling vector
– UID & GIDs (User/Group ID)
– scheduling priorities
– thread(s)

Properties inherited from parent
– fork, execve, clone...

Scheduling




(b) Per CPU red-black tree in the CFS scheduler
MOS4 10-10

- Priorities
- “Real-time” high priority (0-99), FIFO or Round-Robin
 - Timesharing: low priority (100-139)
 - ±20 (nice)
- Completely Fair Scheduler – translates priorities into “speed”
- O(1) in older versions

Delayed jobs:
at 0630 myjob for one-off deferral
cron for regular jobs (hourly, daily, weekly, monthly, etc)

Paging

- Demand-driven
Pages pre-cleaned
- LRU approximation with second-chance
Reference bit (set when page accessed)
Round-robin check pages (& clear ref)




(www.redhat.com/magazine/001nov04/features/vm)

inode


- (see handout 16)
- Owner, last time info, permissions, size, links to the file
- 15 content pointers to disk blocks:
- 12 pointers to direct blocks
 - 1 single-indirect; 1 double-indirect; 1 triple-indirect
- ext2: groups inodes, which point to nearby blocks
- ext3: + journaling
ext4: journaling, +extent

Virtual Memory

- Paged (e.g. 4kB – now also huge pages)
“segments”: Text (code), Data (& heap), Stack
x86: 1GB Kernel, 3GB user; x86-64: 128TB each
can be shared; copy-on-write
memory-mapped files
Page Tables (MOS fig 10-17)
– 2-level: Intel x86
– 3-level: DEC Alpha (43-bit virtual addresses); x86-PAE
– 4-level: Intel x86-64 (48-bit virtual addresses)



mount (MOS4 fig 10-25)



Protection Model

- Process Concepts: UID GID
- File Attributes: RWXRWXRWX
Owner, Group, Other – Read, Write, Execute
- Primarily files, directories
- Same mechanisms used by devices, network connections, ...
- Security:
user particulars in /etc/passwd
Holes: read about the Morris Internet Worm (1988)

I/O and Device Drivers

Drivers are privileged code (not user-supplied)

user-access to devices via special files:
e.g. /dev/fd /dev/tty

Can be character or block devices

Summary of key points

- 1 Background
- 2 Shell
- 3 Components: layers & managers
 - Scheduling
 - Memory
 - Files
 - Input/Output

Your Questions

Exam Questions

In Unix, what is the use of the shell variable PATH? (2 marks)

Briefly explain how a shell implements a pipe between two commands (2 marks)

Briefly explain how Unix implements input redirection in a shell command. (2 marks)

Briefly explain that a Unix shell does to execute the following command /bin/who > myfile (2 marks)

Glossary

I/O redirection
shell
pipe
filter
fork
exec (execve)
waitpid
2-level (high & low) scheduling
character & block devices

Reading

newer OSC/J: Ch. 21

older OSC/J: Ch. 20

MOS: Ch. 10

<https://www.kernel.org/doc/>

David A Rusling, The Linux Kernel
<http://tldp.org/LDP/tlk/tlk-toc.html>
(Linux 2.0.33, 1999)