

COMP25111: Operating Systems

Lecture 3: Computer Architecture – MU0 Control Signals

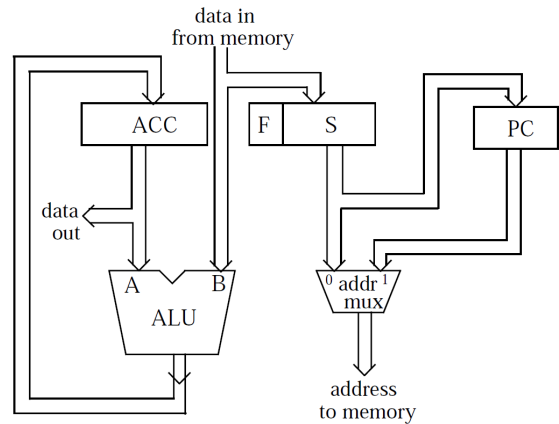
Will Toms

School of Computer Science, University of Manchester

Autumn 2016

From last time – fetch phase

Shade in the path usage for the fetch phase on the MU0 datapath diagram below:



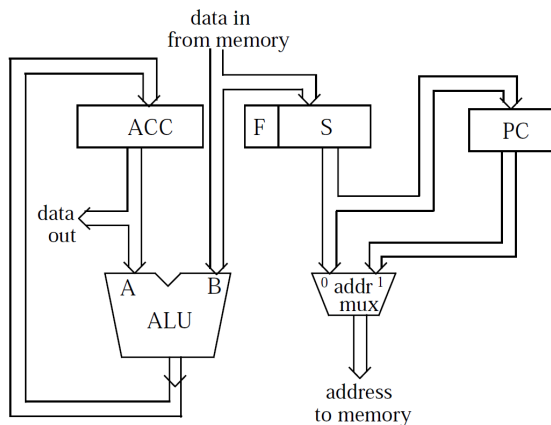
COMP25111 Lecture 3

1/12COMP25111 Lecture 3

2/12

From last time – LDA execute phase

Shade in the path usage for the execute phase for the LDA instructions on the MU0 datapath diagram below:



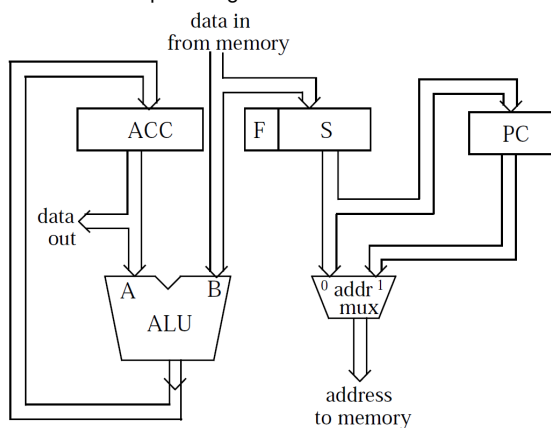
COMP25111 Lecture 3

3/12COMP25111 Lecture 3

4/12

From last time – JMP execute phase

Shade in the path usage for the execute phase for the JMP instruction on the MU0 datapath diagram below:



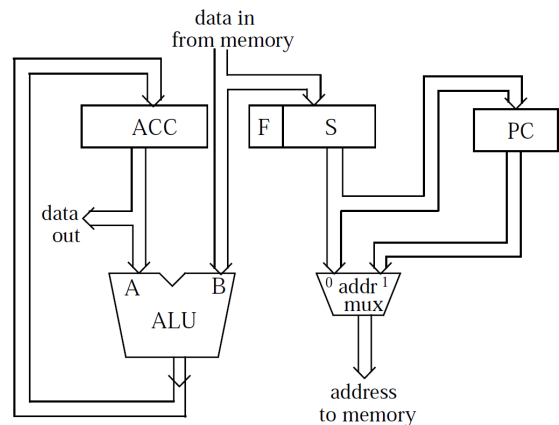
COMP25111 Lecture 3

5/12COMP25111 Lecture 3

6/12

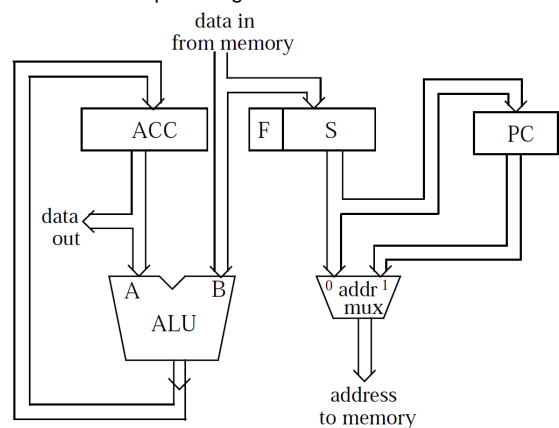
From last time – STA execute phase

Shade in the path usage for the execute phase for the STA instruction on the MU0 datapath diagram below:



From last time – ADD execute phase

Shade in the path usage for the execute phase for the ADD instruction on the MU0 datapath diagram below:



Overview & Learning Outcomes

MU0 Control Signals

Introduction to lab 1

Verilog

What Control Signals?

Push: values from registers & devices always available
Control the actions registers & devices perform on inputs

Each register needs an enable, to allow it to be written to: En_ACC, En_PC, En_IR

Multiplexer needs a signal to select an input

Memory needs actions: Ren (Read Enable) & Wen (Write Enable)

ALU needs actions: add, sub, & byp (bypass)

ALU control signals

We are using three signals (add, sub & byp)

At most one of the three can be set during any phase.

In theory two (encoded) bits would do

Decode either in Control (3 signals) or in ALU (2 signals)

MU0 Control

A “black box”

with these inputs:
external signals: clock reset
from datapaths: F[3:0] N Z
(N & Z come from ACC, needed for JGE & JNE)

and these outputs:
within CPU: En_IR En_PC En_ACC byp add sub addr_Mux
to memory: Ren Wen

Lab 1: What is inside the “black box”?

Must create output signals for each phase of each instruction

These derive from the 6 input bits (F[3:0] N Z)

Some outputs can be “don’t care”
(may simplify the logic circuit)

e.g. ALU action during a JMP

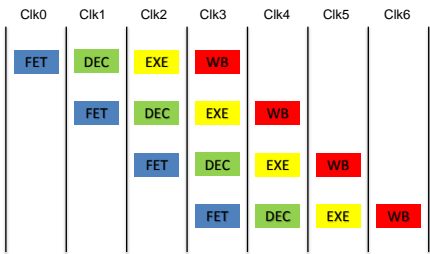
(indicate by “X” instead of “1” or “0”)

Never make register enables or Wen don’t care!



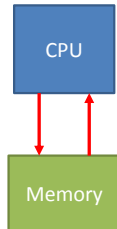
Pipelining

- Speed up execution by operating phases in parallel
- Pentium 4 had 30 pipeline stages
- Power consumption too great
- Multiple simplified cores (Core2 and successors)



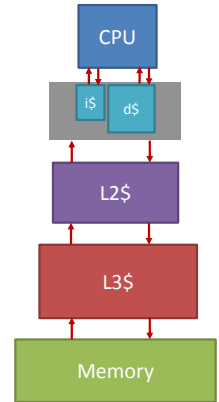
Von Neumann Architecture

- Unified Memory Model
- Instructions and Data reside in the same space
- CPU reads and writes directly to memory



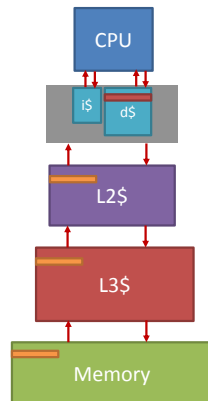
Caches

- Memory much slower than CPU
- Caches exploit
 - Temporal Locality
 - Locations used more than once
 - Spatial Locality
 - Things close together used at the same time



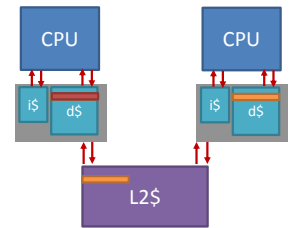
Caches

- Line Size 64 Bytes
- Hit Latency:
 - L1i 1 Cycle
 - L1d 2 Cycles
 - L2 12 Cycles
 - L3 20 Cycles
- Main Memory ~100cycles
- Written (dirty) lines can reside in the cache until they need to be written back to memory



Cache Coherence

- With multiple cores copies can be out of sync
- Must *invalidate* other copies before write
- Different applications can overwrite each others data
- Some applications caching not beneficial



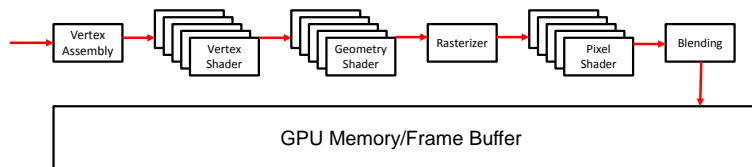
Energy-Efficiency

- CPUs unlikely to go faster
- ICs unlikely to get bigger
- Power Consumption biggest factor
 - Data Centre's account for 2% of global energy usage
 - Battery Life biggest limiting factor in mobile computing
- Technology will become more efficient
- Specialised compute resources for common tasks (*accelerators*)

Energy-Efficiency

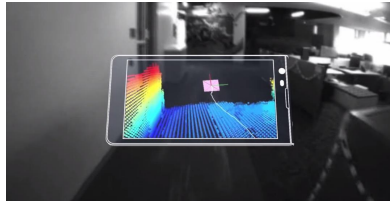
- CPUs unlikely to go faster
- ICs unlikely to get bigger
- Power Consumption biggest factor
 - Data Centre's account for 2% of global energy usage
 - Battery Life biggest limiting factor in mobile computing
- Technology will become more efficient
- Specialised compute resources for common tasks (*accelerators*)

GPUs

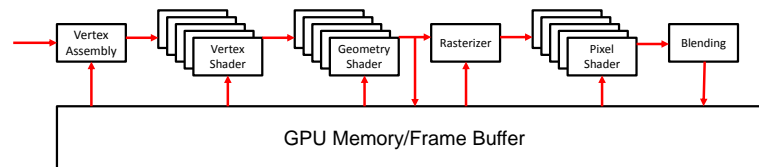


- Graphics Coprocessors have been around since 1980's
- Multi-stage pipelines with many parallel multipliers
 - Great for matrix multiplication tasks

Accelerators



GPGPUs



- Iterative Pipelines allow general purpose computations
- GPUs now coherent
- Difficult to program

Post Von-Neumann?

- Von Neumann model still remain
- Interaction with accelerators via API's/Data-Types and Design Patterns
- Not all designs will map well to heterogeneous architectures
- **Data Movement** most expensive operation