

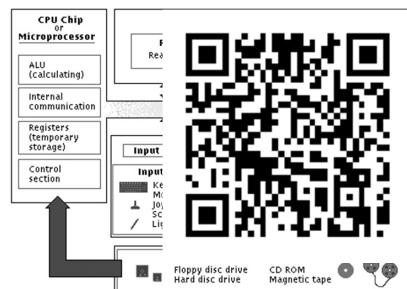


Learning; comprehension; & introspection

COMP25111

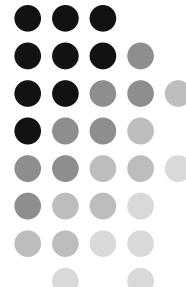
Operating Systems

Lectures 15: Controlling Input and Output 2



Week

9



© Copyright Richard Neville 2009\10

Dr Richard Neville

r.neville@manchester.ac.uk

Room: G12 Kilburn Building,

Bottom floor

NOTE: The up-to-date version of this lecture is kept on the associated web site – available [on-line] @
Blackboard select: COMP25111 Introduction to Computer Systems www.manchester.ac.uk/portal

<http://www.cs.man.ac.uk/~neville/COMP25111/Lecture15/Lecture15.html>

1



Learning; comprehension; & introspection

Where to find this Lecture 10 of the COMP25111 course?

First Go to Blackboard 9; then select: [COMP25111 Operating Systems](#)

Week 7

Then select:

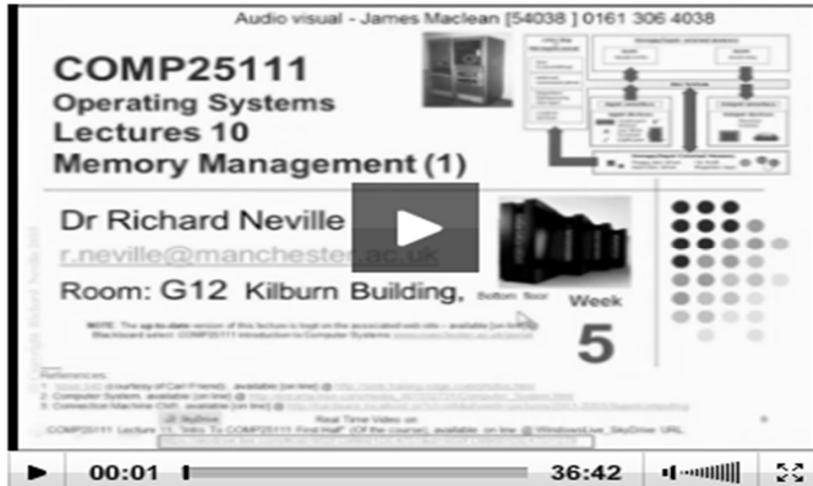
This topic provides...
10: Memory management 1 (Introduction to basics) by RN;
11: Memory management 2 (Virtual Memory (1)) by RN;

Then select: [Lecture 10 Information](#)

Then select: [Real Time Video of Lecture 10](#)

© Copyright Richard Neville 2009\10

Then select:



1. Question

List a number of devices that connect to the system bus, via I/O modules.

ANSWER(S):

- Answer(s):

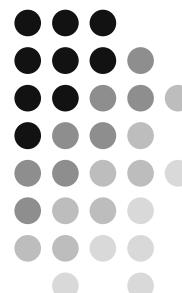
- NOTE: In the exam approximately 2 question are taken from the topics (and program examples) coved in each lecture

COMP25111

Systems Architecture 1

Lectures 15

Controlling Input and Output 2



This material is presented to ensure timely dissemination of Lecture materials. Copyright and all rights therein are retained by authors (i.e. © Copyright Richard Stuart Neville.). All persons copying this information are expected to adhere to the terms and constraints invoked by each author's copyright. These works may not be reposted without the explicit permission of the copyright holder © Copyright Richard Stuart Neville.
Personal use of this material is permitted. However, permission to reprint/republish this material for advertising or promotional purposes or for creating new collective works for resale or redistribution to servers or lists, or to reuse any copyrighted component of this work in other works must be obtained from © Copyright Richard Stuart Neville.
Proper referencing of this material is essential. The expected norm for referencing the material is:
Citation [reference in body of text] (Harvard style): (Neville, 2010).
Reference: Neville, R., (2010). Lecture notes (and all associated materials) for COMP15111 Fundamentals of Computer Architecture; Lecture series, developed and presented by R. Neville.

Learning Outcomes

You should be able to:

D-words

C-words

- 1) Describe what a processor does in response to an **interrupt**;
- 2) Discuss what is meant by interrupt-driven I/O; &
- 3) Explain what is meant by direct memory access and how it is useful in I/O;
- 4) Exams and Revision; and
- 5) L3 Week 12 Hints.

Footnote

1: Describe: aligned to paraphrase, discuss & give example [pictorial or diagrammatic].
2: Discuss: aligned to describe, paraphrase, discuss & give example..
3: Explain: aligned to describe; discuss; give examples.

FOOTNOTE2: D-words: direction words. C-Words: content words.
Ref.: Michael J. Wallace (1980, 2004)
Study Skills in English, ISBN 9780521537520.
D-Words also aligned to: Ref.: *Taxonomy of Educational Objectives: The Classification of Educational Goals*; pp. 201–207; B. S. Bloom (Ed.)
Susan Fauer Company, Inc. 1956.

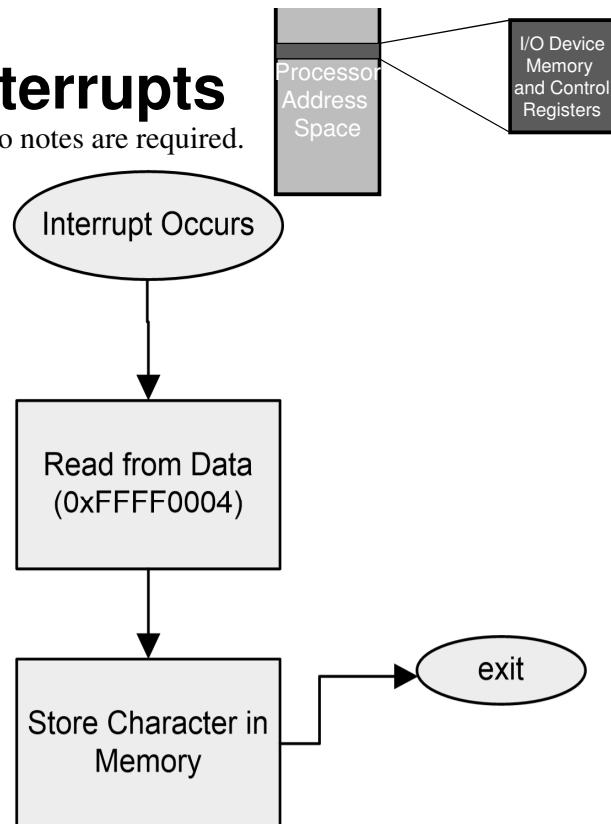
Interrupts

- Interrupts have a wide variety of uses.
- We will focus on a hardware interrupt generated by an external device.
- So, let's give our keyboard a connection to our processor's interrupt line.
- Our keyboard will send an interrupt (change a wire from '0' to '1') when it has a character ready to send.

Flow Chart with Interrupts

A few additions to notes are required.

- There is **no need** to read the **status register**, the interrupt indicates that a character has arrived, so we can simply read from the **data register**.
- Note that every time a character arrives the processor will be 
- There is no time wasted checking the status register when no character has arrived.



References:

1: Lecture 9. Review Question. 2. How is *Memory mapped IO* used in ECE412. The PCMCIA attribute memory and common memory is mapped into the kernel space by the ..., available [on line] @ www.cs.rice.edu/~gw4314/lectures/sram-dram.ppt. (Last date accessed 28-11-2010).

What happens in an interrupt? (1)

Steps the processor takes:

- 1) The I/O device will signal that an interrupt has occurred by using the **interrupt line**;
 - 2) The processor finishes executing its current instruction; and then
 - 3) The processor then sends a special type of memory read (called an interrupt acknowledgement – **IACK**);
- The device that has interrupted the processor will respond with a value that indicates its identity.
 - This is required because several devices may need to interrupt the processor, the processor needs to know **which** device has sent the interrupt – the code that is run in response to an interrupt from the keyboard is different to that to handle an interrupt from a hard disk.

What happens in an interrupt? (2) Cont.

- 4) The processor saves the current value of the Program Counter and the contents of its registers into memory.
- The registers are memory locations inside the processor that it uses for [fast] calculations.
 - It [the registers] are a bit like the use of the output of a calculator but with several rather than just one.
 - The value of the PC and registers are stored so that after the interrupt has been serviced, the program continue from the point at which it was interrupted.

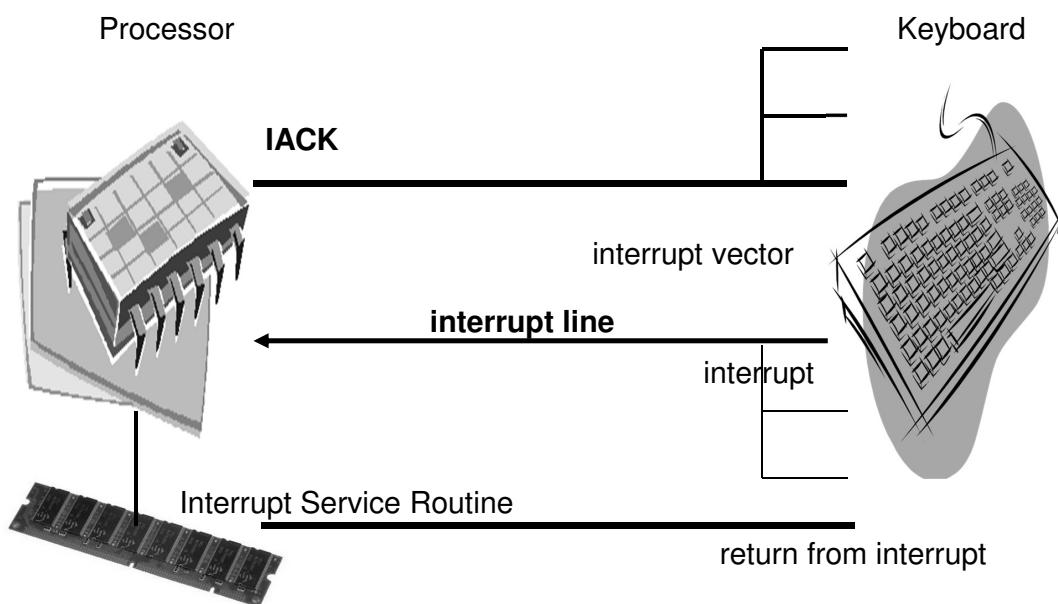
What happens in an interrupt? (3) Cont.

- 5) The value that the processor received from the IACK cycle identifies what device interrupted the processor;
- This is known as the interrupt vector.
- 6) The interrupt vector is used to access a table that holds the starting addresses of all programmes that handle interrupts;
- The interrupt vector is used so that we get the starting address for the appropriate device.
 - This program is called an Interrupt Service Routine (ISR) and contains code that handles the type of interrupt.
- 7) This starting address of the ISR is loaded into the PC.
- 8) This code is now executed until a special instruction (return from interrupt) is executed;
- This marks the end of the Interrupt Service Routine.

What happens in an interrupt? (4) Cont.

- 9) When the return from interrupt instruction is reached:
 - The PC and registers that were stored earlier are loaded back into the processor.
- 10) The program that was executing earlier continues as if nothing has happened.
 - The character that has been placed into memory may be used later – or rather than place it in memory, the ISR might have processed it in some way.

Diagrammatic interrupt sequence:



Interrupt summary

interrupt acknowledgement – IACK

This summary sequence
may be worth
remembering...

- 1) External line ~~interrupts~~ processor;
- 2) IACK cycle identifies the interrupting device;
- 3) Processor accepts interrupt after current instruction;
- 4) The processor stores the information necessary to restart the original program following the interrupt;
- 5) ISR is run for interrupting device until return from interrupt instruction is reached; then finally...
- 6) Stored information is reloaded into the processor, processor continues executing the original program as if nothing had happened.

Interrupt Service Routine (ISR)

13

Interrupts in I/O

- The interrupt is very useful in our keyboard example, rather than periodically checking to see whether a key has been typed by reading the status register, the processor only executes code to get characters when one has been typed;
- This removes the need to periodically check the status register of the I/O device.
- Interrupts are even more useful in the case of mass storage devices.

Interrupts in I/O

- Suppose that a processor wants to write to a disk –this may take as much as 15ms before the sector is located;
 - This process can be speeded up utilising direct memory access (DMA).
- **The 4-steps DMA process is:**
 1. Processor inform the disk DMA I/O - to write data
 - by writing to a **command register** in the DMA I/O device;
 2. DMA controller starts write process;
 3. Process gets on with another work [process] until;
 4. DMA device is finished;
 - which is signalled by an **interrupt**.
 - [This explicitly is where the **interrupt I/O** comes into play...]



Flowcharts

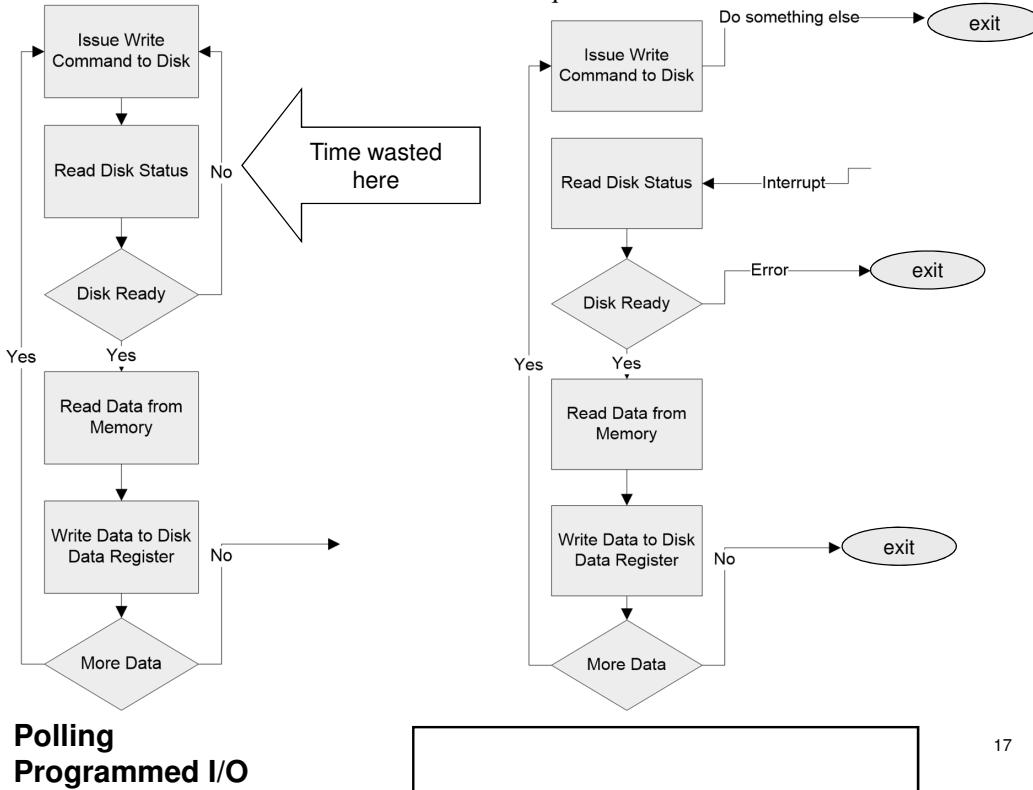
- The next slide shows flowcharts for using programmed and interrupt-driven I/O for a disk where multiple writes are required;
 - i.e. we will be writing several multi-byte words, requiring several data transfers.
- We assume the disk I/O has a:
 - 1) Command register where commands such as:
 - READ, WRITE, FORMAT can be written;
 - 2) Status register that indicates whether the disk is ready for the data or not;
 - This is similar to the status register in the keyboard example.

LC Learning; comprehension; & introspection

Disk writing flowcharts

A few additions to notes are required.

© Copyright Richard Neville 2009\10



17



Learning; comprehension; & introspection

© Copyright Richard Neville 2009\10

Summary

- Interrupt-driven I/O is more efficient than programmed I/O, less time is wasted;
 - But both schemes rely on the processor to transfer the data; &
 - Data rate is limited by the processor's ability to read/write to/from the I/O device.
- For large volumes of data (e.g. disks) Direct Memory Access is performed.

18

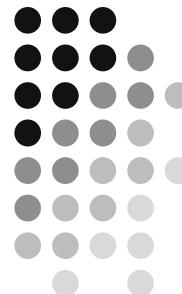
COMP25111

Systems Architecture 1

Lectures 15 Cont ...

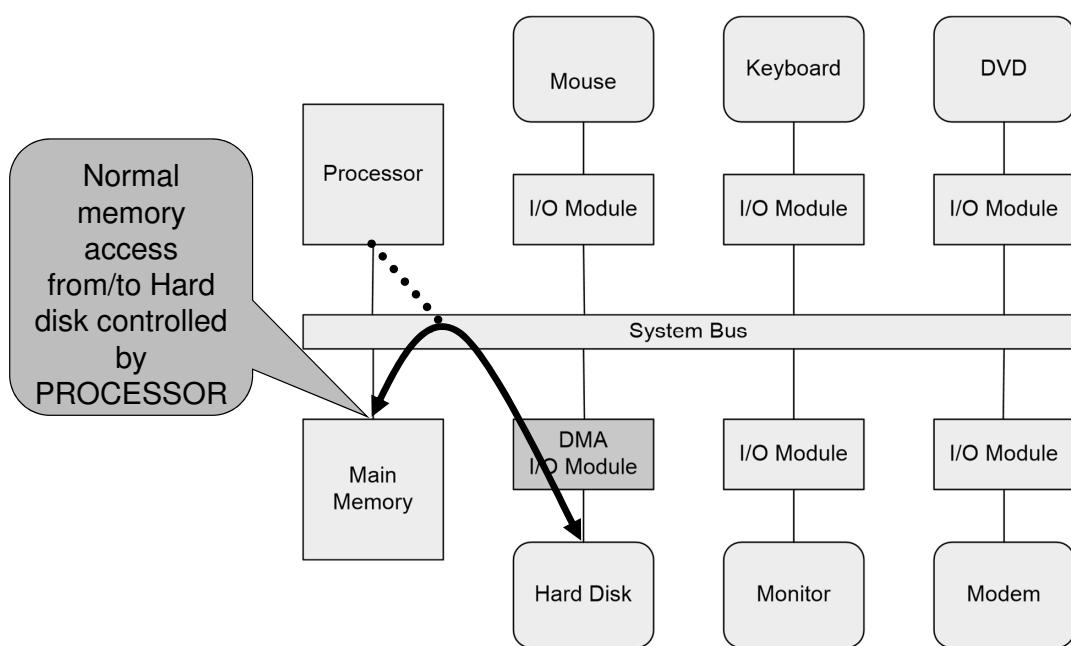
DMA

Direct Memory Access



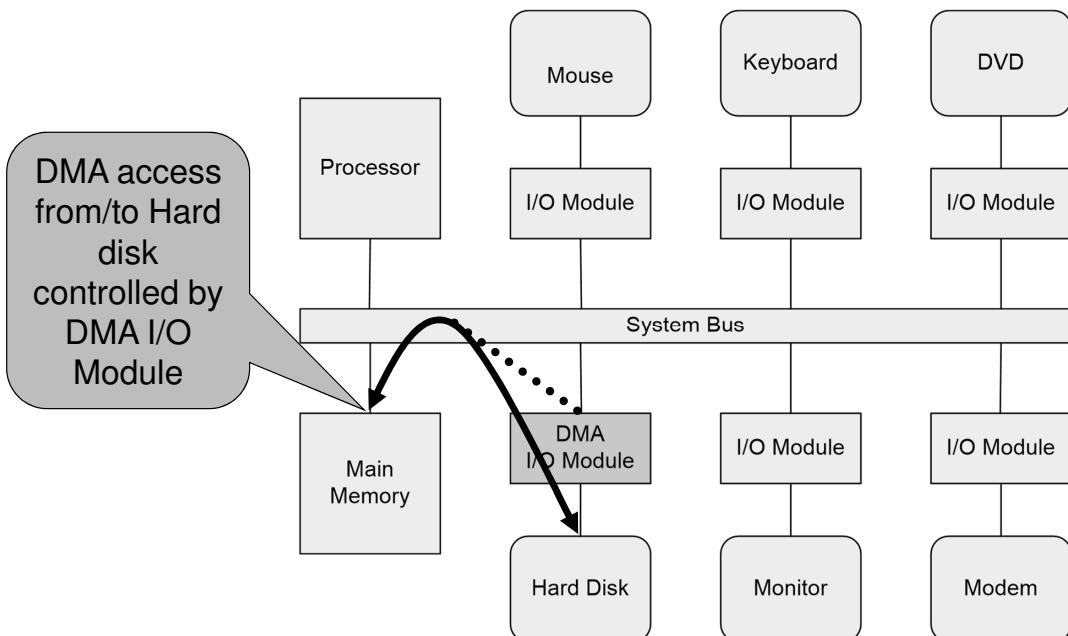
19

Normal I/O



20

DMA I/O module



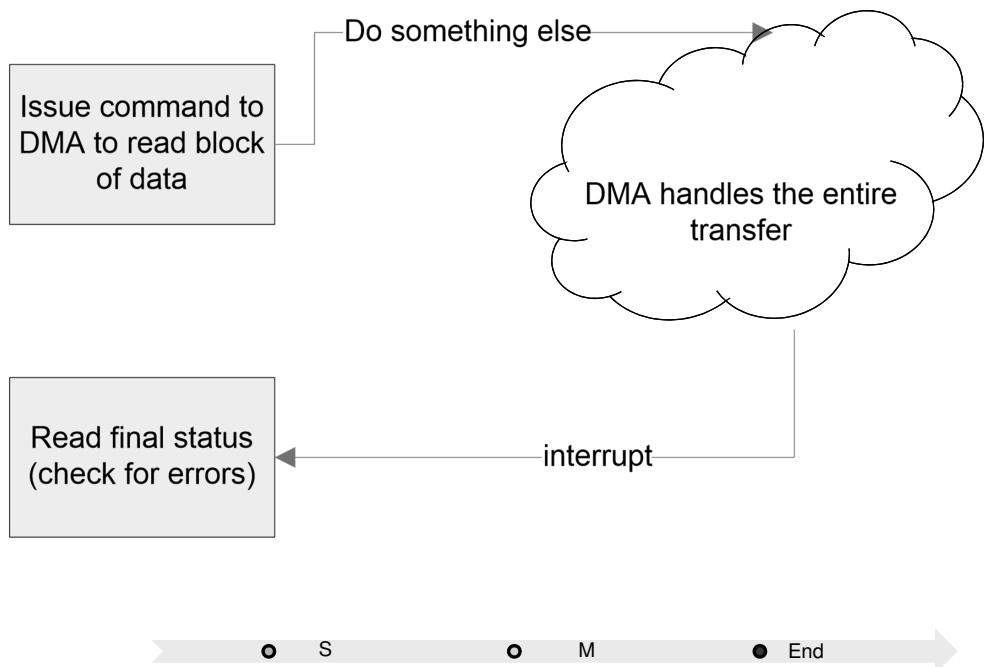
21

DMA

- A DMA device is capable of reading from or writing directly to memory in the same way as a processor does.
- It is optimised for transferring blocks of data into the memory system.
- In our example of reading from a disk, the DMA will handle the transfer of a whole block of data **without** processor intervention.

22

DMA Flowchart



23

END

This material is presented to ensure timely dissemination of Lecture materials. Copyright and all rights therein are retained by authors (i.e. © Copyright Richard Stuart Neville.). All persons copying this information are expected to adhere to the terms and constraints invoked by each author's copyright. These works may not be reposted without the explicit permission of the copyright holder © Copyright Richard Stuart Neville.

Personal use of this material is permitted. However, permission to reprint/republish this material for advertising or promotional purposes or for creating new collective works for resale or redistribution to servers or lists, or to reuse any copyrighted component of this work in other works must be obtained from © Copyright Richard Stuart Neville.

Proper referencing of this material is essential. The expected norm for referencing the material is:

Citation [reference in body of text] (Harvard style): (Neville, 2010).

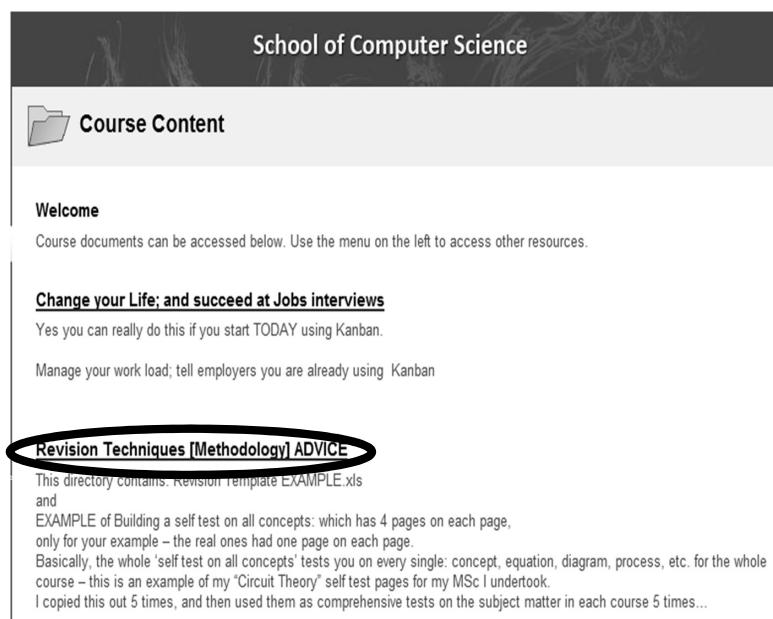
Reference: Neville, R., (2010). Lecture notes (and all associated materials) for COMP17022 Introduction to Computer Systems; Lecture series, developed and presented by R. Neville.

24

•See •Blackboard 9

Li Learning; comprehension; & introspection
Why Didn't Someone Tell Me How To Revise

- If you do not have a comprehensive revision
- Methodology.
- Try the:
- Located on Blackboard 9.
- In the assessment directory.



The screenshot shows a Blackboard 9 course content page. At the top, it says "School of Computer Science". Below that is a "Course Content" section with a folder icon. Underneath is a "Welcome" section with the text: "Course documents can be accessed below. Use the menu on the left to access other resources." There is also a link titled "Change your Life; and succeed at Jobs interviews" with the subtext: "Yes you can really do this if you start TODAY using Kanban." Below that is another link: "Manage your work load; tell employers you are already using Kanban". At the bottom, there is a section titled "Revision Techniques [Methodology] ADVICE" which is highlighted with a black oval. The text in this section reads: "This directory contains: Revision Template EXAMPLE.xls and EXAMPLE of Building a self test on all concepts: which has 4 pages on each page, only for your example – the real ones had one page on each page. Basically, the whole 'self test on all concepts' tests you on every single: concept, equation, diagram, process, etc. for the whole course – this is an example of my 'Circuit Theory' self test pages for my MSc I undertook. I copied this out 5 times, and then used them as comprehensive tests in each course 5 times..."

Why Didn't Someone Tell Me How To Revise

- If you do not have a comprehensive revision
- Methodology.
- Try the:
- Located on Blackboard 9.
- In the assessment directory.

© Copyright Richard Neville 2009\10



Revision Template EXAMPLE.xls

[ExampleOfBuildaSelfTestOnAllConcepts.pdf](#)

COMP19512 - Lecture 13

27

Exams and Revision

• Help a

- Blackboard

In c

Why Didn't Someone Tell Me How To Revise	
Specific revision guidance for COMP1702 Introduction to Computer Systems	
01 Re-listen to all lectures - boot up Audio recording for lecture - then select lecture - click through lecture while listening to Lecture audio...	
11 Re-read lecture notes THEN Undertake Q3's for each lecture	
21 Visit hardest concept(s) THEN Undertake all questions at the end of each lecture	
31 Compile (your own) summary sheets for each lecture	
41 Compile list of keywords for each lecture	
51 compile mind maps for each lecture	
61 Go through past exam - see module guide = past exam in module guide [at end]	
71 Compile a possible future exam (virtual) - using pattern of last - sit un-timed	
81 Sit virtual exam timed	
91 Build self test sheets - these are a set of blank pages that you write QUESTION and ask yourself - the concepts - the equations - the theory - for all concepts and theory of whole course - copy these out 5 time to test yourself on ALL the theory the course covers	

A comprehensive revision schedule is:

- 11 Read & comprehend lecture notes
- 21 Visit hardest subject and reread notes
- 31 Compile summary sheets
- 41 Compile keywords
- 51 Compile mind maps (revise this if you do not find them useful)
- 61 Go through past exam papers
- 71 Compile virtual exam sit un-timed
- 71 Compile virtual exam sit timed
- 81 Build self test on all concepts

Note number 81 is the hardest

- as you write out pages of info. that you should know - with blanks for answers - copy these pages 5 times - make sure you know all the concepts - using them to test - if you cannot answer one - first go to 41 - see if these prompt you

- second: go to 31 - see if these help

- thirdly: go to 11

If you can go through one of the copies with no mistakes

- you have learnt all the concepts - but really these are just an aide-mémoire

I hope this above helps

- I used these techniques when I undertook my MSc
- which had 8 subject per sem.
- all of which I had to comprehend fully.

THE exam times are PROVISIONAL

Monday	Tuesday	Wednesday	Thursday	Friday	REVISION
24th March 1pm-4pm Work hard on lecture notes	25th March 1pm-4pm Work hard on Compile project and summary documents	26th March 1pm-4pm Work hard on Compile project and summary documents	27th March 1pm-4pm Work hard on Compile project and summary documents	28th March 1pm-4pm Work hard on Compile project and summary documents	29th March 1pm-4pm Work hard on Compile project and summary documents
2pm-3pm CT210	2pm-3pm CT292	2pm-3pm CT212	2pm-3pm CT296	2pm-3pm CT218	2pm-3pm CT210

.pdf

Theory

Example

28

Li Learning; comprehension; & introspection

A possible methodology: Stage 1 - Revision schedule:

© Copyright Richard Neville 2009\10

	Monday 26th March	Tuesday 27th March	Wednesday 28th March	Thursday 29th March	Friday 30th March	REVISION
9am-lpm	Scan read Notes Visit hardest subject and read notes	Compile summary sheets/keywords	Scan read Notes Compile summary sheets/keywords	Scan read Notes Compile summary sheets/keywords	Read & Build self test on concepts & keywords	Go through Compile past exam virtual exam sit timed
Break		CT210	CT202	CT203	CT206	CT210
2pm-6pm				CT218		
	Monday 2nd April	Tuesday 3rd April	Wednesday 4th April	Thursday 5th April	Friday 6th April	REVISION
9am-lpm	Read comprehend lec notes	& Build self test on concepts & keywords	Go through past exam sit timed	Read comprehend lec notes	& Build self test on concepts & keywords	Go through Compile virtual exam sit timed
Break			Compile virtual exam sit timed			
2pm-6pm	CT202	CT202	CT213	CT213	CT211	
	Monday 9th April	Tuesday 10th April	Wednesday 11th April	Thursday 12th April	Friday 13th April	REVISION
9am-lpm	Go through past exam exam sit timed papers	Compile virtual exam sit timed	Read comprehend lec notes	Go through past exam sit timed	Read comprehend lec notes	Go through Compile virtual exam sit timed
Break			& Build self test on concepts & keywords	Compile virtual exam sit timed	& Build self test on concepts & keywords	
2pm-6pm	CT211	CT206	CT206	CT218	CT218	
	Monday 16th April	Tuesday 17th April	Wednesday 18th April	Thursday 19th April	Friday 20th April	REVISION
9am-lpm	Revision	half day coursework lectures	Revision CT210	half day coursework lectures	Revision CT202	half day coursework lectures
Break						
2pm-6pm	CT210		CT210		CT202	
	Monday 23rd April	Tuesday 24th April	Wednesday 25th April	Thursday 26th April	Friday 27th April	REVISION
9am-lpm	Revision	half day coursework lectures	Revision CT213	half day coursework lectures	Revision CT213	half day coursework lectures
Break						
2pm-6pm	CT202					
	Monday 30th April	Tuesday 1st May	Wednesday 2nd May	Thursday 3rd May	Friday 4th May	REVISION
9am-lpm	Revision	half day coursework lectures	Revision CT211	half day coursework lectures	Revision CT206	half day coursework lectures
Break						
2pm-6pm	CT211		CT206		CT206	
	Monday 7th May	Tuesday 8th May	Wednesday 9th May	Thursday 10th May	Friday 11th May	REVISION
9am-lpm	Revision	half day coursework lectures	Revision CT218	half day coursework lectures	Final Revision CT204	Final Revision CT204
Break						
2pm-6pm	CT218					

29

Li Learning; comprehension; & introspection

A possible methodology: Stage 1 - Revision schedule: In detail:

Monday 26th March	Tuesday 27th March	Wednesday 28th March	Thursday 29th March	Friday 30th March
Scan read Notes Visit hardest subject and read notes	Compile summary sheets/keywords	Scan read Notes Compile summary sheets/keywords	Compile summary sheets/keywords	Read & Build self test on concepts & keywords
CT210	CT202	CT213	CT206	CT210
				Go through Compile past exam virtual exam sit timed
Monday 2nd April	Tuesday 3rd April	Wednesday 4th April	Thursday 5th April	Friday 6th April
Read comprehend lec notes	& Build self test on concepts & keywords	Go through past exam sit timed	Read comprehend lec notes	& Build self test on concepts & keywords
CT202	CT202	CT213	CT213	CT211
				Go through Compile virtual exam sit timed
Monday 9th April	Tuesday 10th April	Wednesday 11th April	Thursday 12th April	Friday 13th April
Go through past exam exam sit timed papers	Compile virtual exam sit timed	Read comprehend lec concepts keywords	Go through past exam sit timed	Read comprehend lec notes
CT211	CT206	CT206	CT218	CT218
				Go through Compile past exam virtual exam sit timed
Monday 16th April	Tuesday 17th April	Wednesday 18th April	Thursday 19th April	Friday 20th April
half day coursework lectures	Revision CT210	half day coursework lectures	Revision CT210	half day coursework lectures
				Revision CT202
				half day coursework lectures

A possible methodology: Stage 2 - Exam schedule:

	Monday 7th May	Tuesday 8th May	Wednesday 9th May	Thursday 10th May	Friday 11th May	REVISION
9am-1pm Break 2pm-6pm	Revision CT218	half day coursework lectures	Revision CT218	half day coursework lectures	Revision CT218	Final Revision CT204
						REVISION
						REVISION
						REVISION
						REVISION
	Monday 14th May	Tuesday 15th May	Wednesday 16th May	Thursday 17th May	Friday 18th May	EXAMS
9.30am-11.30am Break 2pm-4pm		Final Revision CT218		Final Revision CT206		EXAMS
						EXAMS
						EXAMS
						EXAMS
	Monday 21th May	Tuesday 22nd May	Wednesday 23rd may	Thursday 24th May	Friday 25th May	EXAMS
9.30am-11.30pm Break 2pm-4pm	CT206	EXAM	Final Revision CT211	Final Revision CT213	EXAM CT213	EXAMS
						EXAMS
						EXAMS
						EXAMS
	Monday 28th May	Tuesday 29th May	Wednesday 30th May	Thursday	Friday	EXAMS
9.30am-11.30am Break 2pm-6pm		Final Revision CT202 / CT210	EXAM CT210	EXAM		EXAMS
						EXAMS
						EXAMS
						EXAMS

© Copyr

31

Li Learning; comprehension; & introspection

A possible methodology: Stage 3

Specific revision guidance for COMP15111 Introduction to Computer Systems

0/ Re-listen to all lectures - boot up Audio recording for lecture - then select lecture

- click through lecture while listening to Lecture audio...

1/ Re-read lecture notes THEN Undertake Q3's for each lecture

2/ Visit hardest concept(s) THEN Undertake all questions at the end of each lecture

3/ Compile (your own) summary sheets for each lecture

4/ Compile list of keywords for each lecture

5/ compile mind maps for each lecture

6/ Go through past exam - see module guide -- past exam in module guide [at end]

7/ Compile a possible future exam (virtual) - using pattern of last - sit un-timed

8/ Sit virtual exam timed

9/ Build self test sheets - these are a set of blank pages that you write QUESTION and ask yourself

- the concepts - the equations - the theory - for all concepts and theory of whole course

- copy these out 5 time to test yourself on ALL the theory the course covers

32

Li Learning; comprehension; & introspection

COMP15111

33

Learning; comprehension; & introspection

E Feedback From student whom used the method:

- Hello Richard,
 -
 - I hope you had great Christmas/New Year.
 -
 - XXXX
 -
 - *On a side note, I'm using the 'self-assessment' technique you've recommended (for revision) and have written myself a 30 question test for one of my modules which I do every other day. It's proving very useful!*
 -
 - Thanks,
 - Xxxxxx
 -
 - The point here is:
 - Is that a student has used the 'self-assessment' technique; and said:
“It's proving very useful!”

- Another - Note great feedback from student, before Exam:
- From the above student...
- With regards to #RN3, the self-assessment questions definitely allowed me to retain and recall large amounts of domain knowledge.
- It was especially useful in shorter questions and proved more beneficial than simply re-reading notes in a repetitive manner.
- To build/reinforce this knowledge I found it useful to answer questions/teach others, finding that if I could not sufficiently explain to a **peer** the concept,
- Then I did not truly understand it myself. I'm not a massive fan of group revision, but we organised several Q&A sessions in small groups which worked well.
- I think the final piece of the puzzle and something which is hard to 'revise', is *knowledge application*. For one of my modules I was extremely well prepared with my own self-assessment, past-papers and mock-questions; I therefore had a substantial amount of experience in applying my knowledge come the exam.
- However for another module there was a lack of available material with which to apply my knowledge, due to changes in the syllabus there was only one relevant past paper and no mock questions were provided by the lecturer (they chose not to, for fear of similarity to the actual exam!).
- Ultimately this meant repetition of notes and group sessions were the only real ways to prepare. Which meant I definitely wasn't as confident walking into this exam.
- If you have any suggestions for different techniques for revision given this kind of situation, that would be appreciated.



Important Exam advice, cont...

1. Write your answers on the exam paper in a big font. In previous exams, I have tried to decipher writing that was equivalent to 5 or 6 font – far too small. If you have trouble reading it, so will the examiner.
2. Take care writing in the answer book, so [at least] the answers and diagrams and mathematical working-out is readable, tidy, clear and written in good Plain English.
3. Students need to layout their answers, in their answer book, so each different answer is CLEARLY differentiated from the next. Leave a clear line between each answer.
4. If number conventions such as 2., then 2.a, then 2.a.ii PLEASE use them.
5. Please try to answer your questions in order.
6. Please explain any diagrams [or tables] you draw [up] in full. Full marks cannot be awarded unless full explanations are given to answer each question full and comprehensively.

BIG HINT – diagram really do help:

Remember, good – honours grade answer – in the exam – for a question – to maximise marks – should – or you should think of adding A DIAGRAM;

basic layout:

- 1) Textual answer;
- 2) Diagram supporting answer; &
- 3) Full explanation of diagram...

This sort of answer will maximise your marks...

...practice doing this for the majority of your answer – if you can align a diagram that is...

COMP15111

37

Where to find Lab. L3 Week 12 Hints – Final Lab (Page table)?

First Go to Blackboard 9; then select:  COMP25111 Operating Systems

Lab3_RN

Then select:

Move to this folder, to download ex3:
Click "Lab3.pdf" to obtain (read) Lab3 (ex3) full specifications.
Plus Lab 3 Hint 1: [#H1]; Hint 2: [#H2] etc.
Note read the "Lab Rules" as you must use SUBMIT for Lab 1, 2 & 3
If you use Netbeans it is on all Lab PC :- Where to Find NetBeans on the lab machines:
Boot into Windows 7; 1) Click "Start;" 2) Click "All programs;" 3) Click the "NetBeans" folder [icon]; 4) Click the "NetBeans IDE 7.0.1" icon.

Lab3 Hints

Then select:

Hint 1: [#H1]
Hint 2: [#H2] etc.
If you use Netbeans it is on all Lab PC :- Where to Find NetBeans on the lab machines:
Boot into Windows 7; 1) Click "Start;" 2) Click "All programs;" 3) Click the "NetBeans" folder [icon]; 4) Click the "NetBeans IDE 7.0.1" icon.

Then select: see next page...



Sub directory:

[H1.pdf](#)[H2.pdf](#)[H3.pdf](#)[H3_2.pdf](#)[H4.pdf](#)[traceSingleFetch](#)[traceSingleRead](#)[traceSingleWrite](#)

39

[FIFOvisualExample.pptx](#)

#H1

Sub directory:

'Hint No 1' hints at the requirements you should have data mined from reread section 5.

Stage 1: extracting the requirements of the laboratory from COMP20051 Exercise 3: Simulation of Paging Behaviour laboratory description document.

[#R1] Write a java program to implement a simple MMU.

Extracted from:

... Purpose of the exercise is to write Java code to implement a simple MMU, which contains a page table and code to perform the necessary actions when read or write accesses occur to a virtual memory address. ...

[#R2] Implement a LRU algorithm.

[#R3] [Implement a] 'timestamp' in each page table entry.

Extracted from:

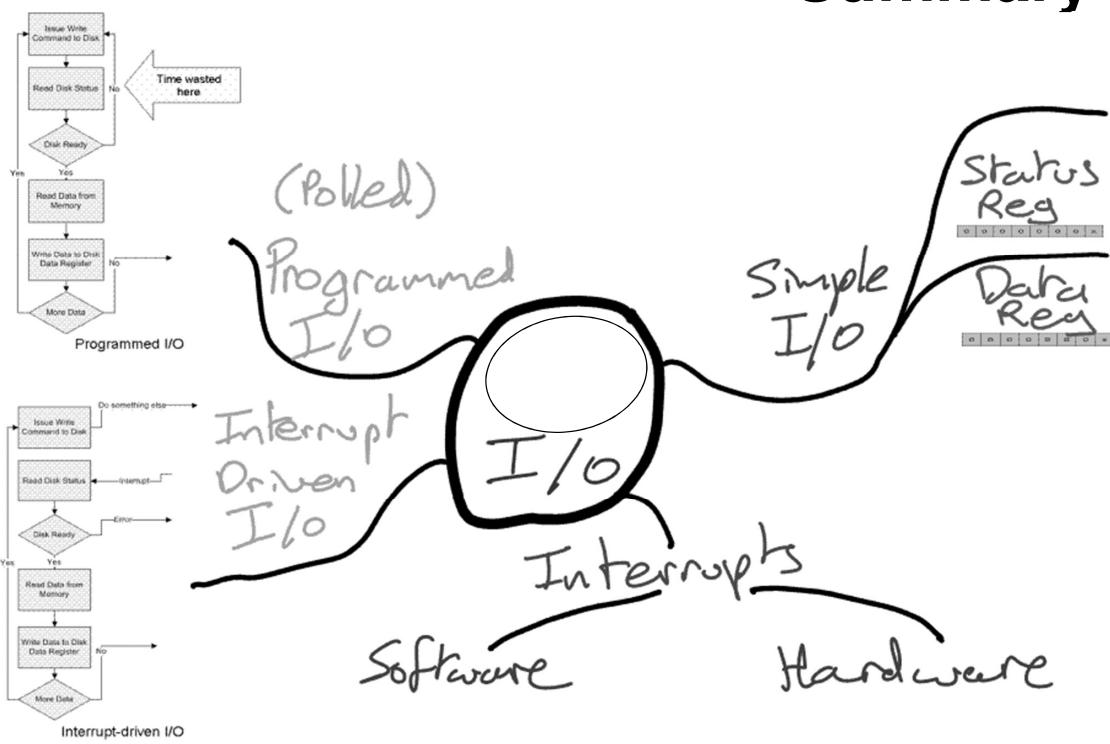
... You should implement a LRU algorithm to determine page rejection. To make this simple, you should have a 'timestamp' in each page table entry, which is updated appropriately when a page is accessed. ...

[#R4] [Implement] a 'dirty' indication in the page table.

Extracted from:

... You should also include a 'dirty' indication in the page table implementation, which is set on a write. This indicates if the page needs to be written back to disk if it is rejected. ...

Summary

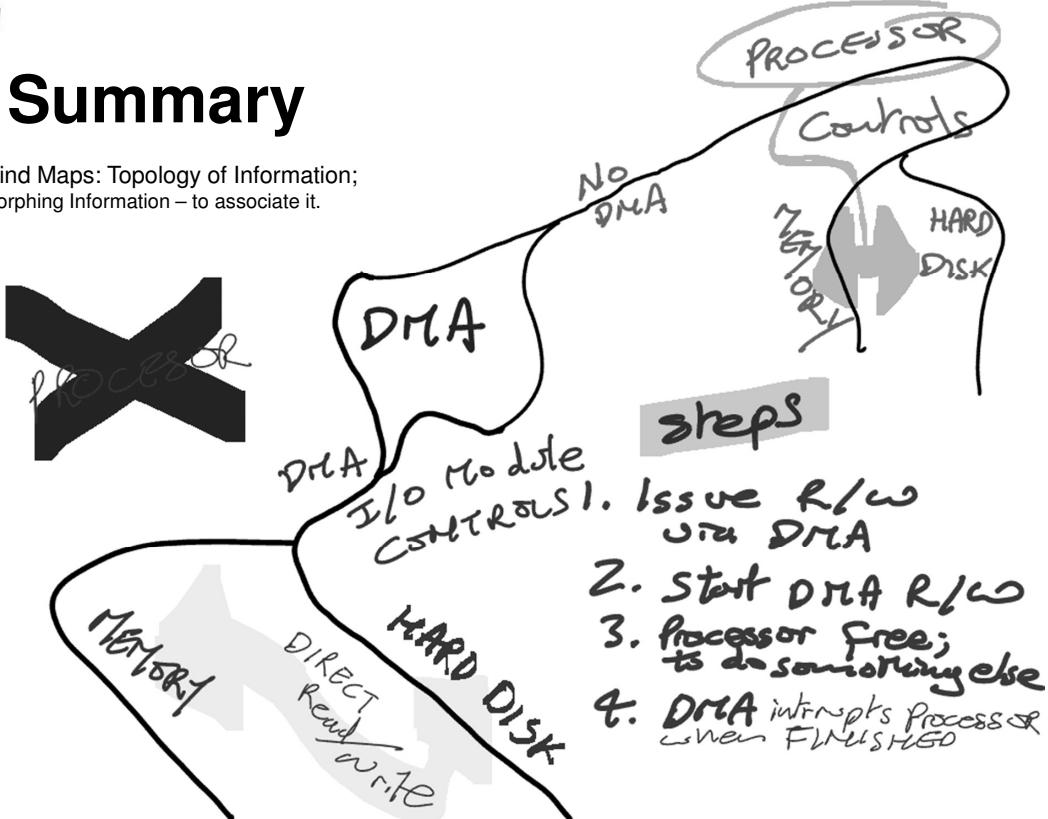


41

Mind Maps: Topology of Information;
Morphing Information – to associate it.

Summary

Mind Maps: Topology of Information;
Morphing Information – to associate it.



42



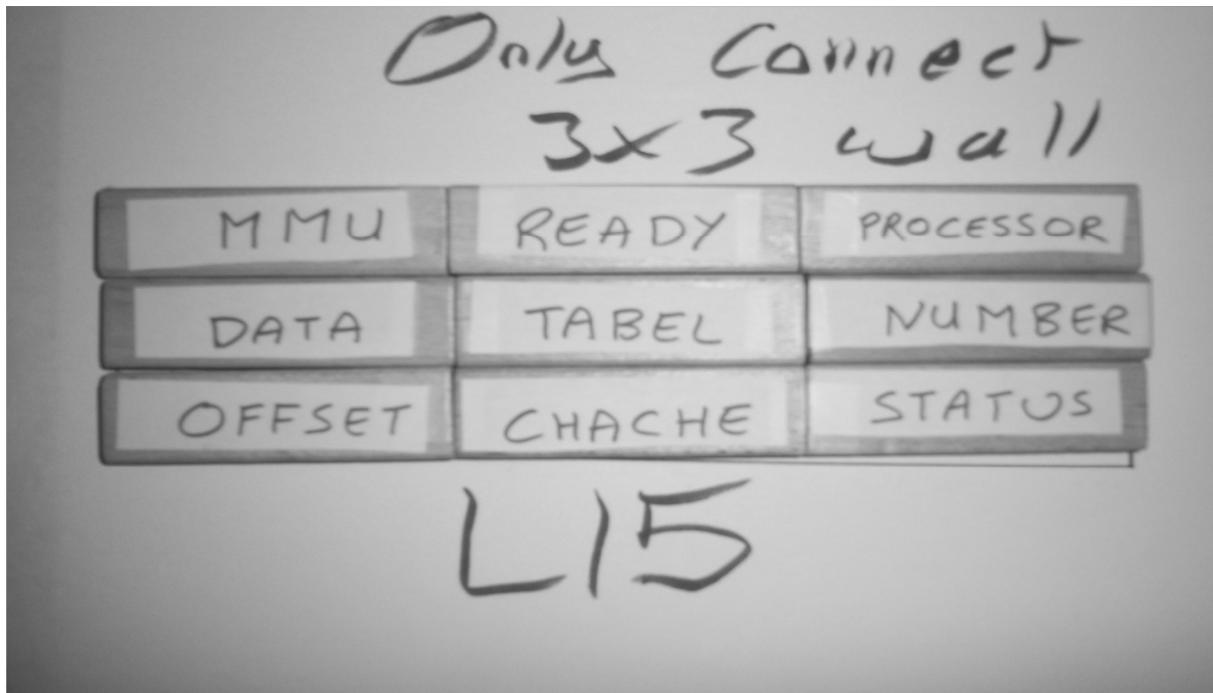
Only
Connect
Quiz



Only
Connect

[The Rules, how to play](#)

<http://youtu.be/59qwTcAceTA>



Quiz



Learning; comprehension; & introspection

List of Questions to ask lecturer

- Before the 9a.m. start lecture the lecturer will be half an hour early and you can ask [any and all] questions in that half hour; before the lecture:

- 1.
- 2.
- 3.
- 4.
- 5.

Getting ready for next week

Do next week's Q3's NOW

- Once you have re-read the lecture notes; and listened to the audio recording [while stepping through the PPT] of the lecture again:
- Please have a think about next week's Q3's
 - on the next page
- If you try to answer the Q3's now you will be in a much better position to recall the information.
- Once you have done this, transfer your answers to next weeks "Student [OWN answers] version" at the start of next weeks lecture.
 - YES this implies bringing the last weeks lecture notes to the next lecture ...

THIS week's

Short Exam Questions

Q3

1. Question
List a number of devices that connect to the system bus, via I/O modules.

Answer(s):

2. Question
Name the two different types of information I/O registers store?

Answer(s):

3. Question
Name the two different types of interrupts.

Answer(s):

- NOTE: In the exam approximately 2 question are taken from the topics (and program examples) coved in each lecture

1. Question

List a number of devices that connect to the system bus, via I/O modules.

Answer(s):

2. Question

Name the two different types of information I/O registers store?

Answer(s):

3. Question

Name the two different types of interrupts.

Answer(s):

- NOTE: In the exam approximately 2 question are taken from the topics (and program examples) coved in each lecture

49

Glossary

• Why build a Glossary for each course unit you undertake?

- It is imperative that the correct terminology [keywords] are utilised in context in your exam answers; this is so important that the lecturer has added glossaries to each paper copy of your lectures. It is of such importance that in your notes [prior to the start of the glossary] the following advice is given:
 - Each module you undertake uses its own jargon.
 - This can be a problem for new students, whom are trying to comprehend the new domain knowledge attached to a particular new module.
 - One way to get to know the new jargon is to build your own GLOSSARIES for each course module.
 - The glossary on the next few pages is a starting point for this module [unit].
 - Please feel free to add to the glossaries throughout the unit...
 - The glossary is full of potential exam questions of the form "define X" or "briefly explain X."
- Please heed the advice in the future; even if your lecturers do not supply a glossary build your own as without knowledge of the appropriate terminology [keywords] when expounding your knowledge you will not be viewed as comprehending the details of any theory.

GLOSSARY

Using the on-line resources and any other resources compile a glossary of the terms below [PIPLINES: memory management]:

- Interrupt →
- Interrupt-driven →
- Acknowledgement →
- Vector →
- Interrupt vector →
- Command register →
- Status register →
- Polling →

51

Learning Resources 1

- **Descriptions [Theory] (in text books)**

- Remember the key issues, highlighted in GREEN, are the concepts to look for in any book:

- Section on Interrupt, Vectored interrupt handling, Command register, Status register, Polling, in chapter 8 the Input-Output in: Chalk BS, Carter AT, Hind RW (2004) Computer Organisation and Architecture: An introduction 2nd Edition, Palgrave, ISBN 1-4039-0164-3 .
- Section on Interrupt, Interrupt-driven, Acknowledgement, and Vectored Interrupt – in chapter 6 the Input/Output in: Computer Organization and Architecture, Fifth Edition by William Stallings.
- Section on Interrupt – in a few chapters in: Structured Computer Organization, 5/E, Andrew S. Tanenbaum, Vrije University, Amsterdam, The Netherlands, ISBN-10: 0131485210, ISBN-13: 9780131485211, Publisher: Prentice.
- Section on I/O hardware, Interrupt, Interrupt-handler, DMA, Software Interrupt vector in chapter 13 I/O systems in Operating System Concepts, A. Silberschatz et al, Wiley.
- Section on Interrupt, Interrupt-driven, Acknowledgement, Interrupt vector, Polling in chapter 5 Input/Output in Modern Operating Systems, A. Tanenbaum, Prentice Hall.

- **Web resources:**

- interrupt; available [on-line] @ <http://www.webopedia.com/TERM/I/interrupt.html>

52

Questions

Introduction to Questions:

The set of questions are based on lecture 15.

Answer Sheet will be given later in year and will contain the answers to these questions.

- Remember to find detailed and comprehensive answer you should [also] reference associated text books in the library.
- A reasonable starting place for associated book titles are:
 - 1) This units 'module guide'; given to you in RN's first lecture – or on the web [Blackboard];
 - 2) Those books mentioned in 'Background Reading,'
 - 3) Those books [and web resources] mentioned in Learning Resources.

Questions

1. Question

- In the event an interrupt occurring; one process has to hand over to another. State the basic steps that have to happen before the second process starts.
- Answer

1. Answer

Answer(s):

2. Question

- Given the overall steps the processor takes to handle an interrupt can be summarised in six basic steps. List the six basic steps (in a interrupt summary).

- **Answer**

2. Answer

1) Answer(s):

2)

3)

4)

5)

6)

3. Question

- Explain how a peripheral communicates with the CPU using interrupts.

- **Answer**

3. Answer

- The following points should be covered to some degree in the answer:

Answer(s):

Revision Exercises

- Scan read Lecture 15's Questions.
 - Answer Lecture 15's Questions
 - Particularly those questions you had difficulties with when you first tried them.

Background Reading

- [1] Computer Organisation and Architecture: An introduction 2nd Edition, Palgrave, ISBN 1-4039-0164-3, Chalk BS, Carter AT, Hind RW (2004).
 - Sections in chapter 8 the Input-Output in: Section on Interrupt, Vectored interrupt handling, Command register, Status register, Polling.
- [2] Computer Organization and Architecture, Fifth Edition by William Stallings.
 - Sections in chapter 6 the Input/Output in: Section on Interrupt, Interrupt-driven, Acknowledgement, and Vectored Interrupt.
- [3] Structured Computer Organization, 5/E, Andrew S. Tanenbaum, Vrije University, Amsterdam, The Netherlands, ISBN-10: 0131485210, ISBN-13: 9780131485211, Publisher: Prentice.
 - Section on Interrupt – in a few chapters.
- [4] Operating System Concepts, A. Silberschatz et al, Wiley.
 - Section on I/O hardware, Interrupt, Interrupt-handler, DMA, Software Interrupt vector in chapter 13 I/O systems.
- [5] Modern Operating Systems, A. Tanenbaum, Prentice Hall.
 - Section on Interrupt, Interrupt-driven, Acknowledgement, Interrupt vector, Polling in chapter 5 Input/Output.

COMP25111 Exercise 3: Simulation of Paging Behaviour

Contains a copy of the “traceA” file.

Before you start this PLEASE read the multimedia PPTs and audios on the Blackboard 9 portal- on the Blackboard 9 Web Site; select: COMP25111 Operating Systems <https://www.portal.manchester.ac.uk/>; then
1/ Go to COMP25111 Operating Systems 2011-12 1st Semester;
2/ Then Lab3 RN
3/ Then Lab3 Hints
4/ Then read through as many hints as you need; remember the methods, processes, naming conventions, and implicit theory outlined in the hints will be useful to you in your final year projects where you should think about using these techniques...

Duration: 1 Session

1. Learning Outcomes

On completion of this exercise, a student will:

- Have implemented the functionality of a simple Memory Management Unit (MMU) for a paged virtual memory system in Java.
- Have exercised the MMU using a trace file of memory accesses using a Java program which initialises the MMU and calls methods within it to perform read and write actions.
- Have produced statistics, which show the behaviour of the MMU with particular reference to page faults and page rejections.

2. Introduction

Virtual memory is a technique for providing a process with apparent access to the whole addressable memory space of a processor in circumstances where the real memory available may be considerably smaller. One advantages of the scheme is the process can use larger code and data sizes than would otherwise be possible. Another is the run-time layout of code and data can be considerably simplified if it does not have to fit into a confined space. This is particularly true for dynamic data whose size is unknown until run-time.

Virtual memory relies on the fact that, in the majority of programs, the use of both code and data exhibits both *spatial locality* and *temporal locality*. That is that at any point in time, the program is usually only using a small fraction of its code and data and that, over a short time period, the usage will not change significantly.

A virtual memory system therefore tries to keep currently used code and data in the (limited size) real memory of the system while the rest (if it exists) is kept on background storage such as magnetic disk. In order to relieve the programmer from the complex task of working out which data should be where at any point in time, the system arranges to move the data between real memory and disk automatically as required.

A piece of hardware known as a Memory Management Unit (MMU), together with system software within the Operating System provide all the functionality needed to

do this. The purpose of this exercise is to implement the functionality of a MMU in order to gain an understanding of the principles.

3. Paged Virtual Memory

Devices such as disks usually store and retrieve data in blocks of hundreds or thousands of bytes. Accessing a block is often slow compared to CPU memory speeds but once accessed; the contents of that block can be read or written very rapidly. In these circumstances, it makes sense to move data in units which are larger than bytes or CPU words. Virtual memory systems therefore operate at the unit of a *page* when transferring data to and from memory and disk. A page does not necessarily correspond to a disk block size but is fixed for a particular MMU typically between 4 and 64 kilobytes. This will usually correspond to several disk blocks.

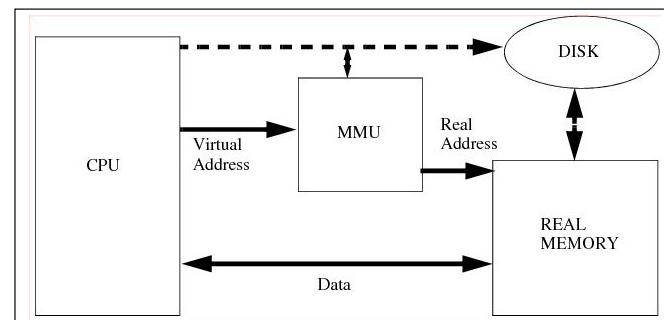


Figure 1 shows the basic structure of a paged virtual memory system.

The addresses issued by the CPU are virtual addresses which are the full width of the processor's addressing capability. In a modern 32 bit processor (for example ARM or x86) these will be a full 32 bits capable of addressing 4Gbytes of memory. The real memory will usually be smaller, typically 512 Mbytes on a laptop.

Each of these memory spaces is considered to be divided into pages. Note that this is not reflected in the internal structure of any memory, it is just an abstract division; the memory addresses are still binary values which cover a linear address space (usually at the unit of bytes). However, we can view any address as being divided into two sections, one which addresses the page and one which is an offset within the page which addresses individual units (bytes). By convention, we usually refer to the virtual address as having *virtual page numbers* (or just page numbers) and the real address as having *page frame numbers*. Figure 2 shows this pictorially.

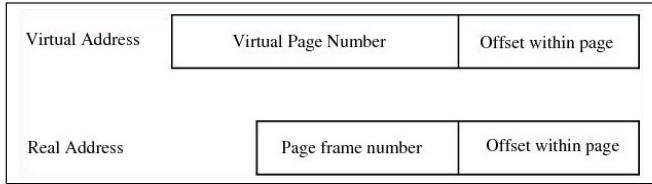


Figure 2 shows this pictorial view of *virtual page numbers* and *page frame numbers*.

The function of the MMU is to keep track of which virtual pages exist in real memory and, which on disk. It does this by keeping a *page table* of every possible virtual page number which contains the page frame numbers of any real copies. It is important to note that the offset is common to both virtual and real addresses. In order to find the data associated with a virtual address it is only necessary to take the virtual page number and replace it by the page frame number (if it exists). The page frame number is looked up in the page table to produce the real address which can then be used to access real memory. This is usually referred to as *address translation*.

If the table indicates that a virtual page does not exist in real memory this is called a *page fault*. It is necessary for the MMU to initiate a transfer of the page's data from the background storage (disk). The exact mechanism for this is not relevant here but it should be noted that this is a relatively slow process and the CPU itself can organise most of the transfer.

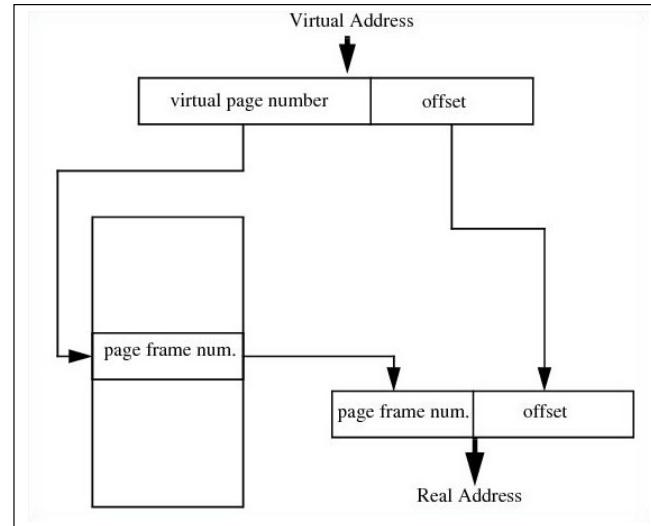


Figure 3 shows more detail of the MMU table and the address translation operation.

The page table is shown here as a single linear structure with the lookup done by the full page number. For large virtual address spaces, this may not be practical and more complex techniques may be required (this is covered in course lectures).

4. Page Rejection

If there is a page fault, a place needs to be found in real memory for the page copy. Initially, assuming all real memory is not in use, it is possible to keep allocating page frames linearly through the address space. However, if all real memory is in use, we must find a page to reject. We must then write its contents back to disk before reallocating its page frame number to the page access, which has caused the fault.

We need to decide which page to choose using a *page rejection algorithm*. The most commonly used is an algorithm called LRU (least recently used). It works on the simple principle that the page, which has survived the longest time without being accessed, is a good candidate for replacement.

One useful optimisation is to note if a page has been written to since it was brought from disk into real memory. If not, there is no need to write it back to disk before reusing the real page frame. This is often true for pages which contain code.

5. The Exercise

The exercise is composed of two parts:

Part One: implements the majority of the code and the LRU page replacement algorithm. The majority of the (theoretical) advice below will aid you in this task;

after you have designed, developed, and implemented the program utilise the two trace files: **traceA** and **traceB**; to test your program.

Part Two: implements the a second algorithm the FIFO page replacement algorithm. The advice below will not aid you in this task. Nevertheless, it is expected that once you have implemented (and tested) the LRU algorithm you will have gained sufficient knowledge to enable you to undertake the FIFO design, development and testing on your own. After you have designed, developed, and implemented the program utilise the trace file: **traceLRUandFIFO**; to test your program. The FIFO algorithm has been covered in your lecture series and it is covered extensively in your course books; also see the Learning Resources (book & Web resources) as well as the Background Reading (book & Web resources) at the end of Lecture COMP25111 Operating Systems Lectures 14: Virtual Memory (3) for further reference material.

Introduction to Exercise

The purpose of the exercise is to write Java code to implement a simple MMU, which contains a page table and code to perform the necessary actions when read, or write accesses occur to a virtual memory address. Because the exercise is using a trace of addresses rather than executing a real program, there is no need to perform any real memory accesses. Instead, it is simply necessary to access the page table to determine if there is a page fault and, if so, find a free page frame to use and place it in the table. If, initially, there are unused page frames these can be allocated linearly until there are none free. At that point it will be necessary to find a page to reject and use its page frame for the newly required page.

You should implement a LRU algorithm to determine page rejection. To make this simple, you should have a ‘timestamp’ in each page table entry, which is updated appropriately, when a page is accessed.

You should also include a ‘dirty’ indication in the page table implementation, which is set on a write. This indicates if the page needs to be written back to disk if it is rejected.

Note that a practical page table will usually also contain page access information which determines the types of access which are allowed to a particular page (for example read only). There is no need to include such information for the purposes of this exercise.

In appropriate Blackboard subfolder, you will find a program harness which, after creating a MMU object, reads the trace of store accesses and makes calls to read and write methods within that object until the trace file ends. It then prints out some statistics:

- Total Accesses
- Total Instruction Fetches
- Total Page Faults
- Total Page Rejections (i.e. ignoring startup page faults)
- Total Page Writebacks (to disk)

The first two are included in the given program; the last three need to be calculated by your MMU implementation. (A skeleton MMU.java file is provided in the appropriate Blackboard subfolder)

The trace file format is a line for each access. The first number indicates the access type (0=read, 1=write, 2=fetch) and the second is a (hex) **24 bit memory address**. Within the program, the address is considered to be a 12-bit page number and a 12 bit offset representing 4k pages each of 4k bytes, i.e. 16Mbytes in total of virtual memory

space. This is deliberately small so that your implementation can use a simple single level page table. Although obviously not representative of modern systems, this is typical of a virtual memory system in computers of the 1960s when virtual memory was invented (in Manchester University on the Atlas computer [1][2]).

An example of ‘traceA’ file is given below:

```
2 400000
0 600000
2 400001
1 600001
2 400002
0 700000
2 500000
0 700001
```

The first line “2 400000” is decoded as follows:

First number	Second number
2	400000
2=fetch	0100 0000 0000 0000 0000 0000 _{HEX}
This is a fetch access	24 bit hexadecimal memory address

The following decode the first number of each of the eight lines in traceA:

First number	Type of access
2	2=fetch
0	0=read
2	2=fetch
1	1=write
2	2=fetch
0	0=read
2	2=fetch
0	0=read

The program given makes a single call to the method which performs the trace simulation using a very very small real memory (2 [real] page frames, i.e. 8k bytes [in MMUSim.java the real memory sizes (in pages) is set to “rmem_pages = 2;”]. This should be run with the file traceA when initially debugging your implementation. This file contains only 8 store accesses which have been devised so that you should be able to predict what statistics your MMU should produce.

When you think this is working, you should modify the program to perform simulations using both 32 and 64 real page frames on traceB. This represents 128k and 256k bytes of real memory, again realistic in the 1960s.

6. Results

You should be prepared to demonstrate the results of the runs on traceA and traceB when your exercise is marked. To help you determine if you have got a correct implementation, the total number of page faults for 32 pages and traceB should be 227.

You should also be prepared to describe your code and draw relevant conclusions from the results.

Assessment

COMP20051 Exercise 3: Developing Simulation of Paging Behaviour

Demonstrators please fill in the student's full detail [below] before starting the assessment; may be the best way to do this is to get the students to fill this in themselves:

Students full name	Student Number	Email address	Course studied

At the deadline, at the end of the Session, the marks awarded are as given in table 1.

Demonstrators please assign marks for each of the eight questions in boxes provided.

At the deadline, at the end of Session 5, the marks awarded are as follows:

Mark awarded for:	Exercise 3
1. Show Correct page table [data] structure; plus comment on how you [the student] developed the data structure.	1 <input type="text"/>
2. Demonstrating an overall [software engineering] correct approach; of the algorithmic and combined data structure: 2 – 5.	
2. State how REQUIREMENT development was undertaken.	1 <input type="text"/>
3. State how [if any] ABSTRACT DESIGN was undertaken.	1 <input type="text"/>
4. State how IMPLEMENTATION was undertaken.	1 <input type="text"/>
5. Demonstrating a methodology for TESTING the program; and comment on: interpretation of results .	1 <input type="text"/>
6. Showing the generation of an appropriate OUTPUT trace file: traceA files [displayed in real-time on command [Prompt] screen] (using LRU).	1 <input type="text"/>
7. Showing the generation of an appropriate OUTPUT trace file: traceB files [displayed in real-time on command [Prompt] screen] (using LRU).	2 <input type="text"/>
8. Showing the generation of an appropriate OUTPUT trace file: traceLRUandFIFO files [displayed in real-time on command [Prompt] screen] (using LRU and FIFO).	2 <input type="text"/>
TOTAL mark	10 <input type="text"/>

Remember to save these exercises in a directory nominally named .../COMP20051/ex3.

Do this exercise [or save this exercise] in a directory named COMP251111/ex3 directory, save your downloaded template [Skeleton(s)] code:-

MMUSim.java; and
MMU.java.

as well as the trace files:

traceA;
traceB; and
traceLRUandFIFO

... in COMP251111/ex3.

Remember to save your trace files of your solution.

The three files traceA; traceB; and traceLRUandFIFO; as well as the template [Skeleton(s)] MMUSim.java & MMU.java are downloadable from Blackboard in Folder Lab3.

References

- [1] One-Level Storage System, T. Kilburn, D.B.G. Edwards, M.J. Lanigan, F.H. Sumner, IRE Trans. Electronic Computers April 1962
- [2] Tom Kilburn (1956). The Atlas, School of Computer Science: information on Tom Kilburn's [computing] effort known as the MUSE (microsecond) computer , available [on-line] @ <http://www.computer50.org/kgill/atlas/atlas.html>, [Last accessed 12/10/09, 11:09].

7. Decoding the COMP20051 Laboratory Exercise 3; and ‘supporting advice’

The previous six pages document is a typical laboratory that you are normally set in the School of Computer Science. The first task is to decode [work out] what explicitly is required. As you have read the document once now reread section 5; this time highlighting or extracting the main requirements of the exercise.

Hint 1: [#H1]

[#H1] If you require a hint, download them from Blackboard. ‘Hint No 1’ hints at the requirements you should have data mined from reread section 5.

After reading document ‘H1.doc’ to validate you have extracted the relevant requirements you next step is to start the design phase that fulfils the requirements [#R1 to #R5] of the exercise. Then once designed the MMU class can be implemented.

Hint 2: [#H2]

[#H2] If you require a hint, download them from Blackboard. After you have compiled a list of requirements, you should now think about designing the MMU class. ‘Hint No 2’ suggests an approach you could take.

Reading document ‘H2.doc’ may help you move forward in the design phase or validate your design.

You should now be able to implement, test, and run MMUSim utilising traceA. Once you have reached this stage you can compare your program’s trace output with a typical trace A output [result] available on Blackboard; the file is named ‘traceAres.’

One of the final requirements of the exercise is to modify the program [MMUSim] to perform simulations using both 32 and 64 real page frames on traceB. To do this you should consider altering the number of real memory pages.

Finally, you should consider the best way to describe your code; when asked. In addition, you should think about drawing relevant conclusions from the results.

The “>>” operator

Final point: what does the operator “>>” actually do?

In the code snippet:

```
int addr = access.addr;
```

int vpage_no = addr >> 24 - vmem_bits; // assumes a 24 bit virt addr given vmem_bits = 12. Hence the actual value of ‘24 - vmem_bits’ is “addr >> 12”

In “doSimul(…)” in class MMUSim.Java. The operator “>>” is utilised. If we look at what happened

Values:
addr = 4194304
vpage_no = 1024

So what has happened?

The java operator “>>”

>> shift bits right with sign extension

It implies [with two objects ‘x’ & ‘y’]:

x >> y

Means: Shift Right - Signed

Explicitly means: Shift **x** to the right by **y** bits. Low order bits are lost. Same bit value as sign (0 for positive numbers, 1 for negative) fills in the left bits.

Given the example **addr** = 4194304 which is specified for **addr** an integer (int) which is base (or radix) 10.

Converting 4194304₁₀:
Converting to hexadecimal:
00400000₁₆
or radix 2 it is:
0000 0000 0100 0000 0000 0000 0000 0000₂

Hence given **addr >> 12**” After shifting right 12 places this is:

0000 0000 0000 0000 0100 0000 0000₂
which is:
00000400₁₆
or radix 2 it is:
1024₁₀

QED