

Before we start

What is RAM fragmentation?  
Why does it happen?

What is the cure?

What is file-store defragmentation?  
Why is it needed?

What (& why) is each of these:

- disk partition
- mounting e.g. a storage device
- hard & soft (symbolic) links in a file-store



COMP25111: Operating Systems  
Lecture 16: The File Manager

Oscar Palomar  
School of Computer Science, University of Manchester  
Autumn 2016

COMP25111 Lecture 16

Overview & Learning Outcomes

- 1 Files & File Systems
- 2 Naming Service
- 3 Storage Service
  - Data Structures
  - Allocation
- 4 File Manager & Virtual Memory

COMP25111 Lecture 16

File system

File Identifiers:  
SFID – System – lifetime of file  
UFID – User – lifetime of process  
(Unix: “file descriptor”, Windows: “file handle”)

Requirements – system calls:

open: file-name → UFID  
read: UFID & count → data  
write: UFID & data →

Multiple OS Layers:

- naming service: e.g. open
- storage service: e.g. read & write (vector of bytes)
- disk driver: access disk sectors

1/35COMP25111 Lecture 16

What is a file?

Collection of related information on secondary storage:  
e.g. data, programs (.java, .c, .h, .class, .o, binary, ...)

Structure: lines of text; or sequence of bytes; or ... ?

Types: should OS recognise/support?  
– in the name: .com .exe etc. (MSDOS)  
– “magic number” at the beginning of some files (Unix)

Attributes: (name?) size, last update, owner, ...  
(try `ls -la`)

Operations: create, open, read, write, close, delete, ...

Access:  
– Sequential: processed in order, from start to end  
– Direct (Random): logical records, processed in any order

3/35COMP25111 Lecture 16

Files & File Systems

File system organisation: Directories


Directory: file-name → SFID  
(SFID gives access to contents & attributes)

Originally:  
1 directory per partition (1-level), or  
1 per user (2-level)

Nowadays: Tree (or Forest ) of directories

Stored on disk just like files but treated differently

Example Directory Tree-Structure



Directories contain files & directories

A normal file is a leaf in the tree

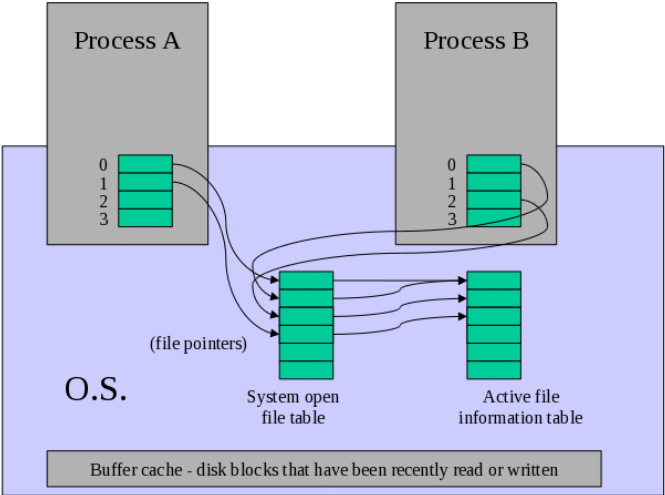
Decoding a path name

- Split path at separators (e.g. / or \)
- Absolute: left-most component = root directory
- Relative: implicitly starts with current working directory
- Each component from left must:
  - identify a directory
  - contain the next component
- Final component = file or directory


Data Structures

- Each process has:
  - working directory (inherited on creation)
  - UFIDs
- Each file has attributes (metadata, File Control Block FCB):
  - file size, permissions, owner, group, dates, ...
  - where to find data on disk
- Open file table in memory:
  - entry = attributes, number of readers & writers.
  - indexed by UFID
- open: create entry in the file table
- (last) close: write attributes to disk
- Implementations vary e.g. maybe also a table per process

Data Structures



Disks (MOS2 fig 6.11)



Physical structure: platters, tracks, sectors, etc.  
Logical structure: blocks

Free Space

- e.g. bit-map, or list of block-no? (no = number)
- e.g. 100GB partition = 25M \* 4kB blocks
- Q: bitmap size (blocks)?
- Q: block-no size (bytes)?
- Q: list size (blocks)?
- list can use free blocks, bitmap needs extra disk space
- Q: search O(?)

File structure – contiguous blocks

(e.g. CD/DVD – ROM/WORM)

file = start-block-no & block-count

e.g. file starts at 3, uses 4 blocks:

1	2	3	4	5	6	7	8	9	: block numbers
		F	I	L	E				: disk blocks

- + simple, fast
- interleaved user requests → seeks
- fragmentation

File structure – list of blocks (1)

“next” within block

e.g. file starts at 3 (next=0 indicates EOF)

1	2	3	4	5	6	7	8	9	: block numbers
		F 4	I 7			L 9	E 0		: disk blocks

- random/direct access very slow

File structure – list of blocks (2)

“next” in separate monolithic table  
(e.g. MSDOS FAT – File Allocation Table)

+ table can also hold free-block info.

e.g. file starts at 3; abc starts at 2 (–1 indicates free)

1	2	3	4	5	6	7	8	9	: block numbers
	a	F	I	c		L	b	E	: disk blocks
–1	8	4	7	0	–1	9	5	0	: FAT

- need to cache table (e.g. 100MB) in memory

File structure – list of blocks (3)

“next” in separate partitioned data-structure

+ one table in RAM (proportional to file size) per open file


e.g. file table in block 1; abc table in block 6

1	2	3	4	5	6	7	8	9
3 4 7 9	a	F	I	c	2 8 5	L	b	E

e.g. UNIX i-node/inode


File structure – list of blocks (4)

i-node = file-attributes + 11 to 15 block-nos  
– first 8 to 12 = first blocks of file  
– last 3 = block of block-nos, block of blocks of block-nos, ...



inodes in separate disk area (earlier slide)

Disks (again)



inodes in separate area from file/directory blocks

Directory structure

- File Name → directory-entry sizes
- max length → fixed-size  
(e.g. MSDOS = 8+3, early UNIX = 14)
  - unlimited → variable-size (including strlen)
  - unlimited → fixed-size (+ “heap” for strings)

- File Attributes:
- in directory entry (e.g. FAT)
  - pointed at by directory entry (e.g. in inode)

- Disk Address:
- e.g. FAT: block-number of start of file
  - e.g. inode: via inode-number

Q: efficiency? – what are the commonest directory operations?

Directory structure (MOS4)

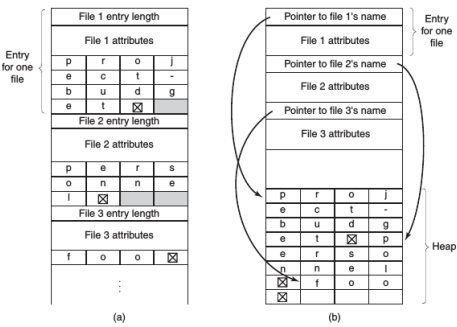


Figure 4-15. Two ways of handling long file names in a directory. (a) In-line. (b) In a heap.

Other Issues

- Concurrency: how should multiple accesses be coordinated?
- allow either 1 writer; or many readers (inflexible?)
  - applications (e.g. database) define specific protocols

- Performance:
- cache
  - efficiency dependent on algorithms/types of data
  - RAID (Redundant Array of Independent Disks): striping

- Access Protection:
- R/W/X permissions (man chmod)
  - access control lists (e.g. man acl)

- Recovery: backups!
- consistency checking (MSDOS:chkdsk, Linux:fsck)
  - partitions
  - journalling
  - RAID: mirroring

Virtual Memory and Storage

Virtual Memory & File managers both copy info RAM ⇔ disk

Unified VM/File Manager: memory-mapped files

File Operation	Equivalent VM Operation
open	map file into Virtual Address Space
read	access virtual address: page fault causes disk read
write	access virtual address: (eventual) page rejection causes disk write [n.b. zero fill on write to unallocated page?]
close	unmap pages

- Pros: programs & libraries; sharing
- Cons: different access patterns?

Free Space Management

How to keep track of free space?

- Keep a bit vector (0 or 1; 1 for a free disk block)
- simple to find first free block  $(bits\_per\_word) \times (0\_words) + (offset\_first\_1\_bit)$
  - 4GB disk with 512 byte blocks needs 1MB space

- Linked list of free blocks
- Need I/O to locate n free blocks
  - Alternatively: keep info in the FAT

Grouping: for every n free blocks record info in 1 about the n-1

Summary of key points

- 1 Files & File Systems
- 2 Naming Service
- 3 Storage Service
  - Data Structures
  - Allocation
- 4 File Manager & Virtual Memory

Your Questions

For next time

A file system uses inodes which contain 8 block-numbers. These are for the first 7 blocks of the file and an indirect block, which just contains block-numbers for the remaining blocks in the file.  
A block-number occupies 2 bytes. Each block is 4k bytes.

What is the maximum size of a file in bytes?

What is the maximum total size of directories and files in a single disk partition?

Exam Questions

Explain the algorithm used to locate the file referred to by a full path name in a hierarchical file system. (5 marks)

Using a FAT16 file system (i.e. each FAT entry occupies 16 bits) how much space would be available on a 160MB disk for directories and files, for block sizes of 2KB and of 4KB? Explain your reasoning. (5 marks)

A disk storing a hierarchical file system will hold three forms of data: directories, file contents, and metadata. Illustrate this for a system using a File Allocation Table (FAT). Explain how this information is used and modified by a process making a new copy of an existing file on the disk. (9 marks)

Glossary

- file
- file attributes
- sequential access
- direct/random access
- magic number
- naming service
- storage service
- directory
- hierarchical (tree-structured) directories
- SFID, UFID, file descriptor/handle
- pathname
- relative v. absolute pathnames
- metadata, FCB
- File Allocation Table, FAT
- free space
- memory-mapped file

Reading

MOS2: 6.1-6.3

MOS3: 4.1-4.3

OSC/J: 10, 11.1-11.5